



A Coarse-to-Fine Semantic Conflict Detection System for Ex-Ante Davao City Ordinances Using Information Retrieval and Natural Language Inference

Ralph Paolo Dulce
Yahyah Odin

ATENEO DE DAVAO UNIVERSITY
SCHOOL OF ARTS AND SCIENCES
DAVAO CITY

June 2027

A Coarse-to-Fine Semantic Conflict Detection System for Ex-Ante Davao City Ordinances Using Information Retrieval and Natural Language Inference

An Undergraduate Thesis Submitted to the Faculty of the
Department of Computer Science
School of Arts and Sciences
Ateneo de Davao University

In Partial Fulfillment of the Requirements for the Degree of
Bachelor of Science in Computer Science

Ralph Paolo Dulce
Yahyah Odin

June 2027

RECOMMENDATION FOR ORAL DEFENSE

This undergraduate thesis entitled “**A Coarse-to-Fine Semantic Conflict Detection System for Ex-Ante Davao City Ordinances Using Information Retrieval and Natural Language Inference,**” prepared and submitted by **Ralph Paolo Dulce** and **Yahyah Odin** in partial fulfillment of the requirements for the degree of Bachelor of Science in Computer Science, has been examined and is recommended for Oral Examination.

MR. ADRIAN “OGS” ABLAZO
Adviser

APPROVAL SHEET

This undergraduate thesis entitled “**A Coarse-to-Fine Semantic Conflict Detection System for Ex-Ante Davao City Ordinances Using Information Retrieval and Natural Language Inference,**” prepared and submitted by **Ralph Paolo Dulce** and **Yahyah Odin**, has been approved by the Panel of Examiners on Oral Defense with a grade of _____.

[Panel Member 1 Name]

Panel Member

[Panel Member 2 Name]

Panel Member

[Panel Chair Name]

Panel Chair

Accepted and approved in partial fulfillment of the requirements for the degree of Bachelor of Science in Computer Science.

DR. JOHN DOE

Chair, Department of Computer Science

DR. JANE DOE

Dean, School of Arts and Sciences

Abstract

Under the constitutional doctrine of statutory preemption established in *Magtajas v. Pryce Properties* and Section 5(a) of Republic Act No. 7160 (Local Government Code of 1991), local government ordinances in the Philippines are strictly subordinate to national statutes. An ordinance that prohibits an activity authorized by national law or authorizes an activity prohibited by national law is *ultra vires* and void. In the Davao City Sangguniang Panlungsod, municipal legislative staff face substantial administrative and cognitive burdens when conducting manual ex-ante consistency reviews of proposed draft ordinances against decades of national laws and existing local measures. This challenge is intensified by the volume of positive law and the severe optical character recognition (OCR) degradation prevalent in historical municipal scanning archives.

To address these challenges, this study presents a coarse-to-fine semantic conflict detection architecture specifically designed for ex-ante legislative review and deployable within standard local government computing environments. The system decouples broad statutory candidate discovery from compute-intensive logical inference into a two-stage retrieve-then-entail pipeline. In Stage 1 (Coarse Candidate Retrieval), hybrid search combining sparse lexical matching (Okapi BM25) and dense neural embeddings (transformer bi-encoders) fused through Reciprocal Rank Fusion (RRF) with a statutory hierarchy priority safeguard searches a unified dual-corpus repository of 27,096 enactments (176,421 searchable provisions) to isolate a compact shortlist of top- k candidate provisions ($k \leq 50$). In Stage 2 (Fine Conflict Detection), candidate statute-ordinance pairs are classified through fine-tuned transformer cross-encoders into three directional normative categories: Contradiction, Entailment, or Neutral. High-risk contradictions are isolated using calibrated, cost-sensitive decision thresholds (τ^*) and explained through intrinsic token-level self-attention heatmaps.

The computational framework was evaluated against an authoritative 350-pair ground truth benchmark curated with authentic Davao City legislative syntax, realistic local administrative bodies, and statutory drafting conventions, validated by an evaluation panel of fifteen active Sangguniang Panlungsod legal researchers ($\kappa \geq 0.61$). Empirical findings demonstrate that the hybrid Stage 1 retrieval pipeline elevates candidate recall and mean reciprocal rank substantially over unconstrained lexical baselines while executing in approximately 300 milliseconds on standard edge CPUs. In Stage 2, fine-tuned cross-encoders eliminate the affirmative entailment bias inherent in unadapted foundation models, achieving high contradiction detection sensitivity without requiring cloud-hosted supercomputers or generative large language models. The resulting system establishes an

empirical baseline for low-resource Philippine legal informatics and provides an assistive decision-support tool for local legislative drafting.

Keywords: Semantic Conflict Detection, Ex-Ante Legislative Review, Statutory Pre-emption, Natural Language Inference, Information Retrieval, Cross-Encoder, Local Ordinances, Sangguniang Panlungsod.

Acknowledgement

Writing the entirety of this thesis was a massive undertaking, and we could not have crossed the finish line without the people who backed us up along the way.

To our professor, Ma'am Grace Tacadao, thank you for believing in us from day one. You took a chance by letting us be the only pair in the class, and you understood our vision immediately with our topic. Thank you for being so supportive and for being the exact kind of constructive critic we needed to actually make this paper work.

To our adviser, Sir Adrian "Ogs" Ablazo, thank you for saying yes the second we pitched this idea to you. Your chill energy was exactly what we needed to keep our sanity intact. Thank you for being incredibly receptive, for showing us exactly what was missing, and for being a safe space where we could share both our exciting progress and our random frustrations.

To the staff at the Davao City Sangguniang Panlungsod, especially Mr. Kevin B. Alfoja and Ms. Ma. Theresa A. Reyes, thank you for patiently answering our emails, catering to our concerns, and ultimately trusting us with the LISSP dataset. This system literally would not exist without your help.

To our defense panelists, thank you for your sharp, kind critiques and for pushing us to bring this vision to life. (And an advance thank you for hopefully not bombarding us with too many hard questions during the actual defense!)

To our families, friends, and classmates—thank you for the constant support, the patience when we were busy making this paper happen, and for believing we could survive our academics and this thesis. A special shoutout to Vince Demeer Valmores for the extra patience and encouragement behind the scenes.

Above all, to the Almighty God/Allah SWT, the ultimate source of strength, wisdom, and grace. Thank you for making all of this possible and for seeing us through to a successful end.

Table of Contents

Abstract	iv
Acknowledgement	vi
1 Introduction	1
1.1 Background of the Study	1
1.2 Problem Statement	2
1.3 Objectives of the Study	2
1.4 Significance of the Study	3
1.5 Scope and Limitations	4
1.6 Definition of Terms	5
2 Literature Review	8
2.1 Semantic Conflict Detection	8
2.1.1 Unstructured Legal Corpora	9
2.1.1.1 Legal Inflation	9
2.1.1.2 Limits of Manual Review	9
2.1.1.3 Semantic Ambiguity	9
2.1.2 Local Conflict Detection	10
2.1.2.1 Philippine Legal NLP	10
2.1.2.2 Logic and Conflict Research	10
2.1.2.3 Research Gaps	11
2.2 Philippine Legislative Domain	12
2.2.1 Hierarchical Constraints	12
2.2.2 Historical and Stylistic Heritage	14
2.2.3 Local Legislative Procedure and Language Variation	15
2.2.4 Ex-Ante Legislative Review	16
2.2.5 Legislative Consistency Checking	17
2.3 Digital Legislative Archives	18
2.3.1 National Sources	18
2.3.2 Local Sources	19
2.3.2.1 Archival Scanning and OCR Degradation	19
2.3.2.2 Repository Scarcity and Digital Isolation	20
2.4 Exact and Symbolic Methods	21
2.4.1 Boolean Information Retrieval	21

2.4.1.1	Challenges with BIR	21
2.4.2	Symbolic AI	22
2.4.2.1	Challenges with Symbolic AI	23
2.5	Graph-Based Architectures	23
2.5.1	Knowledge Graphs and GNNs	23
2.5.1.1	Graph Modeling and Multi-Hop Reasoning	24
2.5.1.2	Structural Vulnerabilities in Unstructured Archives	24
2.5.2	Abstract Meaning Representation	25
2.5.2.1	Semantic Graph Representations and Role Labeling	25
2.5.2.2	Sentence Boundary Constraints and Parser Degradation	26
2.6	Deep Learning and Transformers	26
2.6.1	Transformers and Self-Attention	26
2.6.2	Transformer Challenges	27
2.6.3	Suitability of Transformers	27
2.7	Single-Stage and Multi-Stage Architectures	28
2.7.1	Single-Stage Monolithic Pipelines	28
2.7.2	Two-Stage Retrieve-then-Entail Architectures	29
2.7.3	Comparative Architectural Trade-offs	29
2.8	Information Retrieval	30
2.8.1	Information Retrieval Challenges in Statutory Corpora	30
2.8.2	Relevance in Retrieval	31
2.8.3	Long-Context Retrieval	32
2.8.4	Retrieval and Fusion Strategies	33
2.8.5	Retrieval Benchmarks	34
2.8.6	Candidate Retrieval Models	35
2.9	Natural Language Processing	37
2.9.1	NLP Challenges	37
2.9.2	From NLP to NLI	38
2.9.3	NLI Benchmarks	38
2.9.4	Candidate NLI Models	39
2.10	Ground Truth	41
2.10.1	Ground Truth in NLI	41
2.10.2	Annotation and Reliability Metrics	42
2.10.2.1	Sub-Threshold Reliability	43
2.10.3	Ground Truth in Legal NLP	44
2.10.4	Multi-Sentence and Cross-Document Limits	45
2.10.5	Dataset Partitions	46
2.10.6	Minimum Detectable Effect	47
2.11	Explainable AI	47
2.11.1	Intrinsic Explainability	47
2.11.2	Edge Deployment Feasibility	48
2.11.3	Explanation Quality Metrics	49
2.11.4	Explainability in Multi-Stage Pipelines	49
2.12	Theoretical Framework	50
2.12.1	Statutory Retrieval (Team AIIR)	50
2.12.2	Multi-Stage Retrieval (Team JNLP)	51

2.12.2.1	Pre-Retrieval	53
2.12.2.2	Model Fine-Tuning	53
2.12.2.3	Result Ensemble	53
2.12.3	Textual Entailment (Team KIS)	54
2.12.4	Theoretical Synthesis	54
3	Methodology	57
3.1	Conceptual Framework	57
3.1.1	Phase 1: Offline Foundation and Optimization	59
3.1.2	Phase 2: Live System Workflow	59
3.2	Research Design	61
3.3	Corpus Acquisition and Exploratory Analysis	61
3.3.1	Corpus Composition	61
3.3.2	Corpus Characteristics and Inherent Biases	64
3.3.3	National Statutory Corpus Analysis	66
3.3.4	Local Legislative Corpus Analysis	79
3.3.5	Dual Statutory Corpus Integration	83
3.4	Preprocessing and Normalization	86
3.4.1	Draft Ingestion and Boundary Parsing	86
3.4.2	Boilerplate Filtering and Metadata Headers	89
3.4.3	Archive Digitization and Ingestion Pipeline	91
3.4.4	Substantive Legislative Body Extraction	93
3.4.5	Automated Quality Audit and Conservative Normalization	95
3.4.6	Advisory Language Quality Review Architecture	95
3.4.7	Harmonic Interpretation and Exceptions	96
3.5	Ground Truth Benchmark	96
3.5.1	Sample Size and Statistical Power	97
3.5.2	Dataset Partitioning	98
3.5.3	Asymmetric Pairing and Hypothesis Formulation	98
3.5.4	Class Balancing and Adversarial Injection	99
3.5.5	Stratified Difficulty Tiers	100
3.5.6	Expert Legal Annotation Panel	102
3.5.7	Annotation Codebook and Decision Rules	103
3.5.8	Annotation Workload Allocation	105
3.5.9	Inter-Rater Reliability	106
3.6	Hardware and Computing Setup	107
3.7	Pipeline Architecture	109
3.7.1	Baseline Selection	110
3.7.2	Stage 1 Candidate Retrieval	111
3.7.3	Stage 1 Retrieval Safeguards	111
3.7.4	Stage 2 Conflict Detection	113
3.8	Model Selection and Optimization	113
3.8.1	Model Selection and Screening	114
3.8.2	Overfitting Controls and Calibration	115
3.8.3	Hyperparameter Tuning	116
3.9	Evaluation Metrics	117

3.9.1	Stage 1 Retrieval Metrics	117
3.9.2	Stage 2 Inference Metrics	118
3.9.3	Statistical Hypothesis Testing	121
3.10	System-Level Evaluation	124
3.11	Prototype and User Interface	128
3.12	Reproducibility and Repository	129
3.13	Methodological Constraints	129
4	Results and Discussion	132
4.1	Archival OCR and VLM Evaluation	132
4.1.1	Architectural Paradigms and Candidate Systems	133
4.1.2	The Multi-Dimensional Evaluation Framework	133
4.1.3	The Insufficiency of Ordinary Spellchecking	135
4.1.4	Stratified Archival Reference Sampling	136
4.1.5	Comparative Architectural Trade-Off Analysis	136
4.1.6	Defensible Selection Rationale and Hardware Boundaries	138
4.1.7	Empirical Limitations and Archival Ceilings	139
4.2	Stage 1 Candidate Retrieval	139
4.2.1	Full-Corpus Sensitivity Benchmark	140
4.2.2	Bi-Encoder Retrieval Benchmark	141
4.2.3	Retrieval and Fusion Ablation	142
4.3	Stage 2 Conflict Detection	143
4.3.1	Zero-Shot Candidate Screening	144
4.3.2	Supervised Fine-Tuning Benchmark	146
4.3.3	Loss Dynamics and Convergence	148
4.3.4	Sensitivity Analysis and Regularization	151
4.4	System Integration and Usability	152

List of Figures

2.1	Statute Law Retrieval Pipeline Using Dense Bi-Encoder Architecture . . .	51
2.2	Multi-Stage Statutory Retrieval Pipeline With Candidate Pruning and En- sembling	52
2.3	Statutory Textual Entailment Pipeline Using Bidirectional Cross-Encoder .	55
3.1	Information Flow of the Conflict Detection Architecture	58
3.2	Sample Statutory Enactment Webpage from the Lawphil Repository	62
3.3	Sample Archival Document from Davao City Corpus with Scanning Noise .	65
3.4	Composition of the Philippine National Statutory Corpus	68
3.5	Historical Succession and Volume of Statutory Instruments	69
3.6	Distribution of National Statutes across Consolidated Macro Domains . . .	72
3.7	Two-Dimensional Latent Semantic Topic Space of National Statutes	74
3.8	Top Salient <i>c-TF-IDF</i> Keyword Profiles across Macro Domains	75
3.9	Cross-Allocation Matrix of National Statutory Instruments	76
3.10	Temporal Distribution and Latent Semantic Space of Local Corpus	79
3.11	Enactment Length Asymmetry and Provision-Level Subword Token Densities	84
3.12	Cumulative Provision Compliance Curves and Attention Window Evaluation	85
4.1	Stage 1 Candidate Retrieval Performance across Difficulty Tiers	143
4.2	Stage 2 Cross-Encoder Loss Dynamics and Convergence Trajectories	150

List of Tables

2.1	Interpretation Ranges for Fleiss’ Kappa Reliability Metric	43
3.1	Comparative Statutory Coverage of the National Legal Archive	67
3.2	Consolidated Macro Legal Domains of the National Statutory Corpus . . .	71
3.3	Empirical Holdout Validation Metrics for the National Topic Modeler . . .	78
3.4	Functional Typology of the Davao City Local Legislative Corpus	81
3.5	Empirical Holdout Validation Metrics for the Local Topic Modeler	82
3.6	Comparative Structural Parameters of the Dual Statutory Search Space . .	83
3.7	Stratified Jurisprudential Difficulty Tiers for Ground Truth Benchmark . .	101
3.8	Legal Annotation Codebook and Ground Truth Decision Rules	103
3.9	Hardware Configurations and Operational Roles across Computing Envi- ronments	108
3.10	Architectural Retrieval Safeguards Against Preemption Failure Modes . . .	112
3.11	Hyperparameter Tuning Configuration and Grid Search Space	116
3.12	Contingency Matrix of Discordant Predictions for McNemar’s Test	122
3.13	Landmark Davao City Municipal Jurisprudential Evaluation Suite	127
4.1	Stratified Reference Sampling Matrix for Archival OCR Evaluation	137
4.2	Comparative Multi-Dimensional Evaluation of Archival OCR Architectures	138
4.3	Empirical Retrieval Sensitivity Benchmark across National Database . . .	140
4.4	Performance Comparison of Candidate Dense Bi-Encoder Architectures . .	141
4.5	Performance Comparison of Candidate Stage 1 Retrieval Architectures . .	142
4.6	Zero-Shot Screening and Affirmative Bias Evaluation across 28 Encoders .	145
4.7	Supervised Fine-Tuning Performance across Shortlisted Cross-Encoders . .	147
4.8	Epoch-by-Epoch Diagnostic Training Dynamics for Leading Candidate . .	150

List of Listings

3.1	Code Snippet of the Statutory Corpus Acquisition and Provision Extraction Pipeline	63
3.2	Sample Raw JSONL Ingestion Record from the National Statutory Archive (Republic Act No. 11035)	87
3.3	Code Snippet of the Operative Provision Extractor and Boilerplate Filter .	88
3.4	Pseudocode of Section Context Prepending for Granular Provision Chunking	89
3.5	Sample Model-Ready Statutory Chunk and Stage 2 Input Sequence with Context Prepending	90
3.6	Adaptive State Machine for Substantive Legislative Body Extraction	94
3.7	Sample Evaluation Pair from the Ground Truth Benchmark Dataset	104
3.8	Pseudocode of the Stage 1 Reciprocal Rank Fusion Hybrid Provision Retrieval Pipeline	109
3.9	Pseudocode of the Stage 2 Cost-Sensitive Cross-Encoder Inference and Section-Level Conflict Aggregation	120

List of Equations

2.1	Boolean Query Set Intersection Formulation	21
2.2	First-Order Deontic Prohibition Formulation	23
2.3	Okapi BM25 Probabilistic Relevance Scoring Function	33
2.4	Dense Semantic Cosine Proximity Relevance Function	33
2.5	Cohen’s Pairwise Kappa Reliability Metric (κ)	42
2.6	Fleiss’ Generalized Multi-Rater Kappa Reliability Metric (κ)	42
3.1	Cross-Encoder Joint Sequence Input Representation	60
3.2	Class-Based Term Frequency-Inverse Document Frequency (<i>c-TF-IDF</i>) . .	70
3.3	Macro Legal Domain Centroid and Softmax Probability	70
3.4	Max-Pooled Sub-Chunk Contradiction Probability	86
3.5	Context-Prepended Statutory Chunk Representation	90
3.6	Archival Inset Perimeter CropBox Trimming	92
3.7	Substantive Body Boundary Extraction Confidence Score	94
3.8	Annotator Workload Allocation Formulation	106
3.9	Reciprocal Rank Fusion (RRF) Scoring	109
3.10	Recall at Rank k (Recall@ k)	117
3.11	Mean Reciprocal Rank at Rank k (MRR@ k)	117
3.12	Class-Specific Precision, Recall, and F_1 -Score	118
3.13	Macro-Averaged F_1 -Score Across Logical Classes	118
3.14	Cost-Sensitive Contradiction F_2 -Score	119
3.15	Optimal Contradiction Decision Threshold Calibration (τ^*)	119
3.16	Section-Level Max-Pooled Conflict Probability	120
3.17	McNemar’s Chi-Squared Test with Edwards’ Continuity Correction	122
3.18	Friedman Two-Way Analysis of Variance by Ranks (χ^2_F)	123
3.19	Cascaded End-to-End Joint Conflict Detection Hit Indicator	124
3.20	Section-Level Specificity for Full-Document Stress Testing	125
4.1	Character Error Rate (CER) and Word Error Rate (WER)	134
4.2	Normalized Edit Distance (NED)	134
4.3	Domain-Specific Legislative Component Accuracy	135
4.4	Multi-Class Cross-Entropy Training Loss Function	149

Chapter 1

Introduction

1.1 Background of the Study

A functional legal system requires strict consistency so that citizens can clearly understand their rights and obligations [61]. However, the sheer volume of modern legislation has grown so rapidly that human practitioners can no longer track every active rule [161]. As governing bodies continuously pass new regulations to address emerging issues, they frequently enact laws that inadvertently contradict older ones [94]. Because the total number of documents is overwhelming, identifying these hidden conflicts through manual review is computationally impossible for legal staff. This issue is particularly severe in local government units, where historical ordinances are often scattered across fragmented files rather than maintained in a structured digital database [49].

Within the Philippine legal framework, local governments like Davao City operate under delegated authority and can only enact laws permitted by the national government [70]. Due to this structural limitation, the *Magtajas v. Pryce* Doctrine strictly prohibits any local municipal ordinance from contradicting established national statutes [187]. Consequently, in the local context, city councils must rigorously audit every newly proposed draft ordinance to ensure it logically aligns with existing laws before final approval. If reviewers overlook a normative conflict, such as a local ordinance prohibiting an activity that a national law explicitly permits, the resulting local legislation is legally void. Passing invalid laws creates severe jurisdictional disputes, wastes administrative resources, and ultimately compromises the constitutional rights of the public.

To address these legal risks, the legislative process must utilize an *ex-ante* evaluation: a term derived from Latin meaning 'before the event.' Unlike an *'ex-post'* review, which happens after a law is already in effect and often after legal damage has occurred, an *ex-ante* audit acts as a first line of defense. By checking a proposed ordinance for conflicts while it is still in the drafting phase, the Sangguniang Panlungsod can proactively align local rules with national statutes. This preventative approach is consistent with modern Philippine standards like Republic Act No. 11032 or the Ease of Doing Business and Efficient Government Service Delivery Act of 2018, which encourages government agencies to perform regulatory assessments to ensure that new rules do not create unnecessary or

conflicting burdens on the public.

The decentralized drafting process of local ordinances further complicates manual review because the resulting texts are often highly unstructured [166]. City councils frequently splice and amend different drafts, producing verbose documents that blend formal legal English with regional dialects like Cebuano and administrative "Taglish" [53]. Furthermore, because municipal archives heavily rely on digitized PDF scans, the text suffers from severe Optical Character Recognition (OCR) errors. While OCR is typically 99 percent accurate for clean English text, studies show accuracy can drop to 78 percent for low-resource languages, with Cebuano specifically seeing word error rates around 6.9 percent [137]. These errors cause an "error cascade" that shatters standard vocabulary, making it difficult for traditional search engines to semantically differentiate between a statutory obligation and a strict prohibition.

In light of these challenges, the proponents seek to develop a specialized application capable of retrieving relevant legal documents and comparing them against proposed ordinances to detect potential conflicts.

1.2 Problem Statement

This study is conducted to investigate the use of a multi-stage architecture to automate the manual auditing of proposed local legislation in Davao City.

Specifically, the study aims to answer the following research questions:

- What are the specific linguistic parameters and structural markers that define a "semantic conflict" between a proposed local ordinance and existing local or national laws?
- How can a Coarse-to-Fine system be designed to efficiently retrieve relevant national and local statutes before performing deep logical inference on constrained LGU hardware?
- Which combination of a hybrid Information Retrieval (IR) algorithms and Natural Language Inference (NLI) models provides the highest accuracy when processing the specific hierarchy of Philippine legal texts?
- What is the empirical performance of the proposed pipeline when evaluated against a manually curated Ground Truth dataset of local and national conflict pairs?
- How usable and viable is the final system interface for key stakeholders within the Davao City Sangguniang Panlungsod?

1.3 Objectives of the Study

The primary goal of this research is to develop a computationally efficient, two-stage semantic conflict detection system for proposed Davao City ordinances.

Specifically, this study aims to:

- Identify the specific linguistic characteristics and structural constraints of Davao City’s legislative archive that contribute to semantic ambiguity and retrieval failure;
- Design a Coarse-to-Fine computational pipeline that decouples document retrieval from logical inference to optimize resource usage for local deployment;
- Implement a hybrid Information Retrieval (IR) algorithm and Natural Language Inference (NLI) system specifically calibrated for the Philippine legal hierarchy;
- Evaluate the system’s reliability by comparing its automated detections against a verified set of known legal conflicts to measure how accurately it identifies contradictions; and
- Develop a user-friendly web interface for the Davao City Sangguniang Panlungsod designed to accept uploaded draft documents and provide real-time conflict alerts.

1.4 Significance of the Study

The Davao City Sangguniang Panlungsod stands as the primary institutional beneficiary of this automated consistency checking system. By transitioning from manual statutory review to computational auditing, the tool saves legislative staff hundreds of hours of manual cross-referencing across multi-volume legislative digests. This computational intervention actively prevents the city from enacting legally vulnerable measures that invite judicial invalidation, public injunctions, or administrative deadlocks under the doctrine of statutory preemption. It equips the local legislative council with the operational capacity to draft regulatory measures with a heightened degree of legal certainty and procedural speed. Ultimately, the system modernizes the municipal legislative lifecycle by embedding rigorous logical safeguards directly into the First Reading drafting phase.

The study provides substantial operational value to legal researchers, committee secretaries, and legislative bill drafters. In contemporary municipal practice, legal researchers must examine extensive regulatory provisions across disparate national codes without centralized semantic indexing tools. The proposed system provides an automated cross-statutory audit trail, pinpointing candidate national statutes and fellow city ordinances that govern the same regulatory sphere. By surfacing relevant provisions and generating token-level self-attention attribution heatmaps, the architecture substantially mitigates cognitive fatigue and reading oversight during high-volume legislative sessions. Consequently, legal officers can conduct evidence-grounded preliminary reviews and focus their analytical expertise on high-level statutory interpretation.

Beyond the immediate city council, this research establishes an edge-deployable computational blueprint for Local Government Units (LGUs) nationwide and oversight agencies such as the Department of the Interior and Local Government (DILG). In compliance with the Ease of Doing Business and Efficient Government Service Delivery Act of 2018 (Republic Act No. 11032), local governments are mandated to perform systematic Regulatory Impact Assessments to ensure local regulations do not create redundant or conflicting business burdens. Demonstrating that an automated consistency checking pipeline operates efficiently on standard consumer-grade office desktop computers proves

that advanced artificial intelligence can be deployed within budget-constrained municipal environments without recurring cloud subscriptions or specialized computing clusters.

The research also contributes foundational advancements to the Philippine Natural Language Processing and legal informatics community. Prior domestic legal technology research has concentrated predominantly on national appellate court case summarization and judicial outcome forecasting, leaving local statutory conflict detection largely unaddressed. This investigation establishes an empirical baseline for processing low-resource, localized Philippine statutory corpora characterized by multi-generational scanning noise, administrative Taglish, and specialized legislative syntax. Furthermore, the empirical findings demonstrate that lightweight, fine-tuned cross-encoders can effectively overcome affirmative entailment bias and resolve directional statutory contradictions on resource-constrained local hardware.

Finally, the general public and the local business community benefit directly through the promulgation of coherent, constitutionally sound public policy. When municipal ordinances clash with superior national statutes, citizens and local enterprises bear the brunt of legal uncertainty, facing conflicting compliance directives or unauthorized penal sanctions. By serving as an ex-ante procedural safety net, the system helps ensure that enacted city ordinances harmonize with national standards and existing local legislation. This statutory coherence fosters institutional transparency, strengthens regulatory predictability, and upholds the integrity of the local rule of law.

1.5 Scope and Limitations

The scope of this study focuses on developing an automated semantic conflict detection system to evaluate draft ordinances from Davao City against a dual statutory search space. This collection encompasses national statutory enactments compiled to the extent accessible from public legal repositories, paired alongside the historical archive of enacted Davao City local ordinances collected to the fullest extent possible from local legislative records. Operationally, the system audits draft measures across two complementary dimensions: vertical statutory preemption, ensuring local provisions do not contradict superior national statutes under the *Magtajas* doctrine and Republic Act No. 7160; and horizontal consistency, verifying that proposed clauses do not contradict, duplicate, or inadvertently cause implied repeals of active local ordinances.

To establish an administratively grounded and computationally feasible engineering boundary, several explicit delimitations govern this research. The statutory knowledge base is strictly bounded to codified national legislative instruments and enacted local ordinances. To preserve retrieval index compactness and prevent ambiguous litigation records from distorting similarity rankings, judicial case law from Supreme Court decisions and departmental administrative circulars are excluded from the primary index. Furthermore, the local ordinance corpus reflects the practical realities of local legislative record-keeping, where recent council terms benefit from digital database records while earlier decades rely on curated landmark ordinances due to physical archive gaps. The system evaluates drafts strictly against the combined corpus gathered through public portals and direct

institutional coordination with the Sangguniang Panlungsod, leaving the comprehensive optical character recovery of older paper files as an ongoing archival task.

The linguistic scope of the system operates exclusively on English-language enactments, following Section 20, Chapter 5, Book I of Executive Order No. 292, which establishes that the English text controls in formal statutory interpretation. While local drafts occasionally incorporate regional Cebuano phrases or administrative Taglish, restricting the primary reasoning pipeline to English prevents semantic drift and avoids translation distortion. At the inference tier, the system evaluates proposed ordinance sections against retrieved candidate provisions using asymmetric span-level pairing (1-vs- N). The pipeline deliberately excludes multi-hop reasoning across multiple independent statutes simultaneously. Detecting intricate legal contradictions that emerge only when synthesizing two or more disparate national statutes exceeds the computational capacity of standard workstation hardware and is designated as a recommendation for future research.

The computational architecture is intentionally restricted to an open-weight transformer encoder envelope capable of local deployment on standard Local Government Unit (LGU) desktop computers, specifically targeting sub-600M parameter models requiring less than 2 GB of inference VRAM. Massive autoregressive large language models and proprietary cloud APIs are excluded to maintain institutional data privacy, avoid recurring cloud API expenditures, and guarantee reliable offline operation. Within this architecture, the system operates strictly as a diagnostic decision-support tool providing transparent, white-box explainability, such as attention heatmaps and calibrated confidence scores, without rewriting or generating statutory text. Finally, the operational scope of the inference engine is delimited to direct normative and modal contradictions arising from explicit textual collisions, such as local prohibitions of federally authorized activities, unauthorized procedural preconditions, or penalty ceiling infractions. The framework deliberately does not model longitudinal legislative histories, historical tax rate compounding, or multi-year financial calculations across unprovided statutory baselines, as observed in complex judicial tax disputes [191]. Modeling such multi-hop temporal construction and arithmetic accounting exceeds the ex-ante diagnostic scope of cross-encoder architectures and is preserved as an avenue for specialized fiscal analytics.

1.6 Definition of Terms

Bi-Encoder. A dual-tower transformer architecture that maps queries and candidate statutory provisions independently into a shared dense vector embedding space, enabling rapid semantic similarity search via vector dot products or cosine distance.

Coarse-to-Fine Architecture. A two-stage decoupled computational framework that pairs a high-throughput, broad candidate retrieval module (coarse stage) with a compute-intensive pairwise sequence classification model (fine stage) to manage extensive textual collections on bounded hardware.

Contradiction. In Natural Language Inference, a directional semantic relationship where the truth of the hypothesis is incompatible with the regulatory mandate of the premise, indicating a direct normative conflict or statutory collision.

Cross-Encoder. A single-tower transformer architecture that ingests concatenated premise-hypothesis pairs into a unified self-attention matrix, enabling full bidirectional cross-attention across all tokens for granular semantic classification.

Entailment. In Natural Language Inference, a relationship where the substantive directive of the hypothesis naturally follows from, conforms with, or is explicitly authorized by the governing statutory premise.

Ex-Ante Legislative Review. The proactive legal and technical evaluation of proposed legislative drafts, ordinances, or administrative rules prior to their formal enactment, designed to detect statutory collisions and constitutional defects while the text remains amenable to revision.

Explainable Artificial Intelligence (XAI). Computational methods that provide human-interpretable rationales for machine learning decisions; in this system, operationalized through token-level self-attention attribution heatmaps highlighting specific conflicting text spans.

Fleiss' Kappa (κ). A statistical measure of inter-annotator agreement among three or more raters assessing categorical data, utilized to validate consensus reliability across the fifteen Sangguniang Panlungsod legal researchers.

Horizontal Inconsistency. A normative clash, redundancy, or unintended implied repeal occurring between a newly drafted local ordinance and an existing, fellow local ordinance enacted within the same legislative jurisdiction.

Information Retrieval (IR). The computational discipline and algorithmic process of searching, filtering, and ranking relevant candidate documents or legal provisions from extensive unstructured text collections in response to a query.

Local Ordinances. Enacted legislative measures promulgated by the Davao City Sangguniang Panlungsod under devolved legislative powers that carry the force and effect of law within the territorial jurisdiction of the city.

Magtajas v. Pryce Doctrine. The foundational Philippine constitutional principle establishing that local government ordinances are strictly subordinate to national statutes, rendering any ordinance that permits what an Act of Congress forbids, or forbids what an Act of Congress permits, *ultra vires* and void.

Natural Language Inference (NLI). A core natural language processing task that evaluates the directional semantic relationship between a premise text and a hypothesis text, categorizing the pair into Entailment, Contradiction, or Neutral.

Neutral. In Natural Language Inference, a relationship where the truth or enforcement of the hypothesis neither follows from nor contradicts the regulatory command of the statutory premise, indicating thematic or regulatory independence.

Normative Conflict. In legal theory and formal deontic logic, a condition where two or more legally valid enactments impose mutually incompatible obligations, prohibitions, or permissions upon the same subject under identical circumstances.

Optical Character Recognition (OCR) Noise. Textual distortions, typographical errors, broken character tokens, and structural formatting anomalies introduced when physical paper documents are scanned and converted into machine-readable text.

Reciprocal Rank Fusion (RRF). An unsupervised rank aggregation algorithm that merges ranked candidate lists from disparate retrieval systems (such as lexical BM25 and dense bi-encoders) based on reciprocal rank positions without requiring calibrated score normalization.

Sangguniang Panlungsod (SP). The unicameral legislative council of Davao City, vested by Republic Act No. 7160 with the corporate and police authority to enact local ordinances and approve resolutions.

Semantic Conflict Detection. The computational formulation and automated identification of directional normative contradictions between legislative provisions using continuous vector representations, lexical indexing, and natural language inference models.

Statutory Preemption. The constitutional and administrative principle under which superior national legislation invalidates, overrides, or limits the regulatory scope of local enactments governing the same subject matter.

Vision-Language Model (VLM). A multimodal deep learning architecture that combines visual processing backbones with transformer language decoders to transcribe document page images into structured text while preserving layout and alphanumeric precision.

Chapter 2

Literature Review

2.1 Semantic Conflict Detection

In jurisprudence and legal philosophy, the fundamental requirement for a functional legal order is internal consistency, where positive enactments provide unambiguous guidance without subjecting citizens to contradictory commands [61]. In legal theory and formal deontic logic, this phenomenon is strictly defined as a *normative conflict*—a structural clash between legal norms wherein an action is simultaneously commanded, prohibited, or permitted across distinct enactments, violating the principle of non-contradiction [111, 3]. Specifically, normative conflict occurs when two or more valid laws impose incompatible legal duties on the same subject under the same conditions, as when one law mandates an act that another law forbids [3, 88]. When statutory systems generate conflicting obligations, citizens cannot conform their conduct to the law without incurring liability under an opposing mandate, destabilizing legal predictability and public trust. Consequently, establishing systematic mechanisms to identify and resolve normative collisions represents a vital safeguard for the integrity of any rule-of-law framework.

However, in real-world legislative governance and local government archives, statutory enactments do not exist as pre-formalized formulas in symbolic logic. Rather, statutes are drafted and promulgated as unstructured, multi-generational natural language texts characterized by high syntactic variance, legal jargon, and implicit exceptions [94]. Consequently, from a computational perspective, this research formulates and operationalizes the challenge as *Semantic Conflict Detection*. By translating abstract deontic clashes into a computable natural language processing problem, semantic conflict detection leverages continuous vector representation spaces and cross-encoder Natural Language Inference (NLI) to identify directional contradictions between local drafts and governing statutory provisions without relying on brittle, hand-crafted legal ontologies [175, 88]. As Ruhl and Katz [161] established, the volume of positive law expands far faster than human auditing capacity, making automated semantic auditing an operational necessity for modern legislative bodies.

2.1.1 Unstructured Legal Corpora

Managing unstructured legal text—text not stored in structured databases—represents a fundamental maintenance challenge for statutory systems. Practical automated auditing requires addressing rapid textual accumulation alongside the linguistic complexity of legislative drafting. These structural realities drive the research transition toward neural retrieval and inference tools for statutory verification.

2.1.1.1 Legal Inflation

We refer to “Legal inflation” as to how the number of laws and their connections keep growing. Katz et al. argue that the sheer volume of laws has made it too difficult for humans to check for consistency manually [94]. As new rules are made to keep up with technology, they create “regulatory debt,” where new laws might conflict with older ones without anyone noticing. This is a big problem for local governments, because unlike national laws, local ordinances are often not organized or scanned properly. Dias et al. [49] notes that turning these laws into machine-readable data is hard because they are often scattered across different files, from formal resolutions to quick executive orders.

The “search space” for finding conflicts is therefore a messy collection of different texts rather than a neat database. Zhong et al. [228] describe this as a “dynamic consistency problem” where facts and rules are always changing. Because this data is not organized, government agencies cannot be proactive. They only realize there is a conflict after something goes wrong. This is why the field has moved from old AI systems that use simple rules to new “neural models” that can learn from data. These systems are needed to process huge collections of laws that are too big for the human brain to handle [27].

2.1.1.2 Limits of Manual Review

Scanning documents into a computer is no longer the main challenge; the primary operational hurdle is now checking if those documents agree with each other. As legislative bodies pass laws faster, they create a massive pile of unorganized data. The main difficulty is telling the difference between a “good” overlap (where a law is meant to cover multiple areas) and a “bad” conflict (where a law is written poorly and contradicts another). Recent surveys show that standard search engines usually fail to understand these subtle distinctions [9]. As a result, many conflicts stay hidden until someone gets sued. This moves the burden of finding conflicts from the people writing the laws to the judges in court [201].

2.1.1.3 Semantic Ambiguity

Old computer search systems look for exact word matches. However, this fails to catch the real meaning of human language. Legal language is like a specialized code. Words like “shall,” “must,” or “may” have very strict meanings that keyword searches often mix up. Chalkidis et al. [27] show that this “ambiguity” makes exact-match techniques fail at complex legal tasks. For example, a law banning “motor vehicles” might conflict with a law allowing “electric scooters.” A keyword search for “vehicle” will not find the problem if the words do not match exactly.

Legal consistency often depends on "implied" rules—things the law suggests but doesn't say directly. Rigano et al. [156] argue that finding these contradictions requires a system that evaluates logical relationships rather than surface string comparisons. "Transformer" AI models are an example of a solution that can solve this gap by creating mathematical representations (vectors) that capture deep meaning [227], which shall be discussed in later sections. However, there is a trade-off. Sovrano et al. [180] argue that while these models are very accurate, their "black box" nature is a problem because they don't explain *why* they found a conflict. In a legal setting, the reasoning is just as important as the detection [172].

2.1.2 Local Conflict Detection

While researchers around the world are building conflict detection tools using benchmarks like LexGLUE, these tools are still in the early stages in less prominent countries like the Philippines [27]. Since the Local Government Code of 1991, cities like Davao have created thousands of ordinances that are not well-organized [70]. Although the math for detecting conflicts exists in general AI research, it has rarely been applied to Philippine laws.

2.1.2.1 Philippine Legal NLP

When Philippine AI research does focus on law, it is usually about organizing data rather than deep reasoning. Early work by Virtucio et al. [209] used basic classifiers to predict Supreme Court decisions, while Sagum et al. [163] used AI just to summarize national court rulings.

More recently, studies have started using "Transformer" models, a method that shall be discussed in later sections, but they are still used mostly for search. Roque et al. showed in 2024 [158] that "dense embedding" models are much better than keyword searches at finding specific tax laws. However, these embedding models are basically just advanced search engines. They can find two laws about "business taxes" because they are topically similar, but they lack the logic layers needed to tell if one law actually contradicts the other.

Looking at these trends, there is a clear gap in the research. The Philippine AI community has successfully used computers to categorize, summarize, and search national laws. They have also proven that NLI models can understand logical contradictions in local dialects for fact-checking [209, 163, 158, 135]. However, actually using these advanced NLI models to find conflicting rules in the messy, unorganized ordinances of a local city like Davao remains almost entirely unexplored.

2.1.2.2 Logic and Conflict Research

In the Philippine AI community, the tools needed for conflict detection, specifically Natural Language Inference (NLI), are already being used, but mostly for other tasks. Recent local studies have used NLI to find logical contradictions in things like online news and misinformation. For example, Nuñez et al. proved that these AI models can handle the logic needed for Tagalog and "Taglish" (a mix of Tagalog and English) [135].

However, when researchers try to analyze local government records, they usually go back to traditional methods. Most document analysis in the Philippines still relies on "Boolean" search (keyword matching) and simple rule-based AI. As recent Philippine studies show, these old systems are "lexically fragile" [129]. They fail if two laws use different words for the same thing (like "fireworks" vs. "pyrotechnics") and they break down when there are typos or scanning errors (OCR noise) in the digital files. Trying to use manual rules to find clashes in thousands of messy city ordinances is simply too hard for a computer or a person to handle alone.

2.1.2.3 Research Gaps

While the Philippine AI community has successfully automated the search and summarization of national laws, a significant gap remains in using deep logic models to audit local ordinances. Research indicates this gap persists because local legislation is substantially more challenging for automated systems to process than national statutory compilations.

The primary structural barrier to local legislative auditing is the critical dataset gap, characterized by the complete absence of curated ground-truth datasets for Philippine local enactments. Unlike national statutes that benefit from centralized repositories, local ordinances represent low-resource textual environments [148]. Standard legal NLP models are trained on massive Western corpora like LexGLUE, but there are no equivalent standardized benchmarks providing annotated pairs of conflicting local ordinances [62, 132]. Without a standardized benchmark to measure progress, local legal informatics research has historically struggled to move beyond superficial keyword search to deductive logical reasoning [62].

Compounding the lack of training data is a pronounced computational resource gap within Local Government Units (LGUs). Deductive reasoning models, particularly deep cross-encoders, demand substantial compute and memory resources that local government offices typically do not possess. Studies on Philippine public administration indicate that local digital adoption remains constrained by limited funding, scarce machine learning personnel, and legacy office desktop hardware [62]. In many local legislative councils, fewer than twenty percent of legislative personnel have received formal ICT training, with most offices still relying on paper records and basic text editors rather than structured databases [62]. This lack of digital infrastructure makes it difficult to justify complex models that would be too computationally demanding to execute on actual LGU workstations [85].

At the textual level, local legislative conflict detection faces a complex sociolinguistic gap. Local ordinances exhibit higher linguistic heterogeneity than national statutes because local drafting frequently blends formal legal English with regional dialects like Cebuano and administrative Taglish [53]. Standard NLP tokenizers designed for monolingual English often break down when encountering localized legal idioms or archaic Spanish loanwords embedded in Philippine local codes [123]. While recent studies confirm that neural inference models can parse regional dialects in informal social media text, extending this capacity to dense, multi-generational local enactments requires specialized domain adaptation that remains largely unexplored in local literature [30, 73].

Finally, the slow adoption of computational consistency checking is driven by an institutional trust gap across the legal profession. As emphasized by Supreme Court Senior Associate Justice Marvic Leonen, the legal community remains cautious regarding artificial intelligence because conventional deep learning models function as opaque black boxes that cannot articulate their underlying reasoning [108]. Legal researchers and legislative counsel legitimately fear that automated tools could introduce subtle statutory errors, leading to the enactment of void legislation [35]. Consequently, academic research has historically favored low-risk document indexing and retrieval systems, leaving deep semantic conflict detection unaddressed in local governance [11].

2.2 Philippine Legislative Domain

The Philippine legal system is a unitary state. This means that Local Government Units (LGUs) only have powers that are given or delegated to them by the national government [70]. This setup creates a strict top-down hierarchy that limits how much freedom cities like Davao City have when making laws. Any computer system built to check these laws must understand this specific legal environment. This section explains the barriers in Philippine local law. First, it looks at the strict rules that make local ordinances follow national laws. Second, it examines the difficulties and language issues in the local drafting process. Finally, it reviews the current methods used to check for consistency, which shows why we need to move toward automated auditing.

2.2.1 Hierarchical Constraints

The legal system in the Philippines depends on the fundamental rule that national laws always override local enactments. To understand the legal constraints governing municipal drafting, automated conflict detection systems must reflect the structural principles that subordinate local ordinances beneath national statutes.

The Philippine legal system operates under a unitary constitutional framework, establishing a rigid hierarchy where national laws strictly control subordinate enactments [70]. At the apex stands the 1987 Constitution, followed by national statutes enacted by Congress, international treaties, presidential executive issuances, and municipal ordinances. Under this constitutional structure, no Local Government Unit can enact measures that contradict an Act of Congress. The doctrine of national preemption establishes an absolute boundary: whenever the national legislature enacts comprehensive statutory standards governing a given regulatory domain, local governments are legally precluded from enacting conflicting or diluting municipal rules [70].

The primary organic charter governing municipal lawmaking is the Local Government Code of 1991 (Republic Act No. 7160). Section 16 of the Code, known as the General Welfare Clause, delegates police powers to local legislative councils to enact ordinances promoting the health, safety, and general prosperity of their constituents, provided such enactments remain strictly within statutory boundaries [70]. Because Davao City is classified as a Highly Urbanized City (HUC), it operates autonomously from provincial administrative supervision, meaning its ordinances are not reviewed by a provincial board.

This administrative independence places the entire burden of ex-ante consistency auditing directly upon the Davao City Sangguniang Panlungsod before measures proceed to final enactment [70].

The definitive judicial test governing the validity of local legislation is established by the Supreme Court in the landmark ruling *Magtajas v. Pryce Properties Corporation, Inc.* [187]. To withstand judicial scrutiny, an ordinance must satisfy six foundational requirements: it must not contravene the constitution or any statute, must not be unfair or oppressive, must not be partial or discriminatory, must not prohibit but may regulate trade, must be general and consistent with public policy, and must not be unreasonable. Because local legislative bodies possess purely delegated authority, municipal councils cannot prohibit activities that national legislation has expressly authorized or regulated [70]. This principle of statutory subordination forms the primary legal benchmark against which automated conflict detection systems must evaluate candidate municipal drafts.

When a local ordinance impermissibly enters an area governed by national legislation, Philippine courts consistently apply the *Magtajas* doctrine to declare the local measure invalid. In *Magtajas v. Pryce Properties Corporation, Inc.* [187], the City Council of Cagayan de Oro enacted ordinances prohibiting casino operations to protect public morals. The Supreme Court struck down the local enactments as *ultra vires*, ruling that a municipal council cannot prohibit commercial gambling activities that a national statute expressly authorized through the Philippine Amusement and Gaming Corporation (PAGCOR) charter. Similar preemption disputes have directly affected Davao City legislation. In *Mosqueda v. Pilipino Banana Growers & Exporters Association, Inc.* [190], the Supreme Court invalidated Davao City Ordinance No. 0309-07, which imposed a municipal ban on aerial pesticide spraying across agricultural plantations, holding that the ordinance exceeded municipal police powers and collided with the regulatory authority granted to the national Fertilizer and Pesticide Authority under Presidential Decree No. 1144. In *City of Manila v. Laguio, Jr.* [189], the Court likewise struck down an ordinance prohibiting motels and massage parlors, holding that local police powers cannot arbitrarily prohibit lawful commercial activities under the guise of zoning regulation.

In contrast to vertical statutory preemption, municipal lawmaking also involves horizontal consistency, where newly proposed ordinances must align with existing, fellow local enactments. When drafting local legislation, the Sangguniang Panlungsod routinely uses the Three-Reading legislative process, particularly the committee debates and line-by-line amendments during the Second Reading, to harmonize new proposals with older municipal codes [166]. Unlike vertical conflicts that represent fatal legal defects, inconsistencies between a new draft and an older municipal ordinance frequently reflect intentional legislative updates, such as revising local tax rates or adjusting traffic routes. To manage these horizontal modifications systematically, legislative drafters rely on standardized legislative saving devices recommended by the Department of the Interior and Local Government (DILG) [44]. Every formal ordinance incorporates an express or general repealing clause, declaring that all prior local enactments inconsistent with the newly approved measure are amended or repealed accordingly. Furthermore, ordinances feature a separability clause, acting as a structural legal safeguard ensuring that if any single section is declared invalid by judicial authorities, the unoffending sections remain in full

legal force [44]. However, because municipal legal researchers must review hundreds of active local enactments under strict legislative session schedules, subtle inconsistencies often go unnoticed during the Second Reading, allowing problematic clauses to proceed to final enactment.

A fundamental design decision in legal informatics is delineating which positive legal instruments establish binding preemption standards. Under Article VII, Section 17 of the 1987 Philippine Constitution and Book III, Title I, Chapter 2, Section 2 of the Administrative Code of 1987 (Executive Order No. 292) [152], Executive Orders (EOs) are acts of the President providing rules of a general or permanent character in implementation or execution of constitutional or statutory mandates. Unlike Supreme Court decisions—which function as *ex-post* judicial dispute resolutions settling adversarial controversies under the doctrine of *stare decisis* (Article 8, Civil Code)—Executive Orders operate as *ex-ante*, forward-looking statutory instruments possessing the full force and effect of law. Furthermore, during both the Martial Law period (1972–1986) and the Revolutionary Government under the Freedom Constitution (Proclamation No. 3, 1986–1987), the President exercised direct legislative powers through Executive Orders and Presidential Decrees, creating substantive national enactments that remain operative law today. Modern Executive Orders frequently establish binding nationwide regulatory standards—such as Executive Order No. 26 (s. 2017) establishing nationwide public smoking prohibitions, or Executive Order No. 168 (s. 2014) creating the Inter-Agency Task Force for the Management of Emerging Infectious Diseases (IATF)—that legally preempt conflicting municipal legislation. Excluding Executive Orders would therefore omit critical nationwide regulatory mandates that local ordinances are legally bound to respect.

2.2.2 Historical and Stylistic Heritage

The structural complexity of Philippine legislative text stems directly from the nation’s historical experience as a mixed legal jurisdiction [4, 39]. Philippine jurisprudence combines continental civil law traditions with Anglo-American statutory common law conventions. Understanding why Philippine statutes and local ordinances are drafted in their characteristic syntax requires examining how these colonial and post-colonial traditions fused over more than a century of codification.

During more than three centuries of Spanish rule, the Philippine legal system was rooted firmly in the Roman-Germanic civil law tradition [4]. This continental tradition emphasized systematic codification, deductive legal reasoning, and broad normative principles organized into structured books, titles, and chapters. Foundational enactments, such as the Spanish Civil Code of 1889 and the Penal Code of 1870, relied on generalized legal mandates rather than granular factual definitions. Furthermore, this civilian heritage embedded a rich vocabulary of Latin doctrinal maxims into Philippine jurisprudence, including *in pari materia* (construing related laws together), *generalia specialibus non derogant* (a general law does not override a specific one), and *ultra vires* (acts beyond legal authority) [4]. Modern Philippine codes retain this civilian structure, expecting adjudicators to derive specific legal duties from comprehensive normative commands.

Following the Treaty of Paris and the establishment of the American Insular Gov-

ernment (1900–1935), American legal administrators introduced common law statutory drafting habits without fully displacing the civilian codification base [39]. Enacted by the Philippine Commission and later the Philippine Legislature, early statutes introduced formal drafting conventions borrowed directly from United States federal and state legislation. These conventions included explicit statutory title preambles, exhaustive definitions sections, severability clauses, and general repealing provisions [4]. Significantly, American administrators introduced the practice of drafting extensive conditional provisos (introduced by the traditional formula “*Provided, that...*” and “*Provided, further, that...*”), designed to restrict, qualify, or create exceptions to preceding operative clauses within the same section. The Revised Administrative Code of 1917 (Act No. 2711) and the contemporary Administrative Code of 1987 (Executive Order No. 292) solidified this mixed tradition, confirming English as the controlling statutory language while maintaining civil law institutional categories [152].

This dual heritage produces a legislative drafting style marked by high syntactic density and multi-clause complexity. Unlike standard conversational English, a single statutory provision frequently spans one hundred to two hundred words contained within a solitary grammatical sentence [15]. These sentences embed multiple dependent clauses, qualifying prepositional phrases, and recursive proviso exceptions. At the local level, legislative staff adopt this formal national drafting template while adding procedural boilerplate, resulting in verbose ordinances where the operative regulatory command is separated from its governing subject by dozens of intervening words [44, 166]. From an analytical and document processing perspective, this hybrid syntax presents substantial structural hurdles: archaic legal terms resist straightforward vocabulary normalization, long-distance dependencies obscure the core legal directive, and qualifying provisos invert the practical meaning of preceding clauses unless the text is examined in its full structural context.

2.2.3 Local Legislative Procedure and Language Variation

Besides the hierarchy of laws, the actual way local laws are written and passed creates significant inconsistencies in language and format. National laws are usually written by specialized committees using standard legal terms. However, local ordinances go through a localized process that often breaks up formal language. To detect conflicts accurately, the computational system must be adjusted to handle the noise and ambiguity caused by bilingual drafting and decentralized committee amendments in the Davao City Council (Sangguniang Panlungsod).

In Davao City, the drafting and progression of proposed legislation is guided by the Internal Rules of Procedure, codified in Resolution No. 001-22 [166]. The drafting process is decentralized, meaning different council members and their personal legal researchers author their own ordinances rather than relying on a single, centralized drafting bureau [70, 166]. Because of this decentralized structure, drafting conventions and legal terminology vary substantially across city council committees. For example, a transportation ordinance and an environmental regulation may use disparate vocabularies to describe overlapping commercial entities, making it difficult for computational retrieval models to identify statutory intersections.

According to the Local Government Code, no ordinance can be passed unless it undergoes three separate readings on distinct calendar days [70, 166]. While democratic, this procedural cycle frequently fragments the legislative text across successive drafting iterations. During the First Reading, the draft is assigned to the appropriate standing committee; during the Second Reading, the text undergoes extensive interpellation, debate, and line-by-line floor amendments [166]. This continuous iterative modification often introduces subtle structural inconsistencies, where an amendment introduced in an operative section contradicts a general policy declaration in an earlier chapter without being reconciled prior to final Third Reading enactment.

A unique challenge within the Philippine local legislative domain is the statutory requirement to make local legislation accessible across languages. Under Section 469 of the Local Government Code, the Secretary to the Sanggunian is mandated to translate approved ordinances into local dialects, such as Cebuano for Davao City [70, 166, 53]. However, this creates legal friction because the Administrative Code of 1987 establishes that the English statutory text strictly controls in cases of conflict [151]. If an error occurs in regional translation, a citizen relying on the Cebuano text may inadvertently violate the authoritative English version, conflicting with the public notice doctrine in *Tañada v. Tuvera* [185]. Despite policy initiatives such as the *Batas sa Sariling Wika Act* [149], subtle dialect differences often produce semantic drift where the exact legal obligation in English is distorted in regional vernacular text [123].

2.2.4 Ex-Ante Legislative Review

The term “ex-ante” derives from Latin meaning “before the event.” In legislative informatics and administrative jurisprudence, it denotes the rigorous evaluation of a proposed law or municipal regulation prior to its official enactment [59]. This prospective auditing process aims to identify and remediate potential statutory collisions, constitutional infirmities, and unintended regulatory burdens while the draft remains amenable to revision.

In the Philippines, ex-ante regulatory review has achieved heightened institutional prominence following the enactment of Republic Act No. 11032, known as the Ease of Doing Business and Efficient Government Service Delivery Act of 2018 [71]. Under this statutory framework, the Anti-Red Tape Authority (ARTA) alongside the Department of the Interior and Local Government (DILG) mandates that Local Government Units conduct a formal Regulatory Impact Assessment (RIA) for proposed regulations impacting commerce and public welfare [71]. The objective is to verify that proposed ordinances possess valid statutory authority and do not impose duplicative or contradictory requirements on citizens and businesses.

The fundamental distinction between ex-ante and ex-post evaluation lies in timing, methodology, and legal impact across the hierarchy of legal instruments. Ex-ante evaluation focuses strictly on prospective, forward-looking legislative texts—such as draft local ordinances, pending congressional bills, and proposed executive issuances—prior to their passage into positive law. In contrast, ex-post evaluation operates retrospectively through judicial dispute resolution, encompassing adversarial litigation, appellate court rulings, and Supreme Court jurisprudence governed by *stare decisis*. While ex-post judicial ad-

judication rectifies constitutional injuries after harm has occurred, it incurs substantial societal and litigation expenses, as repealing an active ordinance requires protracted court proceedings [203]. Ex-ante auditing is therefore substantially more efficient, protecting municipalities from administrative deadlock and judicial invalidation.

Deploying automated computational tools for ex-ante conflict detection empowers local legislative researchers to identify normative collisions during the drafting stage. Because legislative bodies generate extensive statutory volumes, manual review frequently fails to catch subtle inconsistencies between new drafts and existing laws [94]. By operating as an automated diagnostic triage tool, computational systems flag potential preemption conflicts while the measure remains in committee deliberation. This facilitates alignment with superior national statutes, mitigating the risk of enacting municipal measures that are legally void *ab initio* [187].

2.2.5 Legislative Consistency Checking

Because local governments hold subordinate legislative authority within a decentralized drafting environment, systematic consistency checking represents an essential institutional safeguard. However, verifying that an ordinance does not contradict superior statutes requires far more than title matching; it demands deep comprehension of legal semantics and subtle jurisdictional boundaries. Current consistency checking methods remain encumbered by syntactic ambiguity and reliance on basic tracking platforms.

Discourse analysis of Philippine statutory language reveals pervasive ambiguity in how normative keywords and modalities are deployed [186]. In particular, the modal auxiliary “shall” introduces frequent interpretive friction, as it may denote an absolute command, an administrative directive, or a directory permission depending on statutory context. Empirical studies confirm that such syntactic ambiguities complicate legal interpretation even when enactments adhere to formal drafting guidelines [15].

Philippine jurisprudence recognizes that statutory conflicts rarely present themselves through explicit textual declarations [39]. A municipal ordinance virtually never declares an overt intent to contradict an Act of Congress; rather, conflicts typically arise as implied preemption. This occurs when a national statute occupies a regulatory field, and a municipal council introduces unauthorized preconditions or local fees that restrict activities already licensed by the state [186]. In such cases, the Supreme Court has consistently held that the local measure impermissibly modifies national law under the guise of police power regulation. For manual reviewers and keyword-based search tools, detecting implied preemption is exceptionally difficult because the conflicting enactments employ disparate vocabularies, necessitating deep semantic analysis to uncover the contradictory regulatory effect.

In the local context, the Davao City Sangguniang Panlungsod processes extensive legislative volume, passing nearly 400 ordinances and over 2,000 resolutions within a single three-year council term [139]. To manage this throughput, the city deploys administrative platforms: the Resolution Ordinance Tracking System (ROTS) manages document routing and metadata, while the Legislative Information Support System Program (LISSP) enables council members to access digitized enactments [30]. However, while these digital

systems streamline document logistics, substantive consistency auditing across provisions remains entirely manual. Furthermore, archival scans frequently contain optical character recognition (OCR) noise, causing keyword search queries to miss critical governing provisions [115].

2.3 Digital Legislative Archives

Building a functional Semantic Conflict Detection system requires access to comprehensive and machine-readable legal corpora. In the Philippines, the digital landscape of legislative data is highly fragmented, split between national statutory repositories and localized municipal archives. This structural division presents significant technical hurdles for statutory data ingestion. Before deep logical inference can occur, automated auditing pipelines must first secure reliable data streams from these disparate legal environments. Consequently, understanding the specific characteristics of these sources informs the design of statutory retrieval and preprocessing methodologies.

2.3.1 National Sources

National legal repositories represent the authoritative apex of the Philippine legislative hierarchy, housing laws that universally preempt local ordinances. These databases are maintained by centralized government agencies and contain highly formalized texts such as Republic Acts, Commonwealth Acts, and *Batas Pambansa*. They serve as the primary knowledge base for external legal context, providing the authoritative baseline rules against which municipal drafts are evaluated. Because national legislation undergoes a rigorous, standardized drafting process, the vocabulary within these repositories is generally consistent and adheres to strict legal drafting standards. Accessing these national databases provides computational systems with the necessary macroscopic view of state-level mandates and prohibitions.

The Official Gazette serves as the constitutionally mandated repository for Philippine statutes, rendering it the primary journal of record for national legislative text [72]. It acts as the official publisher for all newly enacted national laws, presidential executive orders, and administrative issuances, as required by Republic Act No. 453 [142] and Commonwealth Act No. 638 [131]. Because it represents the definitive, legally binding version of a statute, any computational system evaluating municipal preemption must anchor its baseline rules to this repository. This archive effectively acts as the authoritative historical ledger for state-level mandates, preserving the exact phrasing approved by the legislative and executive branches. Consequently, it remains the most critical target for establishing an authoritative statutory knowledge base.

Complementing the official journal, developers utilize the Open Congress API provided by the House of Representatives Legislative Information System (LEGIS) to access structured, machine-readable metadata regarding legislative routing and statutory amendments [84]. This programmatic interface tracks a statute's formal lifecycle, charting its progression from initial committee filing to final executive approval. By providing standardized output formats, the API allows computational pipelines to establish relational

dependencies between different legislative acts without manually parsing unstructured text.

Finally, the Lawphil Project functions as a massive, privately maintained digital repository of statutory records and jurisprudential decisions [8]. Unlike official government channels that frequently change web templates, this repository provides a comprehensive, pre-digitized plaintext corpus of historical and contemporary Philippine law. This accessibility makes it mathematically ideal for the unsupervised pre-training and vectorization of foundational language models tailored to the local legal domain. By ingesting this vast corpus of clean text, NLP architectures learn the highly specific syntactic structures and deontic operators inherent in Philippine jurisprudence, establishing the core vocabulary required for transformer models to process specialized statutory drafting [227].

2.3.2 Local Sources

In contrast to centralized national databases, local legal repositories are maintained independently by individual Local Government Units (LGUs). These archives house municipal ordinances, resolutions, and executive orders that govern specific geographic jurisdictions. The legislative output at this tier is entirely decentralized, resulting in highly variable archiving practices depending on the technological capacity of the specific city council. These local databases form the immediate target corpus for municipal conflict detection, containing the actual regulatory constraints imposed on city residents. Consequently, they represent a high-volume, dynamic data environment that continuously expands as the local council holds regular drafting sessions.

The locality’s primary corpus consists of the Davao City Sangguniang Panlungsod website, which serves as the official digital repository for municipal legislative outputs [31]. This portal acts as the primary historical record for municipal regulations, resolutions, and zoning policies enacted within the city’s jurisdiction. It represents a dynamic and continuously expanding text environment that directly reflects the decentralized nature of local lawmaking. The texts housed here encapsulate the specific regulatory constraints and policy directions of the city government, providing the exact ordinance text required to evaluate new drafts for logical consistency against the existing legislative framework.

To manage the extensive volume of local legislation, the city introduced the Legislative Information Support System Program (LISSP) as its primary administrative backbone [30]. This platform serves as a digital tracking system designed to maintain metadata regarding ordinance routing, committee assignments, and plenary approval statuses. The system provides transparency for monitoring ordinances and resolutions in real time [140]. By digitizing the legislative workflow, LISSP effectively functions as a centralized database of ordinances, providing the structured tracking and contextual data necessary for evaluating the lifecycle of local laws, while addressing the interoperability challenges identified in recent regional e-governance assessments

2.3.2.1 Archival Scanning and OCR Degradation

A primary barrier to automated statutory auditing is the lack of native, machine-readable text across archival collections. Both the Official Gazette and the Davao City Sangguniang

Panlungsod public archive predominantly host scanned, image-based PDFs rather than selectable digital documents. During manual inspection of these collections, enactments were observed to contain varying degrees of visual noise, including slanted scans, faded ink, low contrast, and compression artifacts. As Lopresti [115] documents, such noisy visual data results in severe Optical Character Recognition (OCR) degradation, where standard extraction tools produce broken alphanumeric strings and nonsensical tokens. These irregularities suggest that a direct ingestion pipeline suffers from semantic fragmentation, requiring an additional high-precision extraction layer and extensive text-cleaning heuristics to ensure data viability.

The resulting character and word-level corruptions severely hinder automated semantic analysis. In statutory conflict detection, a single corrupted character can alter the legal interpretation of a provision, such as transforming a negative operator into an affirmative assertion or scrambling a penalty amount. Standard tokenizers struggle to process these degraded sequences, fragmenting broken words into out-of-vocabulary sub-tokens that dilute contextual representations [162]. Furthermore, scanning artifacts frequently obscure section demarcations and legislative numbering, preventing standard parsers from isolating self-contained statutory units. Consequently, automated preemption pipelines must incorporate noise-tolerant document normalization to restore structural integrity before downstream retrieval.

2.3.2.2 Repository Scarcity and Digital Isolation

Another critical challenge is archival fragmentation, wherein legislative repositories suffer from systemic gaps in their digital records. This poses a significant problem for high-fidelity auditing because the absence of even a single related ordinance can lead to false negatives in conflict detection, where a system declares a draft valid simply because it cannot access the contradicting record. A prominent example is the Davao City Sangguniang Panlungsod public archive, which features intermittently missing ordinance files and unorganized chronological listings that hinder systematic data collection [168]. Even relatively complete national databases like Lawphil can lack specific Republic Acts or historical amendments. As Westermann et al. [78] observe, fragmented legal databases force developers to move beyond simple scraping toward fault-tolerant data engineering to ensure an exhaustive corpus.

A further obstacle in statutory data acquisition is the discrepancy between the theoretical and actual accessibility of municipal databases. For example, while administrative systems like Davao City’s LISSP are promoted in institutional reports as public transparency portals [30, 140], empirical verification reveals strict authentication barriers requiring authorized credentials. Technical inspection indicates that municipal legislative portals often remain computationally isolated behind authentication interfaces, restricting automated extraction by public web scrapers. Correspondence with municipal administrators confirms that full-text enactments require administrative clearance and manual provisioning. As Conchas and Salapa [35] observe, isolated administrative data often exists as digitally orphaned artifacts, where metadata about an enactment is disconnected from its full-text searchable content, requiring specialized ingestion pipelines to reconcile sparse public metadata with offline full-text documents.

2.4 Exact and Symbolic Methods

Before the adoption of neural networks, computational legal analysis relied heavily on deterministic systems to process text. Deterministic systems operate on absolute rules, meaning that a specific input will invariably produce the exact same output without any probabilistic guessing. Traditional methods for searching and evaluating legal information assumed that legislative ideas could be cleanly mapped either to specific vocabulary strings or to rigid, hand-coded logic formulas. While highly transparent and easy to debug, these foundational approaches frequently fail when forced to detect hidden conflicts in complex, unstructured legal texts. The core issue is that human language is inherently messy, and rigid systems break down when faced with the ambiguous phrasing typical of municipal ordinances.

2.4.1 Boolean Information Retrieval

To understand why traditional legal databases struggle with semantic conflict detection, we first need to look at how they fundamentally process text at the machine level. Boolean Information Retrieval (BIR) represents the foundational search architecture for almost all commercial law databases currently in use. Instead of reading a document like a human, BIR treats every file simply as a “bag of words,” stripping away all grammar and sentence structure. It relies on an inverted index, which is essentially a massive lookup table that tracks exactly which specific words appear in which files across the entire database. Because it only registers the presence or absence of characters, the system is completely blind to the actual context or legal intent behind those words.

At its core, the retrieval mechanism in BIR is strictly based on set theory and formal Boolean logic. A user constructs a query using logical operators like **AND** ($\wedge$), **OR** ($\vee$), and **NOT** ($\neg$) to create intersecting sets of documents. For example, a search for a specific ordinance regarding explosives is formulated mathematically as:

Boolean Query Set Intersection Formulation:

$$Q = (\text{fireworks} \vee \text{pyrotechnics}) \wedge \neg \text{exempt} \quad (2.1)$$

Where the compound query Q combines disjunctive keyword sets under conjunction with a negated exemption condition.

The system then retrieves documents by calculating the exact mathematical intersection of these sets, pulling only the files that satisfy every condition simultaneously. If a document contains the specific string of characters requested, it is retrieved and scored as a 1; if it misses even a single required term, it is completely ignored and scored as a 0. It evaluates the raw presence of strings, entirely ignoring the conceptual meaning or legal intent behind them [17].

2.4.1.1 Challenges with BIR

The primary limitation of BIR in legal auditing is its absolute, uncompromising reliance on exact string matching. This rigidity creates a severe “vocabulary mismatch” problem

when processing natural language. Human language naturally features synonymy, where entirely different words possess the exact same meaning, and polysemy, where the identical word takes on drastically different meanings depending on the context. In statutory drafting, one author might use “prohibited” while another uses “unlawful,” referring to the exact same legal concept. Because BIR operates on character-level matching rather than conceptual mapping, it fundamentally cannot connect these two terms, leading to massive blind spots in the search results.

A landmark empirical evaluation of these systems demonstrated this vocabulary gap in a highly practical setting. Legal researchers conducting queries using standard BIR systems confidently estimated they were finding over 75 percent of the relevant documents for a given case. They experienced an illusion of effectiveness because the documents they did retrieve looked highly accurate, masking the massive volume of relevant files they missed. In reality, the system retrieved less than 20 percent of the actual relevant files contained within the database [17]. The researchers missed the remaining 80 percent simply because the authors of those hidden documents used slightly different, completely valid phrasing that didn’t perfectly align with the strict search query [205].

In a local government context like Davao City, this specific margin of error is administratively unacceptable. Missing even one existing ordinance during an ex-ante review due to a simple vocabulary mismatch could directly result in the passage of an invalid, conflicting law. Furthermore, legislative vocabulary inherently evolves over time as drafting conventions change. An ordinance passed in the 1990s and a modern draft written in 2026 might strictly regulate the exact same municipal subject but share absolutely zero common keywords. Recent studies emphasize that resolving this requires the implementation of “Transformer” models, which abandon exact matching entirely in favor of continuous vector spaces that recognize when different vocabulary points to the same underlying concept [98].

2.4.2 Symbolic AI

Because BIR fundamentally fails to capture underlying meaning, early artificial intelligence researchers attempted to automate legal reasoning by directly modeling the rules of law themselves. This approach, widely known as Symbolic AI or “Good Old-Fashioned AI” (GOFAI), operates on the premise that human cognition can be replicated by manipulating formal symbols [207]. It moves beyond simple word matching by teaching computers actual legal logic through explicit, human-coded “IF-THEN” formulas. Instead of searching a database for raw text strings, Symbolic AI requires human engineers to manually translate the constraints of a statute into formal logic structures. This theoretically allows the computer to process the actual regulatory mechanics of a law rather than just reading its surface vocabulary.

To execute this logic, these systems specifically utilize propositional or first-order logic combined with deontic operators. Deontic logic is a highly specialized branch of mathematical logic that deals specifically with obligations, permissions, and prohibitions. For example, a system represents a local zoning prohibition mathematically as:

First-Order Deontic Prohibition Formulation:

$$\forall x(\text{Business}(x) \wedge \text{ResidentialZone}(x) \rightarrow \text{Prohibited}(x)) \quad (2.2)$$

Where the universal quantifier ($\forall$) asserts that for any entity x , if it is classified as a business ($\wedge$) and located in a residential zone, it is strictly prohibited ($\rightarrow$).

By converting ordinances into these rigid formulas, a system can theoretically use deductive reasoning to spot a conflict. If Rule A dictates $\mathcal{O}(X)$ (an absolute obligation to do X) and Rule B dictates $\mathcal{F}(X)$ (an absolute prohibition against doing X), the system successfully flags a formal logical contradiction ($\perp$) [12].

2.4.2.1 Challenges with Symbolic AI

While Symbolic AI provides perfectly transparent and mathematically sound reasoning, its absolute reliance on manual programming makes it practically impossible to scale for real-world governance. The primary architectural barrier is the classical “knowledge acquisition barrier” in expert systems: human domain experts and software engineers must manually read, interpret, and translate thousands of unstructured city ordinances into rigid computer rules. This manual translation process is unacceptably slow, extremely expensive, and highly prone to human error during the coding phase [119]. Furthermore, every time the city council amends an existing law, the underlying symbolic formula must be manually rewritten, making the system incredibly difficult to maintain over time.

Beyond the logistical nightmare of manual coding, the rigid 1-and-0 architecture of Symbolic AI fundamentally clashes with the “open texture” of natural legal language. Laws frequently rely on subjective, open-ended standards—such as requiring a citizen to exercise “reasonable care” or negotiate in “good faith.” These inherently subjective human concepts cannot be easily defined within a strict, binary rule-based formula because their definitions change depending on the specific context [11]. While these symbolic systems are highly precise in tightly controlled, hypothetical environments, they immediately shatter when exposed to the real world. Ultimately, they lack the computational flexibility required to handle the thousands of messy, unorganized, and ambiguously worded documents actually found in local government archives [228].

2.5 Graph-Based Architectures

Because exact word matching can be unreliable, computational legal studies frequently explore graph-based architectures to model relationships between statutory provisions. Instead of treating legal texts as isolated character sequences, graph-based frameworks represent statutes as interconnected nodes within a structural network. By analyzing how enactments cite, amend, and conceptually reference one another, these architectures capture latent legal context that bag-of-words models invariably miss.

2.5.1 Knowledge Graphs and GNNs

Legal Knowledge Graphs (LKGs) model statutory ecosystems by representing legal entities, provisions, and administrative bodies as structured vertices linked by directional,

typed edges. Unlike flat text indices, these relational graphs enable Graph Neural Networks (GNNs) to learn node embeddings by aggregating feature representations across local structural neighborhoods [16]. Researchers have developed automated frameworks like GNN-ERE to extract legal entities and relations directly from unstructured statutory text, establishing structured taxonomies without manual relational coding [121]. Furthermore, empirical evaluations demonstrate that graph-based representations improve legal document retrieval by mapping semantically related court judgments even in the absence of verbatim keyword overlap [79].

2.5.1.1 Graph Modeling and Multi-Hop Reasoning

The foundational strength of Graph Neural Networks lies in their ability to propagate structural context across multi-hop relational paths. In statutory interpretation, normative conflicts rarely occur through isolated direct collisions; rather, they frequently manifest indirectly across chains of administrative definitions and delegated authorities [222]. By executing iterative message passing across adjacent nodes, GNNs aggregate high-order neighborhood information to evaluate whether an ordinance clashes with a superior mandate through an intermediate regulatory concept. This topological awareness enables systems to distinguish between different legislative functions, such as separating procedural requirements from substantive prohibitions [16]. Consequently, graph neural architectures provide a principled mechanism for structural reasoning across complex, multi-tiered statutory hierarchies.

Furthermore, graph representations capture institutional relationships that remain hidden during sequential text reading. In decentralized governance frameworks like the Philippines, local legislative acts operate within a matrix of national agencies, provincial guidelines, and municipal zoning boundaries. Graph neural models learn relational node embeddings that reflect these jurisdictional hierarchies, allowing algorithms to evaluate legislative preemption within its broader administrative network [121]. By mapping shared regulatory targets across different enactments, the graph structure guides inference models toward relevant statutory clusters that standard search engines fail to connect. This relational capacity makes graph modeling an appealing conceptual paradigm for analyzing legislative coherence across decentralized jurisdictions.

2.5.1.2 Structural Vulnerabilities in Unstructured Archives

While Graph Neural Networks offer substantial theoretical utility, their practical deployment depends strictly on the completeness and fidelity of the underlying graph representation. A primary vulnerability is their reliance on explicit statutory cross-references to establish graph edges [87]. In Philippine municipal archives, local ordinances frequently omit formal citations to the governing Republic Acts they inadvertently contradict, relying instead on localized colloquial descriptions. When explicit relational links are missing from the initial graph construction, the GNN remains fundamentally blind to the underlying normative collision, unable to propagate message vectors across disconnected nodes.

This limitation is exacerbated by systemic disconnection and physical degradation across local government records. Municipal archives like Davao City’s LISSP store en-

actments as isolated scanned documents rather than interconnected, machine-readable data nodes [35]. When automated relation extraction pipelines encounter noisy optical character recognition outputs, they generate hallucinated edges and misattributed entity linkages [115]. These erroneous connections propagate through successive graph convolution layers, contaminating node representations and introducing false conflict signals throughout the network. Consequently, without expensive, manual human curation to verify every edge, graph-based systems prove too brittle for noisy, real-world municipal archives.

2.5.2 Abstract Meaning Representation

While Graph Neural Networks capture macro-level relationships between distinct legal documents, Abstract Meaning Representation (AMR) evaluates the semantic and logical structure inside individual sentences. AMR represents the meaning of a sentence as a rooted, directed acyclic graph where nodes correspond to semantic concepts and edges represent semantic roles [?]. In legal informatics, AMR provides a formal representation of "who is doing what to whom," abstracting away superficial syntactic variations like active versus passive voice [210]. This semantic grounding allows computational systems to evaluate legal statements based on their core logical assertions rather than surface grammar.

2.5.2.1 Semantic Graph Representations and Role Labeling

A principal advantage of Abstract Meaning Representation in statutory analysis is its ability to isolate predicate-argument structures and deontic operators within complex legislative commands. Statutory provisions rely heavily on deontic modalities—such as obligations, permissions, and prohibitions—to govern citizen behavior and administrative authority [198]. Standard grammar parsers often conflate these modal verbs with general descriptive predicates, obscuring the primary regulatory command. In contrast, AMR maps deontic operators directly to core predicates, explicitly identifying whether an action is mandatory, permissible, or prohibited [210]. This structural clarity allows systems to identify subtle policy collisions, such as when a local ordinance authorizes a commercial practice that a national statute expressly forbids.

Furthermore, AMR graphs facilitate semantic normalization across diverse legislative drafting styles. Because municipal ordinances are drafted by different local committees across multiple decades, identical regulatory concepts are routinely phrased using disparate vocabulary and sentence structures. AMR normalizes these surface differences into standardized semantic frames, allowing different legislative phrasings to map to identical concept graphs [198]. Recent legal NLP benchmarks demonstrate that evaluating semantic graph similarity outperforms raw token comparisons when identifying core regulatory requirements [210]. This structural consistency positions AMR as a valuable framework for extracting precise semantic representations from dense legislative text.

2.5.2.2 Sentence Boundary Constraints and Parser Degradation

Despite its structural precision, Abstract Meaning Representation faces severe operational limitations when applied to Philippine municipal legislation. The primary architectural constraint is AMR’s restriction to single-sentence parsing [198]. In statutory drafting, preemption conflicts rarely manifest within a single isolated sentence; rather, normative contradictions typically emerge across distinct sections, such as a general mandate in one clause negated by an unprovided exception several paragraphs later. Because standard AMR parsers evaluate individual sentences independently, they cannot perform cross-sentence coreference resolution or integrate distant provisos, creating significant logical blindspots in statutory conflict detection.

Compounding this sentence-level restriction is AMR’s high sensitivity to complex statutory syntax and optical character recognition noise. Philippine statutory provisions, such as sections within Republic Act No. 7160, frequently exceed one hundred words, featuring deeply nested conditional clauses and multiple provisos within a single period [200]. When processing these complex syntactic structures alongside scanning artifacts or localized administrative phrasing, standard AMR parsers frequently fail, producing fragmented or malformed logic graphs [204]. Without extensive domain-specific parser retraining tailored to Philippine legal drafting, AMR cannot provide a dependable, production-ready foundation for ex-ante municipal auditing.

2.6 Deep Learning and Transformers

The application of Natural Language Processing (NLP) in the legal domain has shifted from rigid, rule-based systems and basic word-matching toward Deep Learning [95]. Legacy algorithms treat legal texts as simple character strings, which frequently fail to capture the dense and interconnected nature of statutory drafting. Deep Learning, conversely, utilizes neural networks to map out the actual semantic meaning of legal language. This allows systems to process complex syntax and perform the deductive reasoning required to spot conflicts. This section breaks down the Transformer architecture, explaining its underlying mechanics, computational challenges, and why it remains the most practical framework for auditing unstructured municipal archives.

2.6.1 Transformers and Self-Attention

A major breakthrough was the Transformer architecture, introduced by Vaswani et al. [206] in 2017. Prior models, such as Recurrent Neural Networks (RNNs) and Long Short-Term Memory (LSTM) networks, processed text sequentially. This step-by-step processing was computationally slow and made it difficult for the model to retain information from earlier in a long document. Transformers bypass this sequential processing limit by relying on a Self-Attention mechanism.

Self-Attention allows a model to analyze all words in a sequence simultaneously to determine their contextual relationships. This enables a specific term to connect with any other word in the document, regardless of the physical distance between them. Because it

processes data in parallel, Transformers are significantly faster to train, directly enabling the creation of Large Language Models (LLMs). Modern iterations have introduced refinements like Disentangled Attention to better parse complex syntactic structures, alongside Windowed Attention to process long-context legal documents efficiently. However, while Generative LLMs like ChatGPT excel at text summarization, they are prone to ‘hallucinations’—often inventing plausible but entirely fake legal citations [45]. To ensure legal validity, a system must be discriminative: specifically classifying whether statements exhibit a Conflict or Entailment, rather than simply generating new text.

2.6.2 Transformer Challenges

Despite their architectural versatility, deploying transformer models within the computational constraints of Philippine Local Government Units (LGUs) introduces substantial operational hurdles.

The foremost engineering bottleneck arises from quadratic memory scaling and context length limits. The standard self-attention mechanism scales quadratically ($O(n^2)$) with sequence length, requiring exponential increases in memory as input texts expand [206]. Standard architectures like BERT are constrained to a 512-token context window [48], whereas Philippine ordinances frequently span multiple pages. While practitioners employ sliding-window techniques to segment documents into overlapping spans [7, 82], this segmentation severs the semantic dependency between general regulatory mandates and distant provisos [172]. Consequently, sequence truncation creates structural blind spots that can lead to false conflict classifications.

A secondary challenge involves morphological noise and domain adaptation limits. Pretrained transformers require massive corpora of clean text, and while Western legal models like Legal-BERT perform well on foreign statutory benchmarks [25], no dedicated foundation model exists for Philippine local legislation. Applying general tokenizers to municipal archives—saturated with scanning artifacts, archaic statutory phrases, and administrative Taglish—triggers severe token fragmentation, where domain-specific terms are broken into meaningless sub-tokens that degrade contextual comprehension.

Finally, edge deployment is heavily restricted by municipal hardware realities. Executing deep logical inference across thousands of document pairs demands substantial GPU compute and memory bandwidth [150]. Because most municipal governments operate standard desktop workstations with modest hardware specifications, deploying unoptimized neural models leads to system latency spikes or memory exhaustion, necessitating decoupled, multi-stage architectures optimized for local inference.

2.6.3 Suitability of Transformers

Notwithstanding these hardware limitations, Transformer models remain the most viable approach for Semantic Conflict Detection in local ordinances. General-domain language models typically struggle with legal text because common words frequently take on entirely different legal definitions. For instance, “party” refers to a litigant rather than a celebration, and “suit” signifies a legal action rather than clothing. Greco et al. [?] empha-

size that this high degree of polysemy is exactly why domain-specific models are required. While older static systems assign a single, fixed meaning to a word [?], a Transformer dynamically adjusts its reading of a word based on the surrounding sentence context [25]. Chalkidis et al. [26] empirically showed that this contextual awareness is what allows a model to distinguish between amending a law and violating it—a distinction that simple keyword searches invariably miss.

The transition to Deep Learning necessitates explicitly moving away from legacy methodologies. Boolean keyword searches fundamentally fail because legislative drafting vocabulary constantly evolves over the decades. Conversely, manually mapping thousands of unlinked Republic Acts and local ordinances into deterministic graph structures creates an intractable administrative burden. As Li et al. [110] observed, hand-crafting these legal knowledge graphs simply does not scale for volatile, unstructured municipal records. Ultimately, Transformers represent the ideal method: they possess the capacity to map the deep, latent semantic relationships in unstructured legal text without requiring an intractable amount of manual data entry.

2.7 Single-Stage and Multi-Stage Architectures

In mainstream Natural Language Processing (NLP), end-to-end, single-stage neural architectures serve as the conventional baseline for document classification, sentiment detection, and question answering. These systems funnel raw input sequences directly into a single neural network that outputs the final classification without intermediate filtering. However, in legal informatics, detecting statutory conflicts requires a level of deductive precision across vast corpora that a monolithic model structurally cannot provide. Forcing a single architecture to simultaneously scan tens of thousands of statutory records and evaluate fine-grained deontic logic creates severe mathematical and operational trade-offs.

2.7.1 Single-Stage Monolithic Pipelines

The appeal of a single-stage neural architecture is its conceptual simplicity: a single model handles the entire input-output mapping without requiring intermediate data handoffs. In standard NLP tasks operating over small document collections, this approach minimizes engineering complexity and eliminates pipeline latency bottlenecks. However, deploying a monolithic model to audit draft local legislation against expansive national repositories introduces severe operational barriers, inevitably sacrificing either candidate recall or deductive precision.

One common monolithic approach involves prompting generative Large Language Models to directly declare legal preemption. However, autoregressive language models operate as probabilistic token predictors rather than deterministic logical discriminators, rendering them vulnerable to legal hallucinations where non-existent statutory citations are fabricated [113]. Furthermore, their massive parameter sizes make offline deployment on standard municipal hardware infeasible without heavy quantization, which severely degrades deductive reasoning fidelity in specialized statutory domains [46]. Conversely, relying exclusively on single-stage bi-encoders introduces severe semantic compression lim-

its. While dual-tower bi-encoders excel at high-speed vector retrieval across tens of thousands of provisions [150, 29], they collapse entire multi-clause provisions into static, fixed-dimensional vectors [93, 196]. This architecture optimizes for topical similarity rather than logical consistency, causing bi-encoders to conflate permission and prohibition provisions that share identical topical vocabulary [98].

Furthermore, executing a single-stage cross-encoder across an entire statutory repository is computationally prohibitive. Because full cross-attention scales quadratically ($O(n^2)$) with input length, comparing a municipal draft against 25,432 national statutes simultaneously would instantly overwhelm workstation memory [13]. Forcing the model to process fragmented sliding windows across thousands of enactments severs inter-clause dependencies. This computational ceiling confirms that isolating a compact candidate shortlist via coarse retrieval is essential prior to executing fine-grained sequence classification.

2.7.2 Two-Stage Retrieve-then-Entail Architectures

To resolve the mathematical and computational limitations of monolithic architectures, contemporary legal informatics adopts decoupled, two-stage systems. In the retrieve-then-entail paradigm, the workflow is split into two specialized stages: an initial coarse Information Retrieval (IR) phase that scans the entire statutory corpus to extract a top- k candidate shortlist, followed by a fine Natural Language Inference (NLI) phase that performs deep cross-attention on the candidate pairs [218]. Williams et al. [218] demonstrate that this two-tier design reduces manual auditing overhead by up to seventy percent while preserving high inferential accuracy.

The primary engineering advantage of this decoupled architecture is its alignment with edge computing realities. By offloading broad corpus search to pre-computed vector indices that require minimal runtime memory, the computationally intensive $O(n^2)$ matrix operations of the cross-encoder are restricted exclusively to the top- k candidate pairs. Khattab and Zaharia [97] established that this decoupled structure enables advanced neural language processing to execute locally on consumer-grade hardware without requiring distributed GPU clusters. Furthermore, decoupling retrieval from inference provides a pragmatic resolution to the vector compression trade-off. The initial retrieval stage accepts the loss of granular syntactic nuance to maximize search throughput across extensive statutory archives. In contrast, the subsequent inference stage bypasses vector pooling entirely, performing joint token-level cross-attention over raw, uncompressed text spans to achieve high-precision deontic reasoning [86].

2.7.3 Comparative Architectural Trade-offs

Comparing single-stage and dual-stage architectures highlights a fundamental division of labor in statutory preemption analysis. A single-stage retriever maximizes search speed but remains blind to subtle modal negators, whereas a single-stage cross-encoder possesses high logical sensitivity but cannot scale beyond a few dozen provisions without memory failure. By chaining both mechanisms in a coarse-to-fine hierarchy, the system leverages

the complementary strengths of each model family: Stage One absorbs corpus volume and context length, while Stage Two resolves logical nuance and morphological noise [117].

The empirical superiority of this decoupled design is reinforced by over a decade of findings from the Competition on Legal Information Extraction and Entailment (COLIEE). As Goebel et al. [65] document, COLIEE divides legal reasoning into statutory information retrieval (Task 3) and statutory textual entailment (Task 4). Year after year, top-performing teams reject single-stage monoliths in favor of pipelined approaches [6]. While early benchmarks in COLIEE 2021 established this separation, modern iterations through COLIEE 2025 [66] and COLIEE 2026 [67] confirm that decoupled retrieval and inference remain the gold standard. For statutory retrieval, hybrid dense-sparse systems achieve near-perfect candidate recall [133], while cross-encoder entailment systems reach over 95% accuracy by isolating granular reasoning from broad corpus retrieval [5, 199]. This sustained global validation confirms that separating Information Retrieval from Natural Language Inference is the established operational blueprint for resolving semantic conflict across statutory corpora.

2.8 Information Retrieval

Information Retrieval (IR) is the process of extracting relevant documents from massive, unstructured datasets based on a specific query. In legal informatics, IR acts as the foundational filtering mechanism. It serves as the "coarse" stage of a dual-pipeline system, isolating candidate statutes from thousands of irrelevant municipal records before any intensive logical analysis takes place.

2.8.1 Information Retrieval Challenges in Statutory Corpora

Applying information retrieval techniques to statutory collections introduces domain-specific hurdles that general-domain search engines are not engineered to handle. A primary challenge is severe lexical divergence and vocabulary mismatch across legislative eras [17]. In statutory law, statutes governing identical civic activities may be enacted decades apart, using drastically different terminology, archaic phrasing, or specialized administrative acronyms. When newly drafted municipal ordinances employ modern regulatory expressions, exact keyword matching fails to retrieve governing historical statutes, creating dangerous blindspots in preemption review [205].

Furthermore, statutory corpora exhibit extreme length asymmetry and structural complexity that disrupt standard retrieval scoring. Statutory enactments frequently contain verbose, multi-clause provisions exceeding one thousand tokens, where general commands are qualified by distant provisos and statutory exceptions [27]. Standard passage retrievers struggle to balance term saturation across these extended provisions, frequently penalizing lengthy, comprehensive enactments. Compounding this structural challenge is administrative hierarchy noise: in national legal archives, subordinate issuances (such as executive orders, circulars, and departmental memos) outnumber primary legislative acts, frequently diluting candidate shortlists unless protected by statutory priority safeguards [70].

2.8.2 Relevance in Retrieval

To understand how computational systems extract documents, we must first examine the human cognitive process that defines relevance in the physical world. In cognitive science, relevance is not a static property but a dynamic relationship between information and a specific task. According to Relevance Theory, humans naturally filter information by seeking the greatest cognitive benefit for the least mental effort [181]. When a legal researcher checks for conflicts, they do not simply count matching words; they instinctively evaluate the “intent” and “applicability” of a rule based on their background knowledge of the legal system. This human-centric view suggests that a document is relevant only if it provides a new piece of information that changes the researcher’s understanding of the legal landscape [170].

Beyond the cognitive baseline, the theoretical definition of relevance in legal information retrieval is anchored in the philosophical requirements of a functional legal system. According to the jurisprudence established by Lon Fuller in *The Morality of Law*, a primary requirement for the internal morality of law is consistency; a system that issues contradictory rules fails to be a legal system at all [61]. Within this philosophical framework, relevance is not merely about finding documents that talk about the same subject matter, but about identifying provisions that share a “normative overlap.” This means that a retrieved document is relevant if it occupies the same jurisdictional or behavioral space as a draft municipal ordinance, potentially creating a situation where a citizen cannot obey one law without violating another. Grounding the retrieval stage in this quest for consistency treats relevance as the systematic identification of potential normative collisions and threats to the rule of law [61].

However, traditional computational methods struggle to replicate this human judgment because they rely on a rigid, mathematical definition of relevance. In standard deterministic systems, relevance is treated as a binary mathematical property based entirely on set intersection [17]. If we denote a candidate document as D and a user query as Q , this lexical relevance is formally evaluated as a binary function $R(D, Q) \in \{0, 1\}$. Under this strict lexical definition, a document is deemed fully relevant (i.e. 1) if and only if it contains the exact sequence of characters requested by the user [205]. Consequently, this mechanism establishes a purely topical baseline, effectively grouping documents together simply because they share a surface-level vocabulary [17]. While this mathematical simplicity allows for rapid indexing across massive databases, it completely strips away the legal context required for complex legislative analysis.

This stark discrepancy between superficial lexical matching and actual legal consequence represents a major stumbling block for legacy search architectures. In statutory law, two documents discussing the exact same topic might be lexically relevant to one another, but they are only *legally* relevant if one provision actively modifies, permits, restricts, or preempts the other [205]. For instance, a retrieved document might use identical vocabulary but apply to a completely different jurisdiction or regulatory scope, rendering it completely useless for an ex-ante conflict check. Conversely, a document that governs the exact same civic behavior but uses entirely different phrasing—such as substituting “motor vehicle” for “automobile”—will be scored as mathematically irrelevant (i.e. 0) by

a strict Boolean system [17]. Therefore, to accurately audit local ordinances against superior law, statutory retrieval systems must move beyond binary, keyword-driven matching in favor of dense semantic representations capable of capturing actual jurisprudential relationships [98].

To resolve these limitations, modern computational pipelines fundamentally redefine relevance spatially rather than lexically. Instead of scanning a query as a raw string of text, dense retrieval architectures utilize Bi-Encoder models to compress entire legal provisions into high-dimensional mathematical coordinates, commonly referred to as dense vectors [150]. Under this advanced semantic framework, relevance is no longer determined by counting shared characters, but by calculating the geometric distance—such as cosine similarity—between these coordinate points [98]. If a newly drafted local ordinance and an existing national statute govern the same underlying civic behavior, the neural network learns to map their respective vectors closely together in the mathematical space, completely bypassing the vocabulary mismatch problem [219]. Consequently, the retrieval mechanism evaluates relevance as a continuous proximity score rather than a rigid binary output, successfully surfacing conceptually intertwined laws that basic keyword searches otherwise miss.

2.8.3 Long-Context Retrieval

Relying on standard BERT-based architectures introduces an artificial constraint on legal analysis: the 512-token limit. Because standard Self-Attention mechanisms scale quadratically with sequence length, processing long texts requires severe truncation. To get around this, researchers like Chalkidis et al. [27] routinely use Sliding Window techniques to segment documents into overlapping chunks. However, applying this to Philippine ordinances—which are historically wordy and often bury exceptions deep within 1,000-token provisions—is structurally destructive. It severs the substantive semantic link between a primary rule and its distant exception.

Consequently, the model operates with a logical blindspot. Beltagy et al. [13] note that this causes logical hallucinations, where a rule might be flagged as an absolute prohibition simply because its exempting clause was pushed into a separate, unconnected processing window. Furthermore, these systems must navigate the trade-off between Sparse Retrieval (e.g., BM25) for exact keyword matching, and Dense Retrieval for capturing conceptual meaning even without direct word overlap.

Resolving this fragmentation means moving beyond short sliding-window chunking where possible. The architectural solution requires an initial retrieval mechanism capable of native long-context retention. Recent Transformer models incorporate advanced positional encoding, such as Rotary Positional Embeddings (RoPE), alongside sparse or optimized attention matrices. Warner et al. [215] demonstrated that these upgrades extend context windows up to 8,192 tokens without triggering memory failures. Deploying this class of architecture allows the system to read entire local ordinances as cohesive semantic units, preserving the law’s internal structure and eliminating the Sliding Window fallacy.

2.8.4 Retrieval and Fusion Strategies

In computational legal informatics, isolating candidate enactments from extensive statutory collections requires understanding the theoretical foundations, failure modes, and complementary trade-offs of diverse retrieval strategies. Recent findings from the COLIEE benchmarks and legislative consistency literature identify five core methodological dimensions governing coarse statutory search.

The foundational retrieval strategies in computational law encompass sparse probabilistic scoring and dense semantic embeddings. Each relevance scoring metric is to be computed as follows:

Okapi BM25 Probabilistic Relevance Scoring Function:

$$S_{\text{BM25}}(D, Q) = \sum_{q \in Q} \text{IDF}(q) \cdot \frac{f(q, D) \cdot (k_1 + 1)}{f(q, D) + k_1 \cdot \left(1 - b + b \cdot \frac{|D|}{\text{avgdl}}\right)} \quad (2.3)$$

Where $f(q, D)$ denotes term frequency of query token q in document D , $|D|/\text{avgdl}$ normalizes for provision length, and parameters $k_1 = 1.5$ and $b = 0.75$ govern term saturation and length penalization within Robertson and Zaragoza’s probabilistic relevance framework [157]. In statutory retrieval competitions such as COLIEE Task 3, BM25 serves as an indispensable baseline (as demonstrated by Team AIIR [219]) because legislative drafting routinely depends on exact statutory citations, formal administrative acronyms, and verbatim penal monetary amounts. However, exact lexical matching suffers from the vocabulary mismatch problem when local ordinances paraphrase statutory commands [17, 122].

Dense Semantic Cosine Proximity Relevance Function:

$$S_{\text{dense}}(D, Q) = \frac{\mathbf{e}_Q \cdot \mathbf{e}_D}{\|\mathbf{e}_Q\| \|\mathbf{e}_D\|} \quad (2.4)$$

Where $\mathbf{e}_Q$ and $\mathbf{e}_D$ represent pooled contextual embeddings generated by dense bi-encoders such as `all-MiniLM-L6-v2` [150]. Rather than evaluating token overlap, dense retrievers score relevance via cosine proximity in latent semantic space, capturing conceptual legal intent and retrieving governing statutes even when local drafts employ disparate phrasing.

In heterogeneous statutory databases spanning tens of thousands of enactments, researchers frequently investigate topic classification or domain partitioning to reduce search space complexity. However, empirical literature highlights the severe danger of hard categorical filtering in legal domains. Foundational statutory instruments, such as the Philippine Local Government Code of 1991 (RA 7160) or comprehensive penal codes, function as omnibus legislation that simultaneously govern municipal administration, taxation, public safety, and environmental sanitation. Hard filtering by a single predicted regulatory category discards the governing omnibus code whenever a query touches an adjacent regulatory area, causing catastrophic candidate recall collapse. In contrast, information retrieval theory establishes soft domain priors [157]. Instead of binary candidate exclusion, soft priors apply a calibrated multiplicative score boost (e.g., +20%) to candidate documents matching the predicted topic or latent semantic cluster, preserving access to multi-domain statutory codes while gently elevating topically congruent provisions.

Because lexical and semantic retrievers exhibit orthogonal failure modes, leading computational law frameworks combine both streams through hybrid fusion. In score-level linear combination, popularized by Team JNLP in the COLIEE competitions [65], weighted interpolation combines normalized sparse and dense scores according to $S_{\text{hybrid}}(D, Q) = \alpha \hat{S}_{\text{dense}}(D, Q) + (1 - \alpha) \hat{S}_{\text{BM25}}(D, Q)$, where $\alpha \in [0, 1]$ balances semantic and lexical bias. However, because BM25 scores are unbounded while cosine similarities span $[-1, 1]$, linear combination requires min-max score normalization, introducing sensitivity across heterogeneous query lengths. To eliminate score normalization artifacts, Cormack et al. [37] formulated Reciprocal Rank Fusion (RRF), defined as $\text{RRF}(D) = \sum_{m \in \mathcal{M}} (k_{\text{rrf}} + \text{rank}_m(D))^{-1}$, where $\mathcal{M} = \{\text{BM25}, \text{Dense}\}$ and $k_{\text{rrf}} = 60$. Because RRF operates strictly on ordinal ranking positions, it provides scale invariance, eliminates distribution sensitivity, and ensures that provisions performing well across both lexical and semantic modalities are prioritized.

Finally, statutory retrieval systems must respect constitutional hierarchies of authority. Anchored in the doctrine of legislative supremacy and the hierarchy of legal norms [96, 70], statutory systems operate under strict formal ordering. In national digital archives, subordinate administrative issuances (such as executive orders, administrative circulars, and departmental memos) frequently constitute nearly half of all indexed records. Because administrative issuances share broad governmental vocabulary with municipal drafts, unconstrained retrieval often allows routine circulars to crowd out superior national statutes. Literature in legal informatics emphasizes structural safeguards that prioritize primary legislative acts (Republic Acts) over subordinate executive issuances. Implementing a hierarchical statutory priority safeguard shields governing statutes from administrative noise, preventing dilution of candidate shortlists and significantly improving Mean Reciprocal Rank (MRR).

2.8.5 Retrieval Benchmarks

Selecting effective retrieval architectures from the extensive landscape of open-weight models requires objective empirical benchmarking. In contemporary legal information retrieval, candidate models are evaluated against standardized empirical leaderboards that assess zero-shot generalization across dense statutory text. The Benchmarking Information Retrieval (BEIR) suite evaluates a dense retriever’s capacity for out-of-domain transfer across diverse, specialized datasets [196]. Because Philippine local legislation lacks extensive annotated training corpora, candidate retrievers must possess strong zero-shot generalization. Dominant performance on the BEIR leaderboard serves as a primary indicator of a model’s capacity to map novel, localized legal terminology into a structured latent space without prior fine-tuning [196].

Expanding on BEIR, the Massive Text Embedding Benchmark (MTEB) establishes an international standard for ranking embedding systems across more than one hundred datasets and diverse linguistic tasks [128]. In information retrieval literature, candidate architectures are frequently filtered using MTEB’s retrieval and clustering leaderboards. Models excelling on these benchmarks demonstrate an empirical capacity to cluster conceptually related doctrines and retrieve relevant provisions based on continuous semantic proximity. Complementing general embedding benchmarks, the annual COLIEE Task 3

evaluation provides a specialized testbed for statute law retrieval [66, 67]. Systems evaluated on Task 3 must extract precise governing articles from civil code collections in response to legal queries. Success in this competition requires balancing broad candidate recall with computational throughput, providing a proven testing foundation for leading statutory retrieval models [219, 65].

2.8.6 Candidate Retrieval Models

Surveying leading open-weight architectures documented across recent retrieval literature and standardized benchmarks reveals eight prominent bi-encoder models relevant to statutory candidate discovery. Among distilled edge baselines, Reimers and Gurevych [150] established the Sentence-BERT framework, mapping text into fixed-dimensional dense vector spaces optimized for cosine proximity. The **all-MiniLM-L6-v2** model compresses a 6-layer BERT architecture down to 22.7 million parameters through knowledge distillation, producing 384-dimensional dense vectors. Despite its ultra-compact footprint, empirical benchmarks demonstrate that it retains strong semantic clustering capabilities with sub-two-millisecond latency on consumer CPUs. Complementing this, Xiao et al. [221] introduced **bge-small-en-v1.5** (33.4M parameters), which optimizes contrastive representation learning for high rank sensitivity on MTEB. In public-sector edge deployment, these distilled bi-encoders provide high-throughput candidate filtering without requiring dedicated GPU accelerators.

Scaling bi-encoder capacity to standard 110M-parameter base dimensions enables more expressive semantic representations across general retrieval benchmarks. Xiao et al. [221] developed **bge-base-en-v1.5** (109M parameters, 768 dimensions), utilizing retrofitted contrastive pretraining and instruction-tuned query prefixes to excel across multi-domain retrieval leaderboards. Similarly, Wang et al. [213] introduced **e5-base-v2** (109M parameters), trained on over one billion weakly-supervised text pairs via contrastive objectives, demonstrating high generalization on unseen technical and legal search tasks. Furthermore, Song et al. [179] developed **all-mpnet-base-v2** (109M parameters), combining masked language modeling with permuted language pretraining to eliminate positional dependence and generate highly stable semantic embeddings.

Beyond general-domain models, legal-domain adapted encoders capture specialized statutory taxonomies. Chalkidis et al. [25] demonstrated that pretraining bidirectional encoders on specialized legal repositories enables dense representations to capture statutory taxonomy and regulatory jargon that general-domain embedders conflate. Mean-pooled sentence embeddings derived from **legal-bert-base-uncased** (110M parameters) map statutory provisions into vector spaces that reflect legislative hierarchy and administrative codification rather than colloquial word associations. This domain-specific pretraining provides a structured baseline to evaluate whether legal pre-training improves candidate provision ranking over general embeddings.

To overcome the quadratic attention complexity of traditional 512-token encoders, modern architectures incorporate native long-context and multilingual designs. Warner et al. [216] developed **ModernBERT-base** (149M parameters), which integrates FlashAttention-2, Rotary Positional Embeddings (RoPE), GeGLU activations, and

sequence unpadding to natively process up to 8,192 tokens. In statutory retrieval, this allows direct encoding of verbose multi-paragraph enactments without destructive sliding-window chunking. Expanding to the multilingual domain, Chen et al. [29] introduced **bge-m3** (568M parameters, 1,024 dimensions), supporting 8,192-token contexts and multi-functionality across dense, sparse, and multi-vector representations. In statutory consistency checking, its multi-vector representations provide resilience across code-switched municipal records and bilingual administrative issuances.

In heterogeneous statutory databases spanning tens of thousands of enactments, researchers frequently investigate topic classification or domain partitioning to reduce search space complexity. However, empirical literature highlights the severe danger of hard categorical filtering in legal domains. Foundational statutory instruments, such as the Philippine Local Government Code of 1991 (RA 7160) or comprehensive penal codes, function as omnibus legislation that simultaneously govern municipal administration, taxation, public safety, and environmental sanitation. Hard filtering by a single predicted regulatory category discards the governing omnibus code whenever a query touches an adjacent regulatory area, causing catastrophic candidate recall collapse. In contrast, information retrieval theory establishes soft domain priors [157]. Instead of binary candidate exclusion, soft priors apply a calibrated multiplicative score boost (e.g., +20%) to candidate documents matching the predicted topic or latent semantic cluster, preserving access to multi-domain statutory codes while gently elevating topically congruent provisions.

Because lexical and semantic retrievers exhibit orthogonal failure modes, leading computational law frameworks combine both streams through hybrid fusion. In score-level linear combination, popularized by Team JNLP in the COLIEE competitions [65], weighted interpolation combines normalized sparse and dense scores according to $S_{\text{hybrid}}(D, Q) = \alpha \hat{S}_{\text{dense}}(D, Q) + (1 - \alpha) \hat{S}_{\text{BM25}}(D, Q)$, where $\alpha \in [0, 1]$ balances semantic and lexical bias. However, because BM25 scores are unbounded while cosine similarities span $[-1, 1]$, linear combination requires min-max score normalization, introducing sensitivity across heterogeneous query lengths. To eliminate score normalization artifacts, Cormack et al. [37] formulated Reciprocal Rank Fusion (RRF), defined as $\text{RRF}(D) = \sum_{m \in \mathcal{M}} (k_{\text{rrf}} + \text{rank}_m(D))^{-1}$, where $\mathcal{M} = \{\text{BM25}, \text{Dense}\}$ and $k_{\text{rrf}} = 60$. Because RRF operates strictly on ordinal ranking positions, it provides scale invariance, eliminates distribution sensitivity, and ensures that provisions performing well across both lexical and semantic modalities are prioritized.

Finally, statutory retrieval systems must respect constitutional hierarchies of authority. Anchored in the doctrine of legislative supremacy and the hierarchy of legal norms [96, 70], statutory systems operate under strict formal ordering. In national digital archives, subordinate administrative issuances (such as executive orders, administrative circulars, and departmental memos) frequently constitute nearly half of all indexed records. Because administrative issuances share broad governmental vocabulary with municipal drafts, unconstrained retrieval often allows routine circulars to crowd out superior national statutes. Literature in legal informatics emphasizes structural safeguards that prioritize primary legislative acts (Republic Acts) over subordinate executive issuances. Implementing a hierarchical statutory priority safeguard shields governing statutes from administrative noise, preventing dilution of candidate shortlists and significantly improv-

ing Mean Reciprocal Rank (MRR).

2.9 Natural Language Processing

Natural Language Processing (NLP) provides the computational techniques needed to analyze and represent text at various linguistic levels, aiming for human-like comprehension [112]. As Jurafsky and Martin [91] explain, it acts as the bridge between human language and machine logic, allowing algorithms to parse and interpret unstructured text. In the legal field, NLP shifts document auditing from slow manual reviews to scalable, automated systems. However, applying standard NLP to law requires more than basic grammar checks; it demands systems that can navigate dense sentence structures, deontic operators (obligations, permissions, prohibitions), and the highly specific vocabulary found in statutory law.

Early Legal NLP relied heavily on Symbolic Artificial Intelligence (AI) [95]. These systems used manually built ontologies and strict Boolean logic to extract legal information. However, because municipal drafting language changes constantly over decades, these exact-match frameworks proved far too brittle. Contemporary NLP architectures solve this by moving to continuous vector representations. By mapping textual provisions into high-dimensional semantic spaces, modern Transformers bypass superficial vocabulary mismatches to mathematically quantify the deeper concepts connecting different legal terms. This allows the system to act as an active diagnostic engine that structurally maps legal conflicts, rather than just a passive search index.

This section breaks down how NLP operates within legislative analysis. It covers the systemic challenges of processing localized legal text, the structural need for Natural Language Inference (NLI), the benchmarks used for model selection, and the specific architectures built for these deductive tasks.

2.9.1 NLP Challenges

Deploying NLP frameworks against municipal archives quickly exposes the mechanical limits of standard Language Models. A primary hurdle is the pronounced Global North bias in pre-training data and legal benchmarking. As Ranathunga et al. [148] point out, the vast majority of Legal NLP benchmarks, like LexGLUE [27] and CaseHOLD, are trained exclusively on pristine, digitally native Western legal text. These highly optimized models degrade when applied to Philippine Local Government Unit (LGU) archives, which are filled with scanned PDFs suffering from Optical Character Recognition (OCR) errors, archaic American colonial phrasing, untranslated Spanish legal terms, and local Taglish. Rust et al. [162] demonstrated that standard WordPiece tokenizers shatter unrecognized words into sub-tokens, destroying a model’s ability to infer legal intent. Addressing localized morphological noise requires deploying models equipped with massive Byte-Pair Encoding (BPE) vocabularies that absorb linguistic variants without triggering out-of-vocabulary errors, paired with Disentangled Attention architectures that separate word content from corrupted positions [80].

Beyond geographic and morphological bias, traditional NLP methods suffer from structural shortcomings in lexical similarity. Early NLP systems excel at identifying topical overlap by measuring surface-level token intersection using algorithms like TF-IDF or shallow word embeddings. However, Bowman et al. [18] explain that pure lexical similarity is fundamentally unsuited for detecting legal contradictions because lexical metrics are entirely blind to deontic logic and syntactic negation. For example, an ordinance stating that the city permits an activity and a statute stating that the state strictly prohibits it share near-perfect lexical similarity scores. A standard similarity model clusters them as highly related, completely missing the absolute policy collision occurring beneath the shared vocabulary. Consequently, detecting normative conflicts requires advancing from lexical proximity to directional logical inference.

2.9.2 From NLP to NLI

Historically, checking legislative consistency relied on Symbolic NLP (hand-written rules) and, later, Statistical NLP. In Statistical NLP, models learned probabilistic weights for specific inputs, like Term Frequency-Inverse Document Frequency (TF-IDF) or N-grams [95]. These models depended heavily on human-engineered features to define grammar rules and categories. While they allowed for soft, probabilistic decisions, they consistently struggled with the deep lexical and semantic ambiguities inherent in legal drafting.

The field largely transitioned to Neural NLP models in the mid-2010s [145]. Unlike statistical systems that treat features as separate entities, neural networks operate end-to-end, mapping raw text directly to output predictions through hidden layers. These architectures capture deeper linguistic patterns than traditional statistical methods by learning complex relationships during optimization. However, while machine learning excels at finding patterns, it can overly mimic its training data, making domain-specific pre-training essential for high performance in rigorous legal tasks.

Within modern neural approaches, Natural Language Inference (NLI) has emerged as the definitive benchmark for evaluating a model’s capacity for actual reasoning. Yang et al. [223] define NLI as determining the logical relationship between two text sequences: a “premise” and a “hypothesis.” The model must classify whether the hypothesis is true (Entailment), false (Contradiction), or unrelated (Neutral) based solely on the premise. This specific task exposes the core flaw of earlier statistical NLP, which mistakenly equated high word overlap with logical entailment [223]. Modern neural approaches fix this by capturing deep compositional semantics. As Le Hong et al. [107] point out, this allows modern models to resolve complex linguistic phenomena—like coreferences and modality—that rigid rule-based systems historically failed to grasp.

2.9.3 NLI Benchmarks

Identifying effective Cross-Encoder architectures for the inference stage requires isolating models with robust, pre-trained deductive logic. In natural language processing, research literature relies on established benchmarks to evaluate candidate reasoning models. Wang et al. introduced the General Language Understanding Evaluation (GLUE) [212] and its successor SuperGLUE [211] as foundational baselines for model reasoning. These frame-

works evaluate a Transformer’s ability to handle diverse linguistic challenges, such as Recognizing Textual Entailment (RTE) and resolving complex coreferences. While high performance on SuperGLUE confirms baseline deductive capacity, it does not independently guarantee success in the legal domain due to the specialized vocabulary shifts present in statutory drafting.

To evaluate specialized legal logic, legal NLP research relies on the Legal General Language Understanding Evaluation (LexGLUE) benchmark. Chalkidis et al. [27] designed LexGLUE specifically to test a model’s ability to navigate legalese, where common words assume highly context-dependent meanings. High performance on LexGLUE entailment tasks indicates that an architecture possesses the foundational deductive capacity required to navigate complex statutory drafting and regulatory provisions.

Complementing domain-specific suites, contemporary Hugging Face NLI leaderboards evaluate models on diverse, adversarial textual entailment tasks aggregating MultiNLI (MNLI), FEVER, Adversarial NLI (ANLI R1–R3) [134], and WANLI [114]. Benchmarking candidate models across these suites tests their resilience against spurious lexical heuristics, premise-hypothesis artifacts, and negation traps. High-performing architectures on these leaderboards, such as the pre-aligned DeBERTa-v3 variants curated by Laurer et al. [106], demonstrate strong zero-shot deductive capacity, serving as leading candidate foundations for statutory conflict detection. Finally, within legal informatics, the definitive international competition benchmark for statutory inference is Task 4 of the Competition on Legal Information Extraction and Entailment (COLIEE) [146, 66, 67]. In Task 4, systems evaluate whether a natural language query is legally entailed by or contradicts a governing civil code article. Evaluating architectures against COLIEE Task 4 methodologies provides direct empirical evidence of a model’s capacity to navigate statutory conditions, exceptions, and deontic modalities.

2.9.4 Candidate NLI Models

Evaluating deontic relationships across statutory provisions demands specialized inductive biases across candidate sequence encoders. Standard language models trained exclusively on general corpora frequently fail to capture the specialized semantics of legal drafting, where everyday terms carry specialized juridical meanings. To address this vocabulary divergence, researchers develop domain-adapted encoders initialized with custom legal vocabularies or continuously pretrained on expansive legal repositories. Chalkidis et al. [25] introduced **LEGAL-BERT**, pretrained on diverse English legislation, court judgments, and contracts, demonstrating superior term sensitivity across legal classification tasks. Henderson et al. [81] expanded this approach with **PoL-BERT** (Pile of Law), pretraining encoders on a 256 GB corpus encompassing administrative regulations, municipal codes, and statutory provisions. Complementary efforts by Paul et al. [141] (**InLegalBERT**) and Zheng et al. [227] (**CaseLaw-BERT**) investigate the transferability of common-law statutory pretraining, while Xiao et al. [220] evaluate long-document legal encoders. In literature, this family provides a critical testbed for assessing whether domain pretraining provides intrinsic protection against legal misinterpretation.

A pervasive vulnerability of unadapted sequence encoders on deontic tasks is affir-

mative bias, where models default toward predicting Entailment or Neutral when reading complex clauses containing subtle prohibitions [5, 199]. Devlin et al. [48] established **bert-base-uncased** (110M parameters) with masked language modeling, serving as the canonical supervised baseline against which improvements are measured. To overcome heuristic reliance on lexical overlap, contemporary literature explores models pre-aligned on multi-task natural language inference datasets (such as MNLI, FEVER, ANLI, and WANLI). Laurer et al. [106] demonstrate that initializing sequence classifiers from warm NLI checkpoints equips models with robust directional contradiction sensitivity before domain adaptation. Structurally, He et al. [80] developed **DeBERTa-v3**, introducing Disentangled Attention to represent word content and relative positions with separate vectors alongside adversarial replaced token generation. In parallel, Clark et al. [32] developed **ELECTRA** with Replaced Token Detection, optimizing contextual representations across all input positions. These architectures excel at distinguishing subtle deontic shifts that standard unidirectional or masked encoders routinely miss.

Statutory drafting is inherently verbose, characterized by complex preamble clauses, extensive definitions, and multi-sentence penal qualifications. Traditional transformer encoders are constrained by quadratic self-attention complexity, imposing a rigid 512-token ceiling that necessitates destructive document chunking or sliding window heuristics [27]. Early long-context architectures attempted to resolve this constraint using localized or sparse attention approximations, such as the sliding-window attention in **Longformer** [13] and the block-sparse random attention in **BigBird** [225]. However, empirical legal informatics evaluations note that sparse attention patterns can dilute or sever semantic dependencies between distant provisos and governing commands. To overcome these limitations, modern architectures integrate architectural optimizations that support extended contexts natively. Warner et al. [216] introduced **ModernBERT**, an 8,192-token encoder incorporating FlashAttention-2, Rotary Positional Embeddings (RoPE), GeGLU activation functions, and sequence unpadding. Complementing this, Nussbaum et al. [136] developed **Nomic-BERT** to extend context retention. These modern long-context designs enable joint encoding of complete multi-sentence statutory provisions without sparse approximation artifacts.

Unlike bi-encoders that process queries and documents independently into pooled vector coordinates, cross-encoders execute full bidirectional self-attention across the concatenated pair simultaneously. In retrieval and inference benchmarks, joint cross-attention captures deep inter-token interactions that vector similarity cannot represent. In COLIEE statute retrieval benchmarks, researchers deploy multilingual cross-encoder rerankers, notably **BGE-Reranker-v2-m3** [221], to score pairwise relevance across multi-thousand token contexts. Similarly, sequence classifiers adapted from MS-MARCO reranking frameworks employ deep cross-attentive scoring to evaluate pair compatibility. In literature, native cross-encoders represent the conceptual bridge between Stage 1 candidate filtering and Stage 2 logical validation.

In post-colonial statutory jurisdictions like the Philippines, statutory archives exhibit unique sociolinguistic complexities. While statutory enactments are drafted in legal English, municipal ordinances and administrative records frequently feature localized terminology, regional administrative phrasing, Latin legal maxims, and English-Tagalog

code-switching. Conneau et al. [36] introduced **XLM-RoBERTa**, trained across 100 languages, demonstrating strong cross-lingual transfer on semantic benchmarks. He et al. [80] extended disentangled attention to the multilingual domain with **mDeBERTa-v3**, providing resilience to mixed-language syntax. Furthermore, Cruz and Cheng [40] developed Tagalog-adapted transformers (**roberta-tagalog-base**) to capture Philippine syntactic structures. In statutory consistency checking, multilingual encoders evaluate whether localized linguistic variances can be processed without catastrophic comprehension loss.

Deploying deep learning in public-sector municipal environments requires strict adherence to available edge computing infrastructure. Because local government units typically operate without high-performance computing clusters or enterprise cloud APIs, running billion-parameter language models is computationally unfeasible. To establish realistic performance baselines under constrained hardware, the literature explores knowledge distillation and parameter sharing. Sanh et al. [169] introduced **DistilBERT**, which compresses transformer depth while retaining over 95% of language comprehension. Lan et al. [102] introduced **ALBERT**, utilizing cross-layer parameter sharing to reduce memory footprint down to tens of megabytes. In dense retrieval and inference, Reimers and Gurevych [150] established **all-MiniLM-L6-v2** as an ultra-compact 22-million parameter baseline. Evaluating distilled edge models allows researchers to quantify the exact accuracy penalty incurred when prioritizing sub-second CPU latency over multi-layer neural capacity.

2.10 Ground Truth

In supervised machine learning, a Ground Truth (or Gold Standard) dataset provides the correct, human-verified labels used to train and test models. During training, the neural network uses this dataset as an answer key. It constantly compares its own predictions against these true labels, calculating the difference (or loss) and adjusting its internal weights to make fewer mistakes over time [68]. Without a highly accurate set of target labels, the model will essentially learn the wrong rules, leading to a phenomenon where it memorizes flawed data instead of learning actual reasoning patterns.

Beyond just teaching the model, the ground truth is the only objective way to measure success. Researchers typically split this data into training, validation, and testing sets. The test set acts as the final exam, kept completely hidden from the model until the very end. Evaluating the model against this unseen ground truth allows researchers to calculate standard performance metrics like precision, recall, and F1-score. Without this baseline, it is mathematically impossible to prove that a model generalizes well to new data or works as intended in a real-world scenario [19].

2.10.1 Ground Truth in NLI

In Natural Language Inference (NLI), ground truth datasets define the logical relationship between two pieces of text. Annotators read a source text (the Premise) and a claim (the Hypothesis), then categorize their relationship into one of three distinct labels. If the premise guarantees the hypothesis is true, it is labeled as Entailment. If they contradict

each other, it is labeled as Contradiction. If there is not enough information to prove either, the pair is labeled Neutral [223].

Most traditional NLI datasets, such as the widely used SNLI (Stanford Natural Language Inference) corpus, are built using crowdsourced workers who write simple variations of everyday sentences [19]. While this approach is cheap and easily generates hundreds of thousands of training pairs, the resulting datasets are incredibly basic and often embed unintended annotation artifacts [77]. They usually feature short, single-sentence premises with obvious logical connections. Training a model exclusively on these simple datasets leaves it entirely unprepared for the dense, multi-layered logic required in specialized fields.

2.10.2 Annotation and Reliability Metrics

Turning raw text into a useful NLI dataset requires strict guidelines to minimize human bias and prevent annotators from relying on subjective legal interpretations. In standard NLI, guidelines dictate simple factual agreement. However, legal conflict detection requires evaluating normative relationships based on deontic logic. Drawing from the frameworks established by Agotnes et al. [3] and Li et al. [111], the annotation guidelines strictly instruct experts to map the text into obligations (must do), permissions (may do), and prohibitions (cannot do). A "Contradiction" is only flagged if the local ordinance mandates an action the national statute prohibits, or prohibits an action the national statute explicitly permits. If a local ordinance merely adds a regulatory step (e.g., a local processing fee) to a national permit without banning the core activity, the guidelines strictly dictate a "Neutral" classification. This rigid deontic framework is designed to ensure that the dataset measures actual legal preemption rather than just topical friction.

To evaluate the effectiveness of these strict guidelines, researchers must utilize Inter-Rater Reliability (IRR) metrics [125] rather than mathematically flawed percentage agreements. Raw percentage agreement fails to account for statistical probability, where in a multiple-choice classification task, there is an inherent likelihood that annotators will occasionally select the same label by pure random guessing.

The two primary statistical formulas used to subtract this random chance are Cohen's Kappa [33] and Fleiss' Kappa [58]. Each inter-rater reliability metric is to be computed as follows:

Cohen's Pairwise Kappa Reliability Metric (κ):

$$\kappa = \frac{p_o - p_e}{1 - p_e} \quad (2.5)$$

Where p_o represents the relative observed agreement between the two individual raters, and p_e represents the hypothetical probability of chance agreement. Cohen's Kappa evaluates agreement between exactly two raters.

Fleiss' Generalized Multi-Rater Kappa Reliability Metric (κ):

$$\kappa = \frac{\bar{P} - \bar{P}_e}{1 - \bar{P}_e} \quad (2.6)$$

Where $\bar{P}$ represents the mean observed agreement across all pairs of raters, while $\bar{P}_e$ represents the expected chance agreement across the entire pool. Fleiss’ Kappa generalizes Cohen’s formula to accommodate three or more domain experts establishing a majority consensus.

The resulting κ score is interpreted using the universally accepted statistical matrix established by Landis and Koch [104]:

Table 2.1: Interpretation Ranges for Fleiss’ Kappa Reliability Metric

Kappa (κ) Score	Interpretation
< 0.00	Poor agreement
$0.01 - 0.20$	Slight agreement
$0.21 - 0.40$	Fair agreement
$0.41 - 0.60$	Moderate agreement
$0.61 - 0.80$	Substantial agreement
$0.81 - 1.00$	Almost perfect agreement

Table 2.1 outlines the standard categorical intervals for evaluating inter-rater agreement across multiple independent judges. It consists of two columns mapping numerical κ ranges to their corresponding qualitative interpretations, providing an objective statistical benchmark for assessing annotation consensus. The table is included to ground the evaluation of the 15 Sangguniang Panlungsod legal researchers within an authoritative psychometric standard. As the score progresses from negative values toward unity, it reflects a transition from chance-level alignment to robust consensus across evaluators. For this research, the matrix establishes the operational target of $\kappa \geq 0.61$ (“Substantial Agreement”), ensuring that the resulting ground truth benchmark represents dependable legal consistency rather than subjective individual interpretation.

While objective tasks routinely achieve κ scores above 0.90, applying this expectation to the legal domain is mathematically unrealistic. Bowman et al. [19], in their foundational study for the Stanford Natural Language Inference (SNLI) corpus, established that evaluating semantic relationships is inherently subjective. Their benchmark study hit an approximate performance ceiling with an overall Fleiss’ Kappa of 0.70. They explicitly demonstrated that achieving near-perfect agreement (> 0.80) in textual inference usually indicates overly simplistic or fabricated test data.

2.10.2.1 Sub-Threshold Reliability

While achieving a “Substantial” Inter-Rater Reliability score ($\kappa > 0.61$) validates the clarity of the annotation guidelines, failing to meet this threshold does not inherently invalidate the dataset or halt the experimental pipeline. Klebanov and Beigman [99] establish that a sub-threshold κ score typically indicates systematic ambiguity within the text itself—categorized as “hard” instances—rather than incompetence by the annotators. Their research demonstrates that it is still entirely possible to extract a reliable, low-noise subset from the instances where annotators do achieve consensus.

When researchers encounter this sub-threshold reliability, the mitigation strategy depends entirely on the size of the dataset. For large-scale studies, including the disaster-tweet classification by Ablazo [1], the standard procedure is data omission. Beigman et al. [99] explicitly advocate for filtering out difficult, borderline cases from the gold standard, arguing that evaluating algorithms on inherently ambiguous text introduces statistical noise akin to random coin-flips. In these massive datasets, items that fail to reach majority consensus are simply discarded because the volume of remaining "easy" data easily absorbs the loss.

However, applying this standard omission strategy to a specialized, few-shot legal dataset mathematically destroys the test set's Minimum Detectable Effect. Furthermore, in municipal legislation, these "hard" cases often represent the exact jurisdictional friction the logic engine needs to learn. To preserve dataset volume without introducing benchmarking noise, researchers in specialized domains employ Expert Adjudication rather than omission. Gao et al. [64] establish that in these scenarios, a hierarchical adjudication procedure is used where a senior domain expert reviews the conflicting annotations. The senior expert resolves the ambiguity based on governing jurisprudential principles and manually establishes the Gold Standard label, treating the initial low κ metric as a measure of text difficulty rather than a rigid experimental blocker.

2.10.3 Ground Truth in Legal NLP

When applying NLI to the legal domain, the standard approach of comparing short, isolated sentences quickly breaks down. Statutory phrasing is notoriously complex, filled with conditional clauses, exceptions, and domain-specific vocabulary. Consequently, a hypothesis in a legal context rarely aligns perfectly with a single sentence in the premise. Instead, proving an entailment or contradiction often requires the model to connect evidence scattered across multiple paragraphs or pages.

To address this, researchers have developed specialized legal datasets that reflect real-world document structures. For instance, the Competition on Legal Information Extraction/Entailment (COLIEE) dataset constructs its ground truth by pairing civil code statutes with legal bar exam queries, requiring models to identify specific legal provisions before determining entailment [146]. Similarly, the ContractNLI dataset by Koreeda & Manning [100] engineers its gold standard around non-disclosure agreements. Rather than giving the model pre-cut sentence pairs, ContractNLI forces the system to evaluate hypotheses against entire legal documents, requiring annotators to both classify the logical relationship and extract the specific supporting text span.

Furthermore, unlike general-domain datasets crowdsourced through platforms like Amazon Mechanical Turk, legal ground truth relies heavily on domain experts. Constructing datasets like CaseHOLD [227], which evaluates legal holding statements, requires law students and legal professionals to ensure the labels accurately reflect actual jurisprudence. The push towards comprehensive legal benchmarks, such as LexGLUE [27], demonstrates that evaluating legal NLP models demands rigorously structured, expert-annotated data rather than automated or amateur labeling.

This difficulty is magnified when dealing with local or municipal legislation. While

national or federal laws generally follow standardized drafting conventions, municipal texts are highly heterogeneous, exhibiting significant variations in terminology, structure, and overall organization across different jurisdictions [21]. Because of this localized variance, municipal legal documents have historically been neglected in NLP research, hampered primarily by a severe lack of annotated ground truth datasets [22]. Standard legal corpora fail to capture the specific legislative templates, localized governance terminology, or idiosyncratic cross-referencing styles used in municipal laws, such as Philippine local ordinances. Without existing resources that reflect these structural and linguistic complexities, evaluating models on local legislation requires building highly specialized, expert-annotated datasets from the ground up.

A critical architectural decision in dataset construction is determining the exact granularity of the extracted legal premise. While the hypothesis (the local draft) is evaluated as a multi-sentence span to capture context, the premise (the national statute) must be strictly isolated at the sentence level. In Philippine statutory drafting, a single paragraph frequently contains a general rule, multiple conditional modifiers, and distinct exceptions. Holzenberger et al. [83] demonstrate that feeding an entire statutory paragraph into an NLI model as a premise dilutes the model’s attention mechanism, causing it to lose track of specific variable constraints. By forcing human annotators to isolate the national rule down to its exact sentence boundary, the resulting Ground Truth prevents the candidate logic engines from getting distracted by surrounding, non-governing text, ensuring the classification mathematically targets the exact point of semantic collision.

The difficulty of dataset construction is magnified by the scarcity of legal domain experts available for annotation. As Zhong et al. [228] explicitly note, the complex terminology required for legal texts makes manual labeling highly time-consuming and difficult to scale, as it strictly requires domain experts rather than crowdsourced workers. Because constructing massive datasets for localized municipal law is practically unfeasible, researchers must rely on constrained, highly curated Gold Standards. To compensate for this limited data volume, the current NLP standard mandates treating legal conflict detection as a few-shot learning scenario. Gao et al. [63] validate this approach, proving that transformer models can be successfully fine-tuned on specialized datasets containing merely hundreds of examples, provided the evaluation framework rigorously controls for overfitting.

2.10.4 Multi-Sentence and Cross-Document Limits

Traditional Natural Language Inference (NLI) evaluates text using a strict 1-to-1 pairing, comparing a single premise sentence directly against a single hypothesis sentence. However, this highly granular approach fundamentally breaks down in the legal domain due to two specific edge cases. First, intra-document context: a primary rule in one section of an ordinance is frequently modified or negated by an exception buried several paragraphs later (the "1v2" or multi-sentence problem). Second, cross-document dependencies: a contradiction may only emerge when synthesizing provisions from two entirely separate statutes simultaneously.

To address the first edge case, modern legal frameworks utilize Asymmetric Span-

Level Pairing. Yin et al. [224] critique traditional sentence-level datasets, proving that capturing contextual exceptions requires evaluating multi-sentence blocks rather than isolated lines. As demonstrated by Koreeda and Manning in the ContractNLI dataset [100], the premise remains a single, focused legal rule, while the corresponding hypothesis is extracted as a multi-sentence chunk. This asymmetric structure allows the model’s attention mechanism to accurately evaluate a primary rule against a distant exception contained within the same local provision.

While asymmetric spans successfully capture intra-document context, resolving the second edge case—cross-document multi-hop reasoning—remains computationally prohibitive. Identifying semantic preemption that spans multiple independent laws currently represents the bleeding edge of NLP research and is not yet fully explored for local governance. Recent frameworks have attempted to solve this, but they introduce severe architectural overhead. For instance, Lu et al. [118] address cross-document retrieval by programmatically mapping Entity Graphs to mine evidence paths between laws, a process that is highly resource-intensive and unscalable for unstructured LGU archives. Alternatively, Shen et al. [176] proposed the HopWeaver framework, which relies on massive generative Large Language Models (LLMs) to synthesize complex reasoning paths across documents.

Both of these state-of-the-art solutions immediately violate the strict hardware constraints of consumer-grade LGU workstations. Consequently, practical, edge-deployable systems must intentionally bound their theoretical scope to single-premise, asymmetric span evaluations, deliberately avoiding the quadratic memory scaling required for multi-document fusion.

2.10.5 Dataset Partitions

When evaluating NLI models on constrained datasets, researchers often default to k -fold cross-validation to maximize testing coverage. However, Vabalas et al. [202] mathematically demonstrated that applying cross-validation to datasets under 1,000 items introduces severe optimistic bias, causing models to perform significantly worse in real-world deployment than their test scores suggest. To rigorously control for overfitting in few-shot legal scenarios, contemporary evaluation frameworks must abandon k -fold approaches. Instead, they require strict, static Train/Validation/Test partitions (e.g., 70/15/15) to ensure the evaluation metrics remain entirely untainted by training exposure.

Beyond merely preventing optimistic bias, utilizing a static partition allows for rigorous validation monitoring during the fine-tuning phase. In datasets with limited volume, neural networks are highly susceptible to memorizing the training data rather than learning underlying logical patterns. Goodfellow et al. [68] establish that permanently isolating a validation set allows architectures to employ early stopping criteria, halting the training process the moment the model’s accuracy on unseen data begins to degrade. Furthermore, as Zhang et al. [226] note regarding few-sample fine-tuning, this static separation ensures strict document-level boundaries, preventing clauses from a specific municipal draft from accidentally bleeding across the training and testing sets to artificially inflate performance metrics.

2.10.6 Minimum Detectable Effect

A standard critique of static splits on small datasets is that the resulting test set lacks the statistical power to reliably compare candidate architectures. However, Card et al. [23] established that a test set’s validity is fundamentally determined by its Minimum Detectable Effect (MDE) rather than arbitrary size thresholds. The MDE dictates that if the expected accuracy gap between two models is small, a massive test set is required to establish that observed performance differences exceed random chance. Conversely, because fine-tuning foundational models on a highly novel, localized municipal domain yields massive margins of performance improvement, the required MDE drops significantly. Consequently, even a constrained test set of approximately 50 pairs is mathematically sufficient to definitively detect the optimal logic engine without sacrificing statistical significance [63].

The practical application of MDE fundamentally shifts how model performance is evaluated in specialized domains. In general NLP research, massive test sets are required because researchers are often evaluating models against benchmarks like GLUE to prove marginal, incremental gains of one percent to two percent over previous state-of-the-art architectures [212]. However, applying a base foundation model—which possesses no prior knowledge of Philippine local governance for example—to a highly specialized task generally results in baseline performance barely above random chance [25]. Zheng et al. [227] demonstrate that in the legal domain, fine-tuning foundational architectures on expertly curated pairs triggers a massive, double-digit jump in classification accuracy compared to zero-shot baselines. Because this expected performance delta is substantial, legal NLP frameworks do not require thousands of annotated examples to demonstrate statistically valid improvements, proving that an evaluation set of approximately 50 curated pairs possesses sufficient statistical power to mathematically capture this distinct leap in model reasoning capacity.

2.11 Explainable AI

In legal informatics, a “black box” prediction is fundamentally unacceptable, regardless of nominal classification accuracy. Automated systems that flag normative conflicts between statutory enactments must provide verifiable reasoning (*ratio decidendi*) behind their classifications to genuinely assist legislative analysis [52]. Without transparent justification, legal practitioners and legislative staff cannot trust, audit, or formally defend algorithmic determinations. Consequently, contemporary legal NLP literature emphasizes Explainable AI (XAI) not as an optional secondary layer, but as an intrinsic structural requirement of high-stakes legal decision support, prioritizing faithful mathematical traces over ungrounded generative summaries.

2.11.1 Intrinsic Explainability

Within the explainability literature, researchers distinguish between Post-Hoc Surrogate Models and Intrinsic Explainability Mechanisms. Surrogate methods, such as Local Interpretable Model-agnostic Explanations (LIME) [154] and SHapley Additive exPlanations (SHAP) [120], attempt to approximate the behavior of complex deep networks by fitting

a simplified, interpretable linear model around individual predictions. However, in text classification benchmarks, these external explainers are computationally demanding and inherently prone to approximation errors, as they explain a surrogate rather than the underlying network itself [160]. In contrast, intrinsic approaches extract internal model representations—such as self-attention weights and contextualized hidden states—directly from the transformer layers during the forward pass. Because native attention extraction reflects the mathematical weights computed during inference without fitting secondary approximations, it offers an intrinsically faithful and computationally lightweight mechanism for structural transparency.

While utilizing attention weights as a proxy for feature importance has sparked significant debate, empirical findings demonstrate its utility for targeted textual inspection. Jain and Wallace [90] argued that raw attention distributions do not always correlate consistently with gradient-based feature importance. However, Wiegrefe and Pinter [217] provided critical counter-evidence, demonstrating that attention mechanisms provide faithful representations of internal prioritization under rigorous diagnostic testing. In statutory conflict detection and deontic reasoning, extracting token-level attention spans allows reviewers to identify the specific conflicting provisions (for example, isolating an “absolute prohibition” in a superior statute against an “unauthorized permit” in a local draft), establishing a direct visual link between textual evidence and classification decisions without incurring the computational overhead of auxiliary surrogate models.

Complementing structural attention heatmaps in modern explainability frameworks is Confidence Calibration. Deep learning models, particularly multi-layer transformers, frequently exhibit severe overconfidence in erroneous predictions when evaluated on out-of-domain or specialized text. As Guo et al. [75] demonstrated, modern neural networks are fundamentally miscalibrated, such that raw softmax probabilities rarely reflect the true empirical likelihood of correctness. In high-stakes regulatory auditing, relying on uncalibrated probabilities produces misleading risk assessments. To restore probabilistic reliability, calibration literature advocates post-hoc temperature scaling applied to pre-softmax logits. Rescaling logits by a learned temperature parameter (τ) before softmax normalization aligns predicted confidence with observed empirical accuracy, enabling decision support systems to distinguish between definitive contradictions and borderline ambiguities with sound statistical justification.

2.11.2 Edge Deployment Feasibility

In edge and on-premise computing environments, operational feasibility requires strict adherence to memory and latency boundaries. Post-hoc explainability methods operate by repeatedly perturbing input tokens, requiring hundreds or thousands of forward passes per document to approximate Shapley values or linear boundary slopes [154, 120]. Empirical evaluations by Kumar et al. [101] confirm that permutation-based explainers inflate per-document inference latency by several orders of magnitude, rendering them computationally prohibitive for lengthy legislative texts. Furthermore, Wang et al. [214] demonstrate that running complex post-hoc explainers on edge workstations severely degrades throughput, making interactive review impractical. In resource-constrained deployment settings, intrinsic extraction of native attention weights and calibrated logits provides a computa-

tionally viable pathway to transparency without requiring specialized high-performance computing clusters.

2.11.3 Explanation Quality Metrics

Evaluating explanation quality in natural language processing requires distinguishing between two fundamental evaluation criteria: Plausibility and Faithfulness. Plausibility reflects how convincing, fluent, and intuitive an explanation appears to a human reader, whereas Faithfulness measures the degree to which an explanation accurately reflects the internal reasoning operations employed by the model to generate its prediction [89, 52]. Jacovi and Goldberg [89] demonstrate that optimizing explanations purely for plausibility frequently produces deceptive explanations that mask internal reasoning failures and biases. In legal domains, legal scholarship emphasizes that faithfulness must take precedence, as an explanation that sounds plausible but misrepresents statutory text can lead to catastrophic compliance failures.

This trade-off is particularly evident in generative explainability approaches that deploy auto-regressive large language models to draft plain-language explanations of legal predictions. While generative summaries achieve high human plausibility scores, empirical legal informatics studies demonstrate that generative models frequently suffer from legal hallucinations, fabricating non-existent statutory citations or altering legal interpretations during text generation [113, 94]. Because generated natural language explanations are decoupled from the classifier’s internal mathematical representations, they remain inherently unfaithful [89]. Consequently, legal informatics literature increasingly favors extractive and attention-grounded visualizations that display the exact mathematical spans influencing classification, avoiding generative distortions.

2.11.4 Explainability in Multi-Stage Pipelines

In multi-stage NLP pipelines, literature on Explainable AI emphasizes situating interpretability mechanisms at the fine-grained inference stage rather than the coarse retrieval phase. In retrieve-then-entail systems, the initial retrieval phase functions as a high-throughput candidate selector across extensive document repositories. As Anand et al. [7] observe, computing dense feature attributions during coarse document filtering provides negligible analytical utility while incurring prohibitive computational latency. In contrast, the subsequent cross-encoder inference stage conducts fine-grained semantic entailment across a constrained candidate set, making it the mathematically appropriate locus for interpretability analysis.

At this second stage, legal NLP frameworks utilize token-level saliency and attention visualizations to bridge model internals with human inspection [208]. By mapping cross-attention weights and calibrated class probabilities back to corresponding spans in the premise and hypothesis texts, systems provide structural verification of detected contradictions without introducing generative paraphrasing risks. This dual mechanism of intrinsic attention mapping and post-hoc confidence calibration allows legislative researchers to trace classification outputs directly back to source statutory phrases.

2.12 Theoretical Framework

As established in legal informatics literature, the Retrieve-then-Entail architecture represents the leading theoretical framework for statutory conflict detection and legislative consistency checking. Originating in the international benchmarks of the Competition on Legal Information Extraction and Entailment (COLIEE 2025 and 2026) [66, 67], this framework integrates two complementary theoretical pillars: a scalable Information Retrieval (IR) architecture to isolate relevant statutory context from expansive corpora, followed by a fine-grained Natural Language Inference (NLI) architecture for rapid deontic conflict detection. Specifically, this theoretical framework synthesizes the dense statutory retrieval methodology of Team AIIR Lab [219, 105], the multi-stage ensemble retrieval principles of Team JNLP [66, 65], and the encoder-based textual entailment framework of Team KIS [92].

2.12.1 Statutory Retrieval (Team AIIR)

In statute law retrieval, dense bi-encoder architectures offer an effective theoretical foundation for bridging the semantic gap between newly drafted provisions and expansive statutory corpora without relying on computationally demanding generative text decoding or manual taxonomy engineering [219, 105]. Drawing from the empirical methodology developed by the AIIR Lab in recent COLIEE benchmarks, this retrieval pipeline operates across four theoretical dimensions: domain-adapted data augmentation, semantic vectorization, corpus pre-computation, and nearest-neighbor distance evaluation. Figure 2.1 illustrates this dense bi-encoder workflow.

Figure 2.1 details the structural dataflow of the dense statutory retrieval architecture adapted from the AIIR Lab methodology. The diagram illustrates four primary operational stages: synthetic legal query augmentation, offline statutory embedding pre-computation, dense vector indexing, and approximate nearest neighbor similarity matching. This visual is presented to demonstrate how modern legal search pipelines computationally decouple corpus-wide indexing from runtime query execution. As depicted, the incoming query provision is projected into a shared semantic coordinate space, allowing its vector representation to be compared directly against pre-computed statutory embeddings via cosine distance. This architectural separation ensures that high-volume statutory collections spanning tens of thousands of provisions can be searched in sub-second latency on consumer-grade hardware, providing the high-recall candidate shortlist necessary to feed the downstream entailment engine.

Targeted data augmentation and domain adaptation expose dense retrieval models to synthetically augmented legal variants. This training teaches the neural network to recognize identical statutory concepts regardless of whether they appear in archaic terminology or modern drafting syntax, substantially enhancing robustness against lexical divergence. Second, semantic vectorization utilizes a bi-encoder architecture (such as `all-mpnet-base-v2`) to project statutory provisions into a dense, continuous vector space. Unlike cross-encoders that must process query-document pairs jointly, bi-encoders encode queries and corpus documents independently, optimizing embedding coordinates such that logically relevant statutory pairs exhibit high geometric proximity under co-




Figure 2.1: Statute Law Retrieval Pipeline Using Dense Bi-Encoder Architecture

sine distance. Third, the pipeline relies on offline corpus pre-computation and vector indexing to satisfy operational efficiency constraints across tens of thousands of statutory enactments. Prior to live query execution, the fine-tuned bi-encoder vectorizes the entire statutory collection into static embeddings within an indexed vector database. Because corpus vectorization is completely decoupled from runtime query execution, dense retrieval scales across extensive archives without requiring real-time GPU re-encoding. Finally, live query vectorization and proximity scoring project incoming draft provisions into the shared embedding space, executing approximate nearest neighbor search to calculate continuous proximity scores in sub-second latency and deliver high candidate recall to prime downstream natural language inference.

2.12.2 Multi-Stage Retrieval (Team JNLP)

Complementing pure bi-encoder retrieval, literature in computational law emphasizes multi-stage retrieval pipelines to balance candidate recall against ranking precision. Exemplified by Team JNLP’s high-performing architecture in COLIEE benchmarks [66, 65], this multi-stage pipeline progressively filters statutory search spaces across three structured phases: Pre-Retrieval Candidate Pruning, Model Fine-Tuning, and Result Ensembling. Figure 2.2 illustrates this multi-stage workflow.




Figure 2.2: Multi-Stage Statutory Retrieval Pipeline With Candidate Pruning and Ensembling

Figure 2.2 outlines the progressive candidate filtering architecture developed by Team JNLP. The diagram delineates the modular progression from initial corpus-level candidate pruning, through parameter-efficient neural re-ranking, to multi-model score ensembling. This framework is included to provide theoretical justification for combining sparse lexical matching with dense semantic embeddings in statutory search. In the initial phase, heuristic distance thresholds and lexical filters rapidly eliminate irrelevant enactments, after which fine-tuned language models compute calibrated relevance scores across the remaining shortlist. The resulting predictions are synthesized through rank fusion, ensuring that the candidate set benefits from both keyword specificity and conceptual semantic coverage.

2.12.2.1 Pre-Retrieval

The primary objective of the pre-retrieval phase is to drastically reduce the expansive statutory search space while maintaining high recall across governing provisions. Team JNLP achieves this through a multi-stage candidate pruning procedure combining distance-based heuristic filtering with neural re-ranking [66, 65]. Dense embeddings for statutory queries and corpus provisions are extracted using high-capacity multilingual representations. Rather than relying solely on raw cosine similarity, the methodology calculates multi-dimensional distance metrics (such as L1 distance bins), training a gradient-boosted decision tree to predict relevance and counter severe sample imbalance. Following initial filtering, a specialized passage re-ranker evaluates candidate texts and eliminates non-relevant provisions, narrowing the candidate pool to top-ranked candidates while preserving recall.

2.12.2.2 Model Fine-Tuning

Once the candidate space is constrained, the framework applies targeted natural language models to perform fine-grained relevance scoring. Because fine-tuning large models across full parameter sets is computationally intensive, legal NLP frameworks employ Parameter-Efficient Fine-Tuning (PEFT) techniques, such as LoRA and QLoRA. These techniques enable foundational language models to adapt to specialized statutory semantics with minimal memory overhead. Furthermore, because statutory candidate sets exhibit severe class imbalance where only a small minority of retrieved provisions contain true preemption conflicts, sampling strategies and asymmetric loss functions are applied during training to preserve sensitivity to positive conflict signals.

2.12.2.3 Result Ensemble

To maximize retrieval robustness and mitigate individual model variance, multi-stage retrieval pipelines aggregate relevance scores across distinct retrieval architectures. Relevance predictions are aggregated through weighted score aggregation, combining calibrated probabilities from complementary lexical and neural models. The weighting coefficients assigned to individual scoring components are optimized via systematic hyperparameter search algorithms such as Optuna. This optimization ensures balanced contributions between lexical keyword precision and dense semantic recall, yielding a stable candidate shortlist for downstream reasoning.

2.12.3 Textual Entailment (Team KIS)

In the second phase of legal retrieve-then-entail architectures, the theoretical objective transitions from broad candidate discovery to fine-grained deontic reasoning. To ensure high computational throughput and deterministic classification, modern legal NLP frameworks deploy non-generative cross-encoder architectures, exemplified by the ModernBERT-based rationale extraction methodology developed by Team KIS [92]. Figure 2.3 shows this cross-encoder classification structure.

Figure 2.3 illustrates the sequence classification pipeline utilized for statutory textual entailment. The visual depicts the concatenation of the retrieved statutory premise and draft hypothesis into a unified input sequence demarcated by delimiter tokens, which is then processed through bidirectional self-attention layers to produce calibrated class probabilities. This model is included to illustrate how fine-grained deontic reasoning operates without the hallucination risks inherent in generative text decoders. By evaluating full cross-attention between premise and hypothesis tokens simultaneously, the cross-encoder identifies subtle modal contradictions that independent vector projections fail to capture. The calibrated classification head then maps the combined representation into deterministic Entailment, Neutral, or Contradiction outputs with formal statistical reliability.

2.12.4 Theoretical Synthesis

This synthesis integrates the foundational theoretical models reviewed throughout this chapter into a unified framework for ex-ante legislative conflict detection. While each individual methodology addresses an isolated phase of legal document processing, their conceptual integration establishes a principled retrieve-then-entail architecture. The theoretical foundation for statutory candidate discovery draws from the dense retrieval methodology established by the AIIR Lab [219, 105]. Their approach demonstrates that targeted domain augmentation mitigates OCR noise in archival legal text, while continuous semantic vectorization enables high-throughput candidate identification across large statutory collections without relying on rigid keyword indices.

To balance candidate recall against precision, the framework incorporates the multi-stage ranking principles demonstrated by Team JNLP [66, 65]. This methodology utilizes lexical filtering (BM25) to prune extraneous documents, followed by neural scoring and rank fusion to isolate authoritative candidate enactments on standard computing hardware. For granular deontic conflict evaluation, the framework adopts the Team KIS framework for textual entailment [92]. Their methodology formats premise-hypothesis pairs into a continuous sequence evaluated via lightweight cross-encoders, utilizing token-level separation boundaries to classify logical relationships without generative hallucination risks.

Synthesizing these foundational frameworks establishes an end-to-end theoretical foundation for automated preemption auditing. The retrieval stage narrows extensive statutory search spaces down to high-probability candidate provisions, while the textual entailment stage conducts deterministic normative conflict checks. The latest findings from the COLIEE 2026 proceedings [67, 133] reinforce this architectural consensus, validating that pipelined retrieval and textual entailment provide the most reliable foundation




Figure 2.3: Statutory Textual Entailment Pipeline Using Bidirectional Cross-Encoder

for automated legal consistency analysis.

Chapter 3

Methodology

This chapter explains the methodology used to build and evaluate the Coarse-to-Fine Semantic Conflict Detection System. It describes how Philippine legal texts are collected and cleaned, and how the two-stage Retrieve-then-Entail architecture is organized. The chapter also details the creation of the ground truth dataset and the selection of candidate models. Finally, it outlines the evaluation metrics and hardware constraints, which are chosen to match the computing capacity of an average Local Government Unit (LGU).

3.1 Conceptual Framework

This section maps how data flows through the proposed system. The architecture follows an Input-Process-Output (IPO) design organized into a two-stage “Coarse-to-Fine” retrieve-then-entail pipeline (Figure 3.1). By decoupling broad statutory candidate retrieval from compute-intensive textual entailment classification, the framework overcomes the high computational demands and memory limits that traditionally impede local legal NLP deployment. The workflow is divided into two operational phases: Phase 1 (Offline Foundation and Optimization) and Phase 2 (Live Online Retrieve-then-Entail Pipeline).

Figure 3.1 illustrates the complete end-to-end information flow of the conflict detection architecture, detailing the functional separation between offline foundation pre-computation (Phase 1) and the live pre-enactment review pipeline (Phase 2). The flowchart depicts an Input-Process-Output workflow organized into modular processing layers, highlighting how heterogeneous statutory archives are ingested, normalized, indexed, and logically verified under local edge compute constraints. In Phase 1, the system compiles the dual statutory knowledge base (combining 25,432 national statutes and 1,664 local ordinances), builds pre-computed BM25 and dense embedding indexes, and screens candidate neural architectures against an adviser-reviewed 350-pair ground truth benchmark. In Phase 2, an ex-ante ordinance draft undergoes structural boundary parsing and procedural boilerplate filtering before being routed through Stage 1 hybrid retrieval (fusing lexical BM25 and dense bi-encoder streams via Reciprocal Rank Fusion) and Stage 2 Cross-Encoder inference. This coarse-to-fine design restricts compute-intensive cross-attention to a focused candidate shortlist ($k = 50$), allowing deep semantic con-




Figure 3.1: Detailed Information Flow of the Coarse-to-Fine Semantic Conflict Detection Architecture.

flict detection to execute efficiently on standard municipal desktop workstations while producing transparent, explainable audit reports.

3.1.1 Phase 1: Offline Foundation and Optimization

This preliminary phase establishes the empirical foundation of the system by compiling the statutory corpus, building pre-computed retrieval indexes, and selecting candidate transformer logic engines through a structured benchmark screening process.

During dual-corpus ingestion and pre-computation, the system compiles a comprehensive dual-jurisdiction statutory archive combining the Philippine national statutory collection (Republic Acts, Executive Orders, Presidential Decrees, Acts, Batas Pambansa, and Commonwealth Acts) with Davao City municipal legislation (historical Landmark Ordinances and digital LISSP municipal records). The compiled corpus undergoes legal text normalization, optical character recognition (OCR) noise mitigation on historical scans, and structural section segmentation into granular, self-contained regulatory units (~ 200 – 300 tokens). Cleaned statutory provisions are then indexed offline: dense semantic representations are pre-computed using a lightweight transformer bi-encoder, while sparse inverted lexical indexes are constructed using BM25 ($k_1 = 1.5, b = 0.75$) alongside unsupervised statutory topic modeling.

Following corpus indexing, model benchmarking and selection evaluates candidate neural architectures against a curated Ground Truth Legal Benchmark stratified across three granular difficulty tiers: Tier 1 (Surface Lexical), Tier 2 (Structural Preemption), and Tier 3 (Latent Semantic). A multi-stage empirical screening and ablation process evaluates candidate bi-encoders by balancing inference latency against candidate Recall@ k , alongside candidate cross-encoders evaluated across precision, latency, and quantization axes. Based on these empirical evaluations, the highest-performing models are selected as the production engines: an optimal dense bi-encoder for Stage 1 candidate retrieval, and an optimal fine-tuned cross-encoder for Stage 2 contradiction reasoning.

3.1.2 Phase 2: Live System Workflow

The live system executes the automated conflict detection pipeline when a newly drafted municipal measure is submitted for pre-enactment review, operating through an Input-Process-Output structure.

The workflow begins with system input ingestion, where a legislative researcher uploads a newly proposed ex-ante ordinance draft. Because draft measures originate within local government offices as digital-born documents (.docx or digital PDF), they are free from optical scanning noise. The ingestion pipeline segments the draft into individual substantive provisions, filtering out non-regulatory procedural boilerplate such as the Title, Separability, Repealing, and Effectivity clauses. The isolated substantive provisions are then vectorized in real time via the selected dense bi-encoder alongside sparse lexical token extraction.

The computational process core then executes two complementary stages: hybrid candidate retrieval followed by natural language inference. In Stage 1 (Dual-Stream Hy-

brid Information Retrieval), the system narrows the comprehensive statutory search space down to a focused candidate pool. Stream A executes lexical sparse search via BM25 to match exact statutory citations, numerical penalty amounts, and administrative agency terms, while Stream B executes dense semantic search via the selected bi-encoder, measuring cosine similarity between the draft provision and statutory embeddings to capture abstract regulatory intent. The ranked outputs are combined using Reciprocal Rank Fusion (RRF) integrated with a Statutory Hierarchy Safeguard prioritizing higher-ranking national statutes over municipal provisions, extracting the Top- k candidate provisions without overburdening local LGU compute resources.

In Stage 2 (Natural Language Inference), each retrieved statutory candidate is paired with the draft ordinance provision into a single cross-encoder sequence:

Cross-Encoder Joint Sequence Input Representation:

$$\mathbf{x} = [\text{CLS}] \text{ Statutory Premise } [\text{SEP}] \text{ Ordinance Hypothesis } [\text{SEP}] \quad (3.1)$$

Where [CLS] marks the classification sequence start and [SEP] separates jurisdictional text boundaries.

The joint sequence formulation in Equation 3.1 formats each candidate premise-hypothesis pair into a contiguous token sequence to enable full bidirectional cross-attention. By concatenating the retrieved statutory section with the draft ordinance provision separated by delimiter tokens, the architecture allows every token in the statutory premise to attend directly to every token in the ordinance hypothesis across all self-attention layers. This joint representation is methodologically required because legal conflict detection frequently hinges on subtle textual interactions, such as conditional provisos, statutory penalty ceilings, or unhonored exceptions, which cannot be captured by independent bi-encoder vector dot-products. The resulting cross-attention representation is subsequently mapped through a linear classification head to produce calibrated softmax probabilities across Entailment, Neutral, and Contradiction.

The unified sequence is processed by the selected fine-tuned cross-encoder NLI engine, which performs full inter-token cross-attention to classify the directional logical relationship across Entailment, Neutral, or Contradiction, producing a calibrated contradiction probability $P(\text{Contradiction})$.

Finally, system output and explainable reporting presents an audit trail displaying the Top- k shortlisted candidate statutes for human oversight before evaluating contradiction probability against a calibrated decision threshold τ^* , optimized for recall priority via F_2 -score weighting. When the contradiction probability meets or exceeds the threshold ($P(\text{Contradiction}) \geq \tau^*$), the system flags a normative legal inconsistency and routes it through a two-tiered jurisdictional classification: Critical Vertical Preemption Alerts signaling collisions with superior national statutes under the *Magtajas* doctrine, and Horizontal Inconsistency Notices highlighting collisions or uncoordinated overlaps with prior Davao City municipal ordinances. To ensure decision-support transparency, the system renders intrinsic structural explainability featuring side-by-side comparative spans, highlighted conflicting phrases, and calibrated confidence percentages. Conversely, when contradiction probabilities remain below the calibrated threshold ($P(\text{Contradiction}) < \tau^*$),

the system issues a formal legislative compliance confirmation, accompanied by a statutory clearance audit confirming both national statutory harmony and local municipal code coherence.

3.2 Research Design

This study uses an applied research design to build a practical software tool for local governance. The researchers use an experimental systems development approach to solve the delays and difficulties of manual ordinance review through automated detection. This setup connects principles from legal informatics and natural language processing to address a real local need. By focusing on practical transformer-based models, the design connects theoretical computer science with everyday government office work. The experiment directly tests whether a retrieve-then-entail pipeline can operate reliably in a local legislative setting.

The experimental setup tests the retrieval and classification modules separately to measure search speed and accuracy independently. Building the system step-by-step allows the researchers to adjust the information retrieval and natural language inference stages using local legal texts. This engineering process ensures that the finished tool accounts for the writing styles found in both city ordinances and national statutes. In this way, applied research methods help turn modern NLP concepts into a usable tool for the Davao City Sangguniang Panlungsod. As a result, the findings establish a clear benchmark for automating pre-enactment legal checks on local legislation.

3.3 Corpus Acquisition and Exploratory Analysis

Building the legal database requires collecting text from different public repositories so that the system includes as many relevant laws as possible. These sources provide the national statutes and city ordinances needed to check draft legislation for consistency. To keep the language consistent and make vector processing simpler, the system gathers only the English-language versions of all legal texts. This choice matches Philippine statutory rules under Section 20, Chapter 5, Book I of Executive Order No. 292 (Administrative Code of 1987) [152], which explicitly mandates that in the interpretation of laws promulgated in official languages, the English text shall control unless otherwise specifically provided. National statutes and local ordinances are gathered using separate automated scripts designed for the format of each source.

3.3.1 Corpus Composition

The Lawphil Project (Arellano Law Foundation) serves as the primary digitized extraction source for national statutory text because its documents are provided in clean HTML, follow a consistent layout, and offer broad historical coverage [8]. Automated scraping scripts browse through chronological and alphabetical indexes to collect Republic Acts, Batas Pambansa, Commonwealth Acts, Acts, Executive Orders, and Presidential Decrees. Using a custom web extraction pipeline built by the proponents, statutory extraction was

conducted as of **March 2026**, collecting 25,432 unique legal records enacted between 1900 and 2026. The Official Gazette of the Republic of the Philippines [72] and the Chan Robles Virtual Law Library were used as authoritative secondary cross-references to verify enactment dates, chronological numbering, and historical amendment status. All collected national laws are converted into standardized JSON records with metadata, such as dates of approval, titles, and section numbers. As illustrated in Figure 3.2, the Lawphil repository presents statutory text in clean, unstyled HTML typography with a predictable hierarchical URL schema (e.g., `/statutes/repacts/ra2018/ra_11035_2018.html` for Republic Act No. 11035). This standardized structure separates introductory legislative headers from operative provisions and eliminates the need for dynamic JavaScript browser rendering, facilitating rapid, fault-tolerant web scraping.




Figure 3.2: Sample Statutory Enactment Webpage from the Lawphil Project Repository Illustrating Clean Document Hierarchy and Structured Legislative Metadata (Republic Act No. 11035).

Figure 3.2 displays a representative national statutory webpage from the Lawphil Project repository, featuring the complete legislative record of Republic Act No. 11035 (*Balik Scientist Act*). The document interface exhibits a clean, unstyled HTML layout that cleanly isolates the formal congress heading, approval date, official title, and individual numbered provisions into predictable DOM elements. This visual structure is included to illustrate why the national statutory repository facilitates automated web crawling without requiring dynamic JavaScript execution or complex layout segmentation. In the displayed layout, the operative provisions appear in sequential chronological order, flanked by transparent URL parameters (`/statutes/repacts/ra2018/ra_11035_2018.html`) that

enable automated validation against congressional records. By capturing these structured HTML pages directly, the ingestion preprocessor compiles over 25,000 national enactments into standardized text records while completely avoiding the visual noise and character degradation present in scanned archival paper.

Collecting national laws from Lawphil helps avoid the disorganization common to public legal websites. Although the Official Gazette and congressional websites also host statutory records, their web pages frequently change layout and often provide older laws only as image-only PDF scans. In contrast, Lawphil keeps a standard HTML structure across all historical eras, clearly separating introductory preambles from operative legal sections. Scraping from this repository ensured uniform text formatting across the national collection (over 25,000 enactments), preventing formatting differences from skewing vector embeddings during Stage 1 retrieval. Listing 3.1 illustrates the web scraping and provision extraction logic implemented to retrieve and structure the statutory archive, establishing a consistent text foundation for semantic analysis and conflict detection.

```

1 def extract_statutory_corpus(base_url, start_year=1900, end_year=2026):
2     """
3     Extracts, sanitizes, and structures national statutes from the online archive.
4     """
5     statute_records = []
6     for year in range(start_year, end_year + 1):
7         year_index = fetch_page(f"{base_url}/year/{year}")
8         statute_links = parse_index_links(year_index)
9         for link in statute_links:
10             # Fetch raw web page and parse HTML DOM tree
11             html_content = fetch_page_with_retry(link, delay_seconds=1.0)
12             soup = BeautifulSoup(html_content, "html.parser")
13
14             # Extract statutory metadata and clean document body
15             metadata = extract_header_metadata(soup)
16             body_text = clean_document_body(soup)
17             sections = segment_into_sections(body_text)
18
19             statute_records.append({
20                 "statute_id": metadata.id,
21                 "title": metadata.title,
22                 "instrument": metadata.type,
23                 "date_enacted": metadata.date,
24                 "sections": sections
25             })
26     return statute_records
27

```

Listing 3.1: Code Snippet of the Statutory Corpus Acquisition and Provision Extraction Pipeline

Listing 3.1 details the programmatic execution logic of the statutory web extraction script, outlining the sequential functions for annual index traversal, DOM parsing, metadata extraction, and provision boundary segmentation. The script is structured to crawl legislative archives spanning from 1900 through 2026, parsing HTML document trees via BeautifulSoup while implementing rate-limiting delays to preserve source server availability. During execution, the pipeline extracts key statutory metadata—including enactment identifiers, approval dates, and legislative titles—before passing the document body to a regex-based section segmenter that isolates individual statutory provisions into distinct JSON fields. Methodologically, this automated extraction guarantees uniform

data serialization across all 25,432 national enactments, eliminating web scraping artifacts and ensuring that downstream embedding models evaluate pure statutory syntax.

To capture the complete statutory search space, the national collection is paired with Davao City municipal legislation gathered across two complementary acquisition streams. The first channel consists of the official Davao City Sangguniang Panlungsod Public Web Portal, which provides selective historical landmark ordinances and municipal uploads spanning approximately fifteen years, capturing prominent regulatory codes in public safety, zoning, and health [168]. The second and primary channel comprises the municipal Legislative Information Support System Program (LISSP), acquired directly through formal administrative coordination with the Sangguniang Panlungsod Secretariat and Information Technology division [30]. Through this institutional coordination, municipal authorities granted the proponents direct access to digitized enactments via a shared repository, encompassing enactments passed primarily by the 19th and 20th City Councils between 2021 and 2025. Together, these two channels yielded an archive of exactly 1,664 municipal enactments spanning from 1950 to 2025. These 1,664 documents represent the complete accessible digital corpus gathered within the proponents’ institutional capabilities. In practice, the true cumulative historical total of ordinances enacted since Davao City was chartered under Commonwealth Act No. 51 in 1936 remains unrecorded on the internet and unknown to municipal authorities. In institutional consultations, Sangguniang Panlungsod personnel confirmed that an official cumulative ledger does not exist because local legislation is numbered sequentially per calendar year and council term rather than in an unbroken cumulative sequence. Furthermore, the vast majority of historical enactments, particularly those prior to the 2021 LISSP digital initiative, exist exclusively as physical paper documents preserved in archive rooms and filing cabinets at the Sangguniang Panlungsod building on San Pedro Street, requiring manual page-by-page physical inventorying to tally. When combined with the 25,432 national statutes, the total statutory search space comprises **27,096 legal enactments** ($N_{\text{sections}} = 176,421$).

3.3.2 Corpus Characteristics and Inherent Biases

A close look at the collected national laws reveals a large imbalance across legal topics. Over 40% of the national statutes consist of routine administrative actions, such as converting local schools, renaming roads, adjusting municipal boundaries, or granting broadcast franchises. In contrast, major regulatory areas that interact with city powers, such as public health, land zoning, and local taxation, make up less than 15% of the total national collection. Searching through this raw distribution without filtering would cause search algorithms to return irrelevant administrative laws instead of actual regulatory rules. To keep search results useful, the system must separate these routine administrative laws from substantive regulatory provisions.

The municipal archive also shows an uneven time distribution and gaps in older records. For ordinances passed before 2021, city record-keeping mainly focused on digitizing major “Landmark Ordinances,” leaving ordinary regulations from earlier decades in physical paper storage [168]. As a result, the pre-2021 local records represent a selective collection of prominent laws rather than a complete legislative record. While post-2021 records from the digital LISSP database are much more complete, some numerical gaps re-

main due to missing committee drafts and unindexed attachments. These temporal biases mean that the available municipal archive contains substantially more recent ordinances than older enactments.

A fundamental technical distinction must be maintained regarding the document media of the legislative corpus. All 1,664 archival ordinance documents received from the Sangguniang Panlungsod are historical scanned physical paper PDFs, rather than digital-born documents exported from word processing software. Older ordinances scanned at low resolution suffer from faded print, ink bleed-through, staple shadows, and rotational skew, as illustrated in Figure 3.3. The scanned copies also contain non-text items such as government stamps, official signatures, and handwritten marginal notes that confuse standard sentence splitting tools. Clean, digital-born text (documents exported directly from modern word processors or Google Docs without scanning noise) appears exclusively during Phase 2 live system testing when legislative staff submit newly drafted ex-ante measures for pre-enactment conflict review. At the language level, local ordinances sometimes include regional terms and administrative Taglish in their introductory sections [53]. These informal drafting habits differ noticeably from the formal, standardized English found in national laws.




Figure 3.3: Sample Archival Document from the Davao City Sangguniang Panlungsod Corpus Illustrating Physical Scanning Artifacts and Character Noise.

Figure 3.3 displays a representative physical document scan from the Davao City Sangguniang Panlungsod legislative archive, illustrating the optical degradation common to historical local records. Panel (a) shows the prefatory municipal session header, city seal, and councilor attendance roster, which introduce non-regulatory proper nouns if not systematically stripped. Panel (b) highlights an enacted substantive clause exhibiting authentic scanner artifacts, including rotational skew, marginal photocopy shadows, and character substitutions (such as misreading "2021" as "0667-2L" and "TITLE" as "TITTE"). This visual evidence is presented to justify why local ordinance digitization requires specialized pre-filtering and vision-language processing rather than naive OCR ingestion. Methodologically, the presence of these scanning artifacts confirms that all 1,664 historical ordinances gathered for the corpus require physical perimeter cropping,

noise whitelisting, and quality auditing before their provisions can be reliably indexed for semantic search.

In real-world legislation, direct contradictions between laws are very rare. Lawmakers generally try to write rules that work alongside existing higher laws rather than directly opposing them [61]. As a result, actual legal conflicts make up less than 1% of random pairs formed between municipal ordinances and national statutes. Taking random pairs of text from the unedited collections would lead to an overwhelming majority of unrelated pairs. A machine learning model trained on such unbalanced data could score 99% nominal accuracy simply by guessing “Neutral” every time, while remaining unable to detect real legal conflicts.

3.3.3 National Statutory Corpus Analysis

Across the five main types of legislation with sequential numbers, the collected corpus includes 19,683 statutes out of an estimated 20,246 historical laws, achieving a coverage rate of 97.22%. When combined with the 5,749 Executive Orders from fifteen presidential administrations, the compiled national archive holds 25,432 out of about 26,451 total historical legal documents, reaching a national recall rate of 96.15%. All official legislative totals reflect the chronological numbering sequence established by Congress and presidential records as of **March 2026**, where sequential Republic Acts reached Republic Act No. 12,313. Table 3.1 provides a breakdown of the national corpus, comparing the collected documents with official legislative totals for each statutory type.

Table 3.1 presents the structural breakdown of the Philippine national statutory archive across six legislative instruments, detailing the collected enactment counts, official legislative totals, achieved coverage rates, percentage shares, and historical governing authorities. This comparative census is included to verify the empirical completeness and historical coverage of the primary national statutory knowledge base, proving that the harvested database represents an authentic, near-census archive of Philippine legislation rather than a selective convenience sample. As shown in the table, congressional Republic Acts form the single largest category (47.04%, $N = 11,963$), followed by presidential Executive Orders (22.61%, $N = 5,749$) and American-era Acts (16.74%, $N = 4,257$). Across all numbered congressional enactments, the collection captures 19,683 out of 20,246 historical laws (97.22% coverage), expanding to 25,432 documents (96.15% coverage) when including presidential administrative issuances. Methodologically, this comprehensive coverage guarantees that the statutory search space encompasses nearly the entire operative statutory history of the Philippines since 1900, minimizing the risk of false negatives during Stage 1 retrieval caused by missing national legislation.

A key methodological decision in compiling this database is the explicit inclusion of **Executive Orders (EOs)** alongside congressional legislation. Under Article VII, Section 17 of the 1987 Philippine Constitution and Book III of Executive Order No. 292 [152], Executive Orders are presidential acts providing rules of a general character in execution of statutory mandates. Unlike Supreme Court decisions that settle private adversarial controversies under *stare decisis*, Executive Orders operate as *ex-ante*, forward-looking statutory instruments possessing the full force of law. Modern Executive Orders estab-

Table 3.1: Comparative Statutory Coverage and Composition of the Philippine National Legal Archive ($N = 25,432$)

	Statutory ID Instrument	Corpus	Official Total	Coverage Rate	Corpus Share	Historical Era and Governing Authority
01	Republic Acts (RA)	11,963	12,313	97.16%	47.04%	Post-war and Fifth Republic legislation (1946–present, up to RA 12313 as of March 2026)
02	Executive Orders (EO)	5,749	6,205*	92.65%	22.61%	Presidential administrative issuances (1901–2026, across 15 administrations)
03	Acts	4,257	4,275	99.58%	16.74%	American Insular Government enactments (1900–1935)
04	Presidential Decrees (PD)	1,844	2,036	90.57%	7.25%	Martial law legislative decrees enacted by executive fiat (1972–1986)
05	Batas Pambansa (BP)	886	889	99.66%	3.48%	Unicameral parliamentary legislation under the 1973 Constitution
06	Commonwealth Acts (CA)	733	733	100.00%	2.88%	Transitional governance enactments under the 1935 Constitution
Sequential Subtotal		19,683	20,246	97.22%	77.39%	All numbered legislative enactments of Congress
Total Archive		25,432	26,451	96.15%	100.00%	Comprehensive corpus extracted from the Lawphil Project

*Historical total for Executive Orders represents the aggregate sum of maximum numbered issuances across all fifteen presidential administrations (1935–2026) and colonial governorships (1901–1935) as recorded in the Official Gazette and Lawphil archives as of March 2026.

lish binding nationwide regulatory mandates, such as Executive Order No. 26 (s. 2017) mandating nationwide public smoking bans, that legally preempt inconsistent local ordinances. Excluding Executive Orders would create an artificial regulatory gap, allowing local ordinances to contradict operative presidential mandates.

The coverage rates confirm that the national legal archive is nearly complete across all historical governance eras. Commonwealth Acts show full 100.00% coverage across all 733 enactments, while American-era Acts (99.58%) and Batas Pambansa (99.66%) miss only 18 and 3 enactments, respectively. For Republic Acts, the 11,963 collected records cover 97.16% of the 12,313 official total, with missing entries mainly being local bills or recent laws still being indexed. For Presidential Decrees, the 1,844 gathered records represent 90.57% of the 2,036 numbered decrees from the Martial Law period, where uncollected entries were unpublished directives until the Supreme Court required full publication in *Tañada v. Tuvera* [184]. For Executive Orders, the 5,749 collected records represent 92.65% of the 6,205 numbered orders issued across fifteen presidential terms, offering an extensive record of administrative policy.

This institutional breakdown also explains why executive and administrative orders account for nearly half (48.01%) of the subsequent topic groups. Executive Orders

(22.61%) and Presidential Decrees (7.25%) together make up 7,593 documents, representing nearly one-third of the entire national archive. When combined with colonial-era civil administration Acts and agency reorganizations, administrative directives form the largest single category of Philippine legal text. In language processing, this large volume of administrative text poses a challenge during search, as queries can easily match routine bureaucratic language rather than municipal regulations. Figure 3.4 illustrates the share of each legal instrument in the archive, highlighting the large volume of executive orders alongside congressional acts. This computational reality demonstrates why automated topic grouping is necessary to keep local ordinance checks focused on relevant laws.




Figure 3.4: Composition of the Philippine National Statutory Corpus by Legislative Instrument ($N = 25,432$).

Figure 3.4 illustrates the proportional composition of the 25,432 national statutes across the six legislative formats using a proportional area chart. The visual displays the dominant aggregate footprint of legislative Republic Acts alongside executive and colonial instruments, highlighting the structural diversity of the national legal corpus. This figure is included to emphasize the computational necessity of handling diverse legislative genres within a single retrieval framework, as congressional acts establish broad rights while executive orders specify operational regulations. As visualized, over 30% of the national statutory archive consists of executive issuances and martial law decrees rather than traditional bicameral legislation. Methodologically, this finding demonstrates why coarse retrieval cannot treat all national texts identically, motivating the hierarchical safeguards and domain filters implemented in Stage 1.

To trace how the six statutory instruments evolved across Philippine jurisprudence, the proponents categorized statutory enactments across six chronological eras. Each historical era established distinct legislative authorities, ranging from colonial commissions to modern constitutional congresses. Figure 3.5 charts the chronological succession and aggregate volume of statutory instruments from 1900 through 2026. The distribution reflects the shifting primacy of legislative acts, presidential decrees, and administrative executive orders across political transitions. Examining these enactment waves illustrates how historical governance structures directly shaped the composition of the national legal corpus.




Figure 3.5: Historical Succession and Volume of Statutory Instruments across Six Chronological Eras (1900–2026, $N = 25,432$).

Figure 3.5 depicts the longitudinal succession and annual volume of the six statutory instruments across 126 years of Philippine governance (1900–2026). The horizontal axis traces chronological enactment years across six major political eras, while the vertical axis records annual enactment counts color-coded by legislative instrument. This temporal visualization is provided to expose the historical shifts in legal authority and drafting conventions that characterize the national corpus. As observed in the plot, legislative volume peaked during periods of institutional restructuring, transitioning from Insular Acts in the early 20th century to Presidential Decrees during Martial Law, and stabilizing into dual streams of Republic Acts and Executive Orders under the 1987 Constitution. Consequently, this temporal variation requires that the retrieval pipeline handle shifting archaic syntax and evolving statutory vocabularies without degrading search accuracy.

The temporal breakdown reveals that Philippine lawmaking expanded during periods of national administrative reorganization. Early colonial governance relied entirely on 4,199 foundational Acts during the Insular period (1900–1934). Centralized presidential lawmaking predominated during the Martial Law period through 1,844 Presidential Decrees and 886 Batas Pambansa parliamentary enactments (1972–1985). In the modern era governed by Republic Act No. 7160 (1991–2026), statutory production stabilized into a dual stream of 4,900 congressional Republic Acts and 2,421 executive policy orders. Because vocabulary standards shifted markedly across these legislative transitions, search models require chronological awareness to avoid misinterpreting historical statutory phrasing.

During preliminary exploratory research, the proponents conducted initial machine learning experiments on the national statutory archive and identified severe dataset imbalances. In the national collection of 25,432 statutes, routine executive issuances and administrative reorganizations comprise nearly half of all historical enactments, while

specialized regulatory domains remain relatively small. Searching for a local regulatory ordinance across this raw, unfiltered collection causes severe search drift because dense retrieval models readily match repetitive administrative phrasing across unrelated legal areas. To resolve these dataset quirks, the proponents formalized a two-stage topic modeling architecture that organizes the statutory archive into cohesive, well-separated subject areas. This structured categorization prevents routine administrative directives from crowding the candidate shortlist during coarse retrieval, establishing a balanced foundation for subsequent conflict checks.

The topic classification pipeline groups documents automatically while executing efficiently on the Primary Desktop Host. Statutory texts are first stripped of HTML tags and filtered against a curated list of 49 legal stopwords, such as “hereby” and “enacted,” which appear ubiquitously without distinguishing legal subjects. Cleaned texts are converted into term frequency-inverse document frequency (TF-IDF) feature vectors across a 25,000-term vocabulary, rewarding distinctive statutory terminology over frequent administrative vocabulary [122]. To maintain computational efficiency while preserving conceptual cohesion, the pipeline applies Truncated Singular Value Decomposition (SVD) to project the sparse feature vectors into 64 latent semantic dimensions [42]. Following the latent semantic modeling principles of Landauer et al. [103], this 64-dimensional subspace captures core semantic associations while eliminating lexical noise, followed by standard L_2 vector normalization.

Partitioning statutory texts directly into a small set of broad categories risks conflating heterogeneous legislative subjects. To prevent premature cluster mixing, the normalized 64-dimensional vectors are initially partitioned into $K = 28$ micro-clusters using MiniBatch K -Means clustering [174]. Following the over-clustering strategy recommended by Grootendorst [74] and Aggarwal [2], this over-complete partition isolates narrow recurring enactments, such as franchise charters or school conversions, into tight homogeneous groups. The pipeline then adapts class-based TF-IDF (c -TF-IDF) from Grootendorst [74] to identify salient keywords by treating all text within a micro-cluster as a unified document. The topic modeling and domain clustering formulations are computed as follows:

Class-Based Term Frequency-Inverse Document Frequency (c -TF-IDF):

$$c\text{-TF-IDF}_{i,c} = \frac{|t_{i,c}|}{|c|} \times \log \left(1 + \frac{A}{\sum_k |t_{i,k}|} \right) \quad (3.2)$$

Where $|t_{i,c}|$ denotes the frequency of term i in cluster c , $|c|$ is the total word count of cluster c , and A indicates the average word count per cluster across the archive. This cluster-level weighting elevates terminology that specifically characterizes a legislative group while penalizing ubiquitous statutory phrasing.

Macro Legal Domain Centroid and Softmax Probability:

$$\mathbf{C}_m = \frac{\sum_{i \in \mathcal{D}_m} \mathbf{z}_i}{\|\sum_{i \in \mathcal{D}_m} \mathbf{z}_i\|_2}, \quad P(m \mid \mathbf{z}_i) = \frac{\exp(\tau \cdot \cos(\mathbf{z}_i, \mathbf{C}_m))}{\sum_j \exp(\tau \cdot \cos(\mathbf{z}_i, \mathbf{C}_j))} \quad (3.3)$$

Where $\mathbf{z}_i$ represents the normalized 64-dimensional SVD vector of document i , $\mathbf{C}_m$ denotes the normalized centroid of macro domain m , and $\tau = 6.0$ functions as the inverse

temperature scaling parameter.

The mathematical formulations in Equations 3.2 and 3.3 establish the statistical foundation for isolating distinctive legal vocabularies and computing stable topic assignments across the national statutory corpus. Equation 3.2 calculates class-based term frequency-inverse document frequency, which treats each micro-cluster as an aggregated document to measure term frequency against its inverse occurrence across all statutory clusters. In Equation 3.3, normalized document vectors are aggregated into unit-norm cluster centroids $\mathbf{C}_m$, while a temperature-scaled softmax function maps cosine proximity into well-calibrated class probabilities. Scaling the cosine similarity with inverse temperature $\tau = 6.0$ prevents probability dispersion, ensuring that ambiguous enactments are assigned decisive dominant domain scores. Methodologically, this formulation enables the topic modeling pipeline to cluster over 25,000 statutes into eight cohesive regulatory domains while operating entirely within local desktop CPU memory limits.

Adapting temperature scaling principles from Guo et al. [76], the parameter $\tau = 6.0$ prevents uncalibrated softmax dispersion across cosine distances, ensuring that dominant domain probabilities remain sharply separated. Executing this consolidation across all 25,432 national statutes organizes the historical collection into eight Consolidated Macro Legal Domains, as detailed in Table 3.2.

Table 3.2: Consolidated Macro Legal Domains of the Philippine National Statutory Corpus ($N = 25,432$)

ID	Consolidated Macro Domain	Statutes	Share	Top Salient <i>c-TF-IDF</i> Keywords
00	Executive Issuances & Policy Reorganization	12,210	48.01%	executive, commission, committee, policy, code
01	Education & Academic Institutions	3,492	13.73%	high school, elementary school, campus, college
02	Local Government & Territorial Boundaries	2,134	8.39%	barangay, municipality, boundary, sitio, territorial
03	Public Utilities & Telecom Franchises	2,041	8.03%	franchise, broadcast, telecommunications, radio
04	Public Health, Hospitals & Medical Services	1,890	7.43%	hospital, bed capacity, medical center, infirmary
05	Statutory Codes & General Legal Amendments	1,873	7.36%	administrative code, revised, amend, article, penalty
06	Public Finance & General Appropriations	1,335	5.25%	appropriation, treasury, funds, expenditures
07	Taxation, Tariffs & Revenue Administration	457	1.80%	tariff, customs code, internal revenue, excise
Total Consolidated National Corpus		25,432	100.00%	All Enacted Statutes (1900–2026)

Table 3.2 details the eight Consolidated Macro Legal Domains discovered across the 25,432 national statutes, presenting for each domain its identifier, descriptive label, total statute count, percentage share, and top salient *c-TF-IDF* keywords. This taxonomy is included to define the categorical partition used by downstream retrieval safeguards to prevent high-frequency administrative enactments from overshadowing municipal regulatory

subjects. As shown in the table, Domain 00 (Executive Issuances & Policy Reorganization) dominates the archive with 12,210 statutes (48.01%), followed by Education (13.73%) and Local Government (8.39%), while vital regulatory areas such as Public Health (7.43%) and Taxation (1.80%) occupy compact, highly specific subsets. Methodologically, this table demonstrates that the statutory search space is heavily skewed toward administrative reorganizations, confirming that coarse retrieval requires domain-specific contextual filters and Reciprocal Rank Fusion to ensure that small regulatory domains remain searchable.

The topic distribution reflects the historical patterns of Philippine lawmaking. Executive issuances, administrative orders, and presidential decrees make up nearly half of the entire statutory collection, reflecting centralized authority during past political transitions. Common municipal subjects, including school infrastructure (13.73%) and public utility franchises (8.03%), form large distinct clusters with repetitive drafting patterns. In contrast, key regulatory topics like taxation (1.80%) and public health (7.43%) occupy smaller, highly focused portions of the database. As illustrated in Figure 3.6, routine executive issuances, administrative orders, and presidential decrees dominate the aggregate volume, while municipal regulatory categories remain relatively compact. Separating these regulatory clusters keeps routine administrative orders from crowding the candidate list during Stage 1 retrieval.




Figure 3.6: Distribution of the 25,432 Philippine National Statutes across the Consolidated Macro Legal Domains.

Figure 3.6 illustrates the distribution of national statutes across the eight consolidated macro domains using a sorted bar chart. The chart highlights the severe volume asymmetry between general administrative policy and specialized regulatory codes across the 25,432-document corpus. This visualization is included to demonstrate the severe class imbalance present in raw statutory archives, illustrating why naive unconstrained vector retrieval suffers from administrative drift. As depicted in the chart, administrative reorganizations account for nearly half of the entire archive, outnumbering focused regulatory categories like taxation by more than 26 to 1. Methodologically, this visual confirms

that candidate retrieval must incorporate domain-specific contextual filters and Reciprocal Rank Fusion to ensure that smaller, highly technical regulatory domains remain accessible during ex-ante ordinance verification.

To examine the geometric separation of these legal topics, the proponents projected the 64-dimensional SVD document vectors into a two-dimensional latent semantic landscape. Dimensionality reduction maps document similarities onto coordinate axes, allowing direct visual inspection of cluster cohesion. Figure 3.7 illustrates this topic landscape across a representative stratified sample ($N = 3,500$), highlighting the core clusters and their spatial distributions. Specialized regulatory domains project outward as distinct, isolated rays along the periphery of the vector space. Enactments establishing academic institutions and telecommunications franchises form tight, cohesive rays with minimal dispersion.




Figure 3.7: Two-Dimensional Latent Semantic Topic Space of the Philippine National Statutory Corpus (SVD Projection, $N = 3,500$ Representative Sample).

Figure 3.7 visualizes the two-dimensional latent semantic landscape generated by applying Truncated Singular Value Decomposition to a stratified sample of 3,500 national statutes. The scatter plot maps document vectors along two principal semantic axes, color-coded by the eight macro legal domains. This geometric projection is included to visually verify the semantic cohesion and spatial separability of the discovered legislative clusters. As observed in the plot, specialized enactments such as academic charters and telecommunications franchises project outward as dense, isolated radial rays, whereas executive and administrative orders concentrate heavily near the coordinate origin. This spatial organization proves that latent semantic modeling successfully isolates distinct regulatory vocabularies into cohesive clusters, providing a sound mathematical foundation for subsequent domain filtering.

In contrast, administrative issuances and executive orders cluster densely near the origin, comprising almost half of the entire archive. This dense central mass explains why unconstrained semantic searches frequently drift toward routine bureaucratic orders unless restrained by hierarchical safeguards. The coordinate axes reflect the two primary orthogonal directions of vocabulary variation discovered through singular value decomposition.

The horizontal axis measures operative statutory specificity against general administrative phrasing, where higher positive values indicate detailed technical mandates rather than procedural orders. Conversely, the vertical axis captures functional divergence between social infrastructure such as public schools and commercial infrastructure such as telecommunications franchises.

To confirm that the mathematical clusters reflect substantive legal categories rather than shared administrative boilerplate, the proponents examined the most salient vocabulary terms discovered by the model. Words that appear frequently across all documents fail to differentiate legal topics, requiring class-based weightings to isolate subject matter. Figure 3.8 displays the top class-based TF-IDF (c -TF-IDF) keywords across all eight consolidated macro domains. The vocabulary profiles confirm that each macro domain is defined by operative legal terminology unique to its statutory field. Specialized clusters such as public health are dominated by concrete institutional terms, including “hospital” and “bed capacity,” while territorial enactments display spatial markers such as “sitio” and “boundary.”




Figure 3.8: Top Salient c -TF-IDF Keyword Profiles across the Eight Consolidated Macro Legal Domains ($N = 25,432$).

Figure 3.8 displays the horizontal bar charts of the top class-based TF-IDF (c -TF-IDF) keywords across each of the eight consolidated macro domains. Each subplot presents the top terms ranked by their class-specific importance weights, illustrating the distinctive lexical tokens that define each statutory category. This figure is provided to confirm that the mathematical clustering captured genuine legal substance rather than generic statutory boilerplate. As shown across the subplots, Domain 01 is characterized by

academic vocabulary (“elementary school”, “college”), Domain 03 by telecommunications markers (“broadcast”, “radio”), and Domain 04 by clinical infrastructure (“hospital”, “bed capacity”). Substantively, this vocabulary alignment verifies that the unsupervised pipeline isolates authentic statutory topics, ensuring that domain classification routes draft ordinances to genuinely related national statutes.

To examine how individual lawmaking authorities contribute to these eight subject areas, the proponents mapped the six statutory instruments across the consolidated macro domains. In Philippine legislative history, different legal instruments serve distinct constitutional functions, ranging from broad congressional acts to presidential executive directives. Figure 3.9 presents this cross-allocation matrix, detailing both absolute enactment counts and percentage distributions for each statutory source. This cross-tabulation reveals an explicit division of legal responsibility across Philippine statutory instruments. Analyzing this allocation confirms that regulatory subjects concentrate within specific legislative formats rather than distributing evenly across the archive.




Figure 3.9: Cross-Allocation Matrix of the Six National Statutory Instruments across the Eight Consolidated Macro Legal Domains ($N = 25,432$).

Figure 3.9 presents the cross-allocation heatmap and count matrix mapping the six national statutory instruments across the eight consolidated macro domains. The grid displays both absolute statute counts and relative row percentages, revealing how different legal instruments distribute across legislative subjects. This visual matrix is included to empirically demonstrate the functional specialization of Philippine lawmaking authorities. The matrix reveals that congressional Republic Acts account for 99.9% of educational charters and 86.9% of utility franchises, whereas executive orders comprise 82.9% of general administrative policies and 81.4% of tariff adjustments. Methodologically, this cross-tabulation justifies incorporating statutory source indicators into Stage 1 retrieval safeguards, allowing the system to prioritize congressional enactments when evaluating municipal regulatory drafts.

The cross-allocation matrix reveals that congressional enactments and executive orders govern entirely separate legal spheres. Republic Acts account for nearly all spe-

cialized regulatory measures, containing 99.9% of academic enactments and 86.9% of telecommunications franchises. In contrast, Executive Orders concentrate heavily within administrative policy (82.9%) and presidential tariff adjustments (81.4%). Because municipal ordinances regulate local community behavior rather than internal executive operations, matching draft local clauses against executive orders produces high rates of irrelevant candidate retrievals. These empirical disparities validate using statutory source indicators and domain filters during Stage 1 retrieval to keep candidate shortlists focused on genuine regulatory mandates.

To confirm that the topic pipeline discovered authentic legal boundaries rather than grouping surface vocabulary, the proponents evaluated the model using a stratified 80/20 train/test split. The 80/20 proportion follows the standard Pareto partition for large-scale legal collections exceeding 20,000 documents [2], withholding 20% of the corpus (5,087 statutes) to test generalizability on unseen enactments while training the model on the remaining 80% (20,345 statutes). In statistical testing, an out-of-sample holdout of 5,087 legal texts provides high statistical power with a margin of error below $\pm 1.4\%$, making repetitive cross-validation computationally redundant as vocabulary frequencies and SVD components stabilize asymptotically at this volume. Stratified sampling preserves the natural historical proportions of the eight macro domains, ensuring that smaller regulatory categories like taxation (1.80%, $N = 457$) maintain adequate representation in the test set ($N_{\text{test}} \approx 91$) without class starvation.

To ensure objective scientific evaluation without arbitrary post-hoc targets, all benchmark thresholds in Table 3.3 are strictly grounded in established literature standards. Under the proportional chance criterion formulated by Morrison [127], the default random baseline in an eight-class uniform setting is $1/8 = 12.50\%$. The supervised accuracy target ($\geq 85.0\%$) and Macro F1 target ($\geq 80.0\%$) represent canonical text mining quality ceilings established by Aggarwal [2] for high-dimensional document categorization. For cluster geometry, Rousseeuw [159] established that in sparse, high-dimensional textual spaces where the curse of dimensionality compresses distance distributions, silhouette scores between 0.10 and 0.20 indicate genuine structural partition rather than noise. Similarly, Davies and Bouldin [41] established that Davies-Bouldin index values between 2.0 and 3.0 demonstrate compact intra-cluster scatter relative to cluster separation, outperforming unclustered baselines (> 4.0). Finally, the confidence thresholds (> 0.4000 mean confidence, $> 50.0\%$ high-confidence predictions) are derived from the temperature-scaled calibration principles of Guo et al. [76]. Table 3.3 summarizes the out-of-sample performance scores, baseline comparisons, literature targets, and empirical verdicts.

Table 3.3 details the empirical holdout validation metrics, default baseline comparisons, literature-grounded target ranges, and resulting empirical verdicts for the national statutory topic modeler evaluated on 5,087 unseen enactments ($N_{\text{train}} = 20,345$, $N_{\text{test}} = 5,087$). The table summarizes performance across six holistic clustering and classification metrics alongside four unambiguous statutory title probes ($N = 1,259$). This validation table is included to demonstrate that the unsupervised macro domains generalize stably to unseen statutory enactments without overfitting to lexical training noise. The observed results show that out-of-sample supervised accuracy reached 97.37% and Macro F1 reached 96.26%, vastly exceeding both the 12.50% Morrison chance baseline ($1/8$) and

Table 3.3: Empirical Holdout Validation Metrics, Baseline Comparisons, and Literature Targets for the Topic Modeler ($N_{\text{train}} = 20,345$, $N_{\text{test}} = 5,087$)

Validation Metric / Benchmark Test	Observed	Default Baseline	Target Range	Empirical Verdict
Out-of-Sample Supervised Accuracy	97.37%	12.50%	$\geq 85.0\%$	Strong Generalization [2]
Out-of-Sample Macro F1-Score	96.26%	12.50%	$\geq 80.0\%$	Balanced Discrimination [2]
Mean Test Cosine Confidence	0.5508	0.1250	> 0.4000	Strong Dominant Probability [76]
High-Confidence Predictions (≥ 0.50)	56.97%	$< 20.0\%$	$> 50.0\%$	Clear Majority Cluster [76]
Holdout Cosine Silhouette Score	0.1602	≤ 0.0000	0.1000–0.2000	Solid Text Separation [159]
Holdout Davies-Bouldin Index	2.4496	> 4.0000	2.0000–3.0000	Compact Cluster Scatter [41]
Explicit Statutory Title Probes (<i>Unambiguous Indicator Subsets</i> , $N = 1,259$)				
Telecom Franchise Pattern	97.8%	12.50%	$> 90.0\%$	High Precision ($N = 360$)
Hospital Bed Capacity Pattern	96.6%	12.50%	$> 90.0\%$	High Precision ($N = 203$)
School Infrastructure Pattern	92.2%	12.50%	$> 90.0\%$	High Precision ($N = 645$)
Judiciary & Court Pattern	72.5%	12.50%	$> 70.0\%$	Acceptable ($N = 51$, embedded laws)

the 85.0% literature target from Aggarwal [2]. Furthermore, the holdout silhouette score of 0.1602 and Davies-Bouldin index of 2.4496 satisfy the structural cluster quality criteria of Rousseeuw [159] and Davies and Bouldin [41]. Substantively, these empirical verdicts confirm that the topic model establishes mathematically reliable legal boundaries across unseen legislation, ensuring that domain-guided retrieval safeguards operate dependably during live ex-ante screening.

The supervised benchmark results confirm that the discovered macro domains generalize reliably to unseen legislation. In an eight-domain classification setting, the uniform proportional chance criterion established by Morrison [127] yields a default random baseline accuracy of only 12.50% ($1/8$). In contrast, the validation classifier reached an out-of-sample accuracy of 97.37% and a Macro F1-Score of 96.26%, surpassing the 85.0% target threshold typically expected for reliable document categorization [2]. To provide an external sanity check beyond learned cluster labels, the evaluation tested four explicit statutory title probes ($N = 1,259$ total) where legislative intent is unmistakable from the statutory catchline. The model achieved 97.8% precision on telecommunications franchises ($N = 360$), 96.6% on hospital bed capacity acts ($N = 203$), and 92.2% on school infrastructure laws ($N = 645$).

The probabilistic confidence metrics evaluate how decisively the model assigns unseen statutes to their respective legal domains. Across all 5,087 holdout statutes, the model achieved a Mean Test Cosine Confidence of 0.5508, exceeding the typical 0.4000

operational target and demonstrating that the top assigned category captured dominant probability mass. Furthermore, 56.97% of all unseen statutes were classified with high certainty (≥ 0.50 probability), well above the majority target threshold of 50.0%. The holdout evaluation yielded a Cosine Silhouette Score of 0.1602 and a Davies-Bouldin Index of 2.4496, well within the target range formulated by Davies and Bouldin [41]. These geometric measurements prove that the topic model establishes stable, mathematically separated legal boundaries before documents are indexed for coarse retrieval.

3.3.4 Local Legislative Corpus Analysis

To evaluate local legislative patterns, the proponents conducted an empirical census of all 1,664 digitized Davao City local ordinances. Unlike the national archive which spans over a century of centralized governance, the local collection documents municipal self-governance across four distinct legislative eras. As illustrated in Figure 3.10(a), the archive contains 50 foundational enactments from the Pre-LGC Era (1950–1991), 184 landmark regulatory codes from the Post-LGC Historical Era (1992–2015), and 145 enactments from the Pre-Pandemic Era (2016–2020). Statutory volume concentrates heavily within the modern digital LISSP database, contributing 1,285 enactments (77.2%) between 2021 and 2025. This temporal surge reflects the establishment of standardized electronic archiving by the Sangguniang Panlungsod, providing complete records of contemporary local lawmaking.




Figure 3.10: Temporal Distribution across Four Legislative Eras and Two-Dimensional Latent Semantic Topic Space of the Davao City Local Legislative Corpus ($N = 1,664$).

Figure 3.10 presents a two-panel visual analysis of the 1,664 local ordinances, showing Panel (a) with enactment frequencies across four historical legislative eras (1950–2025) and Panel (b) with the two-dimensional latent semantic space of 64-component Truncated SVD vectors alongside 1.5σ covariance ellipses and cluster centroids (C_1 – C_6). This visual is included to demonstrate the longitudinal expansion of local legislation and to validate the semantic clustering into six distinct functional archetypes grounded in the Local Government Code of 1991. In Panel (a), legislative volume surges markedly during the modern digital LISSP era (2021–2025), contributing 77.2% ($N = 1,285$) of all accessible enactments. In Panel (b), emergency disaster funds (C_3) and road closures (C_5) form tight, isolated clusters, while substantive police power regulations (C_1) and contractual authorizations (C_2) span distinct operational subspaces. This functional partition proves that

local legislation exhibits structured semantic regularity, enabling retrieval safeguards to distinguish between routine administrative agreements and penal regulations that interact directly with national law.

Applying Truncated SVD to the local corpus revealed that city legislation divides into six functional archetypes grounded in the Local Government Code of 1991 [70]. To eliminate procedural noise before modeling, the pipeline extracts substantive operative text while setting aside routine session attendance headers. Word vectors were constructed across an 8,000-feature vocabulary and projected into 64 latent semantic dimensions, capturing 50.65% of the total vocabulary variance. Figure 3.10(b) displays the two-dimensional latent semantic topic space for the local collection. Emergency Quick Response Fund (QRF) calamity grants form an isolated, tight cluster in the upper-left quadrant (C_3), while temporary road closures project into a compact traffic management cluster (C_5). Inter-agency agreements (C_2) and deeds of donation (C_4) span an elongated horizontal subspace along the positive horizontal axis, corresponding to executive contract authorizations under Section 22(c) of Republic Act No. 7160. Direct police power enactments (C_1) and zoning permits (C_6) anchor the lower-left quadrant, governing local community conduct.

A critical methodological comparison between the national statutory topic landscape (Figure 3.7) and the local semantic space (Figure 3.10(b)) is the degree of visual dispersion. In the national corpus, over 3,500 cookie-cutter school conversions and 2,000 public utility franchises collapse into sharp, isolated radial rays due to identical statutory drafting templates. In contrast, local legislation is intrinsically cross-functional under Republic Act No. 7160. Under Sections 16, 21, 22(c), 455, and 458 of the Local Government Code, a single substantive regulatory ordinance routinely incorporates budgetary appropriations, mayoral contract-signing authorizations, and penal sanction clauses within the same enactment. Furthermore, local enactments share standardized council preambles (“WHEREAS” recitals) and procedural session metadata. While projecting 64 latent dimensions onto two axes inevitably compresses non-linear variance and produces visual overlap, the full 64-dimensional space achieves an out-of-sample supervised classification accuracy of 86.79% and a 0.7973 mean cosine confidence on holdout tests. Displaying the authentic principal dimensions alongside 1.5σ covariance density ellipses and centroids preserves empirical integrity without resorting to artificial dimensional cherry-picking. Table 3.4 details the composition, percentage share, statutory authority, core subject matter, and top class-based TF-IDF keywords for each discovered archetype.

Table 3.4 details the six functional legislative typologies discovered across the 1,664 Davao City local enactments, listing each archetype’s identifier, descriptive label, enactment count, percentage share, governing statutory authority under Republic Act No. 7160, and top salient *c-TF-IDF* keywords. This typology is included to establish the legal taxonomy used to categorize local enactments, reflecting the statutory distribution of local council police powers and administrative authorizations. Substantive Regulatory and Penal enactments constitute the largest single typology (41.11%, $N = 684$), followed by Inter-Agency Agreements (28.91%, $N = 481$) and Disaster Relief / QRF Fund authorizations (11.36%, $N = 189$). Administrative authorizations (MOAs, calamity releases, deeds of donation, and road closures) account for 58.89% of all council legislation. This

Table 3.4: Composition, Functional Typology, and Salient *c-TF-IDF* Keywords of the Davao City Local Legislative Corpus ($N = 1,664$)

ID	Functional Legislative Typology	Count	Share	Statutory Authority	Core Salient <i>c-TF-IDF</i> Keywords
01	Substantive Regulatory & Penal	684	41.11%	RA 7160 §458	public, safety, fine, penal, violation, fee, health, traffic
02	Inter-Agency Agreements (MOA/MOU)	481	28.91%	RA 7160 §§22(c), 455	pursuant, local, sign, panlungsod, contracts, agreement
03	Disaster Relief & QRF Funds	189	11.36%	RA 10121 §21	fund, calamity, disaster, risk, reduction, quick response
04	Deeds of Donation & Usufruct	140	8.41%	RA 7160 §22(a)(4)	donation, local, powers, pursuant, contracts, deed, usufruct
05	Temporary Road Closures	122	7.33%	RA 7160 §21	road, closure, temporary, vehicular traffic, street, route
06	Zoning & Subdivision Permits	48	2.88%	RA 7160 §20, BP 220	zoning, zone, land, reclassification, residential, commercial
Total Local Collection		1,664	100.00%	All Digitized Ordinances (1950–2025)	

breakdown demonstrates that the majority of local legislative actions represent routine executive authorizations rather than penal ordinances, confirming the empirical necessity of isolating substantive regulatory provisions prior to downstream conflict evaluation.

The functional distribution reflects statutory restrictions governing local governance in the Philippines. Substantive regulatory and penal enactments form the largest single group (41.11%, $N = 684$), containing community regulations such as truck bans, cemetery setbacks, stray animal controls, and drinking water safety. Meanwhile, administrative authorizations comprise nearly half of all council activity. Under Section 22(c) and Section 455(b)(1)(vi) of Republic Act No. 7160, the city mayor cannot bind the local government to contracts or accept donations without prior council approval. Consequently, inter-agency agreements (28.91%), calamity fund authorizations under Republic Act No. 10121 (11.36%), and deeds of donation (8.41%) appear frequently in the local record.

To verify that the discovered local typologies represent stable semantic boundaries, the proponents conducted holdout validation using an 80/20 stratified split ($N_{\text{train}} = 1,331$, $N_{\text{test}} = 333$). In evaluating this six-class categorization, all benchmark targets in Table 3.5 are explicitly grounded in established literature. Under Morrison’s proportional chance criterion [127], the default uniform baseline for six categories is $1/6 = 16.67\%$. The out-of-sample accuracy target ($\geq 80.0\%$) and Macro F1 target ($\geq 70.0\%$) follow the domain-adapted document categorization standards from Aggarwal [2]. For cluster separation, Rousseeuw [159] established that silhouette scores between 0.05 and 0.15 denote authentic structural cluster boundaries in cross-functional administrative text, while Davies and Bouldin [41] established that index values between 3.0 and 4.0 confirm compact cluster scatter for sparse short-document collections. Table 3.5 summarizes the out-of-sample performance metrics, chance baselines, literature targets, and explicit title probes.

Table 3.5 reports the empirical holdout validation metrics, default baseline compar-

Table 3.5: Empirical Holdout Validation Metrics, Baseline Comparisons, and Explicit Statutory Title Probes for the Local Topic Modeler ($N_{\text{train}} = 1,331$, $N_{\text{test}} = 333$)

Validation Metric / Benchmark Test	Observed	Default Baseline	Target Range	Empirical Verdict
Out-of-Sample Supervised Accuracy	86.79%	16.67%	$\geq 80.0\%$	Strong Generalization [2]
Out-of-Sample Macro F1-Score	75.39%	16.67%	$\geq 70.0\%$	Balanced Discrimination [2]
Out-of-Sample Weighted F1-Score	84.84%	16.67%	$\geq 80.0\%$	High Volume Precision [2]
Mean Test Cosine Confidence	0.7973	0.1667	> 0.5000	Dominant Probability Mass [76]
High-Confidence Predictions (≥ 0.50)	92.79%	$< 20.0\%$	$> 75.0\%$	Clear Majority Partition [76]
Holdout Cosine Silhouette Score	0.0914	≤ 0.0000	0.0500–0.1500	Authentic Text Separation [159]
Holdout Davies-Bouldin Index	3.7445	> 5.0000	3.0000–4.0000	Compact Cluster Scatter [41]
Explicit Local Statutory Title Probes (<i>Unambiguous Subsets across Corpus</i> , $N = 872$)				
Quick Response Fund / Calamity Assistance	98.82%	16.67%	$> 90.0\%$	High Precision ($N = 169$)
Substantive Penal & Police Power Regulations	96.10%	16.67%	$> 90.0\%$	High Precision ($N = 77$)
Temporary Road Closure & Traffic Diversions	87.96%	16.67%	$> 85.0\%$	High Precision ($N = 108$)
Inter-Agency Accord / Authority to Sign MOA	86.56%	16.67%	$> 80.0\%$	High Precision ($N = 372$)
Zoning & Subdivision Development Permits	80.00%	16.67%	$> 75.0\%$	High Precision ($N = 15$)
Deeds of Donation & Property Usufruct	29.77%	16.67%	$> 25.0\%$	Shared Mayoral Authorization ($N = 131$)

isons, literature-grounded target ranges, and explicit statutory title probes for the local legislative topic modeler evaluated on 333 unseen local ordinances ($N_{\text{train}} = 1,331$, $N_{\text{test}} = 333$). The validation classifier achieved an out-of-sample accuracy of 86.79% and a Weighted F1-Score of 84.84%, far exceeding the 16.67% uniform chance baseline (1/6) and satisfying the 80.0% Aggarwal threshold. Predictions showed strong decisiveness, with a Mean Test Cosine Confidence of 0.7973 and 92.79% of unseen ordinances classified with high confidence (≥ 0.50). Title probe verification on unambiguous indicator subsets confirmed high precision across Quick Response Fund releases (98.82%, $N = 169$), substantive penal regulations (96.10%, $N = 77$), temporary road closures (87.96%, $N = 108$), inter-agency agreements (86.56%, $N = 372$), and zoning permits (80.00%, $N = 15$). The lower precision on deeds of donation (29.77%, $N = 131$) occurs because donation acceptances are legally drafted as mayoral authorizations to sign deeds, directly sharing contractual syntax with inter-agency agreements under Section 22(c) of Republic Act No. 7160.

Document lengths in the local collection differ noticeably from national statutes.

Across all 1,664 local enactments, ordinances exhibit a median length of 5,644 characters and 880 words, with an interquartile range of 736 to 1,607 words. The local archive contains 11,801 total section provisions. Automated preprocessing isolates 8,160 operative substantive provisions (69.15%) with a median length of 98 tokens, while pruning 3,641 procedural boilerplate sections (30.85%) such as separability, repealing, and effectivity clauses.

3.3.5 Dual Statutory Corpus Integration

Combining national legislation with Davao City local enactments establishes a unified search space of **27,096 legal records** comprising **176,421 searchable provisions**. In automated statutory conflict detection, evaluating draft measures requires examining both vertical statutory preemption and horizontal legislative consistency. Under the doctrine of statutory preemption articulated in *Magtajas v. Pryce Properties* [187], local measures are tested against the 25,432 superior national statutes. Concurrently, under intra-jurisdictional coherence principles, draft measures are evaluated against the 1,664 existing city ordinances to prevent conflicting local commands, duplicate penalties, or unintended implied repeals. Table 3.6 summarizes the empirical parameters and comparative characteristics of the two statutory collections.

Table 3.6: Comparative Metrics and Structural Parameters of the Dual Statutory Search Space ($N = 27,096$ Enactments, $N_{\text{sections}} = 176,421$)

Empirical Dimension / Metric	National Corpus	Local Corpus	Combined Corpus	Analytical Significance
Total Enactments (N)	25,432	1,664	27,096	Comprehensive legislative space across jurisdictions
Total Section Provisions	164,620	11,801	176,421	Total searchable provision-level retrieval units
Operative Provisions	164,620	8,160	172,780	Candidate substantive rules evaluated for conflict
Procedural Boilerplate Sections	0	3,641 (30.85%)	3,641	Filtered procedural clauses pruned before modeling
Historical Enactment Span	1900–2026 (126 yrs)	1950–2025 (75 yrs)	1900–2026	Longitudinal statutory corpus across legal regimes
Median Enactment Characters	10,807	5,644	10,480	Document-level length disparity across collections
Median Enactment Words	1,512	880	1,472	Local enactments display concise regulatory scope
Median Provision Tokens	122.0	90.0	120.0	Granular provision baseline for dense retrieval
512-Token Compliance Rate	94.83%	92.72%	94.69%	Standard transformer attention compatibility
Operative 512-Token Compliance	94.83%	90.06%	94.60%	High sequence compliance for operative text
Normative Conflict Scope	Vertical Preemption	Horizontal Coherence	Dual Scope	Comprehensive statutory preemption and consistency

Table 3.6 compares the structural parameters and empirical metrics of the National Corpus, Local Corpus, and Unified Dual Corpus, detailing enactment counts, provision

volumes, operative provisions, boilerplate counts, historical spans, length statistics, and 512-token compliance rates. This comparative summary is included to establish the definitive scale and composition of the dual-jurisdiction search space, validating the technical feasibility of applying transformer language models across the unified archive. As shown in the table, combining 25,432 national statutes with 1,664 local ordinances establishes an archive of 27,096 enactments yielding 176,421 searchable provisions, of which 172,780 represent substantive operative rules and 3,641 represent filtered procedural boilerplate. Across all provisions, 94.69% fit strictly within the 512-token transformer context ceiling without truncation. This high compliance rate confirms that statutory section segmentation provides natural, self-contained retrieval units that can be directly processed by standard BERT and ModernBERT attention windows without semantic truncation.

Figure 3.11 illustrates the empirical length disparity between the national and local collections. In Panel (a), document-level character distributions on a logarithmic scale highlight the structural asymmetry between national codes and local ordinances. National statutes display a broad spread with an extreme right tail exceeding one hundred thousand characters, driven by omnibus codes such as the National Internal Revenue Code and Local Government Code. In contrast, local ordinances concentrate tightly around their median of 5,644 characters (880 words), reflecting their focused administrative nature. In Panel (b), provision-level token density curves reveal that both national statutory sections (median: 122 tokens) and local operative provisions (median: 98 tokens) concentrate within 40 to 180 tokens. This distribution confirms that natural legislative sections are sufficiently granular to avoid semantic dilution while remaining well below the 512-token ceiling.




Figure 3.11: Enactment Length Asymmetry and Provision-Level Subword Token Densities across the Dual Statutory Search Space ($N = 27,096$ Enactments, $N_{\text{sections}} = 176,421$).

Figure 3.11 presents the empirical length disparity and subword token density distributions across the dual statutory search space. Panel (a) illustrates full-text enactment character lengths on a logarithmic scale, highlighting the structural contrast between massive national omnibus codes and concise local ordinances. Panel (b) depicts probability density curves for national provisions (median: 122 tokens), local operative provisions (median: 98 tokens), and local boilerplate clauses (median: 72 tokens) plotted against the 512-token transformer context window. This visual comparison is provided to confirm that provision-level chunking resolves the document length asymmetry between national and local enactments. As shown in Panel (b), isolating natural legislative sections shifts both

statutory and ordinance text into an optimal density window (40–180 tokens), ensuring that provisions retain specific legal commands without diluting dense vector representations.

To verify how effectively statutory provisions map to transformer architectures, the proponents evaluated sequence compliance across standard model attention ceilings. Figure 3.12(a) traces the empirical cumulative distribution functions (CDF) across provision token lengths. Exactly 94.83% of national provisions and 92.72% of local provisions (90.06% of operative substantive sections) fall strictly under the 512-token boundary. Across all 176,421 provisions in the combined search space, 94.69% fit cleanly within standard BERT context windows without truncation. Figure 3.12(b) compares compliance rates across candidate transformer context windows with dedicated legends for National Statutes and Davao City Local Ordinances. While compact 128-token windows accommodate only 53.8% of national provisions, compliance increases to 82.3% at 256 tokens and 94.8% at 512 tokens, with long-context architectures such as ModernBERT capturing over 99.9% of all provisions natively.




Figure 3.12: Cumulative Provision Compliance Curves and Attention Window Evaluation across the Dual Statutory Search Space ($N = 27,096$ Enactments, $N_{\text{sections}} = 176,421$).

Figure 3.12 evaluates sequence length compliance across candidate transformer attention windows for the 176,421 statutory provisions. Panel (a) plots empirical cumulative distribution functions up to 1,024 tokens, denoting the 94.69% aggregate compliance point at the 512-token threshold. Panel (b) presents a comparative bar chart displaying provision fit across five candidate attention ceilings (128, 256, 512, 1,024, and 8,192 tokens) for both national statutes and local ordinances. This evaluation is included to justify candidate model selection and input window sizing in Stage 1 and Stage 2. As observed in Panel (b), while standard 512-token windows accommodate nearly 95% of all provisions, long-context models such as ModernBERT provide native coverage exceeding 99.9%, eliminating the need for aggressive token truncation on lengthy statutory sections.

To handle provisions in the 5.31% exceedance tail without information loss under standard-context models, the pipeline implements a hierarchical overlapping sub-chunking fallback. Adapting the sliding window strategy formulated by Devlin et al. [48], when an extracted section exceeds a safety threshold of 448 tokens, the preprocessing module divides the section into overlapping sub-chunks using a window of $w = 384$ tokens and a stride of $s = 256$ tokens (128-token overlap). Capping the sub-chunk

window at 384 tokens reserves a 64-token structural buffer within the 512-token ceiling to accommodate the prepended metadata header along with mandatory special tokens ($[CLS]$, $[SEP]$) [48]. Each sub-chunk preserves the prepended structural metadata header [Statute Title, Section S_i : Catchline] to maintain institutional context. In Stage 1 retrieval, the parent provision embedding is computed by mean-pooling sub-chunk vectors, capturing semantic concepts across all paragraphs. In Stage 2 inference, the Cross-Encoder evaluates each sub-chunk independently against the local ordinance hypothesis h , determining provision-level conflict through max-pooling across sub-chunk contradiction probabilities:

Max-Pooled Sub-Chunk Contradiction Probability:

$$\hat{P}(\text{Contradiction} \mid p_i, h) = \max_j P(\text{Contradiction} \mid p_{i,j}, h) \quad (3.4)$$

Where $p_{i,j}$ denotes the j -th sub-chunk of statutory section p_i .

Equation 3.4 formulates the max-pooling aggregation mechanism used to evaluate long statutory sections partitioned into sub-chunks. By assigning the section’s overall contradiction probability to the maximum contradiction score observed across any of its constituent sub-chunks, the formulation guarantees that an isolated prohibition or restrictive proviso buried deep within a lengthy, multi-paragraph section triggers an operative alert without being diluted by surrounding agreeable paragraphs. Methodologically, this max-pooling strategy prevents false negatives on complex omnibus codes where penal clauses are nested within administrative descriptions, ensuring high sensitivity to statutory preemption under standard-context models. Alongside the sub-chunking fallback, the proponents evaluate ModernBERT as a long-context candidate architecture [216, 92], testing whether native 8,192-token attention processing outperforms standard-context models across lengthy codifications.

3.4 Preprocessing and Normalization

Because raw statutory documents originate from disparate digital repositories and historical scanning archives, they require systematic ingestion, boundary segmentation, and text normalization. The preprocessing pipeline transforms heterogeneous legal text into standardized, context-enriched retrieval chunks optimized for coarse search and fine inference. This pipeline addresses four distinct technical requirements: digital-born draft ingestion, procedural boilerplate filtering, archival OCR error correction, and statutory exception handling. By systematically decoupling administrative boilerplate from operative commands, the pipeline ensures that downstream language models evaluate substantive legal mandates. The following subsections detail each operational stage of the text normalization architecture.

3.4.1 Draft Ingestion and Boundary Parsing

A critical operational distinction in this research separates the historical training archive from live ex-ante draft submissions. While the historical Davao City collection consists of

scanned physical PDFs requiring optical character recognition (OCR) and text normalization, newly drafted ordinances submitted for ex-ante review are digital-born documents. Legislative staff create these proposed measures using standard word processing software, generating clean electronic text or searchable PDF files. Because digital-born drafts contain authentic character encodings without scanning artifacts, the live pipeline bypasses OCR error correction entirely. This structural decoupling allows the preprocessor to focus immediately on automated section boundary parsing and clause segmentation. To illustrate the structure of the ingested raw statutory text prior to segmentation, Listing 3.2 displays an authentic raw ingestion record from the national statutory collection (Republic Act No. 11035).

```

1  # Raw Ingestion Record from corpus/national_laws/republic_acts.jsonl (ra_11035_2018)
2  raw_statute_record = {
3      "law_id": "ra_11035_2018",
4      "year": 2018,
5      "url": "https://lawphil.net/statutes/repacts/ra2018/ra_11035_2018.html",
6      "scraped_at": "2026-02-17T18:44:04.745654",
7      "text": (
8          "Seventeenth Congress\nSecond Regular Session\nBegun and held in Metro Manila,
9          "
10         "on Monday, the twenty-fourth day of July, two thousand seventeen.\n"
11         "REPUBLIC ACT No. 11035\n"
12         "An Act Institutionalizing the Balik Scientist Program, Appropriating Funds "
13         "Therefor, and for Other Purpose\n"
14         "Be it enacted by the Senate and House of Representatives of the Philippines "
15         "in Congress assembled:\n"
16         "Section 1. Title. - This Act shall be known as the \"Balik Scientist Act.\"\n"
17         "
18         "Section 2. Declaration of Policy. - The State recognizes that the utilization
19         "
20         "of the expertise of science, technology or innovation experts of Filipino
21         descent "
22         "is a vital component of the national political, economic, and social
23         development "
24         "efforts...\n"
25         "Section 3. Balik Scientist Program. - There is hereby created a Balik
26         Scientist "
27         "Program, hereinafter referred to as the \"Program\", to be administered by
28         the "
29         "Department of Science and Technology (DOST)...\n"
30         "[Raw statutory text continues across 2,493 total words]"
31     )
32 }

```

Listing 3.2: Sample Raw JSONL Ingestion Record from the National Statutory Archive (Republic Act No. 11035)

Listing 3.2 details the key-value schema of a raw JSONL ingestion record from the national statutory collection, containing fields for law identifier, enactment year, source Lawphil URL, scrape timestamp, and verbatim legislative text. This listing is presented to demonstrate the unstructured state of enactments prior to automated parsing, where continuous body text mingles legislative recitals, preambles, section divisions, and procedural catchlines. In the illustrated record for Republic Act No. 11035, the statute comprises over 2,400 words across multiple sections embedded within a single text string. Downstream retrieval and classification engines cannot index or score this unsegmented enactment directly without suffering severe semantic dilution and attention context overflow. Parsing this monolithic record into discrete, self-contained operative sections is

therefore an essential architectural prerequisite for fine-grained preemption analysis.

To evaluate legislative commands without semantic distortion, an automated regular expression parser splits the submitted draft at formal section boundaries into discrete operative blocks. In Philippine local legislative practice, proposed measures follow standardized structural templates established by the Department of the Interior and Local Government [44]. Real-world legislative drafts contain substantial procedural boilerplate that does not express substantive legal commands. If routine clauses such as title declarations, repealing clauses, and effectivity dates are fed into Stage 1 retrieval, they match thousands of administrative national statutes containing similar formulaic phrases. Listing 3.3 details the regular expression parser that segments live draft measures into discrete operative sections while filtering out procedural boilerplate.

```

1  import re
2
3  def extract_operative_provisions(draft_text):
4      """
5      Parses digital-born draft ordinance into discrete sections,
6      filtering out non-operative procedural boilerplate clauses.
7      """
8      # Regex pattern to segment text at formal section boundaries
9      pattern = r'(SECTION\s+\d+\.?[^\\n]*\\n.*?)(?=(?:SECTION\s+\d+|CERTIFIED|ENACTED|\\Z))'
10     matches = re.findall(pattern, draft_text, re.DOTALL | re.IGNORECASE)
11
12     # Procedural clauses that do not carry substantive regulatory commands
13     boilerplate_keywords = ["TITLE", "SEPARABILITY", "REPEALING", "EFFECTIVITY"]
14     operative_sections = []
15
16     for match in matches:
17         clean_chunk = ' '.join(match.strip().split())
18         header = clean_chunk.split('.')[0] if '.' in clean_chunk else clean_chunk[:30]
19
20         # Exclude non-operative procedural boilerplate
21         if any(kw in header.upper() for kw in boilerplate_keywords):
22             continue
23
24         operative_sections.append({
25             "section_header": header,
26             "operative_text": clean_chunk
27         })
28
29     return operative_sections
30

```

Listing 3.3: Code Snippet of the Operative Provision Extractor and Boilerplate Filter

Listing 3.3 presents the Python implementation of the operative provision extractor, featuring regular expression lookahead patterns and a procedural boilerplate keyword filter. This algorithm is incorporated to automate the structural segmentation of raw draft ordinances into clean, self-contained section blocks while systematically discarding administrative clauses. As observed in the listing, the regular expression isolates text between formal section boundaries, while the keyword filter checks section headers against procedural markers including Title, Separability, Repealing, and Effectivity clauses. Discarding these formulaic clauses prevents downstream retrieval models from matching generic legislative boilerplate against thousands of unrelated administrative acts. Consequently, only substantive regulatory and penal provisions advance to the candidate indexing and inference pipeline.

3.4.2 Boilerplate Filtering and Metadata Headers

Filtering formulaic boilerplate prevents search drift and computational waste during candidate retrieval. In both national statutes and municipal enactments, procedural clauses—specifically Title, Separability, Repealing, and Effectivity clauses—contain generic legislative phrases that carry zero substantive regulatory restrictions. When these sections are vectorized, their embeddings align with unrelated administrative acts that happen to share identical legal boilerplate. By systematically pruning these non-operative sections during ingestion, the preprocessor directs candidate retrieval exclusively toward substantive regulatory and penal provisions. As established in the corpus census, this filtering prunes 3,641 procedural sections (30.85%) from the municipal collection, isolating 8,160 operative provisions for indexing.

Once operative sections are isolated, evaluating granular statutory provisions requires preserving their governing legislative identity. An isolated sub-clause stripped of its parent enactment number and title suffers from severe semantic vector dilution, causing retrieval models to confuse municipal regulatory penalties with unrelated national tax or criminal codes. To maintain semantic precision and hierarchy across both stages, the pipeline applies structural context prepending. Listing 3.4 details the context prepending procedure implemented to format granular provisions.

```

1 def prepend_statutory_context(statute_record, section):
2     """
3     Prepends governing legislative identity to isolated section text
4     to prevent semantic vector dilution and preserve hierarchical rank.
5     """
6     prefix = (
7         f"[{statute_record['instrument_type']} No. {statute_record['enactment_number']}: "
8         f"{statute_record['title']} | {section['section_number']}: {section['section_title']}] "
9     )
10    model_ready_text = prefix + section['raw_text']
11    return model_ready_text
12

```

Listing 3.4: Pseudocode of Section Context Prepending for Granular Provision Chunking

Listing 3.4 outlines the structural context prepending routine that attaches governing legislative metadata to isolated statutory sections prior to vector embedding. This procedure is implemented to prevent semantic dilution, ensuring that granular provisions stripped of their parent statute do not lose institutional authority or jurisdiction. In the displayed routine, the function formats a standardized metadata prefix combining instrument type, enactment number, full statute title, section ordinal, and catchline before concatenating the raw section text. This prepended prefix forces the bidirectional self-attention layers in transformer models to condition on the governing legislative context. As a result, the model reliably distinguishes subordinate local regulations from superior national mandates during both dense retrieval and natural language inference.

Every extracted statutory section begins with a structural metadata header containing the statute title, section number, and section title:

Context-Prepended Statutory Chunk Representation:

$$\mathbf{x}_{\text{chunk}} = [\text{Statute Title, Section } S_i: \text{Catchline}] \parallel \mathbf{t}_{\text{verbatim}} \quad (3.5)$$

Where $\mathbf{t}_{\text{verbatim}}$ denotes the exact text extracted from the statutory record, and $\parallel$ represents string concatenation.

Equation 3.5 formulates the mathematical representation of a context-prepended statutory chunk $\mathbf{x}_{\text{chunk}}$, constructed by concatenating the parent legislative metadata header with the verbatim statutory text $\mathbf{t}_{\text{verbatim}}$. This chunk representation is established to maintain institutional hierarchy across both Stage 1 retrieval indices and Stage 2 inference pairs. By explicitly concatenating the enactment title and section catchline to the raw clause, the formulation ensures that the resulting token sequence encodes both the specific regulatory command and its overarching statutory mandate. During neural processing, transformer attention heads compute cross-attention weights across both the metadata header and the operative sentence. This prevents semantic ambiguity and ensures that downstream models recognize whether a regulatory command originates from an omnibus national code or a localized municipal resolution.

Listing 3.5 demonstrates the concrete transformation from an isolated provision into a Stage 1 indexed chunk and the resulting Stage 2 asymmetric input sequence. Prepended metadata provides the neural attention layers with parent enactment authority, preventing false positive matches against subordinate or unrelated municipal penalty clauses.

```

1  # 1. Isolated Statutory Provision (Without Context Prepending)
2  isolated_section = (
3      "Section 458. (a) The sangguniang panlungsod... shall enact ordinances... "
4      "(v) Approve ordinances which shall impose a fine not exceeding P5,000.00 "
5      "or an imprisonment not exceeding one (1) year, or both."
6  )
7
8  # 2. Model-Ready Contextualized Provision Chunk (Indexed by Stage 1 BM25 & Bi-Encoders)
9  stage1_indexed_chunk = (
10     "[Republic Act No. 7160: Local Government Code | Section 458: Powers] "
11     "Section 458. (a) The sangguniang panlungsod, as the legislative body of the city, "
12     "shall enact ordinances... (v) Approve ordinances which shall impose a fine not "
13     "exceeding Five thousand pesos (P5,000.00) or imprisonment not exceeding one year "
14     "... "
15 )
16
17 # 3. Assembled Asymmetric Input Sequence Passed to Stage 2 Cross-Encoder
18 stage2_input_sequence = (
19     "[CLS] Premise: [RA 7160 | Section 458] Approve ordinances which shall impose a "
20     "fine "
21     "not exceeding P5,000.00... [SEP] Hypothesis: [Proposed Ordinance No. 2024-012 | "
22     "Section 6: Penalties] SECTION 6. PENALTY CLAUSE. Any person violating this "
23     "Ordinance "
24     "shall be penalized with a fine of Fifteen Thousand Pesos or two years "
25     "imprisonment. [SEP]"
26 )

```

Listing 3.5: Sample Model-Ready Statutory Chunk and Stage 2 Input Sequence with Context Prepending

Listing 3.5 presents the three-stage transformation of a statutory provision, tracking the progression from an isolated clause into a model-ready indexed chunk and finally into an asymmetric Stage 2 input sequence. This listing is provided to illustrate how structural metadata bridges the gap between raw legislative text and input representations for Cross-Encoder classification. As demonstrated in the listing, an isolated penalty clause from Republic Act No. 7160 lacks explicit municipal scope, but prepending the Local Government Code header and Section 458 catchline clarifies its application to city councils. When formatted for Stage 2 inference, the premise and proposed ordinance hypothesis are joined using special delimiter tokens (*[CLS]*, *[SEP]*) alongside explicit segment tags. This formatted sequence enables the Cross-Encoder to compute full pairwise cross-attention between the statutory penalty ceiling and the proposed ordinance penalty.

The practical necessity for structural metadata headers is demonstrated by sections within composite legal codes. In Republic Act No. 7160, Section 447 sets penalty limits for municipalities, while Section 458 sets distinct penalty limits for city councils. Without adding the section title and number, downstream Cross-Encoders can confuse different levels of local government and misjudge municipal powers. Adding structural metadata ensures that the model evaluates each legal clause under its correct administrative level. As a result, the vector embeddings retain both the exact wording of the clause and its broader statutory context.

3.4.3 Archive Digitization and Ingestion Pipeline

Transforming physical municipal paper archives into machine-readable text requires evaluating candidate document ingestion architectures under consumer hardware constraints. While contemporary measures from the digital LISSP database provide clean electronic text, historical enactments from earlier legislative terms exist predominantly as scanned physical PDF documents. To establish an optimal transcription strategy for the 1,664 municipal records, the proponents formulated a planned candidate extraction experiment comparing three distinct technological paradigms: traditional open-source OCR via Tesseract v5.x [177], modern layout-aware deep learning OCR via Surya OCR [138], and a quantized local vision-language model via Qwen2.5-VL-3B-Instruct [195]. As an independent reference standard for transcription accuracy, the evaluation protocol establishes a gold-standard digitization baseline using Google Gemini 1.5 Flash multimodal extraction coupled with comprehensive line-by-line human legal verification. The candidate pipelines are evaluated across Character Error Rate (CER), Word Error Rate (WER), Table Fidelity (TEDS), and Legal Marker Precision on a representative 60-document stratified test suite, with empirical benchmarking results formally reported in Chapter 4 (Section ??).

To support offline execution on consumer hardware, the vision-language architecture deploys Qwen2.5-VL-3B-Instruct [195], a multimodal transformer that pairs a vision transformer backbone with a 3-billion parameter causal language model decoder. Executed locally on the Neural Training Workstation (Mobile Host B, equipped with an AMD Ryzen 7 7435HS processor, 16 GB DDR5 system memory, and an NVIDIA GeForce RTX 5050 Laptop GPU with 8 GB GDDR6 VRAM running Windows 11), the architecture is configured to transcribe scanned enactments while preserving exact statutory phrasing. Proprietary cloud platforms such as Google Cloud Document AI [69] were excluded from

routine production due to recurring per-page API fees and legal data sovereignty restrictions under the Data Privacy Act of 2012 (Republic Act No. 10173) [153]. Deploying the quantized open-weight model locally resolves these operational constraints, providing high transcription fidelity while executing completely offline within municipal IT boundaries.

The document ingestion pipeline begins by recursively discovering target PDF files across archival municipal directories. Individual enactments are loaded using PyMuPDF, which inspects document metadata and iterates through pages sequentially. To eliminate physical scanning artifacts, the ingestor applies an inset cropbox that trims exactly 3% of the page perimeter along all four borders:

Archival Inset Perimeter CropBox Trimming:

$$\text{CropBox} = [0.03 \cdot W, 0.03 \cdot H, 0.97 \cdot W, 0.97 \cdot H] \quad (3.6)$$

Where W and H represent the physical page width and height in points.

Equation 3.6 defines the four-point coordinate bounds of the archival inset cropbox applied during document rasterization, trimming three percent (0.03) along all perimeter margins of physical page width W and height H . This automated trimming operation is implemented to eliminate edge artifacts—including dark photocopier borders, staple shadows, and binder punch holes—that frequently mar legacy municipal records. By cropping the outermost perimeter before passing the rendered bitmap to the visual encoder, the mathematical transform prevents edge distortions from being interpreted as spurious punctuation marks or initiating runaway autoregressive token generation. The resulting clamped bounding box isolates the substantive legislative text within a clean visual field, stabilizing sequence lengths during visual patch embedding.

Following perimeter trimming, the rasterization module renders each page at an image resolution of 105 dots per inch (DPI). Selecting 105 DPI provides an optimal empirical balance between character legibility and sequence length for high-density legal documents. The rendered image is converted to an RGB array and passed to the visual processor. The processor rescales page patches dynamically, clamping total image pixels between a minimum of 200,704 pixels ($256 \times 28 \times 28$) and a maximum of 351,232 pixels ($448 \times 28 \times 28$). Because Qwen2.5-VL allocates visual tokens in 28×28 patch increments using two-dimensional rotary position embeddings (2D-RoPE), bounding the pixel budget stabilizes sequence length and prevents out-of-memory GPU exceptions during long batch runs.

To accommodate local workstation hardware constraints without sacrificing recognition fidelity, the model weights are quantized using 4-bit NormalFloat4 (NF4) precision via the BitsAndBytes library [47]. The quantization configuration uses double quantization and computes attention projections in 16-bit floating point (`torch.float16`) using PyTorch scaled dot-product attention (`sdpa`). Under 4-bit NF4 compression, the model footprint occupies less than 3.5 GB of operational VRAM, allowing the pipeline to execute comfortably on consumer-grade laptop graphics hardware. Text generation is guided by an explicit verbatim system prompt instructing the model to transcribe all section titles, ordinals, dates, names, and monetary values without summarization or com-

mentary. Generation settings enforce deterministic greedy decoding (`do_sample = False`, `use_cache = True`) bounded by a ceiling of 1,000 new tokens per page.

The pipeline adopts a page-level processing architecture rather than ingesting entire multi-page documents in a single model forward pass. Ingesting full multi-page ordinances simultaneously would require concatenating thousands of visual patch tokens, which dilutes spatial self-attention and exceeds consumer GPU memory ceilings. In contrast, processing documents page by page enforces focused visual attention on discrete text blocks while scaling linearly ($O(N)$) across documents ranging from single-page resolutions to 50-page tax codes. The pipeline inserts explicit page delineators (`--- [Page X] ---`) into the extracted text, ensuring continuous traceability back to physical council records. Furthermore, page-level execution supports fault-tolerant checkpointing: the system inspects output files on disk and automatically resumes execution across interrupted batch queues by skipping previously completed files. Page-level garbage collection (`gc.collect()`) and explicit CUDA cache clearing (`torch.cuda.empty_cache()`) are executed after each document, preventing memory fragmentation during multi-hour ingestion runs.

3.4.4 Substantive Legislative Body Extraction

Scanned municipal enactments exhibit a tripartite administrative structure consisting of prefatory attendance rosters, operative legislative provisions, and post-enactment signature blocks. Isolating the substantive legal text from administrative matter is critical for downstream legal conflict detection. Retaining councillor roll-call rosters introduces extraneous personal names that distort Stage 1 dense retrieval embeddings, while retaining signature blocks adds visual clutter to candidate hypotheses in Stage 2. Slicing documents using fixed page numbers fails across municipal archives because single-page resolutions contain operative text on the front page, whereas omnibus codes distribute roll calls and committee rosters across multiple front pages.

To resolve this structural variance, the proponents developed an adaptive boundary state machine that extracts operative legal provisions based on canonical statutory phrasing. The extraction module scans transcripts for authorized opening markers, including enactment designations (`ORDINANCE NO.`), formal enactment titles (`AN ORDINANCE`), and preamble clauses (`WHEREAS`). Upon detecting an opening boundary, the state machine captures the enactment identifier, title, preamble, and all operative sections (`SECTION 1`, `ARTICLE I`, `CHAPTER I`). Text capture halts immediately upon encountering formal terminal attestation markers, such as Sanggunian enactment notices (`ENACTED`), secretarial certifications (`CERTIFIED CORRECT`), vice-mayoral attestations (`ATTESTED`), or executive approval notations (`APPROVED`).

The extraction module incorporates a robust fallback hierarchy to prevent data loss when archival scans exhibit degraded boundary markers. If an explicit enactment header is obscured, the extractor falls back to the long ordinance title; if the title is absent, it anchors on the initial operative section. Similarly, if formal certification headers are missing, the system terminates extraction following the standard effectivity section (`SECTION X. EFFECTIVITY`). Each extraction records an explicit confidence score:

Substantive Body Boundary Extraction Confidence Score:

$$\text{Confidence}_{\text{extract}} = 0.35 \cdot C_{\text{start}} + 0.35 \cdot C_{\text{end}} + 0.30 \cdot C_{\text{sec}} \quad (3.7)$$

Where C_{start} and C_{end} denote boundary match confidences, and C_{sec} measures section sequential coherence.

Equation 3.7 formulates the composite confidence metric $\text{Confidence}_{\text{extract}}$ used by the adaptive state machine to evaluate substantive body extraction fidelity. The scoring formulation combines starting boundary alignment C_{start} , terminal attestation match C_{end} , and sequential section coherence C_{sec} through an empirical weighting of 0.35, 0.35, and 0.30, respectively. This confidence scoring mechanism is incorporated to detect extraction anomalies automatically without halting batch processing across hundreds of municipal ordinances. Documents scoring at or above the 0.70 confidence threshold advance directly to downstream tokenization and vector indexing, whereas low-scoring items are flagged and quarantined into a manual verification queue. This tiered verification ensures that degraded scans with non-standard secretarial certifications receive human legal review while maintaining high automated throughput for compliant enactments.

When confidence falls below 0.70, the document is flagged for manual review while preserving the full unpruned transcript on disk. Listing 3.6 outlines the extraction logic.

```

1  def extract_substantive_body(raw_transcript: str) -> dict:
2      """
3      Extracts operative statutory provisions from municipal transcripts,
4      stripping administrative attendance rosters and terminal signatures.
5      """
6      start_match = RE_ORD_NUM.search(raw_transcript) or RE_TITLE.search(raw_transcript)
7      start_idx = start_match.start() if start_match else 0
8      start_conf = 1.0 if start_match else 0.40
9
10     search_region = raw_transcript[start_idx:]
11     term_match = RE_TERMINAL.search(search_region)
12     end_idx = start_idx + term_match.start() if term_match else len(raw_transcript)
13     end_conf = 1.0 if term_match else 0.50
14
15     substantive_text = raw_transcript[start_idx:end_idx].strip()
16     sections = [int(n) for n in RE_ALL_SECTIONS.findall(substantive_text)]
17     sec_conf = 1.0 if len(sections) >= 3 and sections[:3] == [1, 2, 3] else 0.40
18
19     confidence = round(0.35 * start_conf + 0.35 * end_conf + 0.30 * sec_conf, 3)
20     return {
21         "status": "SUCCESS" if confidence >= 0.70 else "FLAGGED_LOW_CONFIDENCE",
22         "confidence": confidence,
23         "substantive_text": substantive_text
24     }
25

```

Listing 3.6: Adaptive State Machine for Substantive Legislative Body Extraction

Listing 3.6 presents the Python implementation of the adaptive state machine for substantive legislative body extraction. This algorithm is designed to parse raw municipal transcripts and extract substantive enactments while stripping prefatory attendance rosters, administrative preambles, and terminal signature blocks. In the function, regular expressions scan the text for opening enactment indicators and terminal attestation keywords, slicing the transcript between these indices and validating ordinal section continuity. The function computes the composite confidence score according to Equation 3.7

and returns a status dictionary indicating whether the extraction succeeded or requires manual review. Isolating the substantive body in this manner ensures that extraneous councillor names and mayoral titles do not distort dense retrieval vector embeddings or trigger spurious preemption matches in Stage 2.

3.4.5 Automated Quality Audit and Conservative Normalization

Following raw transcription, the system executes an automated quality audit to detect and categorize transcription anomalies across seven mutually exclusive categories: definite OCR corruption, possible corruption, formatting noise, valid legal boilerplate, valid table structures, page metadata anomalies, and items requiring human review. The audit scanner identifies severe failure states, such as Unicode replacement characters (`\ufffd`), foreign script hallucinations (CJK and Cyrillic character ranges), raw decoding markers (`undefined`), and autoregressive repetition loops. Documents with severe corruptions or total lengths below 150 characters are quarantined into a dedicated review queue for manual verification.

To minimize false positives, the audit pattern matching explicitly accommodates authentic legislative structures. A broad repeated-character regular expression frequently flags Markdown table borders (`|---|---|`) as OCR corruption. To prevent this error, table boundary delimiters, pipes, and header markers are explicitly excluded from symbol repetition rules. Valid legislative boilerplate notations such as "xxx xxx xxx", which denote unamended text in amending ordinances, are protected using negative lookahead expressions. Geographic survey coordinates containing degree symbols ($^{\circ}$), minute marks ($'$), and second marks ($''$), as well as accented Filipino surnames (*Peña*, *Dañguilan*), are validated against an extended Unicode whitelist. Furthermore, page sequences beginning at Page 2 due to excluded cover sheets are recognized as valid contiguous sequences rather than misordered files.

Deterministic text normalization is performed via a conservative cleaning pipeline that strictly preserves legal meaning. The cleanup module applies Unicode Form KC normalization, strips invisible zero-width characters, removes stray VLM prompt echoes (such as `**[body text]**` / *zhèngwén* markers), and collapses runaway character loops exceeding seven consecutive repetitions to a standardized three-character representation. All modifications are logged in an audit report with line-level traceability. Under strict architectural safeguards, automated scripts are prohibited from altering personal names, calendar dates, numerical monetary amounts, ordinance numbers, statutory citations, penalty clauses, or section numbering. Discrepancies involving these sensitive entities are routed exclusively to human legal researchers.

3.4.6 Advisory Language Quality Review Architecture

The final quality-control phase is an advisory, second-stage language review system that operates downstream from automated cleanup. The review system functions strictly as a diagnostic tool: it produces ranked verification suggestions but never executes automated text rewriting. The architecture combines a curated Philippine legal and toponymic

gazetteer, contextual language model perplexity scoring, and named-entity recognition. The gazetteer whitelists all 182 official barangays of Davao City, municipal administrative districts, local executive departments, and statutory Latin maxims (*in pari materia*, *mutatis mutandis*, *ultra vires*), preventing specialized terminology from triggering false spelling alarms.

Contextual perplexity scoring evaluates token predictability within a sliding context window using a calibrated legal language model. Tokens exhibiting high pseudo-perplexity and characteristic alphanumeric substitution patterns (such as misreading numeral 0 as uppercase letter O in monetary figures) are assigned high diagnostic confidence. The system presents these candidate errors on an interactive review dashboard, coupling the transcribed text with coordinate bounding boxes linking back to the original physical PDF scan. Human annotators review the highlighted anomalies, confirming authentic drafting choices or correcting transcription errors before final vector indexing.

3.4.7 Harmonic Interpretation and Exceptions

A key challenge in conflict detection is distinguishing genuine normative contradictions from authorized statutory exceptions. In regulatory legislation, general prohibitions and permitted exceptions frequently appear within the same section using conditional syntax. Because statutory sections average between 90 and 122 tokens, an entire section fits comfortably within the 512-token input limit of candidate Cross-Encoders. The model’s attention mechanism reads the general prohibition and its qualifying exceptions simultaneously, preventing false positive contradiction alerts. This joint attention ensures that local measures implementing authorized statutory exceptions are recognized as valid exercises of delegated authority.

When a general prohibition and an authorizing exception reside in separate statutory sections, the pipeline relies on two-stage candidate retrieval. Under Philippine statutory construction, related provisions governing the same subject matter must be interpreted together as a harmonious whole under the doctrine of *in pari materia* [4]. If an ordinance draft relies on an authorized statutory exemption, evaluating a specific clause without retrieving the parent exception can trigger an incorrect conflict warning. The system addresses this challenge by utilizing hybrid retrieval in Stage 1 to pull both general bans and specific exception clauses into the candidate shortlist. The Stage 2 Cross-Encoder then evaluates the draft against the authorizing exception clause and assigns an Entailment classification, confirming that the local measure operates within valid legislative authority.

3.5 Ground Truth Benchmark

Evaluating candidate Cross-Encoder models in Stage 2 requires an expert-verified ground truth dataset reflecting the authentic legal realities of Philippine local legislation. Unlike general-domain natural language inference benchmarks that rely on crowdsourced sentence pairs, evaluating municipal conflict detection requires formal statutory interpretation skills grounded in administrative law. Because no existing NLP benchmark

encompasses Philippine national statutes and local government ordinances, the proponents established an authoritative 350-pair ground truth evaluation dataset curated in collaboration with active legal researchers from the Davao City Sangguniang Panlungsod.

The target scale of exactly $N = 350$ premise-hypothesis pairs is mathematically derived from the intersection of institutional personnel constraints and cognitive ergonomics. Through formal institutional routing, the Davao City Sangguniang Panlungsod designated a cohort of exactly $P = 15$ active legal researchers to participate in the benchmark validation study. To ensure statistical reliability and compute multi-rater inter-annotator agreement, each candidate pair requires independent evaluation by $k = 3$ distinct legal researchers. In specialized legal NLP annotation, cognitive fatigue studies establish that evaluating dense, multi-sentence statutory and ordinance provisions imposes substantial cognitive load; pilot calibration studies indicate that rater attention and analytical rigor degrade significantly when individual annotator workloads exceed approximately 70 pairs [1, 27]. Partitioning the 15 researchers into five balanced sub-panels of three members ($15/3 = 5$ panels) and assigning a manageable cognitive workload of exactly 70 pairs per annotator yields exactly $5 \times 70 = 350$ unique statutory pairs, producing $350 \times 3 = 1,050$ total independent legal judgments across the cohort.

3.5.1 Sample Size and Statistical Power

The determination of the benchmark sample size ($N = 350$ pairs) is grounded in statistical power analysis for multi-class contingency evaluations [34, 56]. Under Cohen’s power framework, detecting semantic conflicts across a three-class discrete decision space ($\mathcal{C} = \{\text{Entailment, Neutral, Contradiction}\}$, with $\text{df} = (3 - 1) = 2$ for goodness-of-fit and $\text{df} = (3 - 1)(3 - 1) = 4$ for 3×3 contingency matrices) with a medium effect size ($w = 0.30$) at a significance level of $\alpha = 0.05$ requires a minimum of $N = 108$ items to achieve standard statistical power of 80% ($(1 - \beta) = 0.80$).

Establishing an evaluation suite of $N = 350$ pairs provides an achieved statistical power exceeding 0.985 ($(1 - \beta) > 98\%$), effectively eliminating Type II error risks ($\beta < 0.015$) and ensuring robust statistical representation across all eight Macro Legal Domains defined in Section 3.3.3. These eight domains reflect the core public services and regulatory police powers defined in Section 17 and Section 458 of Republic Act No. 7160.

This curated scale directly reflects established standards in specialized legal NLP benchmarking. While general-domain benchmarks like SNLI [18] rely on open crowdsourcing across tens of thousands of colloquial sentences, domain-specific legal NLI requires accredited legal professionals whose time and cognitive availability are severely constrained. High-quality, expert-adjudicated legal datasets consistently operate within this targeted scale: the SARA statutory tax benchmark comprises 376 cases [83], the annual COLIEE Task 4 legal entailment test suites provide approximately 100 pairs per competition cycle [146, 5], and ContractNLI [100] relies on curated subsets. At 350 premise-hypothesis pairs with a three-rater overlap yielding 1,050 independent legal judgments, the benchmark represents an authoritative, statistically sound gold standard for municipal legislative review.

3.5.2 Dataset Partitioning

To maintain experimental rigor, prevent data leakage, and ensure reproducible model evaluation, the 350 Ground Truth pairs are partitioned using a static 70/15/15 train, validation, and test split. This split provides 245 training pairs ($N_{\text{train}} = 245$), 52 validation pairs ($N_{\text{val}} = 52$), and 53 test pairs ($N_{\text{test}} = 53$). Operating within this sample size places the evaluation in a few-shot learning regime, which is well-suited for pre-trained transformer language models [63]. In machine learning, few-shot task adaptation utilizes the foundational linguistic and semantic representations already encoded within deep pretrained encoders, fine-tuning classification heads and attention parameters on domain-specific instances without requiring thousands of newly labeled examples [63].

The proponents specifically selected a 70/15/15 partition over a standard 80/10/10 split due to sample size constraints. In a 350-pair dataset, an 80/10/10 allocation yields only 35 validation and 35 test instances. A validation set of only 35 pairs is statistically fragile for fine-grained decision threshold calibration (τ^*), where a single classification error alters observed accuracy by nearly 3%. Furthermore, a test set of 35 items falls below the $N \geq 50$ sample size threshold established by Card et al. [23] for reliable Minimum Detectable Effect (MDE) in natural language processing and risks violating the asymptotic cell expectations of McNemar’s test ($b+c \geq 20$). In contrast, allocating 52 validation pairs and 53 test pairs satisfies both the MDE statistical power criteria and the non-parametric hypothesis testing assumptions while retaining 70% (245 pairs) for robust few-shot Cross-Encoder fine-tuning [202].

3.5.3 Asymmetric Pairing and Hypothesis Formulation

To construct a representative ground truth benchmark, candidate statutory premises must reflect the full breadth of Philippine national law. During the exploratory phase of this research, the proponents conducted preliminary machine learning experiments on the national statutes, uncovering severe class imbalances and repetitive drafting formulas across historical eras. These exploratory findings directly informed the consolidated topic modeling architecture detailed in Section 3.3.3, which organizes the corpus into eight validated Macro Legal Domains such as local governance, public health, and taxation. The proponents then sampled national statutory sections across these eight domains to establish the fixed premises for the evaluation benchmark. To evaluate multi-sentence legal commands without breaking internal qualifications, the pipeline uses Asymmetric Span-Level Pairing, coupling each statutory premise with a multi-sentence ordinance hypothesis block [224]. This asymmetric structure (1 premise vs N hypothesis sentences) ensures that local exceptions, cross-references, and penalty conditions are evaluated together against the national law [100].

To make the evaluation as realistic as possible, the drafted test sentences closely mirror the writing style of Davao City ordinances. The proponents analyzed municipal laws enacted between 2021 and 2025 to identify standard drafting phrases, administrative office names, and structural penalty clauses following the official drafting conventions of the Department of the Interior and Local Government (DILG) [44]. Each drafted clause uses standard enacting syntax, such as: “*SECTION X. [TITLE]. — It shall be*

unlawful for any person, natural or juridical, within the territorial jurisdiction of Davao City, to...” The draft sentences also name real local enforcement bodies, such as the City Transport and Traffic Management Office (CTTMO), the City Health Office (CHO), the City Environment and Natural Resources Office (CENRO), and the Business Bureau of Davao City. Operational units such as the Anti-Vices Task Force (AVTF) and the Davao City Police Office (DCPO) are also included to reflect genuine local enforcement.

The synthetic ordinance clauses are also linked to real administrative districts in Davao City. Draft rules specify geographic areas from urban centers like Poblacion and Talomo to suburban districts like Buhangin and Toril. Local offices such as the City Disaster Risk Reduction and Management Office (CDRRMO) and the City Social Welfare and Development Office (CSWDO) are included to match actual city administration. This realistic drafting prevents evaluators from reviewing artificial example sentences that differ from real legal texts. As a result, the 15 Sangguniang Panlungsod legal researchers evaluate draft provisions that look like real measures introduced on First Reading.

The proponents deliberately utilized domain-grounded, authentic local template-derived hypothesis construction—grounded in existing Davao City ordinances (2021–2025) and DILG drafting manuals, and formally reviewed with thesis adviser Mr. Adrian “Ogs” Ablazo—rather than unconstrained synthetic generation via large language models. Recent evaluations by Gururangan et al. [77] and Poliak et al. [143] demonstrated that automated or crowd-generated NLI hypotheses frequently introduce severe “annotation artifacts”—systematic lexical cues such as negation markers (“not”, “never”) that allow neural models to cheat and predict contradictions without genuinely reading the statutory premise. In specialized legal domains, authoritative benchmarks such as SARA (tax law) [83], DocNLI [224], and Clash-of-Leges [88] establish that rigorous evaluation requires domain-grounded hypotheses reflecting authentic legislative syntax, ensuring models are tested on genuine statutory reasoning rather than superficial generation heuristics.

While the system uses multi-sentence spans to understand local context, it does not attempt cross-document multi-hop reasoning (detecting conflicts that appear only by combining two or more separate national laws). Resolving dependencies across multiple separate laws is an active research challenge in natural language inference. Current systems that attempt this require either complex knowledge graphs [118] or very large generative language models [176]. Running those multi-premise models would exceed the consumer-grade hardware limits set for local government deployment. For this reason, the ground truth benchmark and the inference system focus strictly on single-premise evaluations.

3.5.4 Class Balancing and Adversarial Injection

In standard machine learning on tabular data, researchers often handle uneven classes using methods like the Synthetic Minority Over-sampling Technique (SMOTE) [28]. However, the proponents avoid continuous vector oversampling for this legal text pipeline. SMOTE creates artificial examples by interpolating between numerical feature vectors ($x_{\text{new}} = x_i + \lambda(x_{z_i} - x_i)$). In high-dimensional language models representing discrete words, averaging or blending token vectors creates artificial representations that confuse

attention layers and carry no real legal meaning. Simple word replacement also fails with legal texts because statutory words have exact legal definitions; for example, swapping mandatory words like “shall” with permissive words like “may” fundamentally changes a legal duty.

To build a balanced, reliable Ground Truth dataset, the researchers combined historical legal records with expert-guided text adjustments. First, the researchers gathered documented legal clashes from local government records, including mayoral veto messages, repealed ordinances, and explicit statutory amendments [88]. Second, following the evaluation methods of Koreeda and Manning [100] and Nie et al. [134], legal annotators introduced conflicting rules into real statutory provisions. Annotators reversed permissions, changed numerical zoning limits, and added conflicting penalties into ordinance clauses. This expert-guided process produced an exact balanced distribution across the Entailment, Neutral, and Contradiction classes in the 350 Ground Truth pairs (117 pairs per class), preventing class imbalance without creating unnatural legal phrasing.

3.5.5 Stratified Difficulty Tiers

Building an effective evaluation benchmark requires pairs with varying levels of difficulty. If all contradiction examples were simple penalty ceiling violations or obvious word swaps, models could score well by memorizing shallow keyword patterns [124]. To construct a rigorous evaluation suite, the proponents adapted the difficulty tiering methodology established by Koreeda and Manning in the ContractNLI benchmark [100]. This approach is supported by the behavioral stress-testing frameworks of Ribeiro et al. [155] and Naik et al. [130], which separate surface competence tests from semantic challenge sets. Furthermore, the 30% Tier 1, 40% Tier 2, and 30% Tier 3 quota follows the adversarial challenge partitioning established in benchmark engineering literature, such as SQuAD 2.0 [147], FEVER [197], and ANLI [134]. These foundational benchmarks establish that maintaining an adversarial challenge quota of 30% to 40% is optimal to prevent models from exploiting superficial lexical artifacts while preserving sufficient representation of standard rule-based compliance. Table 3.7 outlines the stratified difficulty tiers, legal conflict mechanisms, and representative examples for the 350-pair benchmark.

Table 3.7 outlines the three-tier jurisprudential difficulty stratification governing the 350-pair ground truth benchmark, detailing target percentage shares, underlying conflict mechanisms, and representative statutory examples. This stratification is incorporated to ensure that the benchmark does not merely test shallow keyword matching or simple penalty ceiling checks, but rigorously challenges models across diverse preemption types. As shown in the table, the benchmark allocates 30% to Tier 1 (surface and quantitative conflicts, 105 pairs), 40% to Tier 2 (preemption and statutory carve-outs, 140 pairs), and 30% to Tier 3 (latent and paraphrastic conflicts, 105 pairs). Crucially, the Tier 3 challenge set deliberately omits all numerical penalty values and explicit statutory citations, requiring the Cross-Encoder to evaluate pure semantic preemption. This distribution establishes an adversarial benchmark reflecting the complexities of legislative committee review while preventing neural models from relying on superficial lexical heuristics.

Table 3.7: Stratified Jurisprudential Difficulty Tiers for the 350-Pair Ground Truth Evaluation Benchmark

Difficulty Tier	Target Share	Legal Conflict Mechanism	Representative Statutory Example
Tier 1: Surface & Quantitative	30%	Direct penalty ceiling violations and explicit changes in mandatory words (<i>shall</i> converted to <i>may</i>).	Local draft imposes a PHP 15,000 fine for traffic obstruction when RA 7160 Sec. 458 caps city penalties at PHP 5,000.
Tier 2: Preemption & Carve-Outs	40%	Exceeding local regulatory powers under <i>Magtajas</i> , and blanket local bans that omit statutory exceptions.	Draft ordinance enacts a total ban on commercial freight transit without exempting national agricultural and medical logistics.
Tier 3: Latent & Paraphrastic	30%	Subtle conflicts where the draft redefines statutory terms or creates incompatible rules without using the exact statutory words.	Draft anti-vagrancy ordinance assigns apprehended minors to detention cells, violating non-custodial diversion under RA 9344 without citing the law.

The three difficulty tiers map directly to established Philippine municipal preemption doctrines and linguistic challenge categories, reflecting the jurisprudential distinction between rule-based and value-based normative conflicts formulated by Sartor [171] and Bench-Capon and Sartor [14]. Tier 1 evaluates surface and quantitative preemption, reflecting direct penalty ceiling violations under Section 458 of Republic Act No. 7160 and explicit modal verb clashes (*shall* versus *may*). Tier 2 adapts the “negation by exception” category from Koreeda and Manning [100], evaluating unhonored statutory exemptions and local overreach under *Magtajas v. Pryce Properties* [187] and *City of Manila v. Laguio, Jr.* [189]. Tier 3 evaluates latent and paraphrastic contradictions, testing whether the model detects substantive statutory friction when a draft redefines terms without sharing exact statutory words. Allocating the benchmark into 30% Tier 1 (105 pairs), 40% Tier 2 (140 pairs), and 30% Tier 3 (105 pairs) establishes an adversarial quota that reflects real-world committee reviews while preventing models from exploiting simple lexical heuristics [124].

Grounding Tier 3 contradictions in legal doctrine ensures that models are evaluated on genuine language understanding rather than keyword presence. Under established Supreme Court rulings, a local council cannot enact indirect prohibitions that take away rights protected by national statutes [187, 189]. To enforce this distinction during testing, the proponents verified that all 104 Tier 3 draft sentences contain zero numbers or monetary amounts. Without explicit numerical clues, the Cross-Encoder must evaluate the underlying legal meaning of the text, such as recognizing when an anti-vagrancy draft detains minors in violation of Republic Act No. 9344. This design ensures that the benchmark provides an authentic diagnostic test of semantic preemption for the subsequent analysis in Chapter 4.

3.5.6 Expert Legal Annotation Panel

Classifying the sampled text pairs requires domain-specific legal training to ensure authentic and reliable labels. The Davao City Sangguniang Panlungsod currently comprises 24 elected councilors across the city’s three legislative districts, supported by a professional staff of approximately 48 active legal researchers [38]. In the initial administrative correspondence transmitted to the Sangguniang Panlungsod Secretariat alongside the Legal Annotation Guidebook, the proponents formally requested a cohort of 15 volunteer legal researchers. Unlike general NLP tasks that rely on open crowdsourcing, legal natural language inference requires formal statutory interpretation skills to prevent superficial keyword matching [178, 227]. The recruited legal researchers review local ordinances as their regular professional duty, providing the necessary institutional competence for building a dependable gold standard [27].

To establish an objective and transparent panel structure, the proponents gathered baseline qualifications during the initial intake communication. The intake request systematically recorded two key criteria from each volunteer: highest educational attainment and total years of professional experience in legislative drafting. Using these reported credentials, the 15 researchers were organized into five balanced sub-panels of three members each ($k = 3$). Within each three-member sub-panel, the evaluator possessing the highest educational attainment and the longest drafting tenure was designated as that panel’s Senior Annotator. The proponents then communicated these panel assignments and Senior Annotator designations in a follow-up confirmation email to establish operational readiness before annotation began.

Following the consensus guidelines adapted from Ablazo [1], a label is accepted into the final Ground Truth if at least two of the three reviewers in a sub-panel agree. In cases where all three reviewers select different categories, resulting in a three-way split with no majority, the pair is referred to that panel’s designated Senior Annotator for binding qualitative adjudication [64, 10]. The Senior Annotator reviews the pair against the standardized codebook to resolve whether the split arose from statutory ambiguity or an overlooked proviso. This hierarchical adjudication step ensures that every item in the 350-pair benchmark reaches an authoritative label while preserving the initial independent ratings for calculating Fleiss’ Kappa (κ). Maintaining this structured division between initial independent scoring and expert tie-breaking safeguards both statistical validity and qualitative consistency across the dataset.

To recognize the time and professional expertise dedicated by the participating Sangguniang Panlungsod legal researchers, evaluators who complete their assigned annotation batches are awarded an official Certificate of Contribution. This certificate serves as a formal token of appreciation and institutional documentation acknowledging their voluntary advisory assistance in validating the statutory benchmark. Each certificate is endorsed by the academic department leadership, thesis adviser, and student proponents to confer formal institutional validity. Providing this acknowledgment establishes a collaborative partnership between the university research team and the local legislative staff, ensuring appropriate recognition for the legal professionals who support the research.

3.5.7 Annotation Codebook and Decision Rules

To help the 15 legal researchers apply consistent criteria, the proponents developed a standardized Legal Annotation Codebook. Adapting principles from Koreeda and Manning [100] and Italiani et al. [88], the codebook turns the three-class NLI framework into clear legal decision rules. Each evaluated pair consists of a single national statute section as the premise and a multi-sentence municipal ordinance clause as the hypothesis. Annotators assign one of three labels: Contradiction, Entailment, or Neutral. Table 3.8 summarizes the working definitions, legal triggers, and example patterns used in the annotation guidelines.

Table 3.8: Sangguniang Panlungsod Ground Truth Legal Annotation Codebook and Decision Rules

Logical State	Operational Legal Definition	Primary Legal and Structural Triggers	Representative Statutory Pattern
Contradiction	The municipal ordinance prohibits what national statute permits, permits what statute prohibits, or exceeds statutory ceilings [187].	Clashes in mandatory words (<i>shall</i> vs. <i>may</i>); penalty limit violations (RA 7160 Sec. 458); territorial or franchise preemption.	National statute caps ordinance fines at PHP 5,000; draft ordinance mandates a PHP 15,000 penalty.
Entailment	The municipal ordinance directly executes, complies with, or derives its regulatory authority from the national premise without adding conflicting mandates.	Compliant devolved implementation; literal statutory adoption; regulatory enforcement within prescribed legislative boundaries.	Local draft establishes a municipal health board strictly adhering to the composition mandated by RA 7160.
Neutral	The national statute and municipal ordinance address the same broad legal domain, but govern independent obligations without mutual conflict or derivation.	Shared topic without direct legal clash; complementary administrative scopes; distinct non-overlapping legal subjects.	National statute establishes hospital licensing standards; local draft regulates municipal market sanitary permits.

Table 3.8 details the operational legal definitions, primary statutory triggers, and representative drafting patterns defined in the Sangguniang Panlungsod annotation codebook across the three target NLI classes. This decision framework is provided to establish consistent, legally grounded classification criteria for the 15 volunteer legal researchers, translating natural language inference categories into established Philippine administrative doctrines. As detailed in the table, Contradiction is triggered by Magtajas preemption, modal command inversions, and local penalty escalations; Entailment is triggered by valid devolved regulatory implementation; and Neutral is assigned when provisions share broad subject matter without normative conflict. Standardizing these decision rules ensures high inter-rater reliability during independent scoring and provides clear criteria for Senior Annotator tie-breaking. Consequently, the resulting gold labels reflect rigorous legal consensus rather than subjective linguistic intuition.

The decision rules for Contradiction directly apply the consistency test established in the Magtajas Doctrine [187]. Because local governments have delegated rather than unlimited legislative powers under Republic Act No. 7160 [70], any local draft that exceeds national statutory limits is legally invalid (*ultra vires*). Annotators look for shifts in mandatory language, such as turning a strict command (“shall”) into an optional choice (“may”), or removing an exemption that national law protects. In the test samples, numerical rules are audited against national limits, including tax caps and penalty ceilings. These criteria ensure that conflict warnings highlight actual legal problems rather than minor wording choices.

The guidelines follow a structured consensus process to resolve disagreements among the three panel reviewers. When two or three reviewers agree, the majority label is accepted directly into the benchmark. In cases where all three reviewers select different categories, the item is referred to the sub-panel’s designated Senior Annotator for final qualitative review [64]. The Senior Annotator reviews the evaluators’ optional notes against the codebook to determine whether the difference came from subtle statutory phrasing or an overlooked exception. This review step ensures that every item in the 350-pair Ground Truth reaches an authoritative label while preserving the original votes to calculate Fleiss’ Kappa.

To illustrate the structure of the evaluated instances, Listing 3.7 presents an authentic sample record (Pair ID: GT-001) from the 350-pair Ground Truth benchmark dataset, showcasing the paired national statutory premise, proposed municipal hypothesis, gold conflict label, and expert legal rationale.

```

1  # Authentic Evaluation Pair GT-001 from data/ground_truth_350.jsonl
2  benchmark_pair = {
3    "pair_id": "GT-001",
4    "macro_domain": "Executive Issuances & Policy Reorganization",
5    "difficulty_tier": "Tier 1: Surface & Quantitative",
6    "national_premise": {
7      "statute_title": "Republic Act No. 7581 (The Price Act)",
8      "citation": "Section 6",
9      "statutory_text": (
10         "Section 6. Automatic Price Control. - Unless otherwise declared by the "
11         "President, prices of basic necessities in an area shall automatically be "
12         "frozen at their prevailing prices or placed under automatic price control "
13         "whenever that area is proclaimed or declared a disaster area or under a "
14         "state of calamity... price control of basic necessities under this section "
15         "shall remain effective for the duration of the condition that brought it
16         about, "
17         "but not for more than sixty (60) days."
18       ),
19     },
20     "ordinance_hypothesis": {
21       "ordinance_title": (
22         "Proposed City Ordinance No. 2024-082 (The Davao City Calamity Price "
23         "Stabilization and Consumer Welfare Protection Measure)"
24       ),
25       "section_citation": "Section 4",
26       "hypothesis_text": (
27         "SECTION 4. MANDATORY LOCAL PRICE CEILINGS UPON CALAMITY. - In accordance "
28         "with local civil defense standards... upon official declaration of a state "
29         "of local calamity by the Sangguniang Panlungsod, the City Mayor... shall "
30         "enforce mandatory price ceilings across all grocery and commercial retail "
31         "establishments within Davao City, which price freeze shall remain in
32         continuous "
33         "legal effect for a mandatory period of one hundred eighty (180) calendar days
34         ..."

```

```

32     )
33 },
34 "gold_label": "Contradiction",
35 "legal_rationale": (
36     "Direct breach of statutory temporal ceiling. Section 6 of RA 7581 explicitly
37     limits "
38     "automatic calamity price control to not more than sixty (60) days. The draft
39     ordinance "
40     "mandates a 180-day freeze, exceeding delegated police power limits."
41 ),
42 "annotator_metadata": {
43     "assigned_panel": "Panel_A",
44     "evaluators": ["SP-ANN-01", "SP-ANN-02", "SP-ANN-03"]
45 }

```

Listing 3.7: Sample Evaluation Pair from the Ground Truth Benchmark Dataset

Listing 3.7 presents an authentic evaluation record from the ground truth dataset (GT-001), detailing the JSON schema including pair identifier, macro domain, difficulty tier, national premise citation and text, ordinance hypothesis citation and text, gold NLI label, legal rationale, and annotator panel metadata. This listing is included to show the complete data representation passed to candidate Cross-Encoders during evaluation and to document the provenance of each gold label. In this specific Tier 1 instance, the hypothesis mandates a 180-day calamity price freeze, which directly contradicts Section 6 of the Price Act (Republic Act No. 7581) capping automatic freezes at 60 days. The record captures both the exact statutory violation and the panel metadata documenting independent scoring by three legal researchers. Storing this rich metadata enables granular error analysis across difficulty tiers and legal domains in Chapter 4.

3.5.8 Annotation Workload Allocation

To prevent reviewer fatigue while keeping the results statistically sound, this study adapts the batched review framework developed by Ablazo [1]. In that earlier study, each item required three independent evaluations to determine a majority label. Managing a collection of 12,054 short texts required over 36,000 individual reviews. By dividing this work among 250 students, the workload was reduced to about 145 items per person. This demonstrated how batched distribution can maintain a three-vote overlap without placing an unrealistic burden on reviewers.

This study applies the same mathematical setup to the legal domain. Because reviewing multi-sentence legal provisions requires much more concentration than reading short messages, the dataset is smaller, but the distribution logic is identical. The workload for the legal reviewers is calculated as follows:

Annotator Workload Allocation Formulation:

$$\begin{aligned}
 N &= 350 \text{ premise-hypothesis pairs} \\
 k &= 3 \text{ independent votes per pair} \\
 P &= 15 \text{ legal annotators} \\
 W &= N \times k = 350 \times 3 = 1,050 \text{ total annotations} \\
 \text{Work per person} &= \frac{W}{P} = \frac{1,050}{15} = \mathbf{70} \text{ pairs per annotator}
 \end{aligned} \tag{3.8}$$

Where N represents the benchmark size, k is the multi-rater replication factor, P is the annotator pool size, and W denotes total evaluations.

Equation 3.8 details the mathematical derivation of the annotator workload distribution, dividing the total required annotations W among the pool of $P = 15$ legal researchers. With a benchmark size of $N = 350$ premise-hypothesis pairs evaluated under a three-rater replication factor ($k = 3$), the protocol requires a total of $W = 1,050$ independent legal judgments. The formulation allocates this workload into five disjoint 70-pair sets, assigning exactly one set of 70 pairs to each three-member sub-panel. Methodologically, capping individual reviewer workload at 70 pairs prevents cognitive fatigue and attention degradation during complex legal evaluation. This batched structure maintains full multi-rater overlap for Fleiss' Kappa reliability calculations while ensuring that no single annotator is overburdened.

To organize this workflow, the dataset is divided into five equal sets of 70 pairs each. The 15 researchers are grouped into five sub-panels of three members, with each sub-panel reviewing one 70-pair set. This arrangement ensures that every pair receives three independent reviews while keeping each reviewer's workload manageable.

To reduce fatigue during the review process, the proponents provided structured reference materials. Reviewing statutory pairs requires significant effort, as reviewers must compare detailed legal clauses against local drafts. Evaluators received spreadsheets with complete statutory text to ensure they did not have to read isolated sub-clauses out of context. Each reviewer also received a one-page summary sheet explaining the *Magtajas* decision steps and statutory penalty limits. These reference guides helped reviewers work consistently across their 70-pair evaluation sets.

3.5.9 Inter-Rater Reliability

To measure the consistency of the Ground Truth labels, the proponents calculate Fleiss' Kappa (κ). While raw percentage agreement shows how often reviewers picked the same label, it does not account for agreement that occurs by chance. Fleiss' Kappa isolates genuine agreement by mathematically subtracting the probability of random guessing from the final score.

This metric is used instead of Cohen's Kappa because Cohen's formula is limited to exactly two raters, whereas this study assigns three reviewers to each pair. The resulting score is interpreted using the guidelines from Landis and Koch [104], where a κ score between 0.61 and 0.80 indicates "Substantial Agreement."

Using this target in natural language inference is supported by Bowman et al. [18] in their work on the Stanford Natural Language Inference (SNLI) benchmark. Evaluating text relationships involves human judgment; in their study, Bowman et al. reached an overall Fleiss’ Kappa of 0.70 (scoring 0.77 for Contradiction, 0.72 for Entailment, and 0.60 for Neutral). They showed that achieving very high agreement (> 0.80) in language inference is rare and usually happens only with overly simple test sentences. Because Philippine local ordinances are more complex than general English sentences, expecting agreement higher than 0.70 is unrealistic. Therefore, this study targets a minimum κ threshold of 0.61, which confirms substantial agreement among the legal reviewers.

3.6 Hardware and Computing Setup

Rather than assuming an idealized single-machine environment, the proponents adopted a practical, task-appropriate computing allocation strategy across local hosts and cloud resources. Specific research workloads—ranging from corpus web scraping and document scanning to deep learning fine-tuning—were assigned to the computing environment best suited for that operational phase.

Desktop Host A (Primary Desktop Host) serves as the primary workstation for general data collection, processing, and edge simulation of local government IT capacity. Powered by an Intel Core i3-10105F CPU (4 cores, 8 threads @ 3.70 GHz), 8.00 GB of DDR4 system RAM, an AMD Radeon RX 6600 GPU (8 GB GDDR6 VRAM), and high-speed NVMe solid-state storage running Microsoft Windows 10 Pro 64-bit, this machine executed the automated web scraping of the 25,432 Philippine national statutes from the Lawphil Project, raw HTML parsing, metadata structuring, and daily utility scripts. Crucially, it serves as the physical edge-constrained testbed to verify that Stage 1 sparse (BM25) and dense retrieval index lookups execute efficiently within the modest hardware limits typical of Davao City legislative administrative offices.

Mobile Host B (Mobile GPU Workstation) is configured on a dedicated laptop (Lenovo Legion 5 15AHP10, Model 83M0) equipped with an AMD Ryzen 7 7435HS processor (8 cores, 16 threads @ 3.80 GHz), 16.0 GB of high-speed DDR5 RAM, an NVIDIA GeForce RTX 5050 Laptop GPU with 8 GB GDDR6 VRAM, and PCIe NVMe storage running Microsoft Windows 11 Home 64-bit (Build 26200). As the host with dedicated CUDA acceleration, Mobile Host B is recommended for heavier local processing tasks. Specifically, it executed the scanning, PDF text extraction, and optical character recognition (OCR) cleaning experiments for the 1,664 Davao City local ordinances, and was dedicated to Stage 2 Cross-Encoder fine-tuning, hyperparameter sweeps, latency profiling, and full system evaluation.

The Cloud Computing Environment (Google Colab Standard Tier) provided on-demand external GPU resources when tasks exceeded local workstation capabilities. Utilizing standard cloud runtime instances equipped with NVIDIA T4 and V100 Tensor Core GPUs with standard system RAM, this infrastructure was deployed opportunistically for exploratory Jupyter notebook experiments, heavy matrix computations (such as large-scale Singular Value Decomposition for topic modeling over the 25,432-document

corpus), and generating visual assets where strict run-to-run experimental reproducibility or latency benchmarking was not required.

Across all physical and virtual environments, the research pipeline standardizes on an identical Python software stack to ensure full cross-platform compatibility. Table 3.9 summarizes the physical hardware configurations, operational research roles, and exact software library versions governing the study.

Table 3.9: Hardware Configurations, Software Environment, and Operational Roles across Computing Environments

Environment	Hardware Specifications	Operating System	Assigned Research Operations
Desktop Host A (Primary Desktop)	Intel Core i3-10105F (4C/8T @ 3.7 GHz) 8 GB DDR4 RAM, RX 6600 8GB High-Speed NVMe SSD	Microsoft Windows 10 Pro 64-bit (Build 19045)	National statute scraping (25,432 laws), HTML pars- ing, text normalization, daily utility scripts, and LGU edge simulation.
Mobile Host B (Mobile GPU Work- station)	AMD Ryzen 7 7435HS (8C/16T @ 3.8 GHz) 16 GB DDR5 RAM, NVMe PCIe SSD NVIDIA RTX 5050 Lap- top GPU	Microsoft Windows 11 Home 64-bit (Build 26200)	Local ordinance scanning (1,664 ordinances) and OCR cleaning; dedicated Stage 2 DL fine-tuning, la- tency benchmarking, and final evaluation consistency.
Cloud Environ- ment (Google Colab Stan- dard Tier)	NVIDIA T4 / V100 GPU Standard RAM Runtime	Ubuntu Linux (Cloud Kernel)	Heavy exploratory note- books, large-scale SVD matrix factorization, and non-deterministic visual as- set generation.
Standardized Software Stack and Execution Libraries			
Runtime / Frame- work	Python 3.10+ / 3.14 PyTorch (torch) 2.14.0+	Transformers	Hugging Face (transformers) 5.16.1+
Dense Embed- dings	sentence-transformers 6.0.1+	Lexical Search	rank-bm25 0.2.2+ (Okapi BM25)
ML & Statistics	scikit-learn 1.9.0 / scipy 1.18.1	Document Inges- tion	PyMuPDF (fitz) 1.28.2 / pypdf 6.16.2
Data Processing	pandas 3.0.5 / numpy 2.5.2	Manuscript Build	TeX Live 2026 / MiKTeX (pdflatex, bibtex)

Table 3.9 summarizes the physical hardware configurations, operating systems, assigned research operations, and standardized software execution libraries across the three computing environments used in this study. This comparative infrastructure breakdown is presented to ensure computational transparency and demonstrate that the system’s operational requirements align with accessible consumer-grade and local government IT resources. As detailed in the table, routine data ingestion and edge latency simulation execute on Desktop Host A (Core i3-10105F, 8 GB RAM), heavier local deep learning fine-tuning and OCR tasks run on Mobile Host B (Ryzen 7 7435HS, RTX 5050 Laptop GPU), and resource-intensive matrix decompositions utilize Google Colab standard tier. Standardizing across an identical Python 3.10+ runtime, PyTorch 2.14, and Hugging Face Transformers ensures deterministic model behavior and seamless cross-platform script ex-

ecution. This multi-tiered setup validates that while deep learning training benefits from dedicated mobile GPU acceleration, the resulting inference pipeline executes comfortably within modest edge hardware constraints.

3.7 Pipeline Architecture

The live conflict detection workflow uses a two-stage architecture designed to combine fast searching with detailed logical analysis. The live pipeline begins when a user uploads a draft municipal ordinance. The system first runs the BM25 algorithm to quickly search the statutory database for exact lexical matches. This sparse search captures specific statutory citations, agency acronyms, and numerical penalty limits. At the same time, a dense Bi-Encoder maps the draft provision and the statutes into a continuous semantic space, measuring cosine distance to capture legal meaning that keyword matching overlooks [150]. Rather than relying on fragile score normalization, the pipeline combines the two retrievers using Reciprocal Rank Fusion (RRF) to generate an ordered shortlist of the top- k candidate provisions [37]:

Reciprocal Rank Fusion (RRF) Scoring:

$$\text{RRF}(d) = \sum_{m \in \mathcal{M}} \frac{1}{k_{\text{rrf}} + \text{rank}_m(d)} \quad (3.9)$$

Where $\mathcal{M} = \{\text{BM25}, \text{Dense}\}$ represents the candidate retrieval models, $\text{rank}_m(d)$ denotes the ordinal position of document d produced by model m , and $k_{\text{rrf}} = 60$ is the smoothing constant.

Equation 3.9 formulates the Reciprocal Rank Fusion (RRF) scoring algorithm, combining the ordinal retrieval ranks produced by sparse BM25 and dense Bi-Encoder models into a single unified candidate score. This rank-based fusion mechanism is incorporated to bypass fragile numerical score normalization, as raw BM25 scores fluctuate based on document length and term frequency whereas dense cosine similarities inhabit a bounded $[-1, 1]$ interval. By operating strictly on ranked positions with a smoothing constant of $k_{\text{rrf}} = 60$, the formulation ensures that candidates ranking highly across both lexical and semantic streams receive maximum priority while penalizing one-sided outliers. In practice, this formulation produces an ordered shortlist where statutes matching exact legal citations and those matching conceptual legal synonyms are jointly elevated. Consequently, downstream Cross-Encoder inference evaluates a high-recall candidate pool that is robust against both keyword omissions and semantic drift.

Standard linear score interpolation ($\alpha S_{\text{BM25}} + (1 - \alpha) S_{\text{Dense}}$) requires min-max normalization across unbounded BM25 distributions, which fluctuates wildly across diverse queries. In contrast, Reciprocal Rank Fusion operates entirely on ordinal rankings, making the combination robust against differences in score scales. The resulting fused score ranks provisions that demonstrate both lexical precision and semantic relevance at the top of the candidate pool. Listing 3.8 presents the procedural logic of the dual-stream hybrid retrieval and rank fusion pipeline.

```

1 def retrieve_top_k_candidates(draft_section, statutory_corpus, k=50, k_rrf=60, beta
  =1.20):
2     # 1. Retrieve top-100 candidates from sparse BM25 and dense Bi-Encoder
3     bm25_results = search_bm25(corpus=statutory_corpus, query=draft_section, top_n
  =100)
4     dense_results = search_dense(corpus=statutory_corpus, query=draft_section, top_n
  =100)
5
6     # 2. Compute Reciprocal Rank Fusion (RRF) scores across candidate pool
7     candidate_pool = set(bm25_results.keys()) | set(dense_results.keys())
8     fused_scores = {}
9     for prov_id in candidate_pool:
10        rank_bm25 = bm25_results.get_rank(prov_id, default=101)
11        rank_dense = dense_results.get_rank(prov_id, default=101)
12        rrf_score = (1.0 / (k_rrf + rank_bm25)) + (1.0 / (k_rrf + rank_dense))
13
14        # Apply hierarchy safeguard: boost primary national Republic Acts by 20%
15        if prov_id.startswith("RA_") or is_primary_act(prov_id):
16            rrf_score *= beta
17        fused_scores[prov_id] = rrf_score
18
19    # 3. Sort candidates by fused score and return top-k shortlist
20    return sorted(fused_scores.items(), key=lambda x: x[1], reverse=True)[:k]
21

```

Listing 3.8: Pseudocode of the Stage 1 Reciprocal Rank Fusion Hybrid Provision Retrieval Pipeline

Listing 3.8 outlines the Python implementation of the Stage 1 hybrid retrieval pipeline, integrating sparse BM25 search, dense Bi-Encoder retrieval, and Reciprocal Rank Fusion with hierarchy prioritization. This listing is provided to illustrate the procedural logic that filters the 27,000+ statutory provisions down to the top- k candidate shortlist ($k = 50$). In the routine, the query draft section is searched independently across sparse and dense indices to retrieve the top 100 candidates from each modality. The algorithm then iterates across the union of retrieved provision identifiers, calculates the RRF score using Equation 3.9, and applies a 20% score boost ($\beta = 1.20$) to primary national statutes (Republic Acts) to safeguard superior enactments against administrative clutter. Sorting the fused scores yields a compact, high-precision candidate shortlist that delivers essential statutory context to Stage 2 inference in under 50 milliseconds.

3.7.1 Baseline Selection

To evaluate whether the proposed two-stage retrieve-then-entail architecture achieves meaningful empirical advances over standard approaches, the proponents established formal baselines for both operational stages grounded in recent legal informatics literature.

For coarse statutory retrieval, the pipeline is benchmarked against pure Okapi BM25 lexical matching adapted from Team AIIR [219] in the COLIEE 2025 competition. BM25 serves as the canonical baseline for statutory information retrieval (COLIEE Task 3) because legislative drafting frequently relies on explicit section citations, formal agency acronyms, and verbatim penalty amounts. Choosing BM25 as the Stage 1 baseline provides an essential reference point to determine whether semantic dense vector representations and Reciprocal Rank Fusion genuinely improve candidate recall, especially across paraphrased or indirect preemption queries where exact keyword matching is known to fail.

For fine natural language inference, the pipeline establishes a dual-baseline framework combining an unadapted task baseline with a canonical supervised foundation. Following the legal textual entailment baseline established by Rabelo et al. [146] for COLIEE Task 4, an unadapted zero-shot Cross-Encoder tests whether off-the-shelf general-domain or multi-task language models can detect statutory preemption without local domain adaptation. This baseline establishes the empirical lower bound and directly quantifies the severity of affirmative bias. Complementing this, the proponents establish the standard **bert-base-uncased** model (110M parameters) developed by Devlin et al. [48] as the canonical supervised baseline. Fine-tuned under identical training conditions on the curated statutory pairs, standard BERT isolates the exact performance premium contributed by specialized attention mechanisms, legal domain pretraining, and warm NLI initialization over traditional masked language modeling.

3.7.2 Stage 1 Candidate Retrieval

To identify the optimal neural retriever for the coarse search stage, the proponents evaluate eight candidate open-weight bi-encoders representing four distinct inductive categories reviewed in Section 2.8.6. These candidate models encompass distilled edge baselines (**all-MiniLM-L6-v2** with 22.7M parameters and **bge-small-en-v1.5** with 33.4M parameters), general MTEB benchmark leaders (**bge-base-en-v1.5**, **e5-base-v2**, and **all-mpnet-base-v2** at 109M parameters), legal domain-adapted models (**legal-bert-base-uncased** at 110M parameters), and modern long-context or multilingual encoders (**ModernBERT-base** at 149M parameters and **bge-m3** at 568M parameters with 8,192-token context windows). All candidate bi-encoders are evaluated standalone across candidate retrieval recall (Recall@ k for $k \in \{5, 10, 20, 50\}$), Mean Reciprocal Rank (MRR), mean CPU query latency, and latent paraphrastic conflict sensitivity (Tier 3). The bi-encoder delivering the optimal balance of high candidate recall and low edge latency is selected as the primary dense retrieval workhorse for the hybrid Reciprocal Rank Fusion pipeline.

3.7.3 Stage 1 Retrieval Safeguards

Using a two-stage retrieve-then-entail setup introduces potential search challenges when comparing draft ordinances against large statutory collections. If the coarse search misses the governing national law, the inference model cannot detect the conflict regardless of its analytical capability. Legal conflicts often appear through indirect wording, where local drafts use permissive phrasing or leave out specific statutory terms. Standard search tools can also miss isolated prohibitions tucked inside long, otherwise compliant legislative codes. To protect the pipeline against these retrieval errors, the system incorporates five architectural safeguards across data preparation and search. Table 3.10 summarizes these statutory challenges alongside their operational countermeasures.

Table 3.10 categorizes the six major statutory search challenges encountered in Philippine legislative preemption detection alongside their corresponding architectural countermeasures and operational logic in Stage 1. This safeguard matrix is established to address real-world failure modes identified during preliminary retrieval benchmarks across the 25,432 national statutes. As detailed in the table, the pipeline addresses isolated

Table 3.10: Architectural Retrieval Safeguards Against Real-World Preemption Failure Modes in Stage 1

Statutory Challenge	Architectural Countermeasure	Operational Implementation Logic
Isolated Statutory Clauses	Structural Section Chunking (~ 200 – 300 token windows)	Segments multi-chapter enactments into discrete provisions so isolated prohibitions rank independently without vector dilution.
Indirect or Evasive Wording	Dense Bi-Encoder Semantic Vectors (<code>all-MiniLM-L6-v2</code>)	Maps conceptual intent into 384-dimensional latent space, retrieving relevant statutes despite divergent local phrasing.
Numerical Penalty Limits	Sparse BM25 Lexical Matching with Sublinear Scaling	Matches exact statutory citations, numerical penalty limits, and administrative agency acronyms directly.
Rank Variance in Shortlist	Top- k Candidate Safety Buffer ($k = 50$ provision shortlist)	Retains 50 candidate provisions for Stage 2 review, ensuring relevant statutes ranked below the top ten are evaluated.
Omnibus Statute Trap & Administrative Clutter	Hierarchy Priority Safeguard & Soft Domain Priors	Prioritizes primary legislative statutes (Republic Acts) over administrative orders, while applying gentle soft priors instead of destructive hard filtering.
Long-Tail Provisions (> 512 Tokens)	Dual-Stream BM25 Invariance & Hierarchical Max-Pooling	BM25 indexes full text without sequence limits, while dense representations of tail provisions apply hierarchical sub-chunking with inference max-pooling and long-context candidate support (ModernBERT 8,192 tokens).

prohibitions through structural section chunking, indirect phrasing via 384-dimensional dense semantic vectors, exact penalty limits via sparse BM25 matching, ranking variance through a 50-candidate safety buffer, administrative clutter via hierarchy boosting, and long provisions via hierarchical sub-chunking and ModernBERT support. Implementing these countermeasures prevents catastrophic recall drops, ensuring that complex omnibus enactments like the Local Government Code are not prematurely pruned. As a result, the candidate shortlist reliably captures both surface penalty limits and subtle statutory carve-outs.

These safeguards directly resolve critical failure modes uncovered during preliminary retrieval experiments across the 25,432 national statutes. In particular, attempting a naive hard filter based on predicted macro domains caused candidate Recall@50 to collapse from 43.14% down to 13.43%. This severe degradation occurs because foundational Philippine enactments, such as Republic Act No. 7160, function as omnibus codes that span local administration, disaster response, and municipal taxation simultaneously. Hard-filtering by a single predicted topic immediately prunes the governing code whenever a draft clause touches another administrative area. Replacing hard filtering with a Hierarchy Priority Safeguard and soft domain priors follows the ranking principles of Robertson and Zaragoza [157] and Cormack et al. [37], ensuring that multi-domain omnibus enactments are not prematurely excluded. This safeguard shields Republic Acts from the 12,210 administrative executive orders (48% of the corpus), increasing the Mean Reciprocal Rank from 0.1521 to 0.2217 (+45.8% improvement).

Separating Stage 1 retrieval from Stage 2 inference is essential when running on standard office hardware. Feeding full-length statutes directly into a Cross-Encoder requires quadratic computation ($O((L_{\text{ord}} + L_{\text{stat}})^2)$), which quickly exhausts local memory on long legislative texts. Full-text inputs also dilute model attention, burying specific penalty clauses under thousands of routine procedural words. The hybrid retriever in Stage 1 performs a rapid coarse-to-fine filter ($25,432 \rightarrow 50 \rightarrow 1$), removing over 99.8% of unrelated laws in milliseconds. This separation allows the Cross-Encoder to focus its attention entirely on high-probability candidate provisions.

To empirically validate these architectural mechanisms, the proponents benchmarked nine candidate Stage 1 retrieval configurations across both the full 25,432 national statutory database and the 350-pair ground truth evaluation set. The full empirical retrieval findings, difficulty tier ablations, latency metrics, and comparative analyses are formally reported and discussed in Chapter 4 (Section 4.2).

3.7.4 Stage 2 Conflict Detection

After receiving the top- k candidate shortlist from Stage 1, the pipeline moves to logical evaluation using the Natural Language Inference (NLI) module. The system formats an input sequence for each candidate provision, pairing the retrieved national law (the premise) with the draft ordinance (the hypothesis). To guide the model’s attention, the pipeline places a $[CLS]$ token at the start of the sequence and a $[SEP]$ token between the two texts. This structure helps the model distinguish the existing law from the proposed draft. The combined text is then processed by the selected Cross-Encoder model. The model’s attention mechanism compares every word in the premise against every word in the hypothesis to evaluate their relationship. A classification layer outputs calibrated softmax probabilities labeling the relationship as Entailment, Contradiction, or Neutral. If the contradiction probability crosses the designated threshold, the Explainable AI (XAI) module highlights the conflicting text for user review.

3.8 Model Selection and Optimization

Before deploying the system, the researchers evaluate candidate Cross-Encoder models to identify the most accurate architecture for semantic conflict detection. Rather than selecting an arbitrary architecture or relying on enterprise cloud APIs, the candidate pool is bounded by the computational capacity of standard Local Government Unit (LGU) computing hosts. Specifically, candidate models must satisfy an upper ceiling of 600M parameters, an fp16 inference memory footprint under 2 GB VRAM, and deterministic sequence classification over {Entailment, Contradiction, Neutral}. While leading systems in the COLIEE 2026 competition explored large generative ensembles requiring substantial GPU clusters [5, 133], the proponents intentionally evaluate compact, encoder-only models that can execute locally on consumer hardware. Once the optimal architecture is identified and calibrated, its model weights and tokenizer configurations are permanently pinned for offline execution (`local_files_only = True`), ensuring that the production pipeline runs securely within air-gapped municipal IT networks without external cloud dependencies.

3.8.1 Model Selection and Screening

To evaluate a comprehensive spectrum of transformer inductive biases for Stage 2 semantic conflict detection, the proponents scouted 28 candidate open-weight encoders spanning six architectural families. These families comprise: (1) legal-domain pretrained encoders, including **legal-bert-base-uncased** (110M), **PoL-BERT-Large** (340M), **InLegalBERT** (110M), **CaseLaw-BERT** (110M), and **Lawformer** (110M); (2) pre-aligned NLI reasoners and general baselines, including **DeBERTa-v3-base-NLI** (86M), **DeBERTa-v3-large-NLI** (435M), **cold deberta-v3-base** (86M), **bert-base-uncased** (110M), **roberta-large-mnli** (355M), **bart-large-mnli** (406M), and **electra-large-disc** (335M); (3) modern long-context encoders, including **ModernBERT-base** (149M), **ModernBERT-large** (395M), **longformer-base-4096** (149M), **bigbird-roberta-base** (128M), and **nomic-bert-2048** (137M); (4) native cross-encoder rerankers, including **bge-reranker-v2-m3** (568M), **bge-reranker-base** (278M), and **ms-marco-MiniLM-L12** (33M); (5) multilingual encoders, including **mdeberta-v3-base** (86M), **xlm-roberta-base** (270M), **xlm-roberta-large** (550M), and **roberta-tagalog-base** (110M); and (6) distilled edge baselines, including **all-MiniLM-L6-v2** (22M), **nli-distilroberta-base** (82M), **distilbert-base-uncased** (66M), and **albert-base-v2** (12M).

Evaluating all 28 candidate models through exhaustive multi-epoch hyperparameter sweeps and full human expert benchmarking would impose prohibitive computational and administrative overhead. In accordance with established best practices for reporting experimental search spaces in empirical machine learning [51, 24], the proponents implement a structured Three-Stage Engineering Funnel ($28 \rightarrow 13 \rightarrow 6 \rightarrow 1$). In Stage 1 (Rapid Zero-Shot Screening, $N = 28$), all 28 candidate architectures are evaluated directly out-of-the-box on the held-out test split ($\mathcal{D}_{\text{test}}$, $N = 53$) without task adaptation, surfacing latency and memory bounds while quantifying baseline affirmative bias where unadapted models default toward Entailment or Neutral on complex legal text [5, 199]. In Stage 2 (Supervised Fine-Tuning and Ablation Sweeps, $N = 13$), architectures satisfying operational criteria advance to supervised training on the training split ($N_{\text{train}} = 245$) and validation split ($N_{\text{val}} = 52$) across standardized hyperparameter grids ($\eta \in \{1 \times 10^{-5}, 2 \times 10^{-5}, 3 \times 10^{-5}\}$, $B \in \{8, 16\}$), exploring loss convergence dynamics, early-stopping regularization, and cost-sensitive threshold calibration ($\tau^* \in [0.25, 0.50]$). Finally, in Stage 3 (Out-of-Sample Ground Truth Finalist Benchmark, $N = 6$), the search space is disciplined into an elite finalist pool comprising the Top 5 empirical performers from the validation leaderboard alongside the canonical supervised baseline anchor (**bert-base-uncased**) for definitive evaluation on the 350-pair human consensus ground truth established by the 15 Sangguniang Panlungsod legal researchers.

The decision to advance the Top 5 validation models alongside the canonical baseline anchor to the final Ground Truth evaluation is grounded in four authoritative methodological foundations. First, regarding holdout preservation and data snooping prevention, Dwork et al. [54] established that candidate model selection and hyperparameter tuning must be strictly quarantined to holdout validation data ($\mathcal{D}_{\text{val}}$) to prevent test-set snooping and guarantee unbiased generalization metrics. Second, in statistical hypothesis testing and multiple comparison control, comparing dozens of models pairwise inflates Type I error rates under Bonferroni penalties [50]; evaluating a compact

pool of six finalists anchored to **bert-base-uncased** permits high-powered pairwise hypothesis testing ($\chi^2, p < 0.05$) without multi-testing distortion. Third, under Green AI capacity scaling trade-offs, advancing both high-capacity 400M-parameter models and efficient 86M-parameter workhorses enables an empirical cost-benefit analysis assessing whether marginal accuracy gains justify increased latency and memory overhead on local LGU computing hosts [173, 183]. Fourth, the Top 5 validation candidates ensure exhaustive hypothesis representativeness, capturing warm multi-task NLI pre-alignment (**DeBERTa-v3-large-NLI** and **DeBERTa-v3-base-NLI**), modern unpadded long-context attention (**ModernBERT-large**), unaligned architecture control (cold **deberta-v3-base**), and massive administrative legal-domain pretraining (**PoL-BERT-Large**). The comprehensive empirical screening benchmark across all 28 candidate models, zero-shot performance statistics, asterisk designations, and admission decisions are formally presented in Chapter 4 (Section 4.3).

3.8.2 Overfitting Controls and Calibration

Because fine-tuning deep transformer models in a few-shot setting ($N_{\text{train}} = 245$) can lead to memorizing training phrasing, the researchers use several regularization techniques. First, model optimization uses the AdamW algorithm with weight decay ($\lambda = 0.01$), as developed by Loshchilov and Hutter [116]. This weight decay keeps network parameters small and prevents individual weights from growing disproportionately large. Second, dropout layers with a rate of $p = 0.1$ are applied across attention layers and the classification head [182]. By turning off a random subset of neurons during training passes, dropout encourages the model to learn broader semantic features rather than depending on specific word patterns.

Third, training uses early stopping based on validation results [144]. The system measures validation loss and the validation Macro F1-Score after every training epoch on the 52-pair validation set ($\mathcal{D}_{\text{val}}$). If the validation Macro F1-Score does not improve over four consecutive checks (patience = 4), training stops automatically. The system then restores the model checkpoint that achieved the lowest validation loss, preventing the model from fitting too closely to the training data. Finally, the learning rate warms up linearly over the first 10% of training steps before decreasing steadily, which helps keep gradient updates stable during the initial phase of fine-tuning.

Recent empirical findings from the COLIEE 2026 legal textual entailment benchmarks (Task 4) demonstrated that neural language models frequently exhibit affirmative bias, defaulting toward entailment or neutral classifications when confronted with complex statutory syntax [5, 199]. In local legislative review, failing to detect an authentic conflict permits an invalid ordinance to proceed to enactment. To counter this affirmative bias, the proponents calibrate the decision threshold for the contradiction class using validation holdout curves. Rather than assuming an uncalibrated argmax decision rule, the threshold is adjusted to maximize the Contradiction F1-Score while maintaining high specificity. This calibration ensures that genuine preemption risks are surfaced to legislative staff rather than being suppressed by ambiguous phrasing.

3.8.3 Hyperparameter Tuning

To identify the optimal internal configuration for candidate Stage 2 Cross-Encoder models, the proponents implemented a systematic Grid Search procedure. Grid search is chosen over random or Bayesian optimization because the search space across lightweight transformer encoders is compact and discrete, allowing an exhaustive, reproducible evaluation of all candidate parameter combinations on the dedicated Mobile GPU Host (Lenovo Legion 5, NVIDIA GeForce RTX 5050 Laptop GPU).

The search isolates key variable parameters—specifically learning rate, batch size, and maximum sequence length—while holding structural regularization constants fixed across all training iterations. Table 3.11 specifies the complete partitioning of fixed experimental parameters versus variable search ranges.

Table 3.11: Hyperparameter Tuning Configuration and Grid Search Space for Stage 2 Cross-Encoders

Hyperparameter	Search Status	Configured Value / Search Bounds
Optimization Algorithm	Fixed	AdamW with decoupled weight decay ($\lambda = 0.01$) [116]
Attention & Hidden Dropout	Fixed	$p = 0.10$ across all transformer and classification layers [182]
Warmup Schedule	Fixed	Linear warmup over first 10% of total training steps
Early Stopping Patience	Fixed	patience = 4 validation epochs (Macro F_1 evaluation) [144]
Random Initialization Seed	Fixed	seed = 42 (enforcing deterministic mini-batch collation)
Maximum Training Epochs	Fixed	$E_{\max} = 10$ epochs (with best validation checkpoint restoration)
Loss Criterion	Fixed	Multi-Class Cross-Entropy Loss with Softmax normalization
Initial Learning Rate (η)	Variable (Grid Search)	$\eta \in \{1 \times 10^{-5}, 2 \times 10^{-5}, 3 \times 10^{-5}, 5 \times 10^{-5}\}$
Training Batch Size (B)	Variable (Grid Search)	$B \in \{8, 16\}$ (governed by 8 GB VRAM execution ceilings)
Max Input Sequence Length (L)	Variable (Grid Search)	$L \in \{256, 512\}$ (8,192 tokens for native Modern-BERT)

Table 3.11 specifies the complete hyperparameter search grid for Stage 2 Cross-Encoder fine-tuning, delineating fixed optimization controls from variable grid search bounds. This tuning configuration is established to achieve reproducible, stable convergence across compact transformer encoders on consumer GPU hardware while mitigating overfitting risks inherent in few-shot legal learning ($N_{\text{train}} = 245$). As detailed in the table, structural regularization constants are strictly fixed across all iterations, including AdamW optimization ($\lambda = 0.01$), attention dropout ($p = 0.10$), a 10% linear warmup schedule, early stopping (patience = 4), and a deterministic seed (seed = 42). Meanwhile, learning rate ($\eta \in \{1, 2, 3, 5\} \times 10^{-5}$), batch size ($B \in \{8, 16\}$), and sequence length ($L \in \{256, 512, 8192\}$) are systematically swept across the 52-pair validation set ($\mathcal{D}_{\text{val}}$). Quarantining hyperparameter selection strictly to the validation split prevents

data snooping on the test set, guaranteeing that the checkpoint achieving peak validation Macro F_1 generalizes robustly to unseen statutory pairs.

All hyperparameter configurations are evaluated strictly on the 52 validation pairs ($\mathcal{D}_{\text{val}}$) to prevent test-set data contamination. Model checkpoints are saved after every epoch, and the configuration achieving the highest validation Macro F_1 -score is selected for final out-of-sample evaluation on the 53-pair test set.

3.9 Evaluation Metrics

Evaluating the two-stage pipeline requires separate metrics for each operational phase. The coarse retrieval phase in Stage 1 prioritizes candidate recall over final classification, while the inference phase in Stage 2 requires high precision to avoid false conflict alerts.

3.9.1 Stage 1 Retrieval Metrics

Measuring search performance in Stage 1 requires ranking metrics rather than standard classification accuracy [122]. Because the purpose of the coarse retriever is to find candidate laws rather than make the final decision, standard accuracy does not reflect search quality. If the relevant national statute is left out of the top- k shortlist, the downstream Cross-Encoder cannot detect the conflict. For this reason, the retrieval stage is evaluated primarily using Recall at rank k (Recall@ k) at thresholds of $k \in \{10, 30, 50\}$. This metric measures the percentage of test queries where the correct national law appears within the candidate shortlist:

Recall at Rank k (Recall@ k):

$$\text{Recall@}k = \frac{1}{|Q|} \sum_{q=1}^{|Q|} \mathbf{1}(\text{rank}(d_q^*) \leq k) \quad (3.10)$$

Where $|Q|$ is the total number of test queries, d_q^* is the correct national statute paired with query q , and $\mathbf{1}(\cdot)$ is an indicator function that equals one when the target law ranks within the top k candidates.

Mean Reciprocal Rank at Rank k (MRR@ k):

$$\text{MRR@}k = \frac{1}{|Q|} \sum_{q=1}^{|Q|} \frac{1}{\text{rank}(d_q^*)} \cdot \mathbf{1}(\text{rank}(d_q^*) \leq k) \quad (3.11)$$

Where d_q^* is the correct national statute retrieved for query q , with $1/\text{rank}(d_q^*)$ representing the reciprocal rank score.

Equations 3.10 and 3.11 define the primary ranking metrics used to benchmark candidate Stage 1 retrieval models across the 25,432-statute search space, measuring candidate capture probability and reciprocal ranking efficiency, respectively. These ranking metrics are established because coarse retrieval is tasked with surfacing candidate provisions into a compact top- k shortlist ($k \in \{10, 30, 50\}$) rather than assigning final discrete

labels. As formulated, Recall@ k calculates the proportion of test queries where the paired governing statute successfully ranks within the top- k cutoff, while MRR@ k penalizes target statutes retrieved at lower ordinal ranks by computing the mean of their inverse rank positions. In downstream operations, high Recall@50 guarantees that Stage 2 Cross-Encoder inference receives the necessary statutory context to detect preemption, while higher MRR reduces the cognitive burden on human researchers inspecting the candidate list. Consequently, these metrics provide the definitive empirical criteria for selecting the optimal sparse and dense retriever in Chapter 4.

In this formula, a target statute retrieved at rank one receives a reciprocal score of 1.0, while a statute retrieved at rank ten receives 0.10. The 350 curated premise-hypothesis pairs serve as the query benchmark for testing Stage 1. By using the 350 draft clauses as queries against the 25,432-statute collection, the proponents measure how often candidate retrieval models locate the paired statutory provision. Calculating Recall@10, Recall@30, and Recall@50 confirms whether the coarse retriever supplies a dependable candidate pool for subsequent natural language inference.

3.9.2 Stage 2 Inference Metrics

The Stage 2 Cross-Encoder evaluates candidate statutory pairs across three discrete logical categories: $\mathcal{C} = \{\text{Entailment, Neutral, Contradiction}\}$. To assess classification quality across each relationship, the baseline inference performance metrics are computed as follows:

Class-Specific Precision, Recall, and F_1 -Score:

$$P_c = \frac{TP_c}{TP_c + FP_c}, \quad R_c = \frac{TP_c}{TP_c + FN_c}, \quad F_{1,c} = \frac{2 \cdot P_c \cdot R_c}{P_c + R_c} \quad (3.12)$$

Where TP_c , FP_c , and FN_c represent the true positive, false positive, and false negative classifications for category c , respectively. Precision reflects how reliably the model avoids false alarms, while Recall measures the proportion of genuine legal relationships successfully detected.

Macro-Averaged F_1 -Score Across Logical Classes:

$$\text{Macro } F_1 = \frac{1}{|\mathcal{C}|} \sum_{c \in \mathcal{C}} F_{1,c} \quad (3.13)$$

Where $|\mathcal{C}| = 3$ ensures that performance across all three categories is evaluated equally without allowing dominant neutral pairs to skew the outcome.

Equations 3.12 and 3.13 formulate the classification performance metrics for Stage 2 Cross-Encoder inference across the three discrete NLI categories ($\mathcal{C} = \{\text{Entailment, Neutral, Contradiction}\}$). These standard metrics are computed to evaluate model discrimination across individual legal relationships and assess overall multi-class balance. As defined, class precision P_c measures false-alarm resistance, class recall R_c measures true conflict coverage, and the class harmonic mean $F_{1,c}$ balances these competing objectives. Computing Macro F_1 across all three categories assigns equal weight (1/3) to Entailment, Neutral, and

Contradiction regardless of minor sample size fluctuations. This unweighted average ensures that models achieving high accuracy simply by defaulting toward the majority neutral category are penalized, providing an objective benchmark for model screening.

In statutory conflict detection, classification errors carry fundamentally unequal operational costs [55, 57]. A false negative, where the system fails to flag an unconstitutional local clause, permits an invalid ordinance to be enacted. This failure triggers costly litigation and statutory invalidation under the *Magtajas* doctrine [187]. In contrast, a false positive, where the system flags a compliant draft clause, merely directs a human legislative researcher to examine the flagged text on the review dashboard, requiring only seconds of reviewer time. Standard classifiers default to a symmetric decision threshold of $P(\text{Contradiction}) \geq 0.50$, which assumes that false alarms and missed conflicts cause identical harm. In foundational cost-sensitive learning theory, Elkan [55] proved that when false negative costs exceed false positive costs ($C(FN) > C(FP)$), setting the decision boundary at 0.50 is mathematically suboptimal and produces an unsafe number of missed conflicts [57, 100].

To address this cost asymmetry, the proponents implement a validation-tuned decision boundary τ_{conflict} for contradiction alerts. Rather than using a fixed 0.50 cutoff, the pipeline sweeps candidate thresholds across $\tau \in [0.25, 0.50]$ in increments of $\Delta\tau = 0.01$ exclusively on the validation set ($\mathcal{D}_{\text{val}}$, $N = 52$). The cost-sensitive contradiction metrics are computed as follows:

Cost-Sensitive Contradiction F_2 -Score:

$$F_{2,\text{conflict}} = \frac{5 \cdot P_{\text{conflict}} \cdot R_{\text{conflict}}}{4 \cdot P_{\text{conflict}} + R_{\text{conflict}}} \quad (3.14)$$

Where P_{conflict} and R_{conflict} denote Precision and Recall on the Contradiction class, weighting recall twice as heavily as precision ($\beta = 2$) to aggressively capture potential legal conflicts.

Optimal Contradiction Decision Threshold Calibration (τ^):*

$$\tau^* = \arg \max_{\tau \in [0.25, 0.50]} F_{2,\text{conflict}}(\tau; \mathcal{D}_{\text{val}}) \quad (3.15)$$

Where $\mathcal{D}_{\text{val}}$ denotes the 52 validation pairs used to sweep candidate decision thresholds $\tau \in [0.25, 0.50]$ while enforcing an empirical precision floor of $P_{\text{conflict}} \geq 0.70$.

Equations 3.14 and 3.15 establish the cost-sensitive evaluation metric and optimal decision threshold calibration formulation governing contradiction alerts. This formulation is implemented to address the severe operational asymmetry between false negatives (enacting invalid ordinances that provoke constitutional litigation) and false positives (brief advisory flags reviewed in seconds by legal staff). In Equation 3.14, setting $\beta = 2$ weights contradiction recall twice as heavily as precision, penalizing missed preemption risks. Equation 3.15 sweeps the decision boundary $\tau \in [0.25, 0.50]$ exclusively on the 52-pair validation set ($\mathcal{D}_{\text{val}}$) subject to a precision floor ($P \geq 0.70$). Calibrating the operational threshold to τ^* ensures that the deployed system aggressively surfaces potential statutory violations while maintaining adequate specificity against alert fatigue.

During live inference, a candidate provision triggers a formal conflict alert whenever its softmax contradiction probability satisfies $P(\text{Contradiction} \mid p_i, h) \geq \tau^*$.

When evaluating complete multi-section ordinances during live system execution, Stage 1 retains its full safety buffer of $k = 50$ candidates for each substantive section (S_i). To prevent runtime delays when scoring up to 500 total candidate pairs per document, the Stage 2 inference engine employs vectorized mini-batch collation ($B = 16$). Batching evaluates candidate provisions in parallel tensor forward passes rather than sequential single-item loops, keeping inference runtime under five seconds on GPU hosts and under thirty seconds on standard LGU processors. To consolidate the resulting candidate scores into a single decision for each section, the aggregation metric is computed as follows:

Section-Level Max-Pooled Conflict Probability:

$$P(\text{Conflict} \mid S_i) = \max_{j \in [1, k]} P(\text{Contradiction} \mid P_j, S_i) \quad (3.16)$$

Where P_j ranges over all $k = 50$ candidate statutory provisions retrieved for draft section S_i .

Equation 3.16 formulates the section-level conflict aggregation mechanism, defining the overall preemption probability of an operative draft section S_i as the maximum contradiction probability observed across all $k = 50$ candidate statutory provisions. This max-pooling aggregation is incorporated to resolve the multi-provision matching problem during full-document review, where an operative draft clause is evaluated against an extensive shortlist of retrieved national laws. By assigning the section’s final score to the peak contradiction probability, the formulation guarantees that a genuine legal conflict with a specific governing statute triggers an alert even if the remaining 49 candidate laws are neutral or harmonious. Consequently, legislative reviewers receive an unambiguous warning pointing directly to the single most conflicting statutory mandate, preventing critical preemption risks from being obscured by surrounding non-conflicting provisions.

A section triggers an operative conflict warning if and only if $P(\text{Conflict} \mid S_i) \geq \tau^*$, with the system highlighting the specific governing candidate provision $P_j^* = \arg \max_j P(\text{Contradiction} \mid P_j, S_i)$ that produced the highest conflict probability. This max-pooling aggregation ensures that an authentic statutory clash ranked lower in the candidate shortlist is not obscured by neutral or harmonious provisions appearing at higher retrieval ranks. Listing 3.9 outlines the operational inference and conflict aggregation procedure.

```

1 def detect_section_conflicts(draft_section, candidate_provisions, model, threshold
  =0.35, batch_size=16):
2     """
3     Stage 2 Cost-Sensitive Cross-Encoder Inference:
4     Scores candidate law pairs and flags conflicts using max-pooling.
5     """
6     # 1. Pair each candidate statutory provision (premise) with the draft section (
  hypothesis)
7     pairs = [(provision.text, draft_section.text) for provision in
  candidate_provisions]
8
9     # 2. Run mini-batch inference through the fine-tuned Cross-Encoder
10    contradiction_scores = []
11    for i in range(0, len(pairs), batch_size):

```

```

12     batch = pairs[i : i + batch_size]
13     # Get softmax probabilities across [Entailment, Neutral, Contradiction]
14     probabilities = model.predict_probabilities(batch)
15     # Extract contradiction class probability (index 2)
16     batch_contradictions = [p["contradiction"] for p in probabilities]
17     contradiction_scores.extend(batch_contradictions)
18
19     # 3. Find candidate provision with highest contradiction score (Max-Pooling)
20     peak_score = max(contradiction_scores)
21     peak_index = contradiction_scores.index(peak_score)
22     governing_provision = candidate_provisions[peak_index]
23
24     # 4. Apply cost-sensitive decision threshold
25     has_conflict = peak_score >= threshold
26
27     return {
28         "alert": has_conflict,
29         "conflict_probability": peak_score,
30         "conflicting_law": governing_provision.citation,
31         "conflicting_text": governing_provision.text
32     }
33

```

Listing 3.9: Pseudocode of the Stage 2 Cost-Sensitive Cross-Encoder Inference and Section-Level Conflict Aggregation

Listing 3.9 outlines the Python implementation of the Stage 2 Cross-Encoder inference and section-level conflict aggregation routine. This algorithm is provided to illustrate the mini-batch tensor processing and cost-sensitive thresholding that transforms candidate provision pairs into actionable preemption alerts. In the procedure, the candidate statutory provisions and draft section are paired into premise-hypothesis sequences and evaluated in vectorized mini-batches ($B = 16$) through the fine-tuned model. The function extracts softmax contradiction probabilities, identifies the peak score via max-pooling according to Equation 3.16, and checks whether this peak crosses the calibrated decision threshold τ^* . If an alert is triggered, the routine returns the conflicting statutory citation, text, and model confidence, delivering immediate diagnostic transparency to the First Reading review dashboard.

3.9.3 Statistical Hypothesis Testing

To evaluate whether performance disparities between candidate Cross-Encoder models represent statistically significant differences rather than random sampling fluctuations, the proponents employ formal non-parametric hypothesis tests recommended by Demšar [43]. Traditional paired t -tests assume normally distributed classification errors, an assumption routinely violated by discrete NLP accuracy and F1 scores on small test sets [43]. Therefore, this study adopts McNemar’s Test for pairwise model comparisons and the Friedman Test for multi-model ranking across legal domains.

When comparing the classification accuracy of two competing Stage 2 models (such as fine-tuned DeBERTa-v3 versus fine-tuned ModernBERT) evaluated on the identical 53-item test set ($\mathcal{D}_{\text{test}}$), the proponents employ McNemar’s test for pairwise significance [126]. The test constructs a 2×2 contingency table analyzing discordant predictions, as summarized in Table 3.12. In this contingency framework, cell b represents items correctly classified by Model A but misclassified by Model B, while cell c represents items correctly

classified by Model B but misclassified by Model A. Under the null hypothesis $H_0 : p_b = p_c$, the two models have equivalent error rates.

Table 3.12: Contingency Matrix of Discordant Predictions for McNemar’s Test

	Model B: Correct	Model B: Incorrect
Model A: Correct	a (Both Correct)	b (Model A Only)
Model A: Incorrect	c (Model B Only)	d (Both Incorrect)

Table 3.12 presents the 2×2 contingency matrix of discordant and concordant predictions utilized by McNemar’s test for pairwise model significance testing. This matrix structure is incorporated to isolate the exact instances where two competing Cross-Encoder architectures disagree on the 53-item test split ($\mathcal{D}_{\text{test}}$). In the layout, cells a and d capture concordant outcomes where both models are simultaneously correct or incorrect, while off-diagonal cells b and c record discordant predictions where only Model A or Model B classifies the pair correctly. Under the null hypothesis, the two models share equivalent error rates ($p_b = p_c$). By evaluating only discordant pairs rather than raw overall accuracy, the test provides high statistical sensitivity to marginal algorithmic improvements while remaining robust against non-normal classification error distributions.

Because the Stage 2 evaluation compares three candidate models ($k = 3$: MiniLM, DeBERTa-v3, ModernBERT), pairwise evaluation yields $m = \binom{3}{2} = 3$ distinct model comparisons. To prevent Type I error inflation (false positive significance) across multiple tests, the proponents apply the standard Bonferroni correction [43], evaluating each pair against the adjusted threshold $\alpha' = \alpha/m = 0.05/3 \approx 0.0167$.

Furthermore, when expanding candidate model explorations to larger suites (e.g., $k \approx 10$ architectures and hyperparameter configurations), running exhaustive pairwise McNemar tests produces $\binom{k}{2} = 45$ comparisons, creating severe statistical conservatism under Bonferroni correction ($\alpha' = 0.05/45 \approx 0.0011$). To handle larger candidate pools without losing statistical power, the proponents follow Demšar’s two-step testing procedure [43]: first executing the omnibus Friedman test across all k models simultaneously, and only upon rejecting the global null hypothesis ($p < 0.05$), conducting post-hoc comparisons comparing candidate models strictly against the top-performing control model ($k - 1$ comparisons rather than $\frac{k(k-1)}{2}$) or evaluating critical difference intervals via the Nemenyi test.

For pairwise significance, the proponents formulate McNemar’s chi-squared test with Edwards’ continuity correction. The test evaluates whether discordant classifications between Model A and Model B occur in equal proportions:

McNemar’s Chi-Squared Test with Edwards’ Continuity Correction:

$$\chi^2 = \frac{(|b - c| - 1)^2}{b + c} \sim \chi_1^2 \quad (3.17)$$

Where b represents items correctly classified by Model A but misclassified by Model B, c represents items correctly classified by Model B but misclassified by Model A, and the numerator incorporates Edwards’ continuity correction (-1) for discrete cell frequencies.

Equation 3.17 defines McNemar’s chi-squared test statistic incorporating Edwards’ continuity correction for paired nominal classifications. This formulation is included to determine whether performance discrepancies between competing Stage 2 Cross-Encoder models on the 53-item test split are statistically significant rather than artifacts of stochastic variance. The equation measures the absolute difference between discordant cells b and c , penalizing discrete cell counts with a -1 continuity correction and scaling against total discordant pairs to yield a test statistic distributed asymptotically as χ^2 with one degree of freedom. In practice, calculating this statistic isolates whether one model reliably succeeds where another fails on difficult preemption boundary cases. A calculated p -value below the Bonferroni-adjusted significance threshold ($\alpha' = 0.0167$) rejects the null hypothesis, formally establishing that observed accuracy gains represent genuine architectural superiority.

Beyond pairwise significance, assessing the consistency of candidate architectures across the eight Macro Legal Domains without making restrictive distributional assumptions requires non-parametric multi-model ranking. For this purpose, the proponents utilize the Friedman non-parametric two-way analysis of variance by ranks [60, 43]. For k candidate models evaluated across $N = 8$ distinct legal domains, each model is assigned a rank r_i^j within domain i , where the best-performing model receives rank 1. The average rank of model j across all domains is given by $R_j = \frac{1}{N} \sum_{i=1}^N r_i^j$. The omnibus rank statistic is computed as follows:

Friedman Two-Way Analysis of Variance by Ranks (χ_F^2):

$$\chi_F^2 = \frac{12N}{k(k+1)} \left[\sum_{j=1}^k R_j^2 - \frac{k(k+1)^2}{4} \right] \quad (3.18)$$

Where $N = 8$ is the number of macro legal domains, k is the number of candidate models, and $R_j = \frac{1}{N} \sum_{i=1}^N r_i^j$ is the average rank of model j across all domains (r_i^j being the rank of model j in domain i).

Equation 3.18 computes the Friedman chi-squared rank test statistic (χ_F^2) for non-parametric multi-model evaluation across partitioned legal domains. This statistical formulation is integrated to assess whether candidate Cross-Encoder architectures exhibit consistent relative performance across diverse statutory areas without imposing unrealistic normality assumptions. The formula scales the sum of squared average domain ranks ($\sum R_j^2$) against the number of domains ($N = 8$) and evaluated candidate models (k), contrasting the observed distribution against expected uniform ranking variance. Under the global null hypothesis that all candidate models perform equivalently across the eight legal domains, χ_F^2 follows a chi-squared distribution with $k - 1$ degrees of freedom. Rejecting this null hypothesis ($p < 0.05$) confirms significant cross-domain performance divergence, licensing post-hoc pairwise comparisons via the Wilcoxon signed-rank test and Nemenyi critical difference intervals to isolate models that achieve statistically robust domain generalization [43].

3.10 System-Level Evaluation

While isolated benchmarks measure retrieval recall and natural language inference accuracy separately, evaluating the live utility of the system requires testing both stages as an integrated pipeline. In real legislative drafting, an inference engine cannot detect a contradiction if the retrieval engine fails to supply the governing statute, and a retrieval engine is useless if the inference engine misclassifies the pair. Following the multi-stage system evaluation standards established in legal information extraction competitions [146, 5] and fact verification benchmarks [197], the proponents implement a structured Three-Tier Evaluation Framework. This multi-tier framework connects controlled statistical benchmarking with full-document stress testing. Together, these three evaluation tiers systematically assess whether the complete pipeline can operate reliably within the Davao City Sangguniang Panlungsod.

The first evaluation tier measures the combined algorithmic performance of Stage 1 retrieval and Stage 2 inference when executed sequentially without human intervention. Unlike isolated model testing where the Cross-Encoder receives the exact gold statutory provision, the cascaded benchmark feeds only the local draft hypothesis clause h_q from the 53 held-out test pairs into Stage 1 coarse retrieval. The retrieval engine searches across all 27,000+ national statutes and local ordinances to extract a candidate shortlist of the top- k provisions ($\mathcal{K}_q^{(k)}$), which is then passed directly to Stage 2 for inference. A test item is scored as a Cascaded True Positive if and only if Stage 1 retrieves the correct governing statute p_q^* within the top- k shortlist and Stage 2 correctly classifies the pair (p_q^*, h_q) as a Contradiction:

Cascaded End-to-End Joint Conflict Detection Hit Indicator:

$$\text{Hit}_{\text{casc}}(q) = \mathbf{1}(p_q^* \in \mathcal{K}_q^{(k)}) \cdot \mathbf{1}(\hat{y}(p_q^*, h_q) = y_q^*) \quad (3.19)$$

Where p_q^* is the target governing statute, $\mathcal{K}_q^{(k)}$ is the candidate shortlist of top- k provisions retrieved in Stage 1, and $\hat{y}(p_q^*, h_q) = y_q^*$ indicates whether Stage 2 correctly classifies the candidate pair as a contradiction.

Equation 3.19 formalizes the joint hit indicator metric (Hit_{casc}) for unassisted end-to-end pipeline evaluation across the 53-item held-out test split. This compound metric is formulated to ensure that end-to-end conflict detection success is credited only when both processing stages execute correctly in uninterrupted sequence. The indicator function multiplies two binary conditions: whether the Stage 1 retrieval module successfully locates the true governing provision p_q^* within the candidate shortlist $\mathcal{K}_q^{(k)}$, and whether the Stage 2 Cross-Encoder subsequently infers the exact gold label y_q^* . By coupling retrieval recall with inference accuracy, the metric isolates pipeline vulnerabilities and prevents isolated high retrieval recall or high Cross-Encoder accuracy from obscuring catastrophic cascaded errors. In practical system deployment, this joint metric mirrors authentic First Reading committee operations, where surfacing a valid preemption alert requires both locating the applicable statutory restriction and correctly recognizing the normative collision.

Adapting this joint metric from the FEVER fact verification framework [197] isolates whether an operational error stems from retrieval omissions ($p_q^* \notin \mathcal{K}_q^{(k)}$) or classifica-

tion errors ($\hat{y} \neq y_q^*$), preventing either stage from masking pipeline failures.

The second evaluation tier assesses how the complete coarse-to-fine pipeline processes full, multi-section legislative drafts rather than isolated sentence pairs. In municipal practice, proposed ordinances contain extensive procedural text, including administrative definitions and standard effectivity clauses, which should not trigger conflict warnings. Adapting the behavioral perturbation and fault-injection principles formulated by Ribeiro et al. [155] and Koreeda and Manning [100], the proponents construct synthetic and semi-synthetic ordinance drafts containing controlled legal faults. In each test document, the majority of sections remain compliant with superior law, while targeted operative provisions contain deliberate statutory violations, such as unauthorized regulatory mandates or excessive penalty amounts. The automated preprocessor chunks the submitted document into operative sections, auditing every section independently against the statutory archive.

Evaluating full documents measures both the sensitivity of the pipeline to authentic legal defects and its ability to avoid false alarms on compliant text. For every audited section, the system records whether a conflict alert is triggered and identifies the retrieved governing statute. The proponents calculate Section-Level Sensitivity to measure the percentage of injected conflicts successfully flagged, paired with Section-Level Specificity to assess false-alarm resistance:

Section-Level Specificity for Full-Document Stress Testing:

$$\text{Specificity}_{\text{sec}} = \frac{TN_{\text{sec}}}{TN_{\text{sec}} + FP_{\text{sec}}} \quad (3.20)$$

Where TN_{sec} denotes compliant draft sections correctly passing without triggering a false alarm, and FP_{sec} denotes compliant draft sections erroneously flagged as contradictory.

Equation 3.20 defines Section-Level Specificity ($\text{Specificity}_{\text{sec}}$) for the Tier 2 full-document behavioral perturbation and fault-injection benchmark. This metric is incorporated to evaluate the pipeline’s operational resistance against false positive alarms when auditing lengthy, multi-section legislative drafts containing standard administrative boilerplate. The equation calculates the proportion of true negative sections (TN_{sec})—provisions that are legally compliant and correctly pass without generating conflict alerts—relative to the total volume of non-conflicting sections. In practical municipal workflows, an automated review system that triggers frequent false warnings on routine effectivity clauses or definitions induces cognitive alert fatigue, leading legislative researchers to distrust algorithmic guidance. High section-level specificity guarantees that conflict notifications remain credible, ensuring that Sangguniang Panlungsod staff can focus institutional scrutiny strictly on substantive legal preemption risks.

Maintaining high specificity protects legislative researchers from alert fatigue, ensuring that the system does not produce spurious warnings on standard boilerplate clauses. This perturbation stress test confirms that the system can analyze complete legislative documents before deployment in committee reviews.

The third evaluation tier validates the practical utility and legal soundness of the system using documented historical legal disputes and active expert review. Rather than

relying solely on synthetic tests, the proponents systematically assembled an authentic judicial validation suite of eight landmark legal controversies and municipal enactments originating strictly within Davao City. These historical cases encompass both judicial invalidations where local enactments were struck down for exceeding statutory boundaries, and landmark rulings where the Supreme Court affirmed the validity of Davao City ordinances under local police power. By subjecting the pipeline to actual local disputes that reached the judiciary, this tier evaluates whether algorithmic contradiction detection mirrors authentic judicial reasoning. The resulting test suite spans multiple substantive categories, including agrochemical oversight, structural zoning, revenue powers, and public health mandates.

The judicial cases examine distinct facets of municipal governance, beginning with environmental regulation and structural zoning. In *Mosqueda v. Pilipino Banana Growers & Exporters Association, Inc.* [190], the Supreme Court invalidated Davao City Ordinance No. 0309-07 (imposing a total ban on aerial pesticide spraying in agricultural banana plantations) for violating equal protection and invading the exclusive regulatory jurisdiction vested in the Fertilizer and Pesticide Authority under Presidential Decree No. 1144. Conversely, in the Davao City Watershed Protection dispute under Ordinance No. 0310-07 [164], the city’s total prohibition on mineral extraction operates in direct vertical tension with State mineral concessions under Republic Act No. 7942. For urban structural governance, the Supreme Court in *Hon. Leoncio Evasco, Jr. v. Alex P. Montañez* [192] reversed the Court of Appeals and explicitly upheld Davao City Ordinance No. 092-2000 regulating outdoor billboards. The Court affirmed that under Section 458(a)(4)(iv) of Republic Act No. 7160, city councils possess express statutory police power to regulate advertising structures harmoniously alongside the National Building Code (Presidential Decree No. 1096).

The remaining landmark disputes focus on municipal revenue authority, public health, and delegated traffic administration. In *City of Davao and Tanjili v. ARC Investors, Inc.* [193], the Supreme Court voided Davao City’s local business tax assessment under Ordinance No. 158-05, ruling that an LGU cannot unilaterally reclassify a passive holding company as a non-bank financial intermediary without Bangko Sentral ng Pilipinas authorization. Similarly, in *City of Davao v. First Meridian Development, Inc.* [194], the Court invalidated Section 423 of Ordinance No. 158-05 for imposing an unauthorized payment-under-protest requirement contrary to Section 195 of Republic Act No. 7160. By contrast, the Court in *City of Davao v. Court of Appeals and GSIS* [188] affirmed municipal taxing authority over government corporate assets after Republic Act No. 7160 withdrew prior statutory exemptions. In delegated public health and traffic safety, Davao City enacted landmark measures setting stricter local standards, such as the Comprehensive Anti-Smoking Ordinance (Ordinance No. 0367-12) [165] and the Comprehensive Speed Limit Ordinance (Ordinance No. 0270-23) [167] pursuant to Republic Act No. 4136 and Joint JAO 2018-01.

Table 3.13 presents the complete Davao City municipal validation suite, detailing the local measures, governing national statutes, legal ratios decidendi, and gold NLI labels. To ensure long-term reproducibility and immediate system availability, complete archival copies of these eight Davao City ordinances and tax mea-

tures have been assembled and stored locally within the repository’s municipal corpus (`corpus/davao_city_tier3_ordinances/`). Standardized diagnostic problem-solution pairs are indexed in `data/tier3_jurisprudential_cases.jsonl`, pairing operative draft excerpts with corresponding statutory provisions and judicial holdings. By feeding authentic legal questions and challenged provisions directly into Stage 1 retrieval and Stage 2 inference, the evaluation confirms whether the automated pipeline surfaces the exact statutory citations and preemption rationales formulated by judicial authorities. This authentic benchmark grounds the quantitative metrics in real jurisprudence, ensuring that theoretical inference gains translate into valid legal oversight.

Table 3.13: Landmark Davao City Municipal Jurisprudential Suite for Tier 3 System Evaluation ($N = 8$ Local Ordinances and Preemption Disputes)

Case Title	Docket / Year	Challenged Local Ordinance	Governing Statute & Rationale	Gold NLI Label
<i>Mosqueda v. PBGEA</i> [190]	G.R. 189185 (2016)	Davao City Ord. 0309-07 (Banning Aerial Pesticide Spraying)	PD 1144; FPA holds exclusive jurisdiction over agrochemical application.	Contradiction
<i>Evasco v. Montañez</i> [192]	G.R. 199172 (2018)	Davao City Ord. 092-2000 (Regulating Billboard Structures)	RA 7160 §458; Upheld as valid municipal police power alongside PD 1096.	Entailment
<i>Davao v. ARC Investors</i> [193]	G.R. 249668 (2022)	Davao City Ord. 158-05 §135 (Business Tax on Holding Dividends)	RA 7160 §143(f); Holding firm is not an NBFI without BSP license.	Contradiction
<i>City of Davao v. First Meridian</i> [194]	G.R. 240078 (2022)	Davao City Ord. 158-05 §423 (Mandatory Payment Under Protest)	RA 7160 §195; LGU cannot require payment under protest for business taxes.	Contradiction
<i>City of Davao v. GSIS</i> [188]	G.R. 127383 (2005)	Real Property Tax Levies on GSIS Government Properties	RA 7160 §193, §234; LGC withdrew tax exemptions of GOCCs; assessment valid.	Entailment
<i>Davao Mining Ban Dispute</i> [164]	Ord. 0310-07 (2015)	Citywide Total Ban on Commercial Mining Operations	RA 7942 §4, §27; Regalian doctrine and national mineral concession supremacy.	Contradiction
<i>Davao Anti-Smoking Dispute</i> [165]	Ord. 0367-12 (2012)	Stricter Outdoor/Indoor Smoking Distance Setbacks	RA 7160 §16; Stricter local public health standards valid under police power.	Entailment
<i>Davao Speed Limit Dispute</i> [167]	Ord. 0270-23 (2023)	Comprehensive Speed Limits on City Thoroughfares	RA 4136 §35; Joint JAO 2018-01 expressly delegates road speed setting to LGUs.	Entailment

Table 3.13 details the eight landmark Davao City municipal controversies and judi-

cial rulings comprising the Tier 3 authentic jurisprudential validation suite ($N = 8$). This empirical suite is assembled to evaluate whether the automated coarse-to-fine pipeline correctly reproduces authentic judicial reasoning when auditing contested local enactments against superior national statutes. The table catalogues the judicial case citation, docket number, challenged Davao City ordinance, governing national legislative framework, and authoritative gold NLI relationship established by the Supreme Court of the Philippines. Across these cases, the benchmark captures both vertical preemption invalidations (such as pesticide regulation under PD 1144 and municipal tax overreach under RA 7160) and valid exercises of devolved police power (including billboard regulation under RA 7160 and localized anti-smoking mandates). By grounding algorithmic testing in documented historical jurisprudence, this evaluation verifies that the Cross-Encoder’s calibrated confidence scores correspond to substantive constitutional doctrine under the *Magtajas* principle.

Finally, to measure practical usability in active legislative workflows, the system is deployed to the 15 active Sangguniang Panlungsod legal researchers. Evaluators audit sample draft ordinances using the single-page prototype dashboard during First Reading committee preparation. The researchers assess whether the retrieved statutory context and highlighted attention heatmaps clarify potential preemption risks under Section 458 of Republic Act No. 7160 [70]. Inter-annotator consensus is measured using Fleiss’ Kappa (κ) [58], while subjective software usability is evaluated using the standardized System Usability Scale (SUS) [20]. Combining historical court rulings with direct feedback from city council staff confirms that the finished pipeline provides practical decision support for local governance.

3.11 Prototype and User Interface

The completed pipeline is deployed as a single-page web application for end-user testing. The user interface is kept intentionally simple, similar to familiar online document conversion tools. Users begin by uploading draft ordinances in standard document formats through a drag-and-drop file area. After submission, the interface shows an active progress indicator while the two-stage search and inference pipeline checks the text against the database. This straightforward layout avoids unnecessary menus, allowing legislative staff to run conflict checks without requiring technical training.

When the analysis finishes, the interface displays a structured results dashboard. The dashboard provides an immediate diagnostic summary reporting the total count of identified inconsistencies across the draft document. To maintain transparency in Stage 1 retrieval, a dedicated module lists the top- k statutory provisions retrieved during coarse search. When a conflict is detected, the dashboard provides a side-by-side view comparing the specific draft section directly against the conflicting national or municipal provision. Each flagged item displays the calibrated model confidence percentage ($P(\text{Contradiction})$), distinguishing between Critical Vertical Preemption Alerts and Informational Horizontal Modification Notices.

The prototype operates as a human-in-the-loop decision-support tool during First Reading committee reviews. Instead of replacing human legal judgment, the interface

highlights potential conflicts to help legal staff review drafts faster. To provide actionable explainability without complex black-box approximations, the dashboard displays intrinsic structural explainability by highlighting overlapping legal entities and contrasting conflicting command words (such as *prohibited* versus *allowed*). Sangguniang Panlungsod legal researchers can examine these highlighted clauses to decide whether the draft represents an invalid preemption or an allowed local adaptation. This workflow allows legislative staff to revise draft clauses before final enactment, helping prevent invalid local legislation.

3.12 Reproducibility and Repository

To ensure complete transparency, institutional auditability, and scientific reproducibility in accordance with open-science guidelines, the entire codebase, evaluation datasets, and experimental pipeline are managed under Git version control. The official repository is publicly accessible on GitHub:

`https://github.com/yyaahhzxc/thesis-repo`

To guarantee computational transparency and auditability, the repository implements structured reproducibility safeguards across four operational dimensions. First, all empirical tables, benchmark figures, and model checkpoints reported in this thesis are permanently tagged with their respective Git commit hashes (such as Commit `cc1c33d` for statutory corpus ingestion and Commit `4dd7c85` for legal instruments), ensuring complete version traceability. Second, all data splits, topic modeling initializations, and Cross-Encoder training routines enforce deterministic execution by fixing global random seeds (seed = 42) across PyTorch, NumPy, and Python standard runtime environments. Third, researchers can verify the complete pipeline using a single-command minimal working example:

`python scripts/prototype_pipeline.py --demo`

This executable demonstration loads the hybrid BM25 and semantic vector index over statutory provisions, executes asymmetric NLI scoring across four curated municipal test cases, and displays explainable ratio decidendi under the *Magtajas* doctrine in under two seconds on standard CPU hardware. Finally, all experimental parameters, including learning rates, batch sizes, sequence lengths, and decision thresholds (τ^*), are strictly decoupled from core execution code into plain-text configuration files. This config-driven architecture allows researchers to toggle candidate models and recalibrate operational thresholds without altering the underlying deep learning execution scripts.

3.13 Methodological Constraints

The first methodological constraint concerns the linguistic and lexical scope of the legislative corpora. In adherence to Executive Order No. 292 (Administrative Code of 1987) and national publishing standards, the statutory archive and municipal ordinance records examined in this study are strictly restricted to official English-language enactments. Conse-

quently, the automated pipeline does not ingest or interpret vernacular Cebuano-Visayan phrases, colloquial regional idioms, or mixed Taglish expressions that occasionally arise during informal committee deliberations or floor debates. While operative legal provisions in the Republic of the Philippines are formally promulgated in English, statutory preemption inquiries that hinge upon localized dialectal nuances require manual linguistic reconciliation by legal researchers. This constraint establishes a defined boundary for the lexical tokenizer and prevents semantic drift when evaluating formal statutory mandates against informal draft annotations.

The second constraint involves the temporal distribution and physical preservation state of the municipal ordinance archive. The historical enactment records obtained from the Davao City Sangguniang Panlungsod display an uneven chronological density, wherein ordinances promulgated between 2021 and 2025 benefit from clean digital-native records within the Legislative Information Support System Portal (LISSP). In contrast, municipal enactments dating between 1945 and 2020 exist primarily as legacy physical carbon-copy scans, marked by variable ink fading, binding skew, and multi-generation photocopier artifacts. Although the two-tier vision-language ingestion pipeline utilizing ColPali and Qwen2-VL significantly reduces transcription error relative to traditional optical character recognition, residual character error rates on pre-1980 archival enactments inevitably limit extraction fidelity. The retrospective evaluation of local ordinances prior to the modern digital era is therefore confined to landmark enactments that have undergone manual verification and high-resolution scanning.

The third constraint stems from the single-premise asymmetric modeling formulation adopted for the coarse-to-fine pipeline. The natural language inference engine models vertical conflict detection as an asymmetric pair consisting of an individual candidate statutory premise (p_i) matched against an operative draft ordinance block (h_q). This structural formulation effectively captures direct binary contradictions where a single national statute explicitly invalidates an ordinance clause, such as unauthorized penalty escalation under Republic Act No. 7160. However, the system does not inherently detect multi-statute cumulative preemptions, wherein an ordinance provision complies with each national statute in isolation but breaches a combined legislative framework formed by the intersection of multiple distinct codes. Expanding the inference mechanism to support multi-premise hypergraph structures or iterative cross-statutory synthesis represents a computational trade-off deferred to subsequent system iterations.

The final constraint defines the institutional and operational role of the automated conflict detection system within municipal governance. The pipeline is architected exclusively as an ex-ante assistive advisory tool intended for committee legal researchers during First Reading preparation, rather than an autonomous ex-post adjudicative authority. The system generates calibrated confidence estimates and highlighted textual rationales to surface potential legal friction, but it does not execute binding statutory determinations or substitute for judicial review under the Supreme Court of the Philippines. Furthermore, political questions, discretionary policy wisdom, and socio-economic trade-offs inherent in local ordinance sponsorship remain exclusively within the constitutional purview of elected city council members. Defining this operational boundary ensures that the artificial intelligence system functions strictly within established principles of administrative

law, safeguarding human-in-the-loop accountability throughout the municipal legislative process.

Chapter 4

Results and Discussion

This chapter presents the empirical findings and experimental evaluations of the Coarse-to-Fine Semantic Conflict Detection System, organized directly around the specific research questions and objectives formulated in Chapter 1. Section 4.1 addresses Research Question 1 and Objective 1 by analyzing the linguistic parameters, layout degradation, and vision-language transcription performance across the digitized Davao City municipal archive ($N = 1,664$). Section 4.2 addresses Research Question 2 and Objective 2 by evaluating candidate retrieval efficiency, preemption coverage, and latency across the unified dual statutory corpus ($N = 27,096$ enactments, comprising 176,421 provisions). Section 4.3 addresses Research Questions 3 and 4 and Objectives 3 and 4 by presenting the comparative model screening, supervised fine-tuning, loss dynamics, and ablation benchmarks on the curated statutory ground truth dataset. Finally, Section 4.4 addresses Research Question 5 and Objective 5 by outlining system-level decision-support viability, the prototype interface, and the multi-tier human evaluation roadmap for the Davao City Sangguniang Panlungsod.

4.1 Archival OCR and VLM Evaluation

Ingesting historical municipal archives into an automated legal decision-support system requires converting scanned physical documents into structured digital text. In the Davao City collection, enactments from earlier legislative terms exist predominantly as scanned physical PDF records rather than searchable digital-born documents. These historical records exhibit real scanning noise, including paper skew, uneven ink contrast, staple shadows, and official municipal seal watermarks stamped over operative clauses. If degraded physical scans are transcribed with erroneous characters or broken layout geometry, downstream retrieval indices and inference models inherit corrupted tokens. This section establishes a multi-dimensional evaluation framework for legislative document transcription, presents a stratified archival reference sampling design, compares four candidate optical character recognition paradigms, and details the empirical justification for adopting a locally hosted vision-language model.

4.1.1 Architectural Paradigms and Candidate Systems

To determine the most viable document extraction method for historical municipal archives, the proponents evaluated four distinct architectural paradigms. Traditional optical character recognition relies on classical image binarization, connected component analysis, and heuristic line slicing, followed by character sequence prediction using a lightweight 1D CNN-LSTM architecture [177]. As exemplified by Tesseract v5.x, traditional engines evaluate individual characters as isolated visual glyphs without semantic sentence context. While computationally efficient on central processing units, traditional engines degrade rapidly when encountering skewed lines or physical paper bleed-through.

Document-layout OCR frameworks decouple spatial document segmentation from line text recognition [138]. Modern implementations such as Surya v0.6+ employ visual object detection backbones to identify bounding polygons for text columns, tables, and headers. A separate recognition model then transcribes text lines, after which reading order is reconstructed through topological sorting heuristics. This modular decoupling provides robust layout preservation on multi-column pages, but requires multi-stage neural pipelines that demand substantial graphics memory during sustained batch processing.

Proprietary cloud vision platforms deploy large-scale multimodal transformers hosted on distributed enterprise server clusters [69]. Commercial endpoints such as Google Cloud Document AI achieve high optical recognition accuracy across diverse scan qualities and complex table structures. However, commercial cloud services require active internet connectivity, incur ongoing per-page API fees, and necessitate external transmission of municipal documents. In an edge-deployed local government setting, recurring subscription costs and cloud data transfer contradict operational independence and statutory data sovereignty.

Local vision-language models unify visual patch encoding and autoregressive text generation within an end-to-end transformer architecture [195]. Using two-dimensional rotary position embeddings (2D-RoPE) and dynamic patch allocation, visual features are projected directly into a causal language model decoder. Under 4-bit NormalFloat4 (NF4) quantization [47], foundation models such as Qwen2.5-VL-3B-Instruct execute completely offline on consumer workstation hardware. The internal causal language decoder acts as an intrinsic linguistic prior, allowing the model to resolve visual ambiguities by drawing on internal syntactic and statutory representations.

4.1.2 The Multi-Dimensional Evaluation Framework

Evaluating document transcription on municipal legislation exposes vulnerabilities that general-domain OCR benchmarks fail to capture. As Lopresti demonstrated [115], optical character recognition errors trigger an error cascade across subsequent natural language processing layers. In statutory construction, a single mutated character or omitted numeral corrupts governing legal semantics. A comprehensive evaluation framework for legislative transcription must combine classical string edit distances with specialized legal entity accuracy metrics and dense semantic embedding fidelity.

The foundational character and word accuracy metrics are formulated via the

Levenshtein edit distance $\mathcal{L}_{\text{Lev}}$ [109]. To assess baseline typographical fidelity across candidate document parsers, Character Error Rate and Word Error Rate are computed as follows:

Character Error Rate (CER) and Word Error Rate (WER):

$$\text{CER} = \frac{\mathcal{L}_{\text{Lev}}(\mathbf{c}_{\text{ref}}, \mathbf{c}_{\text{hyp}})}{|\mathbf{c}_{\text{ref}}|} = \frac{S_c + D_c + I_c}{N_c}, \quad \text{WER} = \frac{\mathcal{L}_{\text{Lev}}(\mathbf{w}_{\text{ref}}, \mathbf{w}_{\text{hyp}})}{|\mathbf{w}_{\text{ref}}|} = \frac{S_w + D_w + I_w}{N_w} \quad (4.1)$$

Where S , D , and I denote character or word substitutions, deletions, and insertions relative to the gold-standard reference sequence of length N .

Equation 4.1 formulates Character Error Rate (CER) and Word Error Rate (WER) using normalized Levenshtein string edit distances. These foundational metrics are incorporated to establish an initial, length-invariant baseline of typographical fidelity when transcribing historical physical enactments. The equations quantify the sum of character or word substitutions (S), deletions (D), and insertions (I) relative to the total token count of the gold reference transcript (N). In legislative document ingestion, low global CER values can obscure catastrophic localized errors, as omitting a single numerical digit in a penalty clause alters the legal meaning entirely. Consequently, while CER and WER quantify coarse optical recognition noise, they treat all tokens symmetrically, necessitating complementary domain-specific accuracy metrics for statutory entities.

To evaluate overall document similarity on a bounded scale regardless of clause length, the proponents compute Normalized Edit Distance:

Normalized Edit Distance (NED):

$$\text{NED}(\mathbf{s}_{\text{ref}}, \mathbf{s}_{\text{hyp}}) = 1.0 - \frac{\mathcal{L}_{\text{Lev}}(\mathbf{s}_{\text{ref}}, \mathbf{s}_{\text{hyp}})}{\max(|\mathbf{s}_{\text{ref}}|, |\mathbf{s}_{\text{hyp}}|)} \quad (4.2)$$

Where $\mathbf{s}_{\text{ref}}$ and $\mathbf{s}_{\text{hyp}}$ denote the reference and hypothesis string sequences, bounded between 0.0 (maximum dissimilarity) and 1.0 (exact match) to provide a length-invariant evaluation baseline.

Equation 4.2 defines Normalized Edit Distance (NED) as an inverted, length-bounded metric of string similarity. This formulation is included to provide an intuitive, normalized score ranging from 0.0 to 1.0 that remains stable across varying document lengths and section sizes. The equation scales the raw Levenshtein distance against the maximum string length between reference and hypothesis texts, subtracting the resulting ratio from unity so that an exact transcription yields a score of 1.0. In benchmarking candidate OCR engines across multi-page municipal ordinances, NED prevents lengthy boilerplate provisions from disproportionately skewing aggregate error measurements. This metric serves as a standardized indicator of overall string alignment before evaluating structural entity retention.

Beyond general string alignment, legislative conflict detection requires exact-match fidelity across critical statutory tokens:

Domain-Specific Legislative Component Accuracy:

$$\text{Metric Accuracy} = \frac{1}{|K|} \sum_{k \in K} \mathbb{I}(\mathbf{v}_{\text{hyp}}^{(k)} \equiv \mathbf{v}_{\text{ref}}^{(k)}) \quad (4.3)$$

Where K represents the set of extracted legal entity instances, and $\mathbb{I}(\cdot)$ is an exact-match indicator function comparing hypothesis values $\mathbf{v}_{\text{hyp}}^{(k)}$ against gold references $\mathbf{v}_{\text{ref}}^{(k)}$.

Equation 4.3 formulates Domain-Specific Legislative Component Accuracy across critical statutory entity classes. This targeted evaluation metric is integrated to ensure that high-stakes legal indicators—including statutory citations, enactment numbers, penalty figures, and effectivity dates—are transcribed with verbatim alphanumeric precision. The formulation computes the proportion of extracted entity instances where the hypothesis string matches the gold standard reference exactly according to the indicator function $\mathbb{I}(\cdot)$. In local ordinance digitization, a minor character corruption in a Republic Act citation breaks downstream BM25 inverted index lookups and misdirects dense vector retrieval. Enforcing exact-match component accuracy verifies whether candidate transcription engines preserve the governing legal tokens required for reliable Stage 1 candidate retrieval.

These domain metrics evaluate exact string equality across critical legislative components: Legal-Entity Accuracy (LEA) for local administrative offices and private contracting parties; Ordinance-Number Accuracy (ONA) for municipal enactment numbers and series suffixes; Date Accuracy (DA) for enactment, approval, and effectivity dates; Monetary-Value Accuracy (MVPA) for fines, fees, and penalty ceilings; Section-Heading Integrity (SHI) for operative structural catchlines; and Page Completeness (PC) for sequential page preservation. Furthermore, tabular schedule fidelity is evaluated using Tree-Edit-Distance-based Similarity (TEDS) [229], measuring row-column cell adjacency across multi-tier tax and zoning grids. Finally, Semantic Embedding Cosine Fidelity (SECF) measures dense vector retention by calculating the cosine similarity between dense embeddings of the reference and hypothesis texts.

4.1.3 The Insufficiency of Ordinary Spellchecking

Standard lexical spellcheckers and unconstrained dictionary autocorrection tools are fundamentally incapable of evaluating legal document transcription. In statutory conflict analysis, the cost of semantic errors is highly asymmetric. If an OCR engine transcribes “Section” as “Secton”, the edit distance is one and an ordinary spellchecker repairs the typographical slip easily. Conversely, if an engine transcribes a penalty of “P5,000.00” as “P50,000.00”, or mutates “shall” into “may”, or alters “Ordinance No. 0667-21” to “Ordinance No. 0661-21”, ordinary spellcheckers detect zero lexical errors because all generated tokens represent valid dictionary words. In municipal jurisprudence under the *Magtajas v. Pryce Properties* doctrine and Section 458 of Republic Act No. 7160, mutating a P5,000.00 penalty ceiling into P50,000.00 converts a valid local ordinance into an unconstitutional, *ultra vires* act.

Furthermore, Philippine legislative enactments contain high concentrations of specialized terminology that trigger massive false-positive rates in generic spellcheckers.

Municipal ordinances routinely incorporate statutory Latin maxims (*in pari materia*, *mutatis mutandis*, *ultra vires*, *ejusdem generis*), official administrative acronyms (CTTMO, CENRO, CPDO, LISSP), and indigenous Visayan and Mindanaoan toponyms. Davao City alone encompasses 182 administrative barangays, including localities such as Buhangin, Bunawan, Calinan, Marilog, Paquibato, Talomo, and Toril. Standard spellchecking algorithms aggressively flag these authentic names as misspellings. When automated autocorrection is enabled, generic tools destructively alter specialized local nomenclature into common English words, permanently corrupting statutory text.

4.1.4 Stratified Archival Reference Sampling

To establish an authoritative ground-truth benchmark without requiring manual transcription of all 1,664 archival documents, the proponents designed a stratified reference sample of $N = 60$ representative municipal enactments, comprising approximately 180 to 240 physical pages. Table 4.1 outlines the stratification matrix across eight legislative and physical axes.

Table 4.1 outlines the multi-dimensional stratification matrix across eight legislative and physical axes for evaluating document transcription across $N = 60$ representative municipal enactments. This stratified sampling design is incorporated to establish an empirically representative, unbiased benchmark of approximately 180 to 240 physical pages without requiring manual transcription of the entire 1,664-document archive. The matrix partitions enactments across four legislative eras, three document length tiers, three scan degradation categories, and six functional typologies, while explicitly tracking table density, mayoral seals, Filipino/Moro proper names, and statutory citation frequencies. In the distribution, historical typewriter scans from 1950–1991 and photocopied enactments from 1992–2015 comprise 50% of the sample, deliberately confronting models with severe physical noise such as carbon bleed-through, staple shadows, and uneven ink contrast. Establishing this balanced reference suite ensures that OCR and vision-language evaluation reliably measures model generalization across the full spectrum of degradation encountered in the Davao City archives.

4.1.5 Comparative Architectural Trade-Off Analysis

Table 4.2 summarizes the multi-dimensional comparative evaluation of the four candidate document parsing paradigms across accuracy metrics, hardware boundaries, legal text retention, data privacy, and operational costs.

Table 4.2 summarizes the multi-dimensional comparative evaluation of four candidate document parsing paradigms across twelve performance, operational, and financial dimensions. This comparative matrix is formulated to provide an empirical basis for selecting the municipal ingestion pipeline, weighing optical character accuracy against local edge hardware ceilings, financial sustainability, and institutional data privacy. The results demonstrate that the locally hosted vision-language model (Qwen2.5-VL-3B-Instruct in 4-bit NF4) closely rivals commercial cloud AI, achieving 3.12% CER, 96.7% monetary penalty accuracy, 91.8% table structure fidelity (TEDS), and 98.6% Semantic Cosine Fidelity. In sharp contrast, traditional scanner OCR (Tesseract v5.x) yields a prohibitive

Table 4.1: Stratified Reference Sampling Matrix for Archival Legislative Document Evaluation ($N = 60$ Enactments)

Stratification Axis	Strata Partition	Sample ($N = 60$)	Legislative & Archival Rationale
Legislative Era	Pre-LGC (1950–1991)	12 (20.0%)	Captures physical paper yellowing, manual type-writer fonts, early digital prints, and modern clean scans.
	Post-LGC Historical (1992–2015)	18 (30.0%)	
	Pre-Pandemic (2016–2020)	15 (25.0%)	
	LISSP Digital (2021–2025)	15 (25.0%)	
Document Length	Short (1–2 pages)	24 (40.0%)	Verifies model stability across single-page resolutions and multi-page omnibus regulatory codes.
	Medium (3–7 pages)	24 (40.0%)	
	Long (≥ 8 pages / Codes)	12 (20.0%)	
Scan Degradation	High Quality (Clean digital)	20 (33.3%)	Tests noise robustness against dark photocopy borders, staple shadows, and carbon-copy bleed-through.
	Moderate (Photocopy, minor skew)	25 (41.7%)	
	Severe (Bleed-through, dark edges)	15 (25.0%)	
Functional Typology	1. Substantive Regulatory / Penal	15 (25.0%)	Reflects the complete functional typology identified across the 1,664 Davao City municipal enactments.
	2. Inter-Agency MOAs / MOUs	10 (16.7%)	
	3. Disaster Risk & Calamity Funds	7 (11.7%)	
	4. Deeds of Donation & Usufruct	8 (13.3%)	
	5. Road Closures & Traffic Rules	10 (16.7%)	
	6. Zoning Reclassifications	10 (16.7%)	
Table Density	Zero Tabular Schedules	30 (50.0%)	Evaluates Markdown grid retention in revenue schedules, tax assessments, and traffic rerouting vectors.
	Simple Tables (2–4 columns)	18 (30.0%)	
	Dense Tables (≥ 5 columns)	12 (20.0%)	
Signatures & Seals	Clean Endings (Minimal)	20 (33.3%)	Evaluates whether mayoral signatures and dry seals trigger hallucinatory generation loops.
	Moderate (Official initials)	25 (41.7%)	
	Heavy (Overlapping seals & stamps)	15 (25.0%)	
Filipino Names	High Local / Moro / Indigenous	30 (50.0%)	Validates out-of-vocabulary character accuracy on Mindanaoan and Visayan personal and place names.
	Standard Anglophone / Spanish	30 (50.0%)	
Statutory Citations	High Citation Density (≥ 5 laws)	25 (41.7%)	Measures recognition fidelity on referenced Republic Acts, Presidential Decrees, and Batas Pambansa.
	Low Citation Density (< 5 laws)	35 (58.3%)	

14.82% CER and collapses to 63.3% accuracy on monetary amounts, introducing fatal errors for downstream statutory analysis. While Google Document AI achieves slightly lower raw error rates (2.15% CER), its mandatory external API transmission and recurring per-page fees conflict with Local Government Unit budget constraints and data sovereignty mandates under Republic Act No. 10173. Consequently, running Qwen2.5-VL-3B-Instruct offline on consumer GPU hardware delivers the optimal balance of enterprise-grade transcription precision and complete institutional independence.

To quantify the downstream consequences of transcription noise, the proponents conducted a controlled validation run on authentic archival municipal legislation (exemplified by Davao City Ordinance No. 000667-21). Traditional scanner OCR produced a Word Error Rate (WER) of 3.11% and a Character Error Rate (CER) of 0.45%. While these

Table 4.2: Multi-Dimensional Comparative Evaluation of Candidate Archival Document Architectures for Municipal Legislation

Evaluation Dimension	Tesseract v5.x	Surya v0.6+	Google Document AI	Qwen2.5-VL-3B (4-bit NF4)
Architectural Family	Traditional CNN-LSTM	Layout-Aware Modular	Cloud Multimodal Net	Local Vision-Language
Character Error Rate (CER)	14.82%	8.45%	2.15%	3.12%
Word Error Rate (WER)	22.30%	12.10%	4.05%	5.24%
Normalized Edit Distance	0.852	0.915	0.978	0.969
Legal-Entity Accuracy	74.2%	86.5%	97.2%	95.8%
Ordinance-Number Accuracy	68.3%	83.3%	98.3%	96.7%
Monetary-Value Accuracy	63.3%	80.0%	98.3%	96.7%
Table Fidelity (TEDS)	41.2%	76.4%	94.1%	91.8%
Semantic Cosine Fidelity	88.4%	93.8%	99.1%	98.6%
Execution Hardware	Minimal CPU (< 500 MB)	Mid-tier GPU (4 GB)	Client HTTPS Interface	Consumer GPU (RTX 5050)
Network Dependency	100% Offline	100% Offline	Mandatory Internet API	100% Offline
Data Privacy Compliance	Fully Sovereign	Fully Sovereign	Third-Party Cloud Route	Fully Sovereign (RA 10173)
Financial Model	Free Open-Source (\$0)	Free Open-Source (\$0)	Paywall per 1,000 pages	Free Open-Source (\$0)
Mean Page Latency	Fast (0.3–0.8 s, CPU)	Moderate (1.2–2.5 s)	Fast (0.5–1.2 s roundtrip)	Moderate (4.0–8.5 s, GPU)

Note: Evaluation reflects architectural properties and empirical validation runs on historical Davao City municipal legislation. Qwen2.5-VL-3B-Instruct evaluated in 4-bit NormalFloat4 (NF4) precision on Mobile Host B (NVIDIA RTX 5050 Laptop GPU). Metric values represent aggregate benchmark scores across evaluated archival reference transcripts.

global error rates appear small on a character level, the traditional baseline achieved only 54.55% Legal Marker Precision. The engine corrupted the explicit citation to Republic Act No. 7160, failed to isolate Section 1 and Section 2 headers, and dropped formatting in the monetary penalty clause.

In contrast, the locally hosted vision-language model achieved 0.00% Word Error Rate, 0.02% Character Error Rate, and 100.0% Legal Marker Precision across all eleven target statutory indicators. Furthermore, vectorizing the extracted text via TF-IDF n-gram embeddings confirmed that the VLM transcription achieved 100.0% Semantic Cosine Fidelity relative to verified reference text, compared to 99.21% for traditional OCR. The underlying language model decoder functions as an intrinsic linguistic prior: when visual glyphs are partially obscured by faint ink or scanner noise, the model’s internal language representations successfully reconstruct canonical statutory phrases.

4.1.6 Defensible Selection Rationale and Hardware Boundaries

Selecting Qwen2.5-VL-3B-Instruct as the primary archival ingestion engine resolves the technical trade-offs between extraction accuracy, operational expense, and statutory data sovereignty. While traditional OCR provides rapid CPU throughput, its inability to maintain section header geometry and its susceptibility to numeral substitution directly undermine Stage 1 lexical retrieval and Section-Level Specificity. Although commercial cloud platforms deliver strong extraction accuracy, deploying commercial APIs across

1,664 multi-page ordinances entails recurring financial expenses. Furthermore, transmitting unpromulgated local government drafts and internal committee working texts to external commercial cloud servers introduces legal compliance hurdles under Republic Act No. 10173 (Data Privacy Act of 2012) [153]. An edge-deployable local system simulating Local Government Unit IT constraints must remain independent of external subscription paywalls.

Deploying Qwen2.5-VL-3B-Instruct in 4-bit NF4 precision delivers enterprise-grade transcription accuracy while running entirely offline. The model processes pages deterministically, respects local data sovereignty, and operates within the 8 GB VRAM boundary of the Neural Training Workstation. The primary operational trade-off of the vision-language approach is computational latency. On the NVIDIA RTX 5050 Laptop GPU, transcribing a single scanned page requires approximately 4.0 to 8.5 seconds, compared to less than one second for traditional OCR. However, because archival corpus digitization is an offline, one-time preprocessing task, this latency expenditure is technically acceptable. Bounding page generation at 1,000 new tokens and isolating document execution within fault-tolerant resume checkpoints ensures dependable batch processing across the complete municipal collection.

4.1.7 Empirical Limitations and Archival Ceilings

The empirical evaluation of archival transcription reveals three concrete technical limitations that establish the performance boundaries of the system. First, physical document degradation imposes an asymptotic ceiling on optical character recognition. On historical carbon copies from the 1950s and 1960s where physical ink bleed-through overlaps characters from opposing page sides, even human legal annotators disagree on obscured numerals, indicating that recognition ambiguity stems from physical artifact loss rather than algorithmic insufficiency.

Second, complex multi-tier tabular schedules in historical municipal tax ordinances feature non-standard visual grids and irregular cell mergers. While the vision-language model successfully captures the numerical tax rates and cell contents in Markdown format, cell alignment variations in irregular tables depress automated Tree Edit Distance (TEDS) metrics despite verbatim preservation of substantive monetary numbers. Third, high-capacity dense embedding models can exhibit metric saturation, where a corrupted digit in a penalty clause results in only a fractional cosine similarity drop, masking a severe statutory preemption conflict. Addressing these edge cases reinforces the core architectural principle that automated ingestion serves as an assistive decision-support tool rather than an autonomous substitute for human legislative review.

4.2 Stage 1 Candidate Retrieval

In a two-stage legal conflict detection pipeline, the primary objective of Stage 1 coarse retrieval is to rapidly filter an extensive statutory collection of 27,096 enactments (176,421 searchable provisions), capturing all potentially preemptive provisions within a compact top- k candidate shortlist ($k \in \{5, 10, 20, 50\}$). In automated legal informatics, coarse

retrieval performance is judged primarily on candidate recall (Recall@ k) and Mean Reciprocal Rank (MRR), rather than classification accuracy [122]. Candidate recall (Recall@ k) quantifies the proportion of queries where the governing statutory provision is successfully retrieved within the top k candidates, while Mean Reciprocal Rank (MRR) measures the reciprocal position of the first relevant statute (1/rank), rewarding models that place governing laws near the very top of the candidate pool. If the governing statutory enactment is omitted from the top- k shortlist during Stage 1, downstream inference models in Stage 2 cannot detect the conflict regardless of their analytical reasoning capacity. Simultaneously, the shortlist must remain sufficiently compact ($k \leq 50$) to avoid diluting the cross-attention mechanism or exceeding local workstation memory ceilings during subsequent pair-wise inference.

4.2.1 Full-Corpus Sensitivity Benchmark

To evaluate retrieval robustness in an unconstrained production environment, the proponents benchmarked baseline retrieval configurations against the entire national statutory database ($N = 25,432$). In this evaluation, all 350 curated municipal queries were submitted directly to the full archive without manual pre-filtering or provision isolation. Table 4.3 reports the empirical retrieval sensitivity, rank distributions, and latency metrics across the full corpus.

Table 4.3: Empirical Retrieval Sensitivity Benchmark across the Full National Statutory Database ($N = 25,432$)

Stage 1 Retrieval Strategy	Latency	R@5	R@10	R@20	R@50	MRR@50
1. Baseline BM25 (Unconstrained)	329.6 ms	21.43%	29.71%	35.71%	43.14%	0.1521
2. BM25 + Hierarchy Priority Safeguard	300.1 ms	29.14%	34.86%	39.43%	43.14%	0.2217
3. BM25 + Hard Domain Filter	297.5 ms	4.86%	7.71%	10.29%	13.43%	0.0467

Note: Evaluated on all 350 ground truth queries searching the entire 25,432 national statute corpus on the Primary Desktop Host (Intel Core i3-10105F CPU @ 3.70 GHz). Latency values represent mean per-query retrieval runtime.

Table 4.3 presents the empirical candidate retrieval sensitivity, rank distributions, and per-query CPU latency across the entire national statutory database ($N = 25,432$) for three Stage 1 lexical strategies. This full-corpus benchmark is incorporated to evaluate whether sparse keyword search can withstand the noise of thousands of heterogeneous administrative issuances without manual pre-filtering. The results show that unconstrained BM25 achieves only 29.71% Recall@10 and 43.14% Recall@50 with a Mean Reciprocal Rank (MRR@50) of 0.1521, causing more than half of governing national statutes to be omitted from downstream inference. Attempting to eliminate administrative noise via a hard domain filter triggers catastrophic retrieval collapse, dropping Recall@50 to 13.43% and MRR to 0.0467 because omnibus statutory codes span multiple regulatory areas. In contrast, incorporating the Hierarchy Priority Safeguard elevates MRR to 0.2217 (+45.8% relative gain) by prioritizing primary Republic Acts over subordinate administrative orders. These findings establish that unassisted lexical search is fundamentally insufficient for large-scale statutory preemption retrieval, necessitating dense semantic representations

and hierarchical ranking safeguards.

4.2.2 Bi-Encoder Retrieval Benchmark

To identify the most accurate and computationally efficient dense semantic retriever for Stage 1 coarse filtering, the proponents benchmarked the eight candidate bi-encoders defined in Section 2.8.6 and Section 3.7.2 across all 350 curated ground truth municipal queries. In this evaluation, each bi-encoder generates continuous provision embeddings (384 to 1,024 dimensions), and candidate provisions are ranked strictly by cosine similarity against the draft query vector without sparse keyword assistance. Table 4.4 summarizes the empirical retrieval sensitivity, latency, and difficulty tier performance across candidate bi-encoders.

Table 4.4: Empirical Performance Comparison of Candidate Stage 1 Dense Bi-Encoder Architectures ($N = 350$)

Category	Candidate Bi-Encoder	Params	Context	Latency	R@5	R@10	R@20	R@50	MRR	Tier 3
										R@5
Distilled Edge	all-MiniLM-L6-v2*	22.7M	512	1.33 ms	87.1%	97.1%	100.0%	100.0%	0.7388	91.3%
	bge-small-en-v1.5	33.4M	512	1.85 ms	87.4%	97.1%	100.0%	100.0%	0.7431	91.3%
General MTEB	bge-base-en-v1.5	109.0M	512	4.82 ms	88.0%	97.7%	100.0%	100.0%	0.7512	92.3%
	e5-base-v2	109.0M	512	4.95 ms	87.7%	97.4%	100.0%	100.0%	0.7480	91.3%
	all-mpnet-base-v2	109.0M	512	5.12 ms	87.7%	97.4%	100.0%	100.0%	0.7495	92.3%
Legal-Domain	legal-bert-base-uncased	110.0M	512	4.88 ms	86.6%	96.6%	99.7%	100.0%	0.7310	89.4%
Long-Context	ModernBERT-base	149.0M	8,192	6.20 ms	88.6%	98.0%	100.0%	100.0%	0.7584	93.3%
Multilingual	bge-m3	568.0M	8,192	24.50 ms	88.9%	98.3%	100.0%	100.0%	0.7621	93.3%

Note: *Selected model integrated into the production Stage 1 hybrid pipeline. Evaluated standalone using normalized cosine similarity on the Primary Desktop Host (Intel Core i3-10105F CPU @ 3.70 GHz) without sparse keyword assistance. Latencies represent mean per-query CPU runtime.

Table 4.4 summarizes the empirical performance of eight candidate dense bi-encoder architectures evaluated across parameter scales, context lengths, CPU latencies, candidate recall thresholds ($k \in \{5, 10, 20, 50\}$), Mean Reciprocal Rank, and Tier 3 latent sensitivity. This evaluation is included to identify the optimal dense semantic retriever for Stage 1 coarse filtering, comparing lightweight edge encoders against large-context foundation models. The results demonstrate that all eight candidate bi-encoders capture continuous legal semantics effectively, each reaching at least 96.6% Recall@10 and 100.0% Recall@50 across the 350 curated queries. While scaled models such as **bge-m3** (MRR = 0.7621) and **ModernBERT-base** (MRR = 0.7584) attain the highest raw precision on paraphrastic Tier 3 queries (93.3% Recall@5), their 8,192-token architectures incur substantial latency penalties (24.50ms and 6.20ms per query on CPU). In contrast, **all-MiniLM-L6-v2** attains 87.1% Recall@5, 97.1% Recall@10, and 100.0% Recall@20 within an ultra-fast 1.33 milliseconds and a 22.7M parameter footprint. Because section-level statutory provisions average under 300 tokens, **all-MiniLM-L6-v2** provides the optimal trade-off between semantic sensitivity and CPU edge execution speed, leading to its selection as the core dense retriever for the hybrid pipeline.

4.2.3 Retrieval and Fusion Ablation

In empirical machine learning, an ablation study systematically isolates, removes, or modifies individual algorithmic components to determine their exact contribution to overall system performance. To determine the optimal combination of lexical and semantic retrieval mechanisms, the proponents conducted an ablation study evaluating nine candidate architectures against the 350 curated ground truth queries. Table 4.5 summarizes empirical performance across candidate architectures across latency, Recall@ k cutoffs, Mean Reciprocal Rank, and Tier 3 latent retrieval sensitivity.

Table 4.5: Empirical Performance Comparison of Candidate Stage 1 Retrieval Architectures ($N = 350$)

Candidate Retrieval Architecture	Latency	R@5	R@10	R@20	R@50	MRR	Tier 3 R@5
1. Pure BM25 Baseline (AIIR)	1.39 ms	79.7%	88.6%	97.4%	99.7%	0.6113	74.0%
2. BM25 + Hard Domain Filter	1.23 ms	99.1%*	100.0%*	100.0%*	100.0%*	0.7921*	99.0%*
3. BM25 + Soft Domain Prior (+20%)	1.43 ms	87.7%	94.6%	99.7%	100.0%	0.7122	79.8%
4. Pure Dense Bi-Encoder (MiniLM)	1.33 ms	87.1%	97.1%	100.0%	100.0%	0.7388	91.3%
5. Hybrid JNLP (Weighted Sum $\alpha = 0.5$)	1.23 ms	82.6%	91.7%	98.9%	100.0%	0.6409	77.9%
6. Hybrid Dense-Biased ($\alpha = 0.7$)	1.20 ms	84.0%	94.0%	99.4%	100.0%	0.6656	78.8%
7. Hybrid BM25-Biased ($\alpha = 0.3$)	1.22 ms	80.9%	90.6%	98.3%	100.0%	0.6202	76.0%
8. Reciprocal Rank Fusion (RRF $k = 60$)	1.26 ms	88.6%	96.0%	99.7%	100.0%	0.7040	88.5%
9. Hybrid + Soft Domain Prior ($\alpha = 0.5$)	1.21 ms	89.7%	95.7%	99.7%	100.0%	0.7328	84.6%

Note: *Configuration 2 reflects the artificial isolation effect; when evaluated across the full 25,432 corpus, hard filtering collapses overall Recall@50 to 13.43%. Latencies represent mean per-query CPU runtime.

Table 4.5 details the empirical performance of nine candidate Stage 1 retrieval architectures across mean per-query CPU latency, Recall@ k thresholds ($k \in \{5, 10, 20, 50\}$), Mean Reciprocal Rank (MRR), and Tier 3 latent paraphrastic sensitivity. This ablation benchmark is designed to isolate the exact empirical contributions of sparse lexical search (BM25), dense vector embeddings (all-MiniLM-L6-v2), domain-specific classification filters, and reciprocal rank fusion. The data show that while standalone BM25 achieves 79.7% Recall@5 and 99.7% Recall@50, its sensitivity drops to 74.0% on Tier 3 queries that lack verbatim keyword matches. Conversely, pure dense retrieval achieves 87.1% Recall@5 and 91.3% on Tier 3, demonstrating high conceptual sensitivity at the expense of exact alphanumeric penalty matching. Combining both streams through Reciprocal Rank Fusion (RRF $k = 60$) achieves 88.6% Recall@5, 96.0% Recall@10, and a perfect 100.0% Recall@50 with a mean CPU latency of only 1.26 milliseconds. Achieving 100.0% candidate recall at $k = 50$ guarantees zero provision attrition when shortlisting candidate statutes for downstream Stage 2 inference.

To visualize the trajectory of candidate retrieval recall as the shortlist expands and inspect how different architectures handle latent paraphrasing, Figure 4.1 illustrates the candidate Recall@ k progression curves alongside the performance breakdown across jurisprudential difficulty tiers.




Figure 4.1: Stage 1 Candidate Retrieval Performance across Shortlist Thresholds and Jurisprudential Difficulty Tiers ($N = 350$). Panel A traces Recall@ k progression curves. Panel B highlights the critical semantic gap in Tier 3 (Latent and Paraphrastic), where dense vectors achieve a +17.3% absolute gain over sparse BM25.

Figure 4.1 illustrates candidate retrieval recall trajectories and difficulty tier breakdowns for Stage 1 architectures evaluated across the 350 ground truth queries. This visualization is included to highlight the operational trade-offs across expanding shortlist sizes and expose the semantic gap between lexical and neural search. Panel A traces the Recall@ k progression curves from $k = 5$ to $k = 50$, demonstrating that hybrid fusion strategies converge rapidly toward 100% recall while standalone lexical search plateaus early. Panel B highlights the critical semantic gap in Tier 3 (Latent and Paraphrastic), where dense vector representations achieve a +17.3% absolute advantage over sparse BM25 (91.3% versus 74.0%). In practical legislative drafting, municipal proposals frequently reword regulatory concepts using localized phrasing rather than quoting statutory acts verbatim. The visual evidence confirms that combining dense semantic vectors with lexical keyword matching under Reciprocal Rank Fusion provides necessary robustness against vocabulary mismatch while maintaining ultra-low CPU latency.

4.3 Stage 2 Conflict Detection

Following candidate retrieval filtering, Stage 2 performs fine-grained semantic conflict detection using specialized Transformer Cross-Encoders. General pretrained language models possess broad linguistic capabilities but are not calibrated for the formal syntax, hierarchical qualifications, and strict preemption doctrines of municipal law. As observed in recent COLIEE benchmarks, zero-shot models routinely suffer from affirmative bias—the systematic cognitive tendency of pretrained language models to presume pairs are harmonious or logically non-conflicting unless prompted by superficial lexical nega-

tion markers—defaulting toward neutral or entailment classifications when confronted with legal text [5, 199]. Fine-tuning adjusts the cross-attention weights between statutory premises and local hypotheses, training the model to recognize when local clauses impermissibly expand, restrict, or contradict superior statutory mandates.

It is important to emphasize that the evaluations reported in this section represent initial model exploration, screening, and hyperparameter optimization experiments conducted on the initial 350-pair ground truth benchmark (expert-curated and adviser-reviewed) to establish baseline inductive properties, optimize internal configurations, and calibrate decision thresholds prior to the final human evaluation deployed by the 15 Sangguniang Panlungsod legal researchers. This section details the initial screening evaluations across 28 scouted models, supervised fine-tuning sweeps, candidate ablations, and convergence dynamics of candidate models.

4.3.1 Zero-Shot Candidate Screening

In accordance with Step 1 of the Two-Step Model Screening Process formulated in Section 3.8.1, the proponents conducted a rapid zero-shot screening evaluation across 28 candidate models representing the six architectural families defined in Section 2.9.4. In this evaluation, all 28 models were tested out-of-the-box on the 53 held-out test pairs ($\mathcal{D}_{\text{test}}$) without prior task-specific fine-tuning. This screening tests the intrinsic deontic sensitivity of each architecture and exposes the severity of baseline affirmative bias before investing substantial GPU training hours. Furthermore, it establishes the empirical *Zero-Shot Baseline* specified in Section 3.7.1, quantifying the lower-bound performance of unadapted language models. Table 4.6 summarizes the empirical zero-shot accuracy, macro F_1 , and conflict F_1 scores across all 28 candidate architectures, indicating with an asterisk (*) the thirteen models that satisfied the admission criteria and advanced to the supervised fine-tuning and ablation suite.

Table 4.6 summarizes the empirical zero-shot screening results and affirmative bias assessments across 28 scouted encoder architectures evaluated on the held-out test split ($\mathcal{D}_{\text{test}}, N = 53$). This initial screening is conducted to benchmark the intrinsic deontic reasoning capabilities of diverse model families prior to GPU fine-tuning, identifying architectures that can resist affirmative bias without exceeding edge hardware ceilings. The data show that pre-aligned NLI reasoners, specifically **DeBERTa-v3-large-NLI** (Conflict $F_1 = 0.7200$) and **DeBERTa-v3-base-NLI** (Conflict $F_1 = 0.6957$), achieve the highest zero-shot contradiction sensitivity, effectively disentangling statutory permissions from local municipal prohibitions. In contrast, unadapted general-domain encoders and distilled edge models suffer from severe affirmative bias, with architectures like **all-MiniLM-L6-v2** (Conflict $F_1 = 0.3846$) and **distilbert-base** (Conflict $F_1 = 0.3478$) systematically collapsing into Entailment or Neutral predictions on over 50% of genuine conflicts. Following the operational screening protocol, exactly thirteen candidate architectures satisfied the admission criteria and were shortlisted for supervised fine-tuning, while fifteen models were eliminated due to prohibitive affirmative bias, excessive parameter footprints, or autoregressive generation risks.

Conversely, fifteen candidate architectures were disqualified based on distinct em-

Table 4.6: Empirical Zero-Shot Screening and Affirmative Bias Evaluation across 28 Scouted Candidate Encoders ($N_{\text{test}} = 53$)

Family	Candidate Architecture	Params	Context	Acc	Macro F_1	Conflict F_1	Affirmative Bias
Legal-Domain	legal-bert-base-uncased*	110M	512	66.04%	0.6380	0.6087	Low (Term sensitivity)
	PoL-BERT-Large*	340M	512	69.81%	0.6845	0.6667	Low (Regulatory fit)
	InLegalBERT*	110M	512	62.26%	0.5980	0.5652	Moderate (Court bias)
	CaseLaw-BERT	110M	512	62.26%	0.5890	0.5455	Moderate (Case bias)
	Lawformer	110M	4,096	58.49%	0.5412	0.4762	High (Sliding window)
Pre-Aligned NLI	DeBERTa-v3-base-NLI*	86M	512	71.70%	0.7042	0.6957	Low (Strong negation)
	DeBERTa-v3-large-NLI*	435M	512	73.58%	0.7251	0.7200	Low (Robust balance)
	deberta-v3-base*	86M	512	62.26%	0.5891	0.5217	Moderate (Cold MLM)
	roberta-large-mnli	355M	512	60.38%	0.5420	0.4444	High (Default neutral)
	bart-large-mnli	406M	1,024	58.49%	0.5310	0.4348	High (Seq2seq lag)
	electra-large-disc	335M	512	62.26%	0.5982	0.5385	Moderate (RTD sample)
	bert-base-uncased*	110M	512	56.60%	0.5240	0.4545	High (General MLM baseline)
Long-Context	ModernBERT-base*	149M	8,192	64.15%	0.6120	0.5600	Moderate (Fast throughput)
	ModernBERT-large*	395M	8,192	67.92%	0.6654	0.6400	Low (Consistent recall)
	longformer-base-4096	149M	4,096	58.49%	0.5385	0.4615	High (Diluted attention)
	bigbird-roberta-base	128M	4,096	56.60%	0.5190	0.4286	High (Sparse block loss)
	nomic-bert-2048	137M	2,048	58.49%	0.5450	0.4783	Moderate (Short window)
Cross-Encoders	bge-reranker-v2-m3*	568M	8,192	64.15%	0.6210	0.5833	Moderate (Relevance head)
	bge-reranker-base	278M	512	60.38%	0.5620	0.5000	Moderate (Redundant)
	ms-marco-MiniLM-L12	33M	512	52.83%	0.4720	0.3636	Severe (Relevance bias)
Multilingual	mdeberta-v3-base*	86M	512	60.38%	0.5694	0.4800	High (Loanword robust)
	xlm-roberta-base	270M	512	54.72%	0.4980	0.3913	High (Heavy cross-lingual)
	xlm-roberta-large	550M	512	56.60%	0.5180	0.4286	High (Excessive VRAM)
	roberta-tagalog-base	110M	512	47.17%	0.4120	0.3077	Severe (Collapses neutral)
Distillation/Edge	all-MiniLM-L6-v2*	22M	512	54.72%	0.4812	0.3846	Severe (Entailment bias)
	nli-distilroberta-base*	82M	512	56.60%	0.5124	0.4167	Severe (Entailment bias)
	distilbert-base-uncased	66M	512	52.83%	0.4680	0.3478	Severe (Collapses neutral)
	albert-base-v2	12M	512	49.06%	0.4350	0.3182	Severe (Entailment bias)

Note: Evaluated zero-shot out-of-the-box on the 53 held-out test pairs ($\mathcal{D}_{\text{test}}$) without task-specific fine-tuning. Models marked with an asterisk (*) successfully satisfied the admission criteria and advanced to the supervised fine-tuning and ablation suite. Affirmative bias denotes the empirical tendency of unadapted models to default toward Entailment or Neutral on statutory conflicts.

pirical failure modes. First, models such as DistilBERT-base, ALBERT-base-v2, and roberta-tagalog-base displayed severe affirmative bias, yielding Conflict F_1 scores below 0.35 due to a near-total collapse into affirmative predictions on contradictory provisions. Second, heavy multilingual and sequence-to-sequence models, including XLM-RoBERTa-large (550M parameters, 72.4 ms latency) and BART-large-mnli (406M parameters, 68.2 ms latency), exceeded the 2 GB edge VRAM ceiling without offering accuracy advantages over compact DeBERTa-v3 architectures. Third, encoders such as ELECTRA-large-disc and BGE-Reranker-base proved redundant relative to standard BERT and DeBERTa models, while older sparse long-context encoders (Longformer-base and BigBird-RoBERTa) were superseded by ModernBERT’s unpadded FlashAttention-2 design. Finally, sequence-to-sequence models like BART-large-mnli were excluded to eliminate autoregressive decoding overhead and non-deterministic text generation risks, preserving the deterministic classification mandate of the screening engine.

4.3.2 Supervised Fine-Tuning Benchmark

In rigorous machine learning research, an *ablation study* systematically modifies one specific model component or inductive bias at a time while holding all other training conditions constant. This controlled approach isolates the exact causal factor responsible for performance improvements, rather than treating competing neural models as opaque black boxes. To investigate these architectural factors, the thirteen shortlisted candidate architectures were fine-tuned on the 245 training pairs ($\mathcal{D}_{\text{train}}$) and evaluated on the 52 validation pairs ($\mathcal{D}_{\text{val}}$) across initial learning rates $\eta \in \{1 \times 10^{-5}, 2 \times 10^{-5}, 3 \times 10^{-5}\}$ and batch sizes $B \in \{8, 16\}$ under deterministic execution (seed = 42).

Table 4.7 documents the validation performance of the thirteen shortlisted architectures, grouped cleanly by architectural family without duplicate entries. Each model was evaluated under its optimal learning rate and batch size configuration ($B = 8$).

Table 4.7 presents the supervised fine-tuning validation performance and hyperparameter convergence metrics across the thirteen shortlisted candidate Cross-Encoders evaluated on the validation split ($N_{\text{val}} = 52$). This controlled benchmark is designed to systematically measure the performance returns across four architectural ablation axes: parameter scaling, NLI pre-alignment, legal domain specialization, and compact edge execution. The results reveal that pre-aligned NLI models achieve the highest overall predictive accuracy, with **DeBERTa-v3-large-NLI** attaining minimum validation loss (0.4320) and peak Conflict F_1 (0.9000), closely followed by its base variant **DeBERTa-v3-base-NLI** (Val Loss = 0.4410, Conflict $F_1 = 0.8947$). Scaling to larger context windows via **ModernBERT-large** yields strong performance (Conflict $F_1 = 0.8750$) supported by native FlashAttention-2 speed, while administrative legal pretraining in **PoL-BERT-Large** (Conflict $F_1 = 0.8696$) outperforms general BERT baselines. Task-specific fine-tuning across all thirteen encoders substantially improves upon both the unadapted zero-shot baseline ($F_1 = 0.4545$) and canonical supervised BERT ($F_1 = 0.8519$), confirming that specialized adaptation is indispensable for resolving complex statutory collisions.

The empirical findings from Table 4.7 systematically resolve the research questions formulated across the four candidate ablation axes. Regarding parameter and capacity

Table 4.7: Supervised Fine-Tuning Performance and Hyperparameter Convergence across the Thirteen Shortlisted Cross-Encoders ($N_{\text{val}} = 52$)

Candidate Model & Inductive Bias	Params	LR (η)	Batch	Val Loss	Val Acc	Conflict F_1
Baseline Models (Zero-Shot & Supervised Foundations)						
Zero-Shot Cross-Encoder (Un-adapted Task Baseline)	110M	—	—	1.0820	56.60%	0.4545
bert-base-uncased (Canonical Supervised General Baseline)	110M	2×10^{-5}	8	0.5124	82.69%	0.8519
Pre-Aligned NLI Reasoners						
DeBERTa-v3-large-NLI (WANLI Scaled Reasoning Ceiling)	435M	2×10^{-5}	8	0.4320	86.54%	0.9000
DeBERTa-v3-base-NLI (Disentangled Warm Workhorse)	86M	2×10^{-5}	8	0.4410	86.54%	0.8947
deberta-v3-base (Cold MLM Pre-trained Baseline)	86M	2×10^{-5}	8	0.4812	84.62%	0.8750
Modern Long-Context Encoders						
ModernBERT-large (FlashAttention-2 Scaled 8K)	395M	2×10^{-5}	8	0.4625	84.62%	0.8750
ModernBERT-base (FlashAttention-2 Base 8K)	149M	2×10^{-5}	8	0.5231	82.69%	0.8500
Legal Domain-Adapted Encoders						
Pol-BERT-Large (Administrative Codification Pretraining)	340M	2×10^{-5}	8	0.4712	84.62%	0.8696
legal-bert-base-uncased (Judicial Domain Pretraining)	110M	2×10^{-5}	8	0.4950	82.69%	0.8571
InLegalBERT (Common Law Statutory Pretraining)	110M	2×10^{-5}	8	0.5180	80.77%	0.8333
Cross-Encoder Rerankers						
bge-reranker-v2-m3 (Multilingual Cross-Encoder Reranker)	568M	2×10^{-5}	8	0.4890	82.69%	0.8571
Multilingual Encoders						
mdeberta-v3-base (Multilingual Disentangled Baseline)	86M	2×10^{-5}	8	0.5020	82.69%	0.8500
Compact Edge Baselines						
nli-distilroberta-base (Distilled NLI Cross-Encoder)	82M	2×10^{-5}	8	0.5690	80.77%	0.8182
all-MiniLM-L6-v2 (Ultra-Compact Edge Baseline)	22M	3×10^{-5}	8	0.5784	80.77%	0.8108

Note: Evaluated on the 52 validation pairs ($\mathcal{D}_{\text{val}}$) across hyperparameter grid searches under deterministic execution (seed = 42). Exactly 13 unique architectures qualified for fine-tuning. Macro F_1 averages across Entailment, Neutral, and Contradiction.

scaling (Axis 1), comparing base encoders against their 400M-parameter counterparts directly measures the performance return on computational scaling. In the ModernBERT family, scaling from base (149M) to large (395M) increased Conflict F_1 from 0.8500 to 0.8750 (+2.50%), lowering validation loss from 0.5231 to 0.4625. In the DeBERTa-v3 family, scaling from base (86M) to large (435M) lifted Conflict F_1 from 0.8947 to 0.9000, establishing the empirical ceiling for the study. However, because DeBERTa-v3-large requires nearly $4\times$ higher inference latency (48.7 ms vs. 14.8 ms), the base variant provides

a superior efficiency-accuracy balance for real-time municipal drafting.

Analyzing natural language inference pre-alignment and baseline contrasts (Axis 2) isolates the scientific value of initializing neural weights from multi-task NLI pretraining (MNLI, FEVER, ANLI) versus cold masked language modeling, while establishing an empirical contrast against the canonical supervised baseline (**bert-base-uncased**). Standard **bert-base-uncased** establishes the foundational supervised baseline at Conflict $F_1 = 0.8519$ (Val Loss = 0.5124). Upgrading the encoder architecture to modern disentangled attention without NLI pretraining (**deberta-v3-base**) improves Conflict F_1 to 0.8750 (+2.31% over base BERT). Initializing from warm multi-task NLI checkpoints (**DeBERTa-v3-base-NLI**) further lifts Conflict F_1 to 0.8947 (+4.28% over base BERT and +1.97% over cold DeBERTa), while the scaled **DeBERTa-v3-large-NLI** achieves 0.9000 (+4.81% over base BERT). Furthermore, all fine-tuned models substantially outperform the unadapted zero-shot baseline ($F_1 = 0.4545$), confirming that task-specific adaptation provides substantial performance benefits for statutory preemption detection.

Evaluating legal-domain pretraining adaptation (Axis 3) indicates that administrative regulatory pretraining (PoL-BERT-Large, Conflict $F_1 = 0.8696$) out-performs judicial case-law pretraining (Legal-BERT, Conflict $F_1 = 0.8571$). Both domestic models outperform Indian statutory pretraining (InLegalBERT, Conflict $F_1 = 0.8333$), confirming that vocabulary adaptation to regulatory codifications is more effective for local municipal ordinances than litigation-centric case text.

Finally, evaluating compact edge baselines (Axis 4) reveals that distilled ultra-compact encoders, such as all-MiniLM-L6-v2 and nli-distilroberta-base, achieved over 81% Conflict F_1 following supervised fine-tuning. Although these models trail larger cross-encoders by approximately 8% in conflict detection accuracy, they execute in under 10 ms per pair on standard CPU hardware, establishing viable lightweight alternatives for municipal departments operating without dedicated graphics accelerators.

To evaluate statistical significance across the candidate models, the proponents conducted a non-parametric Friedman test across the eight legal domains ($\chi_F^2 = 18.42, p < 0.05$), confirming that performance differences between architectural families are statistically significant. Subsequent post-hoc pairwise comparisons using McNemar’s test with Edwards’ continuity correction confirmed that the top-performing candidate (DeBERTa-v3-base-NLI) significantly outperforms both the unadapted zero-shot baseline ($\chi^2 = 7.14, p = 0.0075 < 0.01$) and the canonical supervised BERT baseline ($\chi^2 = 4.27, p = 0.0388 < 0.05$) on discordant conflict pairs.

4.3.3 Loss Dynamics and Convergence

In supervised machine learning, the loss function mathematically quantifies the penalty incurred when the model’s predicted probability distribution diverges from the actual ground truth labels. For multi-class natural language inference, the training objective optimizes Multi-Class Cross-Entropy Loss:

Multi-Class Cross-Entropy Training Loss Function:

$$\mathcal{L}_{\text{CE}} = - \sum_{c \in \mathcal{C}} y_c \log \hat{P}(c) \quad (4.4)$$

Where $y_c \in \{0, 1\}$ denotes the binary ground truth indicator for class $c \in \mathcal{C}$, and $\hat{P}(c)$ represents the calibrated softmax probability output by the model.

Equation 4.4 defines the Multi-Class Cross-Entropy loss objective utilized for supervised fine-tuning of candidate Cross-Encoder architectures. This loss function is incorporated to penalize divergence between predicted class probabilities and ground truth NLI relationships (Entailment, Neutral, Contradiction). The formulation computes the negative log-likelihood of the ground truth label y_c , scaling penalization asymptotically as predicted confidence $\hat{P}(c)$ strays from certainty. During optimization over the 245 training pairs, backpropagating gradients from this loss iteratively updates transformer cross-attention weights, aligning internal representations with formal legal preemption structures. Minimizing cross-entropy loss drives the model toward confident, well-separated class assignments necessary for calibrated thresholding.

An epoch represents one complete training pass across all 245 training pairs. Tracking the trajectory of training loss alongside validation loss across successive epochs serves as the foundational diagnostic method to monitor convergence and detect the boundary between underfitting and overfitting. In healthy training, both training loss and validation loss decrease synchronously. When training loss continues to fall while validation loss stabilizes and subsequently turns upward (the inflection point), the network has exhausted generalizable patterns and begun memorizing idiosyncrasies in the training sample.

The fine-tuning and convergence results show that the Top 5 performing architectures (**DeBERTa-v3-large-NLI**, **DeBERTa-v3-base-NLI**, **ModernBERT-large**, **cold deberta-v3-base**, and **PoL-BERT-Large**) alongside the canonical **bert-base-uncased** baseline achieve the lowest holdout validation losses and stable generalization at Epoch 3. Therefore, these six specific architectures are advanced as the candidate finalists for the subsequent out-of-sample benchmark evaluation and final system deployment roadmap. To evaluate convergence stability across this finalist pool without presuming a predetermined winner, Figure 4.2 illustrates the comparative holdout validation loss dynamics and Macro F_1 trajectories across all six candidate finalists identified from the supervised fine-tuning leaderboard (Table 4.7) under their optimal tuning configurations ($\eta = 2 \times 10^{-5}$, $B = 8$).

Figure 4.2 illustrates the comparative holdout validation loss dynamics and Macro F_1 score trajectories across the six candidate finalists over five fine-tuning epochs ($\eta = 2 \times 10^{-5}$, $B = 8$). This empirical visualization is presented to diagnose convergence behavior, verify gradient stability, and identify the universal bias-variance inflection point across competing transformer backbones. Panel A tracks holdout validation loss, demonstrating that all candidate encoders achieve their global minimum at Epoch 3 before displaying overfitting divergence as training loss continues its monotonic decline. Panel B traces holdout Macro F_1 progression, illustrating that pre-aligned NLI reasoners (**DeBERTa-v3-large-NLI** and **DeBERTa-v3-base-NLI**) consistently sustain superior classification balance throughout training compared to general baselines (**bert-base-uncased**).




Figure 4.2: Stage 2 Cross-Encoder Comparative Loss Dynamics and Convergence Trajectories across the Six Candidate Finalists ($\eta = 2 \times 10^{-5}$, $B = 8$). Panel A traces holdout validation loss across epochs, demonstrating the universal bias-variance inflection point at Epoch 3 alongside the training trajectory. Panel B compares the validation Macro F_1 -score progression, illustrating that pre-aligned NLI reasoners and modern scaled encoders consistently outperform unaligned and general baselines throughout optimization.

Across both panels, the trajectories confirm that task-specific adaptation converges rapidly within three full passes over the 245 training pairs without exhibiting loss oscillations or gradient collapse.

While Figure 4.2 visualizes comparative holdout trajectories across the entire finalist pool, Table 4.8 provides a detailed diagnostic breakdown of the epoch-by-epoch optimization progression for the leading candidate architecture emerging from the validation grid search (DeBERTa-v3-large-NLI).

Table 4.8: Epoch-by-Epoch Diagnostic Training Dynamics, Holdout Loss, and Generalization Progression for the Leading Validation Candidate (DeBERTa-v3-large-NLI, $\eta = 2 \times 10^{-5}$, $B = 8$)

Epoch (E)	Training Loss	Validation Loss	Val Acc	Macro F_1	Optimization Status / Event
1	1.0210	0.7240	71.15%	0.7120	Initial linear learning rate warmup
2	0.6540	0.5120	80.77%	0.8140	Steady convergence; rapid preemption learning
3	0.3840	0.4320	86.54%	0.8580	Optimal Checkpoint (Minimum Validation Loss)
4	0.2520	0.4480	84.62%	0.8420	Early signs of holdout loss divergence (+0.0160)
5	0.1710	0.4850	82.69%	0.8310	Early stopping triggered (patience = 4 check)

Table 4.8 documents the epoch-by-epoch training and validation loss progression, holdout accuracy, Macro F_1 , and optimization events for the leading candidate architecture (DeBERTa-v3-large-NLI, $\eta = 2 \times 10^{-5}$, $B = 8$). This diagnostic case study is included to inspect the exact mechanics of the bias-variance trade-off on a specialized legal classification split ($N_{\text{train}} = 245$, $N_{\text{val}} = 52$) and verify the automated early-stopping

trigger. The table records training loss dropping steadily from 1.0210 at Epoch 1 to 0.1710 at Epoch 5, while holdout validation loss reaches an optimal minimum of 0.4320 at Epoch 3, yielding peak validation accuracy of 86.54% and Macro F_1 of 0.8580. Beyond Epoch 3, validation loss reverses direction, rising to 0.4480 at Epoch 4 (+0.0160) and 0.4850 at Epoch 5 (+0.0530), indicating that the high-capacity encoder begins memorizing phrasing idiosyncrasies within the small training set. The empirical log confirms that the automated patience trigger successfully halted training at Epoch 5 and restored the optimal Epoch 3 checkpoint, preventing catastrophic over-specialization prior to out-of-sample testing.

4.3.4 Sensitivity Analysis and Regularization

The potential for model overfitting represents a critical vulnerability when adapting high-capacity transformer architectures (such as 86-million to 435-million parameter models) to a specialized, compact training set ($N_{\text{train}} = 245$). Left unregularized, a neural network readily memorizes idiosyncratic phrasing artifacts rather than learning broader statutory conflict principles, producing near-zero training error while failing on unseen municipal measures. The empirical training logs confirm that the tripartite regularization strategy—combining decoupled AdamW weight decay ($\lambda = 0.01$) [116], attention and hidden dropout ($p = 0.10$) [182], and early stopping validation checkpoints [144]—effectively curbed parameter over-specialization. At the optimal Epoch 3 checkpoint, the generalization gap between training accuracy (92.4%) and holdout validation accuracy (84.6%) remained tightly bounded at 7.8%, indicating that the cross-encoders learned generalizable statutory preemption patterns rather than superficial keyword heuristics.

Candidate encoders also exhibited sharp sensitivity to the initial learning rate during hyperparameter grid exploration. At a conservative rate of $\eta = 1 \times 10^{-5}$, models displayed noticeable underfitting, as the optimization trajectory progressed too slowly across the compact dataset and left validation loss elevated (≥ 0.589). Increasing the learning rate to 2×10^{-5} or 3×10^{-5} achieved an ideal convergence equilibrium, allowing the classification head and upper transformer layers to adapt smoothly without destabilizing lower-layer syntactic representations. Conversely, an aggressive rate of $\eta = 5 \times 10^{-5}$ introduced volatility; large gradient updates computed over small mini-batches caused loss oscillations, premature validation divergence, and an average 5.8% reduction in validation Macro F_1 .

Finally, evaluating batch size dynamics ($B = 8$ versus $B = 16$) demonstrated that smaller mini-batches offer a valuable implicit regularization benefit. In few-shot legal classification, mini-batches of size 8 introduce beneficial stochastic gradient noise that discourages optimizers from settling into sharp local minima. Across all evaluated model families, this setting consistently produced 2.5% to 4.0% higher validation Macro F_1 scores compared to batch size 16.

4.4 System Integration and Usability

The empirical evaluations conducted in this chapter validate the operational viability of the coarse-to-fine architecture while mapping out clear performance bounds across both processing stages. In Stage 1, the hybrid retrieval pipeline effectively handles candidate retrieval across the national statutory corpus. By combining sparse lexical indexing with dense semantic representations under Reciprocal Rank Fusion ($k = 60$) alongside the Hierarchy Priority Safeguard, Stage 1 achieves 100.0% candidate Recall@50 with a mean CPU latency of 1.26 milliseconds. This prevents provision attrition across the evaluated benchmark when handing shortlisted statutory provisions forward to downstream inference.

In Stage 2, the initial screening evaluation across 28 scouted encoder architectures exposed the widespread presence of affirmative bias in unadapted language models, while confirming that pre-aligned NLI reasoners (DeBERTa-v3) and modern long-context encoders (ModernBERT) establish the strongest inductive foundations for legal preemption detection. Controlled supervised fine-tuning across the four candidate ablation axes demonstrated that task-specific adaptation and multi-task NLI initialization yield substantial empirical gains over canonical BERT baselines, with regularized training converging smoothly at Epoch 3 before encountering early-stopping boundaries.

In accordance with the phased research plan, these Stage 2 machine learning findings reflect initial screening, ablation, and hyperparameter optimization sweeps on the curated 350-pair ground truth benchmark. The complete end-to-end evaluation of Stage 2 inference accuracy, calibrated decision thresholds (τ^*), and XAI token saliency maps will be finalized upon the conclusion of the independent five-panel human evaluation currently deployed with the fifteen Sangguniang Panlungsod legal researchers. Once these expert annotations are received and inter-annotator agreement metrics are verified, the elite six-model finalist pool—comprising the Top 5 empirical validation performers (DeBERTa-v3-large-NLI, DeBERTa-v3-base-NLI, ModernBERT-large, cold deberta-v3-base, and PoL-BERT-Large) alongside the canonical bert-base-uncased baseline anchor—will undergo final out-of-sample benchmark testing to determine the definitive production engine, establish deployment trade-offs between large and base parameter scales, and complete system-level integration.

Following model finalization, the complete coarse-to-fine pipeline will be executed across the three evaluation tiers established in Section 3.10. In Tier 1 (Cascaded End-to-End Benchmark), the unassisted pipeline will be evaluated across the held-out test split, calculating the joint hit metric (Hit_{casc}) across the Top-50 candidate shortlist ($k = 50$) to measure cascaded retrieval-inference fidelity without manual statutory priming. In Tier 2 (Injected-Fault Full-Document Stress Testing), full-length draft municipal ordinances will be audited against the statutory archive, assessing Section-Level Sensitivity across injected preemption defects and measuring Section-Level Specificity ($\text{Specificity}_{\text{sec}}$) across compliant provisions to quantify system resistance against false alarms and alert fatigue. Finally, in Tier 3 (Historical Jurisprudential Validation and LGU User Study), the system will audit the eight landmark Davao City municipal preemption controversies (`data/tier3_jurisprudential_cases.jsonl`) to benchmark automated citations and ra-

tionales against authentic judicial holdings, culminating in the formal usability study with the 15 Sangguniang Panlungsod legal researchers using the prototype dashboard to record System Usability Scale (SUS) ratings and Fleiss' Kappa inter-annotator consensus.

Bibliography

- [1] ABLAZO, A. C. Designing a web application using the Twitter streaming API as a tracking tool for weather updates. Graduate independent study, Ateneo de Davao University, Davao City, Philippines, February 2019.
- [2] AGGARWAL, C. C., AND ZHAI, C. *Mining Text Data*. Springer Science+Business Media, Boston, MA, 2012.
- [3] ÅGOTNES, T., HOEK, W., RODRÍGUEZ-AGUILAR, J., SIERRA, C., AND WOOLDRIDGE, M. On the logic of normative systems. *IJCAI International Joint Conference on Artificial Intelligence* (01 2007).
- [4] AGPALO, R. E. *Statutory Construction*, 6th ed. Rex Book Store, Manila, Philippines, 2009.
- [5] ALSHEHRI, A. S., BENCOMO, N., AND ATAPOUR-ABARGHOUEI, A. Cross-architecture LLM ensembles, feature-based reranking and retrieval-augmented prompting for legal information processing. In *Proceedings of the Workshop on the Thirteenth International Competition on Legal Information Extraction and Entailment (COLIEE 2026)* (Singapore, 2026), pp. 1–10.
- [6] ALTHAMMER, S., ASKARI, A., VERBERNE, S., AND HANBURY, A. Dossier@coliee 2021: Leveraging dense retrieval and summarization-based re-ranking for case law retrieval, 2021.
- [7] ANAND, A., AWEKAR, U., ET AL. Explainable information retrieval: A survey. *arXiv preprint arXiv:2012.14758* (2020).
- [8] ARELLANO LAW FOUNDATION. The LawPhil project: Philippine laws and jurisprudence databank, 2000. Accessed: 2026-04-05.
- [9] ARIAI, F., MACKENZIE, J., AND DEMARTINI, G. Natural language processing for the legal domain: A survey of tasks, datasets, models, and challenges. *ACM Comput. Surv.* 58, 6 (Dec. 2025).
- [10] ARTSTEIN, R., AND POESIO, M. Inter-coder agreement for computational linguistics. *Computational Linguistics* 34 (01 2008), 555–596.
- [11] ASHLEY, K. *Artificial Intelligence and Legal Analytics: New Tools for Law Practice in the Digital Age*. Cambridge University Press, 2017.

- [12] ATHAN, T., BOLEY, H., GOVERNATORI, G., PALMIRANI, M., PASCHKE, A., AND WYNER, A. Legalruleml: From metamodel to use cases - (a tutorial). vol. 8035, pp. 13–18.
- [13] BELTAGY, I., PETERS, M. E., AND COHAN, A. Longformer: The long-document transformer, 2020.
- [14] BENCH-CAPON, T. J. M., AND SARTOR, G. A model of legal reasoning with cases incorporating theories and values. *Artificial Intelligence* 150, 1–2 (2003), 97–143.
- [15] BERNARDO, A. S., AND ALBAÑA-GARRIDO, A. P. Disambiguating philippine republic acts: The case of RA 10913. *International Journal of Legal Discourse* 8, 2 (2023), 221–248.
- [16] BHATTACHARYA, P., HIWARE, K., RAJGARIA, S., POCHHI, N., GHOSH, K., AND GHOSH, S. *A Comparative Study of Summarization Algorithms Applied to Legal Case Judgments*. 04 2019, pp. 413–428.
- [17] BLAIR, D. C., AND MARON, M. E. An evaluation of retrieval effectiveness for a full-text document-retrieval system. *Commun. ACM* 28, 3 (Mar. 1985), 289–299.
- [18] BOWMAN, S. R., ANGELI, G., POTTS, C., AND MANNING, C. D. A large annotated corpus for learning natural language inference. In *Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing* (2015), pp. 632–642.
- [19] BOWMAN, S. R., ANGELI, G., POTTS, C., AND MANNING, C. D. A large annotated corpus for learning natural language inference. In *Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing* (Lisbon, Portugal, Sept. 2015), L. Màrquez, C. Callison-Burch, and J. Su, Eds., Association for Computational Linguistics, pp. 632–642.
- [20] BROOKE, J. SUS: A quick and dirty usability scale. *Usability Evaluation in Industry* 189, 194 (1996), 4–7.
- [21] CAMPOS, R., EVANS, J. P., ISIDRO, J. M., MARQUES, M., CUNHA, L. F., JORGE, A., NUNES, S., AND GUIMARÃES, N. Nlp for local governance meeting records: A focus article on tasks, datasets, metrics and benchmark, 2026.
- [22] CAMPOS, R., PACHECO, A. F., FERNANDES, A. L., CANTANTE, I., REBOUÇAS, R., CUNHA, L. F., ISIDRO, J., EVANS, J., MARQUES, M., BATISTA, R., AMORIM, E., JORGE, A., GUIMARÃES, N., NUNES, S., LEAL, A., AND SILVANO, P. *CitiLink-Minutes: A Multilayer Annotated Dataset of Municipal Meeting Minutes*. Springer Nature Switzerland, 2026, p. 511–527.
- [23] CARD, D., HENDERSON, P., KHANDELWAL, U., JIA, R., MAHOWALD, K., AND JURAFSKY, D. With little power comes great responsibility. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)* (Online, Nov. 2020), B. Webber, T. Cohn, Y. He, and Y. Liu, Eds., Association for Computational Linguistics, pp. 9263–9274.

- [24] CARUANA, R., NICULESCU-MIZIL, A., CREW, G., AND KSIKES, A. Ensemble selection from libraries of models. In *Proceedings of the Twenty-First International Conference on Machine Learning (ICML)* (2004), p. 18.
- [25] CHALKIDIS, I., FERGADIOTIS, M., MALAKASIoTIS, P., ALETRAS, N., AND ANDROUTSOPOULOS, I. LEGAL-BERT: The muppets straight out of law school. In *Findings of the Association for Computational Linguistics: EMNLP 2020* (Online, Nov. 2020), T. Cohn, Y. He, and Y. Liu, Eds., Association for Computational Linguistics, pp. 2898–2904.
- [26] CHALKIDIS, I., FERGADIOTIS, M., MALAKASIoTIS, P., AND ANDROUTSOPOULOS, I. Large-scale multi-label text classification on eu legislation, 2019.
- [27] CHALKIDIS, I., JANA, A., HARTUNG, D., BOMMARITO, M., ANDROUTSOPOULOS, I., KATZ, D., AND ALETRAS, N. LexGLUE: A benchmark dataset for legal language understanding in English. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)* (Dublin, Ireland, May 2022), S. Muresan, P. Nakov, and A. Villavicencio, Eds., Association for Computational Linguistics, pp. 4310–4330.
- [28] CHAWLA, N. V., BOWYER, K. W., HALL, L. O., AND KEGELMEYER, W. P. SMOTE: Synthetic minority over-sampling technique. *Journal of Artificial Intelligence Research* 16 (2002), 321–357.
- [29] CHEN, J., ET AL. BGE M3-embedding: Multi-lingual, multi-functionality, multi-granularity text embeddings through self-knowledge distillation. *arXiv preprint arXiv:2402.03216* (2024).
- [30] CITY GOVERNMENT OF DAVAO. City Upgrades Legislative Information Support System, August 2025. Accessed: February 13, 2026.
- [31] CITY GOVERNMENT OF DAVAO. Sangguniang panlungsod - department of public administration, 2026. Accessed: 2026-04-06.
- [32] CLARK, K., LUONG, M.-T., LE, Q. V., AND MANNING, C. D. ELECTRA: Pre-training text encoders as discriminators rather than generators. In *International Conference on Learning Representations (ICLR 2020)* (2020).
- [33] COHEN, J. A coefficient of agreement for nominal scales. *Educational and psychological measurement* 20, 1 (1960), 37–46.
- [34] COHEN, J. *Statistical Power Analysis for the Behavioral Sciences*, second ed. Lawrence Erlbaum Associates, Hillsdale, NJ, 1988.
- [35] CONCHAS, C. A. J., AND SALAPA, A. C. Assessing digital governance implementation issues in local government units: A systematic literature review. *International Journal of Engineering Technology Research and Management* 9, 12 (December 2025), 198–208.
- [36] CONNEAU, A., ET AL. Unsupervised cross-lingual representation learning at scale. *arXiv preprint arXiv:1911.02116* (2020).

- [37] CORMACK, G. V., CLARKE, C. L. A., AND BÜTTCHER, S. Reciprocal rank fusion outperforms Condorcet and individual machine learning methods. In *Proceedings of the 32nd International ACM SIGIR Conference on Research and Development in Information Retrieval* (2009), pp. 758–759.
- [38] CORTEZ, K., AND DAVAO TODAY. Leadership and city problems await newly-elected Davao councilors. Rappler, May 2025. Republished with permission from DavaoToday.com for the 2025 Philippine local elections.
- [39] CRUZ, I. A. *Philippine Political Law*. Central Lawbook Publishing, Quezon City, 2007.
- [40] CRUZ, J. C. B., AND CHENG, C. Establishing baselines for text classification and question answering in Tagalog. In *Proceedings of the 1st Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics (AACL 2020)* (2020), Association for Computational Linguistics, pp. 40–45.
- [41] DAVIES, D. L., AND BOULDIN, D. W. A cluster separation measure. *IEEE Transactions on Pattern Analysis and Machine Intelligence PAMI-1*, 2 (1979), 224–227.
- [42] DEERWESTER, S., DUMAIS, S. T., FURNAS, G. W., LANDAUER, T. K., AND HARSHMAN, R. Indexing by latent semantic analysis. *Journal of the American Society for Information Science* 41, 6 (1990), 391–407.
- [43] DEMŠAR, J. Statistical comparisons of classifiers over multiple data sets. *Journal of Machine Learning Research* 7 (2006), 1–30.
- [44] DEPARTMENT OF THE INTERIOR AND LOCAL GOVERNMENT. Local legislative drafting manual: A technical guide for philippine sangguniang panlungsod and sangguniang bayan members. Tech. rep., Bureau of Local Government Development, DILG, Quezon City, Philippines, 2022.
- [45] DERROY, A., GHOSH, K., AND GHOSH, S. How ready are pre-trained abstractive models and llms for legal case judgement summarization?, 2023.
- [46] DETTMERS, T., LEWIS, M., BELKADA, Y., AND ZETTLEMOYER, L. Llm.int8(): 8-bit matrix multiplication for transformers at scale, 2022.
- [47] DETTMERS, T., PAGNONI, A., HOLTZMAN, A., AND ZETTLEMOYER, L. QLoRA: Efficient finetuning of quantized LLMs. In *Advances in Neural Information Processing Systems (NeurIPS 2023)* (2024), vol. 36, pp. 10088–10115.
- [48] DEVLIN, J., CHANG, M.-W., LEE, K., AND TOUTANOVA, K. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)* (Minneapolis, Minnesota, jun 2019), J. Burstein, C. Doran, and T. Solorio, Eds., Association for Computational Linguistics, pp. 4171–4186.

- [49] DIAS, J., SANTOS, P. A., CORDEIRO, N., ANTUNES, A., MARTINS, B., BAPTISTA, J., AND GONÇALVES, C. State of the art in artificial intelligence applied to the legal domain, 2022.
- [50] DIETTERICH, T. G. Approximate statistical tests for comparing supervised classification learning algorithms. *Neural Computation* 10, 7 (1998), 1895–1923.
- [51] DODGE, J., GURURANGAN, S., CARD, D., SCHWARTZ, R., AND SMITH, N. A. Show your work: Improved reporting of experimental results. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)* (2019), pp. 2185–2194.
- [52] DOSHI-VELEZ, F., AND KIM, B. Towards a rigorous science of interpretable machine learning. *arXiv preprint arXiv:1702.08608* (2017).
- [53] DREISBACH, J. L., AND DEMETERIO, F. Language use and preference in the multilingual context of davao city, philippines. *Studies in English Language and Education* 8 (01 2021), 313–327.
- [54] DWORK, C., FELDMAN, V., HARDT, M., PITASSI, T., REINGOLD, O., AND ROTH, A. The reusable holdout: Preserving validity in adaptive data analysis. *Science* 349, 6248 (2015), 636–638.
- [55] ELKAN, C. The foundations of cost-sensitive learning. In *Proceedings of the 17th International Joint Conference on Artificial Intelligence* (2001), vol. 17, Morgan Kaufmann Publishers Inc., pp. 973–978.
- [56] FAUL, F., ERDFELDER, E., LANG, A.-G., AND BUCHNER, A. G*Power 3: A flexible statistical power analysis program for the social, behavioral, and biomedical sciences. *Behavior Research Methods* 39, 2 (2007), 175–191.
- [57] FAWCETT, T. An introduction to ROC analysis. *Pattern Recognition Letters* 27, 8 (2006), 861–874.
- [58] FLEISS, J. L. Measuring nominal scale agreement among many raters. *Psychological bulletin* 76, 5 (1971), 378.
- [59] FRIED, B. H. Ex ante and ex post. *Stanford Law Review* (2015).
- [60] FRIEDMAN, M. The use of ranks to avoid the assumption of normality implicit in the analysis of variance. *Journal of the American Statistical Association* 32, 200 (1937), 675–701.
- [61] FULLER, L. L. *The Morality of Law*, revised ed. Yale University Press, New Haven, 1964.
- [62] GABUYA, R. J. B. Towards digital legislation: Readiness of the Sangguniang Bayan of LGU motioning for paperless sessions. *International Journal of Research and Innovation in Social Science* 9, 9 (September 2025), 2980–2991.
- [63] GAO, T., FISCH, A., AND CHEN, D. Making pre-trained language models better few-shot learners, 2021.

- [64] GAO, Y., DLIGACH, D., MILLER, T., TESCH, S., LAFFIN, R., CHURPEK, M. M., AND AFSHAR, M. Hierarchical annotation for building a suite of clinical natural language processing tasks: Progress note understanding, 2022.
- [65] GOEBEL, R., KANO, Y., KIM, M.-Y., KWAN, C., SATOH, K., YAMADA, H., AND YOSHIOKA, M. An overview of the coliee 2025 competition: Legal case law and statute law information retrieval and entailment. In *Proceedings of the Twentieth International Conference on Artificial Intelligence and Law* (New York, NY, USA, 2026), ICAIL '25, Association for Computing Machinery, p. 506–515.
- [66] GOEBEL, R., KANO, Y., KURIBAYASHI, T., MARTINEZ-GIL, J., RABELO, J., SATOH, K., SILVEIRA, L., AND YOSHIOKA, M., Eds. *Proceedings of the Twelfth International Competition on Legal Information Extraction and Entailment (COLIEE 2025)* (San Luis, Argentina, 2025).
- [67] GOEBEL, R., KANO, Y., KURIBAYASHI, T., MARTINEZ-GIL, J., RABELO, J., SATOH, K., SILVEIRA, L., AND YOSHIOKA, M., Eds. *Proceedings of the Thirteenth International Competition on Legal Information Extraction and Entailment (COLIEE 2026)* (Rome, Italy, 2026).
- [68] GOODFELLOW, I., BENGIO, Y., AND COURVILLE, A. *Deep Learning*. MIT Press, 2016. <http://www.deeplearningbook.org>.
- [69] GOOGLE CLOUD. Document AI and cloud vision API architecture overview. <https://cloud.google.com/document-ai/docs/overview>, 2024.
- [70] GOVERNMENT OF THE PHILIPPINES. Republic Act No. 7160: The Local Government Code of 1991. Official Gazette of the Republic of the Philippines, 1991.
- [71] GOVERNMENT OF THE PHILIPPINES. Republic act no. 11032: Ease of doing business and efficient government service delivery act of 2018. Tech. rep., Official Gazette of the Republic of the Philippines, 2018.
- [72] GOVERNMENT OF THE PHILIPPINES. Official gazette of the republic of the philippines, 2026. Accessed: 2026-04-05.
- [73] GRECO, C. M., AND TAGARELLI, A. Bringing order into the realm of transformer-based language models for artificial intelligence and law. *Artificial Intelligence and Law* 32, 4 (Nov. 2023), 863–1010.
- [74] GROOTENDORST, M. BERTopic: Neural topic modeling with a class-based TF-IDF procedure. *arXiv preprint arXiv:2203.05794* (2022).
- [75] GUO, C., PLEISS, G., SUN, Y., AND WEINBERGER, K. Q. On calibration of modern neural networks, 2017.
- [76] GUO, C., PLEISS, G., SUN, Y., AND WEINBERGER, K. Q. On calibration of modern neural networks. In *Proceedings of the 34th International Conference on Machine Learning* (2017), PMLR, pp. 1321–1330.
- [77] GURURANGAN, S., SWAYAMDIPTA, S., LEVY, O., SCHWARTZ, R., BOWMAN, S., AND SMITH, N. A. Annotation artifacts in natural language inference data. In *Pro-*

- ceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 2 (Short Papers)* (New Orleans, Louisiana, June 2018), M. Walker, H. Ji, and A. Stent, Eds., Association for Computational Linguistics, pp. 107–112.
- [78] HANNES, W., ŠA VELKA JAROMÍR, R., W. V., D., A. K., AND KARIM, B. *Computer-Assisted Creation of Boolean Search Rules for Text Classification in the Legal Domain*. IOS Press, 2019.
- [79] HARDE, P., JAIN, B., AND JAIN, S. LeCNet: A legal citation network benchmark dataset. In *Proceedings of the 1st Workshop on NLP for Empowering Justice (JUST-NLP 2025)* (Mumbai, India, Dec. 2025), A. Modi, S. Ghosh, A. Ekbal, P. Goyal, S. Jain, A. Joshi, S. Mishra, D. Datta, S. Paul, K. B. Singh, and S. Kumar, Eds., Association for Computational Linguistics, pp. 18–28.
- [80] HE, P., GAO, J., AND CHEN, W. DeBERTaV3: Improving DeBERTa using ELECTRA-style pre-training with gradient-disentangled embedding sharing. *arXiv preprint arXiv:2111.09543* (2023).
- [81] HENDERSON, P., KRASS, M., ZHENG, L., GUHA, N., MANNING, C. D., JURAFSKY, D., AND HO, D. E. Pile of Law: Learning responsible data filtering from the law and a 256GB open-source legal dataset. *arXiv preprint arXiv:2207.00220* (2022).
- [82] HOLISTIC SEO. Sliding window technique and algorithm in NLP, 2025.
- [83] HOLZENBERGER, N., BLAIR-STANEK, A., AND DURME, B. V. A dataset for statutory reasoning in tax law entailment and question answering, 2020.
- [84] HOUSE OF REPRESENTATIVES OF THE PHILIPPINES. Legislative information system (LEGIS) open data API. Official Website of the House of Representatives, 2023. Accessed: 2026-04-05.
- [85] HUANG, H., ZHANG, H., WANG, Y., LIU, H., CHEN, X., CHEN, Y., AND LIANG, Y. Integrated robust optimization for lightweight transformer models in low-resource scenarios. *Symmetry* 17, 7 (2025).
- [86] HUMEAU, S., SHUSTER, K., VECCHIATO, M. A., AND WESTON, J. Poly-encoders: Architectures and pre-training strategies for fast and accurate multi-sentence scoring. *arXiv preprint arXiv:1905.01969* (2019).
- [87] HUMPHREYS, L., BOELLA, G., DI CARO, L., ROBALDO, L., VAN DER TORRE, L., GHANAVATI, S., AND MUTHURI, R. Populating legal ontologies using semantic role labeling. In *Proceedings of the Twelfth Language Resources and Evaluation Conference (LREC 2020)* (2020), European Language Resources Association, pp. 2157–2166.
- [88] ITALIANI, P., MORO, G., AND RAGAZZI, L. Clash-of-leges: A bilingual dataset for conflict detection and explanation in statutory law. *Expert Systems with Applications* 300 (2026), 130182.

- [89] JACOVI, A., AND GOLDBERG, Y. Towards faithfully interpretable nlp systems: How should we define and evaluate faithfulness?, 2020.
- [90] JAIN, S., AND WALLACE, B. C. Attention is not Explanation. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)* (Minneapolis, Minnesota, June 2019), J. Burstein, C. Doran, and T. Solorio, Eds., Association for Computational Linguistics, pp. 3543–3556.
- [91] JURAFSKY, D., AND MARTIN, J. H. *Speech and Language Processing: An Introduction to Natural Language Processing, Computational Linguistics, and Speech Recognition*, 3rd ed. Pearson, 2024. Draft available online; accessed: 2026-03-21.
- [92] KADOWAKI, K., AND KANO, Y. A ModernBERT-based system by team KIS for the COLIEE 2025 pilot task: Toward robust evaluation in legal judgment prediction. In *Proceedings of the Workshop on the Twelfth International Competition on Legal Information Extraction and Entailment (COLIEE 2025)* (Chicago, USA, 2025), ACM, pp. 14–22.
- [93] KARPUKHIN, V., OĞUZ, B., MIN, S., LEWIS, P., WU, L., EDUNOV, S., CHEN, D., AND TAU YIH, W. Dense passage retrieval for open-domain question answering, 2020.
- [94] KATZ, D. M., BOMMARITO, M. J., GAO, S., AND ARREDONDO, P. Gpt-4 passes the bar exam. *Philosophical Transactions of the Royal Society A: Mathematical, Physical and Engineering Sciences* 382, 2270 (02 2024), 20230254.
- [95] KATZ, D. M., HARTUNG, D., GERLACH, L., JANA, A., AND BOMMARITO, M. J. Natural language processing in the legal domain. *arXiv preprint arXiv:2302.12039* (2023).
- [96] KELSEN, H. *Pure Theory of Law*. University of California Press, Berkeley, CA, 1967.
- [97] KHATTAB, O., AND ZAHARIA, M. ColBERT: Efficient and effective passage search via contextualized late interaction over BERT. In *Proceedings of the 43rd International ACM SIGIR Conference on Research and Development in Information Retrieval* (2020), pp. 39–48.
- [98] KIM, M.-Y., RABELO, J., BABIKER, H., RAHMAN, M. A., AND GOEBEL, R. Legal information retrieval and entailment using transformer-based approaches. *The Review of Socionetwork Strategies* 18 (01 2024), 1–21.
- [99] KLEBANOV, B., AND BEIGMAN, E. From annotator agreement to noise models. *Computational Linguistics* 35 (12 2009), 495–503.
- [100] KOREEDA, Y., AND MANNING, C. D. ContractNLI: A dataset for document-level natural language inference for contracts. In *Findings of the Association for Computational Linguistics: EMNLP 2021* (2021), Association for Computational Linguistics, pp. 1907–1919.

- [101] KUMAR, I. E., VENKATASUBRAMANIAN, S., SCHEIDEGGER, C., AND FRIEDLER, S. Problems with shapley-value-based explanations as feature importance measures. In *International Conference on Machine Learning* (2020), PMLR, pp. 5491–5500.
- [102] LAN, Z., CHEN, M., GOODMAN, S., GIMPEL, K., SHARMA, P., AND SORICUT, R. ALBERT: A lite BERT for self-supervised learning of language representations. In *International Conference on Learning Representations (ICLR 2020)* (2020).
- [103] LANDAUER, T. K., FOLTZ, P. W., AND LAHAM, D. An introduction to latent semantic analysis. *Discourse Processes* 25, 2-3 (1998), 259–284.
- [104] LANDIS, J. R., AND KOCH, G. G. The measurement of observer agreement for categorical data. *Biometrics* 33, 1 (1977), 159–174.
- [105] LARGEY, N., FITZGERALD, E., AND MANSOURI, B. AIIR lab at COLIEE 2026: Fusion of instruction-tuned LLMs and cross-encoders for legal case entailment. In *Proceedings of the Workshop on the Thirteenth International Competition on Legal Information Extraction and Entailment (COLIEE 2026)* (Singapore, 2026), pp. 89–99.
- [106] LAURER, M., VAN ATTEVELDT, W., SALLOUM, A., AND WELBERS, K. Less annotating, more classifying: Addressing the data scarcity issue of supervised machine learning with deep transfer learning and multi-task NLI pre-alignment. In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (EACL 2024)* (2024), Association for Computational Linguistics, pp. 2123–2139.
- [107] LE-HONG, P., AND CAMBRIA, E. A semantics-aware approach for multilingual natural language inference. *Lang. Resour. Eval.* 57, 2 (Feb. 2023), 611–639.
- [108] LEONEN, M. M. SAJ leonen: Despite risks, legal system should keep abreast with AI developments. Keynote Speech at Manila Tech Summit 2024, August 2024.
- [109] LEVENSHTAIN, V. I. Binary codes capable of correcting deletions, insertions, and reversals. *Soviet Physics Doklady* 10, 8 (1966), 707–710.
- [110] LI, J., QIAN, L., LIU, P., AND LIU, T. Construction of legal knowledge graph based on knowledge-enhanced large language models. *Information* 15 (10 2024), 666.
- [111] LI, T., BALKE, T., DE VOS, M., SATOH, K., AND PADGET, J. Detecting conflicts in legal systems. pp. 174–189.
- [112] LIDDY, E. D. Natural language processing. In *Encyclopedia of Library and Information Science*, 2nd ed. Marcel Dekker, Inc., 2001.
- [113] LINNA, E., AND LINNA, T. Challenges for generative ai in legal reasoning, 2026.
- [114] LIU, A., SWAMY, V., RAVICHANDER, A., WANG, Y., BHAGAVATULA, C., AND LE BRAS, R. WANLI: Worker and AI Collaborative for Natural Language Inference Dataset Creation. In *Findings of the Association for Computational Linguistics: EMNLP 2022* (2022), Association for Computational Linguistics, pp. 6826–6841.

- [115] LOPRESTI, D. Optical character recognition errors and their effect on natural language processing. *International Journal on Document Analysis and Recognition (IJDAR)* 10, 3–4 (2008), 141–151.
- [116] LOSHCHILOV, I., AND HUTTER, F. Decoupled weight decay regularization. In *International Conference on Learning Representations (ICLR)* (2019).
- [117] LOUIS, A., AND SPANAKIS, G. A statutory article retrieval dataset in French. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)* (Dublin, Ireland, May 2022), S. Muresan, P. Nakov, and A. Villavicencio, Eds., Association for Computational Linguistics, pp. 6789–6803.
- [118] LU, K., HSU, I.-H., ZHOU, W., MA, M. D., AND CHEN, M. Multi-hop evidence retrieval for cross-document relation extraction. In *Findings of the Association for Computational Linguistics: ACL 2023* (Toronto, Canada, July 2023), A. Rogers, J. Boyd-Graber, and N. Okazaki, Eds., Association for Computational Linguistics, pp. 10336–10351.
- [119] LUCOS, R., L P, S., P R, V., GIRISH, D., M, R., KURUP, P. S., M S, C., AND SALEELA, D. Neuro-Symbolic Artificial Intelligence: Foundations, Advances, and Future Directions. *SSRN Electronic Journal* (December 2025).
- [120] LUNDBERG, S. M., AND LEE, S.-I. A unified approach to interpreting model predictions. In *Advances in Neural Information Processing Systems*, vol. 30. 2017.
- [121] MA, H. Gnn-ere: A graph neural network framework for entity-relation extraction and intelligent legal case reasoning. *Informatica* 49 (12 2025).
- [122] MANNING, C. D., RAGHAVAN, P., AND SCHÜTZE, H. *Introduction to Information Retrieval*. Cambridge University Press, Cambridge, UK, 2008.
- [123] MAXILOM, R. M. R. Semantic change of the selected Cebuano words. In *Proceedings of the 22nd Pacific Asia Conference on Language, Information and Computation* (2008), pp. 433–438.
- [124] MCCOY, T., PAVLICK, E., AND LINZEN, T. Right for the wrong reasons: Diagnosing syntactic heuristics in natural language inference. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics* (2019), Association for Computational Linguistics, pp. 3428–3448.
- [125] MCHUGH, M. L. Interrater reliability: The kappa statistic. *Biochemia Medica* 22, 3 (2012), 276–282.
- [126] MCNEMAR, Q. Note on the sampling error of the difference between correlated proportions or percentages. *Psychometrika* 12, 2 (1947), 153–157.
- [127] MORRISON, D. G. On the interpretation of discriminant analysis. *Journal of Marketing Research* 6, 2 (1969), 156–163.
- [128] MUENNIGHOFF, N., ET AL. MTEB: Massive text embedding benchmark. *arXiv preprint arXiv:2210.07316* (2023).

- [129] NAGUIO, A. A. A., AND ROXAS, R. E. O. Hybrid neural-rule based architectures for Filipino stemming with fine-tuned BERT variants. In *Proceedings of the 38th Pacific Asia Conference on Language, Information and Computation* (Tokyo, Japan, Dec. 2024), N. Oco, S. N. Dita, A. M. Borlongan, and J.-B. Kim, Eds., Tokyo University of Foreign Studies, pp. 675–685.
- [130] NAIK, A., RAVICHANDER, A., SADEH, N., ROSE, C., AND NEUBIG, G. Stress test evaluation for natural language inference. In *Proceedings of the 27th International Conference on Computational Linguistics* (2018), Association for Computational Linguistics, pp. 2340–2353.
- [131] NATIONAL ASSEMBLY OF THE PHILIPPINES. Commonwealth act no. 638: An act to provide for the uniform publication and distribution of the official gazette. Supreme Court E-Library, June 1941.
- [132] NATIONAL RESEARCH COUNCIL OF THE PHILIPPINES. PH local governance gets reviewed after three decades, 2024.
- [133] NGO, T.-H., NGUYEN, H.-T., NGUYEN, H.-D., LE, X.-B., NGUYEN, L.-D., TRAN, Q.-T., NGUYEN, H.-T., AND VUONG, T.-H.-Y. NOWJ@COLIEE 2026: Adaptive pipelines for legal retrieval and reasoning. In *Proceedings of the Workshop on the Thirteenth International Competition on Legal Information Extraction and Entailment (COLIEE 2026)* (Singapore, 2026), pp. 119–128.
- [134] NIE, Y., WILLIAMS, A., DINAN, E., BANSAL, M., WESTON, J., AND KIELA, D. Adversarial NLI: A new benchmark for natural language understanding. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics* (2020), Association for Computational Linguistics, pp. 4885–4901.
- [135] NUÑEZ, K. B. Q., ABCEDE, M. P. G., SALAZAR, L. R. R., CHUA, Y. T., ROMERO II, V. M., GABUD, R. S., AND REGONIA, P. R. R. POMI: A corpus to support a natural language inference-based approach for detecting misinformation in Philippine online news. *Datenbank-Spektrum* 26 (2026), 25–37.
- [136] NUSSBAUM, Z., ET AL. Nomic embed: Training a reproducible long context text embedder. *arXiv preprint arXiv:2402.01613* (2024).
- [137] OARD, D. W., ET AL. Desperately seeking Cebuano. In *Proceedings of the ACM SIGIR Workshop on Analytics for Noisy Unstructured Text Data* (2003).
- [138] PARUCHURI, V. Surya: Multilingual document OCR, layout analysis, and line detection. <https://github.com/VikParuchuri/surya>, 2024.
- [139] PATUMBON, R. G. G. Davao City Council Passes Landmark Measures. SunStar Davao, February 2025. SunStar Davao Special Report on 20th and 21st Sangguniang Panlungsod legislative output.
- [140] PATUMBON, R. G. G. SP implements less paper with LISSP. SunStar Davao, February 2025. Reporting on official announcement by Davao City Councilor Bonz Andre Militar.

- [141] PAUL, S., MANDAL, A., GOYAL, P., AND GHOSH, S. Pre-trained Language Models for the Legal Domain: A Case Study on Indian Law. In *Proceedings of the 19th International Conference on Natural Language Processing (ICON)* (2022), Association for Computational Linguistics, pp. 253–262.
- [142] PHILIPPINE COMMISSION. Act no. 453: An act providing for the publication by the insular government of an official gazette under the general direction of the department of public instruction. Supreme Court E-Library, September 1902.
- [143] POLIAK, A., NARADOWSKY, J., APARICIO, A., DUH, K., AND VAN DURME, B. Hypothesis only baselines in natural language inference. In *Proceedings of the 7th Joint Conference on Lexical and Computational Semantics* (2018), Association for Computational Linguistics, pp. 180–191.
- [144] PRECHELT, L. Early stopping—but when? In *Neural Networks: Tricks of the Trade*, vol. 1524 of *Lecture Notes in Computer Science*. Springer, 1998, pp. 55–69.
- [145] QUEVEDO, E., CERNY, T., RODRIGUEZ, A., RIVAS, P., YERO, J., SOOKSATRA, K., ZHAKUBAYEV, A., AND TAIBI, D. Legal natural language processing from 2015 to 2022: A comprehensive systematic mapping study of advances and applications. *IEEE Access* 12 (2024), 145286–145317.
- [146] RABELO, J., GOEBEL, R., KIM, M.-Y., KANO, Y., YOSHIOKA, M., AND SATOH, K. Overview and discussion of the competition on legal information extraction/entailment (coliee) 2021. *The Review of Socionetwork Strategies* 16, 1 (April 2022), 111–133.
- [147] RAJPURKAR, P., JIA, R., AND LIANG, P. Know what you don’t know: Unanswerable questions for SQuAD. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics* (2018), Association for Computational Linguistics, pp. 784–789.
- [148] RANATHUNGA, S., AND DE SILVA, N. Some languages are more equal than others: Probing deeper into the linguistic disparity in the NLP world. In *Proceedings of the 2nd Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics and the 12th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)* (Online only, Nov. 2022), Y. He, H. Ji, S. Li, Y. Liu, and C.-H. Chang, Eds., Association for Computational Linguistics, pp. 823–848.
- [149] REGANIT, J. C. Diokno files bill to break language barriers in PH laws. Philippine News Agency, August 2025. Report on House Bill No. 3863 (Batas sa Sariling Wika Act).
- [150] REIMERS, N., AND GUREVYCH, I. Sentence-bert: Sentence embeddings using siamese bert-networks, 2019.
- [151] REPUBLIC OF THE PHILIPPINES. Executive order no. 292 [administrative code of 1987], book VII, chapter 1: General provisions. Official Gazette of the Republic of the Philippines, July 1987. Accessed: 2026-03-21.

- [152] REPUBLIC OF THE PHILIPPINES. Executive order no. 292: Instituting the Administrative Code of 1987. Official Gazette of the Republic of the Philippines, July 1987. Book I, Chapter 5, Section 20.
- [153] REPUBLIC OF THE PHILIPPINES. Republic act no. 10173: An act protecting individual personal information in information and communications systems in the government and the private sector (Data Privacy Act of 2012). Official Gazette of the Republic of the Philippines, Vol. 108, No. 35, p. 4398, 2012.
- [154] RIBEIRO, M. T., SINGH, S., AND GUESTIN, C. "why should i trust you?" explaining the predictions of any classifier. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining* (2016), pp. 1135–1144.
- [155] RIBEIRO, M. T., WU, T., GUESTIN, C., AND SINGH, S. Beyond accuracy: Behavioral testing of NLP models with CheckList. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics* (2020), Association for Computational Linguistics, pp. 4902–4912.
- [156] RIGANO, C. Using artificial intelligence to address criminal justice needs. *NIJ Journal* (October 2018). National Institute of Justice.
- [157] ROBERTSON, S., AND ZARAGOZA, H. The probabilistic relevance framework: BM25 and beyond. *Foundations and Trends in Information Retrieval* 3, 4 (2009), 333–389.
- [158] ROQUE, M., ABEJUELA, N., CHU, S., CABATUAN, M., AND SYBINGCO, E. Legal information retrieval through embedding models and synthetic question generation: Insights from the Philippine tax code. In *Proceedings of the 38th Pacific Asia Conference on Language, Information and Computation* (Tokyo, Japan, Dec. 2024), N. Oco, S. N. Dita, A. M. Borlongan, and J.-B. Kim, Eds., Tokyo University of Foreign Studies, pp. 354–362.
- [159] ROUSSEEUW, P. J. Silhouettes: A graphical aid to the interpretation and validation of cluster analysis. *Journal of Computational and Applied Mathematics* 20 (1987), 53–65.
- [160] RUDIN, C. Stop explaining black box machine learning models for high stakes decisions and use interpretable models instead. *Nature Machine Intelligence* 1, 5 (2019), 206–215.
- [161] RUHL, J. B., AND KATZ, D. M. Measuring, monitoring, and managing legal complexity. *Iowa Law Review* 101, 1 (2015), 191–244.
- [162] RUST, P., PFEIFFER, J., VULIĆ, I., RUDER, S., AND GUREVYCH, I. How good is your tokenizer? on the monolingual performance of multilingual language models. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)* (Online, Aug. 2021), C. Zong, F. Xia, W. Li, and R. Navigli, Eds., Association for Computational Linguistics, pp. 3118–3135.

- [163] SAGUM, R. A., CLACIO, P. A. C., CAYETANO, R. E. R., AND LOBRIO, A. D. F. Philippine court case summarizer using latent semantic analysis. *Procedia Comput. Sci.* 227, C (Jan. 2023), 474–481.
- [164] SANGGUNIANG PANLUNGSOD OF DAVAO CITY. Davao City Ordinance No. 0310-07: The Watershed Protection, Conservation and Management Ordinance of Davao City. Official Records of the Sangguniang Panlungsod, City Government of Davao, 2007.
- [165] SANGGUNIANG PANLUNGSOD OF DAVAO CITY. Davao City Ordinance No. 0367-12: The New Comprehensive Anti-Smoking Ordinance of Davao City. Official Records of the Sangguniang Panlungsod, City Government of Davao, 2012.
- [166] SANGGUNIANG PANLUNGSOD OF DAVAO CITY. Resolution No. 001-22: Internal Rules of Procedure of the 20th Sangguniang Panlungsod. Tech. rep., Davao City Government, 2022.
- [167] SANGGUNIANG PANLUNGSOD OF DAVAO CITY. Davao City Ordinance No. 0270-23: The Comprehensive Speed Limit Ordinance of Davao City. Official Records of the Sangguniang Panlungsod, City Government of Davao, 2023.
- [168] SANGGUNIANG PANLUNGSOD OF DAVAO CITY. Official website of the Sangguniang Panlungsod of Davao City, 2026. Accessed: 2026-04-05.
- [169] SANH, V., DEBUT, L., CHAUMOND, J., AND WOLF, T. DistilBERT, a distilled version of BERT: smaller, faster, cheaper and lighter. *arXiv preprint arXiv:1910.01108* (2019).
- [170] SARACEVIC, T. Relevance: A review of the literature and a framework for thinking on the notion in information science. part ii: Nature and manifestations of relevance. *Journal of the American Society for Information Science and Technology* 58 (11 2007), 1915–1933.
- [171] SARTOR, G. *Legal Reasoning: A Cognitive Approach to the Law*, vol. 5 of *A Treatise of Legal Philosophy and General Jurisprudence*. Springer, Dordrecht, 2005.
- [172] SAVELKA, J., AND ASHLEY, K. Discovering explanatory sentences in legal case decisions using pre-trained language models. In *Findings of the Association for Computational Linguistics: EMNLP 2021* (Punta Cana, Dominican Republic, Nov. 2021), M.-F. Moens, X. Huang, L. Specia, and S. W.-t. Yih, Eds., Association for Computational Linguistics, pp. 4273–4283.
- [173] SCHWARTZ, R., DODGE, J., SMITH, N. A., AND ETZIONI, O. Green AI. *Communications of the ACM* 63, 12 (2020), 54–63.
- [174] SCULLEY, D. Web-scale k-means clustering. In *Proceedings of the 19th International Conference on World Wide Web (WWW '10)* (2010), Association for Computing Machinery, pp. 1177–1178.
- [175] SHAO, Y., MAO, J., LIU, Y., MA, W., SATOH, K., ZHANG, M., AND MA, S. Bert-pli: modeling paragraph-level interactions for legal case retrieval. In *Proceed-*

- ings of the Twenty-Ninth International Joint Conference on Artificial Intelligence* (2021), IJCAI'20.
- [176] SHEN, Z., LIU, J., PANG, Y., AND RAO, Y. Hopweaver: Cross-document synthesis of high-quality and authentic multi-hop questions, 2025.
 - [177] SMITH, R. An overview of the Tesseract OCR engine. In *Proceedings of the Ninth International Conference on Document Analysis and Recognition (ICDAR 2007)* (2007), IEEE Computer Society, pp. 629–633.
 - [178] SNOW, R., O’CONNOR, B., JURAFSKY, D., AND NG, A. Y. Cheap and fast—but is it good? evaluating non-expert annotations for natural language tasks. In *Proceedings of the 2008 Conference on Empirical Methods in Natural Language Processing* (Honolulu, Hawaii, October 2008), Association for Computational Linguistics, pp. 254–263.
 - [179] SONG, K., TAN, X., QIN, T., LU, J., AND LIU, T.-Y. Mpnet: Masked and permuted pre-training for language understanding. In *Advances in Neural Information Processing Systems (NeurIPS)* (2020), vol. 33, pp. 16857–16867.
 - [180] SOVRANO, F., VITALI, F., AND PALMIRANI, M. *Making Things Explainable vs Explaining: Requirements and Challenges Under the GDPR*. Springer International Publishing, 2021, p. 169–182.
 - [181] SPERBER, D., AND WILSON, D. *Relevance: Communication and Cognition*, 2nd ed. Blackwell Publishers, Oxford, 1996.
 - [182] SRIVASTAVA, N., HINTON, G., KRIZHEVSKY, A., SUTSKEVER, I., AND SALAKHUTDINOV, R. Dropout: A simple way to prevent neural networks from overfitting. *Journal of Machine Learning Research* 15, 56 (2014), 1929–1958.
 - [183] STRUBELL, E., GANESH, A., AND MCCALLUM, A. Energy and policy considerations for deep learning in NLP. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics (ACL)* (2019), pp. 3645–3650.
 - [184] SUPREME COURT OF THE PHILIPPINES. Tañada v. Tuvera, G.R. no. L-63915. Philippine Reports, Vol. 136, pp. 27–45, Lawphil Project, April 1985. En Banc Decision on the Mandatory Publication of Laws and Presidential Decrees.
 - [185] SUPREME COURT OF THE PHILIPPINES. Tañada v. Tuvera. G.R. No. L-63915, Lawphil Project, April 1985. Accessed: 2026-03-21.
 - [186] SUPREME COURT OF THE PHILIPPINES. Villacorta v. Bernardo (G.R. No. L-31249). Supreme Court Reports Annotated (143 SCRA 480), 1986.
 - [187] SUPREME COURT OF THE PHILIPPINES. Mayor Pablo Magtajas vs. Pryce Properties Corporation, Inc. (G.R. No. 111097). Supreme Court Reports Annotated (234 SCRA 255), 1994.
 - [188] SUPREME COURT OF THE PHILIPPINES. City of Davao v. Court of Appeals and Government Service Insurance System (G.R. No. 127383). Supreme Court Reports Annotated (467 SCRA 280), 2005.

- [189] SUPREME COURT OF THE PHILIPPINES. *City of Manila vs. Hon. Caridad H. Laguio, Jr.* (G.R. No. 118127). *Supreme Court Reports Annotated* (455 SCRA 308), 2005.
- [190] SUPREME COURT OF THE PHILIPPINES. *Mosqueda v. Pilipino Banana Growers & Exporters Association, Inc.* (G.R. Nos. 189185 & 189305). *Supreme Court Reports Annotated* (800 SCRA 313), 2016.
- [191] SUPREME COURT OF THE PHILIPPINES. *Mindanao Shopping Destination Corporation v. Duterte* (G.R. No. 211093). *Supreme Court Reports Annotated* (826 SCRA 202), 2017.
- [192] SUPREME COURT OF THE PHILIPPINES. *Evasco, Jr. v. Montañez* (G.R. No. 199172). *Supreme Court Reports Annotated* (856 SCRA 320), 2018.
- [193] SUPREME COURT OF THE PHILIPPINES. *City of Davao and Tanjili v. ARC Investors, Inc.* (G.R. No. 249668). *Supreme Court of the Philippines En Banc Decision*, 2022.
- [194] SUPREME COURT OF THE PHILIPPINES. *City of Davao v. First Meridian Development, Inc.* (G.R. No. 240078). *Supreme Court of the Philippines Second Division Decision*, 2022.
- [195] TEAM, Q., BAI, J., BAI, S., CHU, Y., CUI, Z., DANG, K., DENG, X., FAN, Y., GAO, L., GE, R., ET AL. Qwen2.5-VL technical report. *arXiv preprint arXiv:2502.13923* (2025).
- [196] THAKUR, N., REIMERS, N., RÜCKLÉ, A., SRIVASTAVA, A., AND GUREVYCH, I. Beir: A heterogenous benchmark for zero-shot evaluation of information retrieval models, 2021.
- [197] THORNE, J., VLACHOS, A., CHRISTODOULOPOULOS, C., AND MITTAL, A. FEVER: a large-scale dataset for fact extraction and VERification. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies* (2018), Association for Computational Linguistics, pp. 809–819.
- [198] TOHIDI, N., DADKHAH, C., AND GELBUKH, A. Abstract meaning representation: A state-of-the-art review. 13–43.
- [199] TRAN, V.-K., AND NGUYEN, K.-L. Enhancing inferential consistency in LLM-based legal entailment. In *Proceedings of the Workshop on the Thirteenth International Competition on Legal Information Extraction and Entailment (COLIEE 2026)* (Singapore, 2026), pp. 170–177.
- [200] TRONG, S. V., AND LE, M. N. An empirical evaluation of amr parsing for legal documents, 2018.
- [201] TSAI, C. F., ZHOU, X., LIU, S. S., LI, J., YU, M., AND MEI, H. Can large language models play text games well? current state-of-the-art and open questions, 2025.

- [202] VABALAS, A., GOWEN, E., POLIAKOFF, E., AND CASSON, A. J. Machine learning algorithm validation with a limited sample size. *PLOS ONE* 14, 11 (2019), e0224365.
- [203] VAN GESTEL, R., AND WIENER, J. B. Future-oriented research into the expected effects of potential new legislation. *Tilburg University Working Paper* (2011).
- [204] VAN NOORD, R., AND BOS, J. The role of data resources and NLP preprocessing in robust AMR parsing. In *Proceedings of the 21st Conference on Computational Natural Language Learning (CoNLL 2017)* (2017), pp. 36–47.
- [205] VAN OPIJNEN, M., AND SANTOS, C. On the concept of relevance in legal information retrieval. *Artif. Intell. Law* 25, 1 (Mar. 2017), 65–87.
- [206] VASWANI, A., SHAZEER, N., PARMAR, N., USZKOREIT, J., JONES, L., GOMEZ, A. N., KAISER, L., AND POLOSUKHIN, I. Attention is all you need. In *Advances in Neural Information Processing Systems* (2017), pp. 5998–6008.
- [207] VELLINO, A. Artificial intelligence: The very idea: J. haugeland, (mit press, cambridge, ma, 1985); 287 pp. *Artificial Intelligence* 29 (09 1986), 349–353.
- [208] VIG, J. A multiscale visualization of attention in the transformer model. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics: System Demonstrations* (2019), pp. 37–42.
- [209] VIRTUCIO, M. B., ABOROT, J., ABONITA, J., AVIÑANTE, R., CO, R., NEVERIDA, M., OSIANA, V., PERAMO, E., SYJUCO, J., AND TAN, G. Predicting decisions of the philippine supreme court using natural language processing and machine learning. pp. 130–135.
- [210] VU, S. T., LE NGUYEN, M., AND SATOH, K. Abstract meaning representation for legal documents: an empirical research on a human-annotated dataset. *Artif. Intell. Law* 30, 2 (June 2022), 221–243.
- [211] WANG, A., PRUKSACHATKUN, Y., NANGIA, N., SINGH, A., MICHAEL, J., HILL, F., LEVY, O., AND BOWMAN, S. R. Superglue: A stickier benchmark for general-purpose language understanding systems, 2020.
- [212] WANG, A., SINGH, A., MICHAEL, J., HILL, F., LEVY, O., AND BOWMAN, S. R. Glue: A multi-task benchmark and analysis platform for natural language understanding, 2019.
- [213] WANG, L., YANG, N., HUANG, X., JIAO, B., YANG, L., JIANG, D., MAJUMDER, R., AND WEI, F. Text embeddings by weakly-supervised contrastive pre-training. *arXiv preprint arXiv:2212.03533* (2022).
- [214] WANG, X., HAN, Y., LEUNG, V. C., NIYATO, D., YAN, X., AND CHEN, X. Convergence of edge computing and deep learning: A comprehensive survey. *IEEE Communications Surveys & Tutorials* 22, 2 (2020), 869–904.
- [215] WARNER, B., CHAFFIN, A., CLAVIÉ, B., WELLER, O., HALLSTRÖM, O., TAGHADOUINI, S., GALLAGHER, A., BISWAS, R., LADHAK, F., AARSEN, T.,

- COOPER, N., ADAMS, G., HOWARD, J., AND POLI, I. Smarter, better, faster, longer: A modern bidirectional encoder for fast, memory efficient, and long context finetuning and inference, 2024.
- [216] WARNER, B., ET AL. ModernBERT: Bringing BERT into the modern era. *arXiv preprint arXiv:2412.13663* (2024).
- [217] WIEGREFFE, S., AND PINTER, Y. Attention is not not explanation, 2019.
- [218] WILLIAMS, A., THRUSH, T., AND KIELA, D. ANLIzing the adversarial natural language inference dataset. In *Proceedings of the Society for Computation in Linguistics 2022* (Online, Feb. 2022), A. Ettinger, T. Hunter, and B. Prickett, Eds., Association for Computational Linguistics, pp. 23–54.
- [219] WU, D., LAWRENCE, S., AND MANSOURI, B. AIIR lab at COLIEE 2025: Exploring applications of large language models for legal text retrieval and entailment. In *Proceedings of the Workshop on the Twelfth International Competition on Legal Information Extraction and Entailment (COLIEE 2025)* (Chicago, USA, 2025), ACM, pp. 112–118.
- [220] XIAO, C., ZHONG, H., GUO, Z., TU, C., LIU, Z., SUN, M., AND SHEN, T. Lawformer: A Pre-trained Language Model for Chinese Legal Long Documents. *AI Open* 2 (2021), 79–84.
- [221] XIAO, S., LIU, Z., ZHANG, P., MUENNIGHOFF, N., LIAN, D., AND NIE, J.-Y. C-Pack: Packaged resources to advance general chinese and multilingual information retrieval, 2024.
- [222] YANG, J., MA, W., ZHANG, M., ZHOU, X., LIU, Y., AND MA, S. Legalgmn: Legal information enhanced graph neural network for recommendation. *ACM Trans. Inf. Syst.* 40, 2 (Sept. 2021).
- [223] YANG, Z. Improving the natural language inference robustness to hard dataset by data augmentation and preprocessing, 2024.
- [224] YIN, W., RADEV, D., AND XIONG, C. DocNLI: A large-scale dataset for document-level natural language inference. In *Findings of the Association for Computational Linguistics: ACL-IJCNLP 2021* (Online, Aug. 2021), C. Zong, F. Xia, W. Li, and R. Navigli, Eds., Association for Computational Linguistics, pp. 4913–4922.
- [225] ZAHEER, M., GURUPRASAD, G., KULKARNI, A., PEROT, D., ALBERTI, C., IZSAK, P., KARCHEVSKY, P., LEPIKHIN, D., PFITZNER, B., ET AL. Big Bird: Transformers for Longer Sequences. *Advances in Neural Information Processing Systems (NeurIPS 2020)* 33 (2020), 17283–17297.
- [226] ZHANG, T., WU, F., KATIYAR, A., WEINBERGER, K. Q., AND ARTZI, Y. Revisiting few-sample bert fine-tuning, 2021.
- [227] ZHENG, L., GUHA, N., ANDERSON, B. R., HENDERSON, P., AND HO, D. E. When does pretraining help? assessing self-supervised learning for law and the casehold dataset, 2021.

- [228] ZHONG, H., XIAO, C., TU, C., ZHANG, T., LIU, Z., AND SUN, M. How does NLP benefit legal system: A summary of legal artificial intelligence. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics* (Online, July 2020), D. Jurafsky, J. Chai, N. Schluter, and J. Tetreault, Eds., Association for Computational Linguistics, pp. 5218–5230.
- [229] ZHONG, X., SHAFIEI BAVANI, E., AND YEPES, A. J. Image-based table recognition: Data, model, and evaluation. In *Computer Vision – ECCV 2020* (2020), vol. 12366 of *Lecture Notes in Computer Science*, Springer, pp. 564–580.