

Neuro-Symbolic Artificial Intelligence: Foundations, Advances, and Future Directions

Reshma Lucos², Supriya L P², Vinod P R³, Dhanya Girish⁴, Reshma M⁵, Parvathy S Kurup⁶, Chinchu M S⁷, Divya Saleela*⁸

reshma.l@esl.ac.uk¹, supriyabinnyb@gmail.com², prvinodmail@gmail.com³, dhanyagirish@skcet.ac.in⁴, reshmamcse24@gmail.com⁵, parvathykurups@gmail.com⁶, chinchums18@gmail.com⁷, d.saleela@soton.ac.uk⁸

¹Lecturer, Elizabeth School of London, United Kingdom

²Assistant Professor, Sree Buddha College of Engineering, India

³Assistant Professor, College of Engineering Adoor, India

⁴Assistant Professor, Sri Krishna College of Engineering and Technology, India

⁵Assistant Professor, College of Engineering Adoor, India

⁶Assistant Professor, Sree Buddha College of Engineering, India

⁷Assistant Professor, Sree Buddha College of Engineering, India

⁸Postdoctoral Research Fellow, University of Southampton, United Kingdom

Abstract: *Neuro-symbolic artificial intelligence (AI) has emerged as a promising paradigm that combines the statistical learning strengths of neural networks with the transparency and logical rigour of symbolic reasoning. Deep learning has achieved state-of-the-art results in perception and pattern recognition tasks, yet it suffers from challenges such as limited explainability, data inefficiency, and difficulties with compositional generalisation. Conversely, symbolic AI offers interpretability and deductive reasoning but fails to scale effectively in noisy and dynamic environments. By integrating these two paradigms, neuro-symbolic systems aim to produce intelligent agents capable of robust perception, logical inference, and human-aligned reasoning. This review explores the historical roots of neuro-symbolic AI, surveys the dominant methodological approaches, and analyses recent advances across natural language processing, robotics, and scientific discovery. In addition, we highlight the major limitations, ethical considerations, and future research directions necessary for achieving scalable and trustworthy neuro-symbolic systems.*

Keywords

Neuro-symbolic AI, deep learning, symbolic reasoning, explainable AI, hybrid intelligence.

I. INTRODUCTION

AI has become a cornerstone of contemporary technological progress, influencing research, industry, and society at scale. Its impact is increasingly visible in domains such as healthcare, finance, transportation, and education, where intelligent systems support decision-making and automation. This expansion has been fuelled primarily by advances in deep learning, which has enabled human-level or superhuman performance in perception, reasoning, and control tasks (Abbas Khan et al., 2025). Large-scale transformer architectures and foundation models have

further accelerated this transformation, underpinning applications in natural language generation, multimodal reasoning, and scientific discovery (Z. Chen et al., 2024). Despite these achievements, significant concerns remain regarding trust, robustness, and accountability, especially in safety-critical sectors (Yigit et al., 2025).

Deep learning models excel at recognising patterns in vast datasets, yet they are inherently opaque. Their decision-making processes are difficult to interpret, often leading to the characterisation of neural networks as “black boxes” (Hassija et al., 2024). This lack of transparency poses obstacles in domains where explanations are essential for adoption, such as medicine, law, and governance. Furthermore, these models are heavily dependent on large volumes of annotated data, are prone to overfitting, and frequently exhibit brittleness when exposed to out-of-distribution or adversarial inputs. While the capacity to approximate complex functions has enabled notable advances, their limitations highlight the insufficiency of deep learning as a complete model of intelligence (Ahmed et al., 2023).

Symbolic approaches to AI offer a contrasting set of strengths and weaknesses. Rooted in logic, rules, and structured knowledge representations, symbolic AI systems provide transparency, compositional reasoning, and the ability to manipulate abstract knowledge. These attributes made symbolic methods dominant in the early decades of AI research, particularly in expert systems for diagnosis and decision support. However, symbolic systems have struggled with scalability, adaptability, and robustness in dynamic, noisy environments. The knowledge acquisition bottleneck, which requires significant manual effort to construct and maintain symbolic rules and ontologies, further restricted their practical deployment (Eriksson et al., 1994).

The complementary nature of these paradigms has motivated renewed interest in hybrid approaches that can combine their

strengths. Neuro-symbolic AI aims to integrate the perceptual adaptability of neural networks with the interpretability and logical rigour of symbolic reasoning (Bhuyan et al., 2024). By leveraging both data-driven learning and structured knowledge, such systems promise enhanced data efficiency, improved robustness, and principled reasoning. This paradigm is increasingly viewed as a pathway towards more trustworthy and human-aligned AI.

Recent advances have revitalised the field, with differentiable logic programming, neural program induction, and graph-based reasoning now being explored as mechanisms to embed symbolic knowledge into deep architectures. Surveys categorise these approaches into key streams, including differentiable reasoning, abductive learning, and symbolic planning driven by large language models (Liang et al., 2025). Concrete demonstrations underscore this progress. For example, AlphaGeometry combined a large language model with a symbolic geometry engine to solve Olympiad-level mathematics problems, illustrating how symbolic reasoning can augment neural perception. Similarly, NeSyGPT integrated foundation models with symbolic logic, reducing manual engineering while improving interpretability. These developments indicate a shift from theoretical discussions towards scalable, practical frameworks.

The relevance of neuro-symbolic methods extends beyond research laboratories. In natural language processing, they have been applied to question answering and semantic parsing, where symbolic reasoning reduces hallucinations and improves factual accuracy. In robotics, perception modules powered by neural networks are increasingly combined with symbolic planners, enabling robots to perform tasks under explicit logical constraints. In healthcare, symbolic medical ontologies are being merged with neural imaging models to develop interpretable diagnostic systems (Vidal et al., 2025). These applications highlight the versatility of hybrid systems across diverse sectors. At the same time, societal and regulatory pressures are accelerating interest in interpretable AI. The European Union's proposed AI Act and similar global initiatives emphasise explainability, accountability, and safety as prerequisites for responsible AI deployment. Neuro-symbolic AI, by embedding symbolic reasoning within statistical learning systems, provides a natural avenue for satisfying these requirements. Industry analysts, have positioned neuro-symbolic AI within their 2025 AI Hype Cycle as a promising technology for enterprise contexts where regulation and accountability are paramount (Raza et al., 2025).

Nevertheless, open challenges remain. Current approaches face difficulties in scaling symbolic reasoning to large, unstructured datasets, and in balancing predictive accuracy with interpretability. The absence of standardised benchmarks impedes objective evaluation and comparison of competing methods. Furthermore, the computational cost of integrating neural and symbolic components raises questions regarding

energy efficiency and sustainability. Addressing these issues will require advances in methodology, interdisciplinary collaboration, and greater emphasis on reproducible research.

This review provides a structured and comprehensive survey of neuro-symbolic AI. It traces the historical development of the field, outlines its core methodological approaches, and synthesises recent advances in applications spanning natural language processing, robotics, healthcare, and scientific discovery. In addition, it identifies the limitations of current methods and highlights promising research directions, particularly in relation to scalability, efficiency, interpretability, and governance. By consolidating developments in this rapidly evolving area, the review aims to contribute to the advancement of AI systems that are both powerful and trustworthy.

The remainder of this paper is organised as follows. Section II reviews relevant literature, including foundational theories, current methodological approaches, and emerging frameworks for neuro-symbolic integration. Section III discusses recent advances and applications, drawing on case studies across multiple domains. Section IV outlines the challenges and open problems that remain unresolved, with emphasis on scalability, standardisation, and ethical considerations. Section V highlights future research directions that may shape the evolution of the field. Finally, Section VI concludes the paper with a summary of key insights.

II. LITERATURE SURVEY

The field of neuro-symbolic AI has developed through distinct intellectual, methodological, and application-driven phases. To provide a structured account of this evolution, the literature is reviewed across five interrelated strands. First, the foundational theories of symbolic and neural integration are revisited, highlighting early hybrid frameworks that attempted to reconcile logical reasoning with adaptive learning. Second, methodological innovations are examined, focusing on differentiable reasoning, probabilistic frameworks, and graph-based approaches that underpin contemporary research. Third, applications are surveyed across domains such as natural language processing, robotics, and healthcare, where hybrid architectures have been deployed in practice. Fourth, recent breakthrough frameworks, introduced between 2024 and 2025, are discussed as exemplars of large-scale, scalable neuro-symbolic systems. Finally, persistent challenges and open directions are considered, encompassing issues of scalability, benchmarking, interpretability, and governance.

A. Foundational Theories of Symbolic and Neural Integration

Early attempts at integrating symbolic and neural paradigms were motivated by recognition of their complementary strengths and weaknesses. Symbolic AI, exemplified by rule-based expert systems, offered transparency and logical rigour but was limited by the knowledge acquisition bottleneck and poor adaptability in dynamic environments (Peralta et al., 2025). In contrast, neural

networks provided adaptive pattern recognition and generalisation but were criticised for their opacity and lack of compositional reasoning (Hupkes et al., 2020). Garcez and colleagues showed that symbolic rules could be encoded within connectionist systems, preserving logical soundness and completeness. This was further advanced by Logic Tensor Networks, which embedded first-order logic into differentiable architectures. Such frameworks demonstrated that hybrid approaches could unify adaptive learning with principled reasoning. However, these contributions were constrained by limited computational resources and small-scale datasets. The trade-off between adaptability and interpretability thus remained unresolved, motivating subsequent methodological innovations.

B. Methodological Innovations in Neuro-Symbolic AI

With the resurgence of deep learning, new methodologies for hybrid reasoning were developed. Probabilistic logic frameworks combined symbolic rules with differentiable inference, enabling reasoning under uncertainty (Institute for Data Science, ETH Zürich, Switzerland et al., 2025). Logic Tensor Networks extended symbolic constraints to gradient-based optimisation, while abductive learning introduced mechanisms for handling incomplete or noisy data by reconciling observed evidence with logical consistency. Parallel advances in graph-based reasoning highlighted the importance of structured relational information. Graph neural networks (GNNs) have been shown to approximate logical inference, and hybrid models now exploit knowledge graphs to ground symbolic reasoning in data-rich contexts. Differentiable inductive logic programming (ILP) further expanded this frontier by enabling models to learn symbolic rules from examples. These methodological innovations reveal a field moving beyond proof-of-concept to scalable, domain-relevant frameworks. Nonetheless, challenges persist. While such models improve interpretability and robustness compared with purely neural architectures, they often incur significant computational overhead and lack standardised evaluation benchmarks. This limits reproducibility and complicates fair comparison across systems. The methodological literature thus underscores a persistent tension between interpretability, efficiency, and scalability.

C. Applications Across Domains

The practical relevance of neuro-symbolic AI is evident across several application domains. In natural language processing, symbolic constraints have been integrated into large language models to reduce hallucinations and improve factual consistency in question answering and semantic parsing (Huang et al., 2025). In robotics, hybrid frameworks combine neural perception with symbolic planning, enabling robots to execute compositional tasks under explicit logical constraints. In healthcare, symbolic ontologies have been merged with neural imaging systems to

support interpretable diagnostic frameworks, thereby increasing trust in machine-generated recommendations.

These applications highlight the versatility of hybrid methods and their ability to enhance trustworthiness in safety-critical domains. However, they also reveal persistent weaknesses: domain-specific symbolic engineering remains labour-intensive, and systems often fail to generalise across tasks. The application literature thus illustrates both the promise and the fragility of hybrid deployments, suggesting that broader adoption will depend on the development of domain-agnostic frameworks.

D. Recent Breakthrough Frameworks (2024–2025)

Recent years have witnessed breakthroughs that have elevated neuro-symbolic AI from experimental settings to high-profile demonstrations. AlphaGeometry integrated a large language model with a symbolic geometry engine to solve Olympiad-level problems, achieving results previously unattainable by either paradigm alone (C. Zhang et al., 2024). Its successor, AlphaGeometry2, achieved gold-medallist performance, underscoring the scalability of the hybrid approach. Similarly, NeSyGPT introduced a framework for combining foundation models with symbolic reasoning, reducing the reliance on handcrafted rules while improving interpretability. Energy-based neuro-symbolic models have also been developed for reasoning over knowledge graphs, further diversifying the methodological toolkit.

These breakthroughs demonstrate the feasibility of large-scale neuro-symbolic integration and provide benchmarks that validate theoretical claims. Yet they also reveal limitations: systems such as AlphaGeometry remain highly resource-intensive, and their applicability to heterogeneous real-world problems is uncertain. The literature thus signals a transition from theory to practice, but one that is not yet accompanied by sustainable deployment strategies.

E. Challenges and Open Directions

Despite substantial progress, neuro-symbolic AI faces unresolved challenges. Scaling symbolic reasoning to unstructured, large-scale data remains a central bottleneck (Yang et al., 2024). The absence of standardised benchmarks impedes comparative evaluation, risking fragmentation of the field. Energy efficiency and sustainability have emerged as additional concerns, as hybrid systems often require significant computational resources. From a governance perspective, policy frameworks such as the European Union's AI Act stress accountability, transparency, and explainability, yet current architectures lack systematic auditing mechanisms.

To summarise, while neuro-symbolic AI has progressed from foundational theories to practical frameworks, it continues to face unresolved challenges in scalability, reproducibility, and regulatory alignment. These limitations highlight the necessity of evaluating hybrid systems not only in theoretical settings but

also in realistic, domain-specific deployments. Building on this foundation, the following discussion turns to recent advances and applications, examining case studies from natural language processing, robotics, healthcare, and other domains to illustrate both the promise and the boundaries of current approaches.

III. RECENT ADVANCES AND APPLICATIONS

Neuro-symbolic AI has progressed from theoretical explorations to concrete demonstrations across diverse domains. These advances illustrate how hybrid frameworks can address limitations of purely neural or symbolic paradigms while also exposing new challenges. This section reviews recent applications in natural language processing (NLP), robotics and embodied intelligence, healthcare and biomedical research, and scientific discovery, each of which highlights the benefits and constraints of neuro-symbolic integration.

A. Natural Language Processing

One of the most prominent application areas of neuro-symbolic AI is NLP, where hybrid methods have been employed to reduce factual inconsistency and improve interpretability. Neural language models, particularly large-scale transformers, are powerful at text generation but prone to hallucinations and logical errors (Ye et al., 2023). To mitigate this, symbolic reasoning has been integrated into neural pipelines for tasks such as semantic parsing, question answering, and knowledge-grounded dialogue. For example, augmenting language models with first-order logic constraints has been shown to reduce hallucination rates and improve factual alignment in QA benchmarks. Similarly, neuro-symbolic semantic parsers that combine neural sequence-to-sequence models with symbolic grammars have demonstrated improved compositional generalisation.

Despite these gains, trade-offs remain. Hybrid NLP systems often require domain-specific symbolic resources such as ontologies or logical rules, which limit scalability across diverse languages and contexts. Furthermore, symbolic integration may slow inference compared to purely neural systems, raising questions about efficiency in real-world deployments. The literature therefore indicates a promising but incomplete transition towards interpretable, trustworthy language models.

B. Robotics and Embodied Intelligence

In robotics, the integration of symbolic reasoning with neural perception has enabled advances in planning, manipulation, and autonomous navigation. Neural networks excel at perceptual tasks such as object recognition and scene understanding, while symbolic planners provide compositional reasoning and task-level decision-making (YanfangZhou et al., 2025). Systems combining symbolic knowledge bases with perceptual modules allow robots to reason about goals, constraints, and object affordances. More recent work has embedded neuro-symbolic

reasoning into reinforcement learning pipelines, allowing robots to learn policies that respect explicit logical constraints.

The key advantage of such approaches lies in their ability to combine low-level perception with high-level reasoning, making robotic behaviour more reliable in structured environments. However, scalability remains a limitation, as symbolic reasoning engines often struggle with real-time performance in complex, dynamic contexts. Additionally, extensive symbolic engineering is typically required, restricting generalisation beyond the training environment.

C. Healthcare and Biomedical Research

Healthcare has become a critical proving ground for neuro-symbolic AI, particularly due to the demand for interpretability in clinical decision-making. Symbolic ontologies, such as the Unified Medical Language System (UMLS), have been integrated with neural imaging models to develop hybrid diagnostic tools that offer transparency alongside predictive accuracy (Almarshadi et al., 2024). For instance, neuro-symbolic diagnostic frameworks combining convolutional neural networks with symbolic medical knowledge graphs have achieved higher interpretability in radiological decision support. Similar approaches have been applied to drug discovery, where symbolic reasoning enforces biochemical constraints on generative molecular models.

The strengths of these systems lie in their ability to improve trustworthiness and regulatory compliance by providing human-understandable explanations. Nevertheless, challenges persist in the form of integration costs, data heterogeneity, and computational overhead. Moreover, clinical adoption requires rigorous validation under real-world conditions, which remains limited. Thus, while neuro-symbolic frameworks hold substantial promise for healthcare, translation into clinical practice is still at an early stage.

D. Scientific Discovery

Beyond applied domains, neuro-symbolic AI has demonstrated potential in accelerating scientific discovery. Hybrid frameworks have been used to formalise hypotheses, guide simulations, and generate interpretable explanations of complex phenomena (Vidal et al., 2025). Notably, AlphaGeometry integrated a large language model with a symbolic reasoning engine to solve Olympiad-level geometry problems, achieving a milestone in formal mathematical reasoning. Similarly, NeSyGPT combined foundation models with symbolic logic programming to enable interpretable scientific inference across chemistry and physics datasets.

These advances illustrate the capability of hybrid systems to extend beyond data-fitting and into hypothesis-driven reasoning. Yet they also highlight resource-intensive demands, as systems such as AlphaGeometry require substantial computational infrastructure. Moreover, the reliance on benchmark-driven

evaluation leaves open questions about applicability to real-world scientific challenges.

Taken together, recent advances demonstrate that neuro-symbolic AI can enhance robustness, interpretability, and domain-specific reliability across multiple domains. At the same time, these systems reveal persistent trade-offs between generalisability, computational efficiency, and symbolic engineering, which constrain their broader adoption. To understand these limitations more fully, the following discussion turns to the key challenges and open problems that continue to shape the trajectory of neuro-symbolic AI.

IV. FUTURE RESEARCH DIRECTIONS

The trajectory of neuro-symbolic AI suggests that its long-term impact will depend on addressing existing limitations while seizing emerging opportunities. Several directions appear particularly salient: scaling symbolic reasoning to large, unstructured data; improving computational efficiency and sustainability; developing multimodal and general-purpose frameworks; strengthening human-AI collaboration; and aligning hybrid systems with ethical and governance requirements.

A. Scalability and Data Efficiency

A central challenge remains the scaling of symbolic reasoning to large, heterogeneous datasets. Current hybrid models struggle with unstructured, high-dimensional inputs such as video, speech, and sensor data (Fitas, 2025). Future research must therefore explore frameworks that enable symbolic abstractions to be derived automatically from raw inputs, reducing the reliance on manual rule engineering. Advances in differentiable logic programming and graph-based learning suggest promising avenues for scalable abstraction. Moreover, transfer learning techniques may allow symbolic constraints to be generalised across domains, thereby enhancing reusability and reducing development costs.

B. Computational Efficiency and Sustainability

The computational cost of hybrid systems remains a pressing concern, particularly in light of the growing energy demands of large-scale neural models (W. Chen et al., 2024). Neuro-symbolic frameworks, which introduce additional reasoning layers, risk exacerbating this burden. Future work must therefore emphasise energy-efficient architectures and hardware-software co-optimisation. Emerging research in neuromorphic computing and energy-aware symbolic reasoning holds potential to reconcile hybrid performance with sustainability goals. Establishing benchmarks that explicitly incorporate energy metrics would also help steer the field towards more responsible computational practices.

C. Multimodality and General-Purpose Reasoning

Most current applications of neuro-symbolic AI are domain-specific, with limited transferability. An important direction lies in the integration of multimodal inputs: text, vision, audio, and structured knowledge, into unified reasoning frameworks (Bhuyan et al., 2025). By grounding symbolic reasoning in multimodal embeddings, systems could move towards more general-purpose intelligence. Recent advances in foundation models provide a natural basis for such integration, but symbolic layers must be adapted to handle multimodal uncertainty and compositionality. Achieving this would bring hybrid systems closer to the goal of robust, human-like reasoning.

D. Human-AI Collaboration and Interactivity

Another promising avenue is the design of neuro-symbolic systems that enhance human-AI collaboration. Symbolic reasoning enables the communication of explanations and justifications in human-interpretable terms, which is particularly relevant for decision-making in domains such as medicine, law, and engineering (X. Zhang & Sheng, 2024). Future research should explore interactive learning frameworks in which humans can dynamically update symbolic rules or provide feedback to guide neural learning. Such human-in-the-loop designs could not only improve interpretability but also accelerate trust and adoption in sensitive domains.

E. Governance, Ethics, and Accountability

Finally, the future of neuro-symbolic AI will be shaped by societal and regulatory demands for transparency, accountability, and fairness. As governments adopt AI-specific regulations, hybrid systems must align with requirements for explainability and auditability (Du, 2025). Embedding symbolic reasoning into neural architectures offers a pathway to meet these obligations, but robust mechanisms for verification, bias detection, and compliance monitoring are still lacking. Future research must therefore engage interdisciplinary perspectives, drawing from law, philosophy, and policy studies to ensure that neuro-symbolic systems evolve responsibly.

Taken together, these prospective directions underline that the future of neuro-symbolic AI rests not only on technical progress but also on its capacity to align with societal expectations of trustworthiness and responsibility. By addressing scalability, efficiency, multimodality, and governance in an integrated manner, the field may progress from niche demonstrations to mainstream adoption. With these opportunities and challenges in mind, the final section concludes this review by synthesising the key insights and outlining their broader implications for the evolution of AI.

V. CONCLUSION

This review has traced the trajectory of neuro-symbolic AI, from early theoretical attempts at combining symbolic logic and neural learning to recent advances in large-scale, domain-specific applications. The literature reveals that hybrid

approaches offer unique strengths in interpretability, compositionality, and trustworthiness, while also exposing persistent weaknesses in scalability, reproducibility, and computational efficiency.

The discussion of applications in natural language processing, robotics, healthcare, and scientific discovery illustrates both the promise and the fragility of current systems. Breakthrough frameworks such as AlphaGeometry and NeSyGPT demonstrate that hybrid reasoning can scale to complex tasks, yet they remain resource-intensive and domain-constrained. These observations emphasise the importance of methodological innovation and standardised evaluation if neuro-symbolic AI is to progress beyond proof-of-concept.

Looking forward, the field must balance technical ambition with societal responsibility. Scalability, energy efficiency, multimodal integration, human–AI collaboration, and governance emerge as critical priorities. By consolidating knowledge across foundational theories, methodologies, applications, and future directions, this review seeks to provide a roadmap for advancing neuro-symbolic AI towards systems that are not only powerful but also transparent, trustworthy, and aligned with human values.

REFERENCES

- Abbas Khan, M., Khan, H., Omer, M. F., Ullah, I., & Yasir, M. (2025). Impact of Artificial Intelligence on the Global Economy and Technology Advancements. In S. El Hajjami, K. Kaushik, & I. U. Khan (Eds.), *Artificial General Intelligence (AGI) Security* (pp. 147–180). Springer Nature Singapore. https://doi.org/10.1007/978-981-97-3222-7_7
- Ahmed, S. F., Alam, Md. S. B., Hassan, M., Rozbu, M. R., Ishtiak, T., Rafa, N., Mofijur, M., Shawkat Ali, A. B. M., & Gandomi, A. H. (2023). Deep learning modelling techniques: Current progress, applications, advantages, and challenges. *Artificial Intelligence Review*, 56(11), 13521–13617. <https://doi.org/10.1007/s10462-023-10466-8>
- Almarshadi, N. F., Alharbi, M. M. A., Alharbi, M. S., Alwaked, W. A., Alsayer, H. M., & Alhussain, S. M. (2024). Clinical and Medical Coding: A New Pathway for Automation-An Updated Review. *Journal of Medical and Life Science*, 6(4), 633–651. <https://doi.org/10.21608/jmals.2024.413890>
- Bhuyan, B. P., Ramdane-Cherif, A., Singh, T. P., & Tomar, R. (2025). Neuro-Symbolic AI in Various Domains. In B. P. Bhuyan, A. Ramdane-Cherif, T. P. Singh, & R. Tomar, *Neuro-Symbolic Artificial Intelligence* (Vol. 1176, pp. 311–324). Springer Nature Singapore. https://doi.org/10.1007/978-981-97-8171-3_17
- Bhuyan, B. P., Ramdane-Cherif, A., Tomar, R., & Singh, T. P. (2024). Neuro-symbolic artificial intelligence: A survey. *Neural Computing and Applications*, 36(21), 12809–12844. <https://doi.org/10.1007/s00521-024-09960-z>
- Chen, W., Ma, X., Wang, Z., Li, W., Fan, C., Zhang, J., Que, X., & Li, C. (2024). Exploring neuro-symbolic AI applications in geoscience: Implications and future directions for mineral prediction. *Earth Science Informatics*, 17(3), 1819–1835. <https://doi.org/10.1007/s12145-024-01278-7>
- Chen, Z., Xu, L., Zheng, H., Chen, L., Tolba, A., Zhao, L., Yu, K., & Feng, H. (2024). Evolution and Prospects of Foundation Models: From Large Language Models to Large Multimodal Models. *Computers, Materials & Continua*, 80(2), 1753–1808. <https://doi.org/10.32604/cmc.2024.052618>
- Du, J. (2025). *Toward Responsible And Beneficial Ai: Comparing Regulatory And Guidance-Based Approaches* (Version 1). arXiv. <https://doi.org/10.48550/ARXIV.2508.00868>
- Eriksson, H., Puerta, A. R., & Musen, M. A. (1994). Generation of knowledge-acquisition tools from domain ontologies. *International Journal of Human-Computer Studies*, 41(3), 425–453. <https://doi.org/10.1006/ijhc.1994.1067>
- Fitas, R. (2025). Neuro-Symbolic AI for Advanced Signal and Image Processing: A Review of Recent Trends and Future Directions. *IEEE Access*, 13, 143360–143376. <https://doi.org/10.1109/ACCESS.2025.3598909>
- Hassija, V., Chamola, V., Mahapatra, A., Singal, A., Goel, D., Huang, K., Scardapane, S., Spinelli, I., Mahmud, M., & Hussain, A. (2024). Interpreting Black-Box Models: A Review on Explainable Artificial Intelligence. *Cognitive Computation*, 16(1), 45–74. <https://doi.org/10.1007/s12559-023-10179-8>
- Huang, L., Yu, W., Ma, W., Zhong, W., Feng, Z., Wang, H., Chen, Q., Peng, W., Feng, X., Qin, B., & Liu, T. (2025). A Survey on Hallucination in Large Language Models: Principles, Taxonomy, Challenges, and Open Questions. *ACM Transactions on Information Systems*, 43(2), 1–55. <https://doi.org/10.1145/3703155>
- Hupkes, D., Dankers, V., Mul, M., & Bruni, E. (2020). Compositionality Decomposed: How do Neural Networks Generalise? *Journal of Artificial Intelligence Research*, 67, 757–795. <https://doi.org/10.1613/jair.1.11674>
- Institute for Data Science, ETH Zürich, Switzerland, Rossi, S., Johnson, S., & Department of Electrical Engineering and Computer Sciences, UC Berkeley, USA. (2025). NEUROSymbolic AI: MERGING DEEP LEARNING AND LOGICAL REASONING FOR ENHANCED EXPLAINABILITY. *International Journal of Advanced Artificial Intelligence Research*, 2(06), 1–7. <https://doi.org/10.55640/ijaair-v02i06-01>

- Liang, B., Wang, Y., & Tong, C. (2025). AI Reasoning in Deep Learning Era: From Symbolic AI to Neural-Symbolic AI. *Mathematics*, 13(11), 1707. <https://doi.org/10.3390/math13111707>
- Peralta, A., Olivas, J. A., Romero, F. P., & Navarro, P. (2025). Integration of Fuzzy Techniques and Formal Representation of Domain and Expert Knowledge in AI Systems: A Comprehensive Review. *Contemporary Mathematics*, 1660–1681. <https://doi.org/10.37256/cm.6220256231>
- Raza, S., Sapkota, R., Karkee, M., & Emmanouilidis, C. (2025). *TRiSM for Agentic AI: A Review of Trust, Risk, and Security Management in LLM-based Agentic Multi-Agent Systems* (Version 3). arXiv. <https://doi.org/10.48550/ARXIV.2506.04133>
- Vidal, M.-E., Chudasama, Y., Huang, H., Purohit, D., & Torrente, M. (2025). Integrating Knowledge Graphs with Symbolic AI: The Path to Interpretable Hybrid AI Systems in Medicine. *Journal of Web Semantics*, 84, 100856. <https://doi.org/10.1016/j.websem.2024.100856>
- Yanfang Zhou, Y., Liu, Y., Li, X., Zhao, Y., Wang, X., Tian, J., Li, Z., & Xu, X. (2025). Metagent-P: A Neuro-Symbolic Planning Agent with Metacognition for Open Worlds. *Findings of the Association for Computational Linguistics: ACL 2025*, 22747–22764. <https://doi.org/10.18653/v1/2025.findings-acl.1169>
- Yang, K., Poesia, G., He, J., Li, W., Lauter, K., Chaudhuri, S., & Song, D. (2024). *Formal Mathematical Reasoning: A New Frontier in AI* (Version 1). arXiv. <https://doi.org/10.48550/ARXIV.2412.16075>
- Ye, H., Liu, T., Zhang, A., Hua, W., & Jia, W. (2023). *Cognitive Mirage: A Review of Hallucinations in Large Language Models* (Version 1). arXiv. <https://doi.org/10.48550/ARXIV.2309.06794>
- Yigit, Y., Ferrag, M. A., Ghanem, M. C., Sarker, I. H., Maglaras, L. A., Chrysoulas, C., Moradpoor, N., Tihanyi, N., & Janicke, H. (2025). Generative AI and LLMs for Critical Infrastructure Protection: Evaluation Benchmarks, Agentic AI, Challenges, and Opportunities. *Sensors*, 25(6), 1666. <https://doi.org/10.3390/s25061666>
- Zhang, C., Song, J., Li, S., Liang, Y., Ma, Y., Wang, W., Zhu, Y., & Zhu, S.-C. (2024). *Proposing and solving olympiad geometry with guided tree search* (Version 1). arXiv. <https://doi.org/10.48550/ARXIV.2412.10673>
- Zhang, X., & Sheng, V. S. (2024). *Neuro-Symbolic AI: Explainability, Challenges, and Future Trends* (Version 1). arXiv. <https://doi.org/10.48550/ARXIV.2411.04383>